



ORACLE®

MySQL Performance & Reporting



MySQL Performance Tuning for Embedded Systems

- Query tuning is the most critical part of MySQL performance tuning
- Can give 100% to 1000% (or more!) performance improvement
- Critical for reporting and embedded systems
 - Reporting queries are usually complex
 - Embedded systems runs “as is”, sometimes without a DBA



Agenda

- Understanding Query Performance
- Indexes and Temporary Tables in MySQL
- GROUP BY / ORDER BY Optimizations
- Subqueries
- Reporting Queries and Summary Tables
- Full Text Search Queries
- What's New in MySQL 5.6 Beta



How to Deal with Slow Performance

- Find slow queries
- Profile/Explain
- Fix queries
 - Add indexes
 - Re-write queries
 - Add summary tables
 - Cache queries
- = Better performance!



Main Query Performance Problems

- Full table scans (no index)
- Temporary tables
- Filesort



Using EXPLAIN to Analyze Queries

- Find out how the query optimizer would improve a

SELECT query

```
mysql> EXPLAIN select * from City where Name = 'London'\G
      id: 1
    select_type: SIMPLE
        table: City
         type: ALL
possible_keys: NULL
         key: NULL
      key_len: NULL
         ref: NULL
        rows: 4079
      Extra: Using where
```

Query Profiling

```
mysql> set profiling =1;  
mysql> show profiles\G
```

```
***** 1. row *****
```

```
Query_ID: 4
```

```
Duration: 0.00069300
```

```
Query: select * from sbtest where k >0 order by  
pad desc limit 10
```

```
***** 2. row *****
```

```
Query_ID: 5
```

```
Duration: 0.00044800
```

```
Query: explain select * from sbtest where k >0  
order by pad desc limit 10
```

Using Query Profiling

```
mysql> set profiling =1;
mysql> select pad, c, count(*) from sbtest where k = 0
      group by pad, c limit 10;
... (22.90 sec)
```

```
mysql> show profile;
```

Status	Duration
starting	0.000017
init	0.000086
checking permissions	0.000016
Opening tables	0.000037
System lock	0.000021
init	0.000034
optimizing	0.000014
statistics	0.000126
preparing	0.000027
Creating tmp table	0.000420
executing	0.000005

| Copying to tmp table | 22.904328 |

Sorting result	0.000053
Sending data	0.000028
end	0.000005
removing tmp table	0.000210
end	0.000008
query end	0.000008
closing tables	0.000018
freeing items	0.000030
logging slow query	0.000005
cleaning up	0.000015

Temporary table

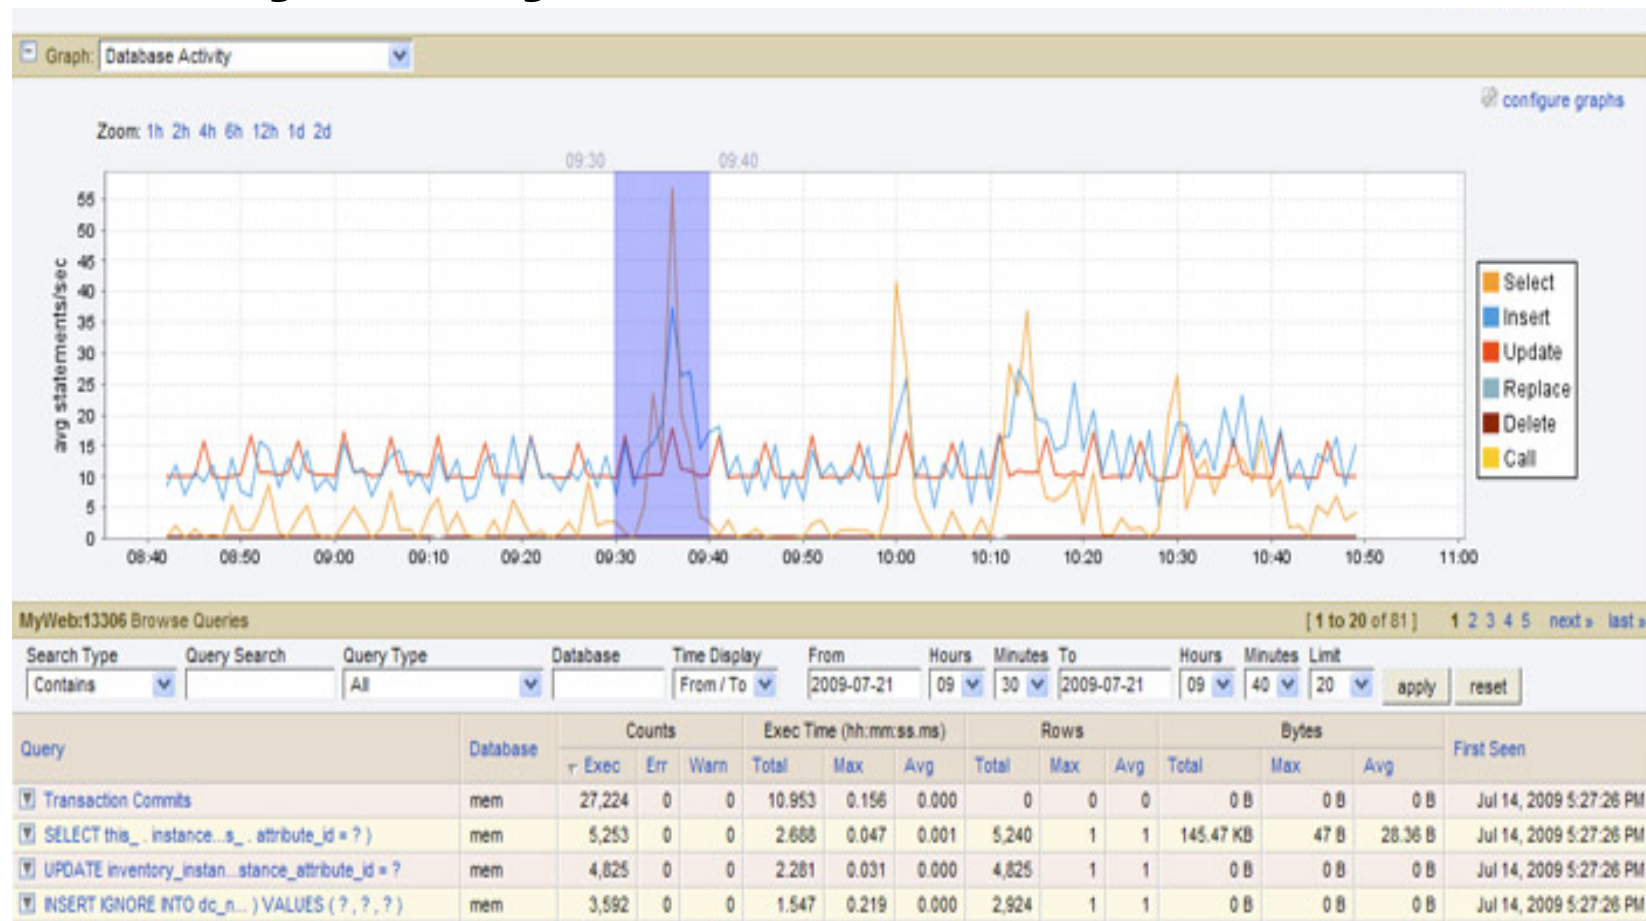
Query Profiling: CPU

```
mysql> select * from sbtest where k = 0 order by pad desc
      limit 10;
... (22.90 sec)
mysql> show profile cpu;
```

Status	Duration	CPU_user	System_cpu
starting	0.000019	0.000000	0.000000
init	0.000079	0.000000	0.000000
checking permissions	0.000016	0.000000	0.000000
Opening tables	0.000039	0.000000	0.000000
System lock	0.000021	0.000000	0.000000
init	0.000036	0.000000	0.000000
optimizing	0.000015	0.000000	0.000000
statistics	0.000117	0.000000	0.000000
preparing	0.000028	0.000000	0.000000
executing	0.000007	0.000000	0.000000
Sorting result	3.431279	3.430478	0.000000
Sending data	0.000130	0.001000	0.000000
end	0.000009	0.000000	0.000000
query end	0.000008	0.000000	0.000000
closing tables	0.000020	0.000000	0.000000
freeing items	0.000413	0.000000	0.000000
cleaning up	0.000015	0.000000	0.000000

Sort is CPU-
intensive

MySQL Enterprise Monitor: Query Analyzer, I



<http://www.mysql.com/products/enterprise/monitor.html>

MySQL Enterprise Monitor: Query Analyzer, II

Canonical Query Example Query Explain Query 908 B 165 B 1

Overview of information collected and aggregated for queries of this form.

Alias
None specified.

Canonical Form
[truncated](#) | [full](#) | [formatted](#)

```
SELECT
  COUNT( message_id ) AS cnt,
  DATE_FORMAT(FROM_UNIXTIME(datetime), ?) AS
    fmtdate
FROM
  logs_amavis
WHERE
  datetime >= ? AND datetime <= ?
GROUP BY
  process_mode ASC, process_type ASC, fmtdate
  ASC
```

Execution Time Statistics

Max Time	Min Time	Avg Time	Total Time	Standard Deviation
11.123	0.274	2.140	12.838	

Row Statistics

Max Rows	Min Rows	Avg Rows	Total Rows	Standard Deviation	Total Size	Max Size
408	37	129	776	82	24.21 KB	12.8 KB

Number of Executions
6

Time Span
From Oct 27, 2008 10:16:07 AM to Oct 27, 2008 10:46:07 AM.

[hide](#)

[expand »](#)

Aggregated Form

Canonical Query

Execution and Row
Statistics

MySQL Enterprise Monitor: Query Analyzer, III

Canonical Query Example Query Explain Query 908 B 165 B 134 B

The query with the longest execution time during the Time Span (usually the slowest but not always).

Sampled Query
[truncated](#) | [full](#) | [formatted](#)

```
SELECT
  COUNT( message_id ) AS cnt,
  date_format(from_unixtime(datetime), "%Y-%m-%d")
  AS fmtdate
FROM
  logs_amavis
WHERE
  datetime >= 1224498843 and datetime <= 1225080000
GROUP BY
  process_mode ASC, process_type ASC, fmtdate ASC
```

Execution Time
11,122 ms

Date
Oct 27, 2008 10:35:01 AM

User
root

Thread ID
298107

From Host
192.168.0.2:22717

To Host
127.0.0.1:3306

Comments

hide

expand »

Example Query

Execution Time for this
Example

MySQL Enterprise Monitor: Query Analyzer, IV

MySQL Explain Plan

Canonical Query

Example Query


Explain Query

908 B

165 B

134.67 B

10:37:01 AM

 alias

Explain of a query that occurred during the Time Span (usually the slowest but not always).

Explain

id	select_type	table	type	possible_keys	key	key_len	ref	rows	extra
1	SIMPLE	logs_amavis	index	null	mdatetimepmpt	47	null	276087	Using where; Using index; Using temporary; Using filesort

hide

expand »

MySQL Query Analyzer Documentation:

<http://dev.mysql.com/doc/mysql-monitor/2.0/en/mem-query-analyzer-queries.html>



Composite Indexes in MySQL



Composite Indexes

- MySQL chooses 1 (best) index per table
 - With some exceptions...
- The more unique values the better
 - Do not index status, gender, etc
- Order of fields inside index matters
 - (in most cases)
- **"Where region = 'US' and date_added > '2010-05-01' "**
 - Index on (region, date_added) preferred

Composite Indexes

- `"Where region = 'US' and date_added > '2010-05-01' "`
- Index (region, date_added):
 1. MySQL will "jump" to index leaf where Region='US'
 2. Scan date_added range starting with the leaf
- Constant + range: Put constant first, range second

GROUP BY Queries



GROUP BY and Temporary Tables

- How many cities in each country?

```
mysql> explain select CountryCode, count(*) from City
        group by CountryCode\G
```

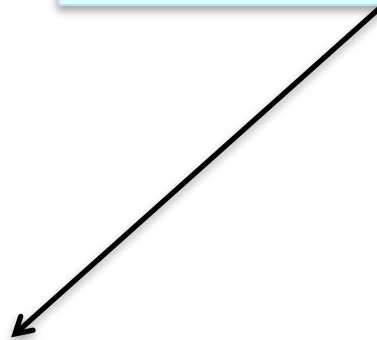
```
***** 1. row
```

```
      id: 1
  select_type: SIMPLE
        table: City
         type: ALL
possible_keys: NULL
         key: NULL
      key_len: NULL
         ref: NULL
        rows: 4079
```

```
Extra: Using temporary; Using filesort
```

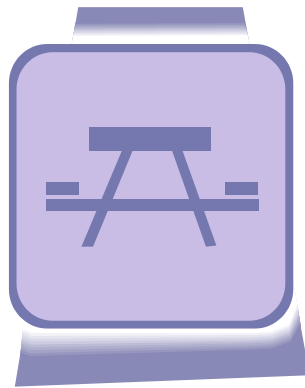
```
1 row in set (0.00 sec)
```

Temporary tables
are slow!





Temporary Tables: Theory





Temporary Tables

- Main performance issues
- MySQL can create temporary tables when query uses:
 - GROUP BY
 - Range + ORDER BY
 - Some other expressions
- 2 types of temporary tables
 - MEMORY
 - On-disk



Temporary Tables

- First, MySQL tries to create temporary table in memory
 - **tmp_table_size**
 - maximum size for in Memory temporary tables
 - **max_heap_table_size**
 - Sets the maximum size for **MEMORY** tables
- If ($\text{tmp_table} > \text{tmp_table_size}$ OR $\text{tmp_table} > \text{max_heap_table_size}$)
- { convert to MyISAM temporary table on disk }



Temporary Tables

- MEMORY engine does not support BLOB/TEXT
- `select blob_field from table group by field1`
- `select concat(...string>512 chars) group by field1`
 - Create on-disk temporary table right away



Temporary Tables: Profiling

- Watch:
 - **Created_tmp_tables** – number of temporary table MySQL created in both *RAM and DISK*
 - **Created_tmp_disk_tables** - number of temporary table MySQL created on *DISK*

```
mysql> show session status like 'created%';
```

Variable_name	Value
Created_tmp_disk_tables	1
Created_tmp_files	0
Created_tmp_tables	10

```
3 rows in set (0.00 sec)
```



Temporary Tables: Practice





Air Traffic Statistics Table for Testing

5M rows, ~2G in size

```
CREATE TABLE `ontime_2010` (  
  `YearD` int(11) DEFAULT NULL,  
  `MonthD` tinyint(4) DEFAULT NULL,  
  `DayofMonth` tinyint(4) DEFAULT NULL,  
  `DayOfWeek` tinyint(4) DEFAULT NULL,  
  `Carrier` char(2) DEFAULT NULL,  
  `Origin` char(5) DEFAULT NULL,  
  `DepDelayMinutes` int(11) DEFAULT NULL,  
  ...  
) ENGINE=InnoDB DEFAULT CHARSET=latin1
```

[http://www.transtats.bts.gov/DL_SelectFields.asp?
Table_ID=236&DB_Short_Name=On-Time](http://www.transtats.bts.gov/DL_SelectFields.asp?Table_ID=236&DB_Short_Name=On-Time)

ORACLE

The World's Most Popular Open Source Database



GROUP BY Query Example

- Find maximum delay for flights on Sunday
- Group by airline

```
select max(DepDelayMinutes) ,  
       carrier, dayofweek  
from ontime_2010  
where dayofweek = 7  
group by Carrier, dayofweek
```

GROUP BY Query Example

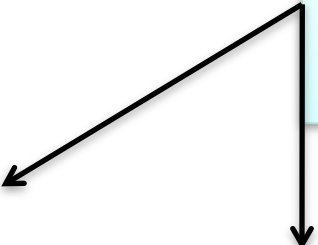
```
select max(DepDelayMinutes), carrier, dayofweek
from ontime_2010
where dayofweek = 7
group by Carrier, dayofweek
```

```
type: ALL
possible_keys: NULL
key: NULL
key_len: NULL
ref: NULL
```

```
rows: 4833086
```

```
Extra: Using where; Using temporary; Using
filesort
```

Full table scan!
Temporary table!



Fixing Full Table Scan

```
mysql> alter table ontime_2010 add key (dayofweek);
```

```
mysql> explain select max(DepDelayMinutes), Carrier,  
    dayofweek from ontime_2010  
where dayofweek =7 group by Carrier, dayofweek\G
```

```
      type: ref  
possible_keys: DayOfWeek  
      key: DayOfWeek  
    key_len: 2  
      ref: const  
     rows: 817258
```

Many rows scanned!
Temporary table!

Extra: Using where; Using temporary; Using filesort

ORACLE

GROUP BY: Adding Covered Index

```
mysql> alter table ontime_2010 add key covered  
      (dayofweek, Carrier, DepDelayMinutes);
```

```
mysql> explain select max(DepDelayMinutes), Carrier,  
      dayofweek from ontime_2010  
where dayofweek =7 group by Carrier, dayofweek\G
```

...

```
possible_keys: DayOfWeek,covered
```

```
      key: covered
```

```
key_len: 2
```

```
      ref: const
```

```
rows: 905138
```

```
Extra: Using where; Using index
```

No temporary table!



When Covered Indexes Aren't Good ...

```
mysql> explain select max(DepDelayMinutes), Carrier,  
    dayofweek from ontime_2010  
where dayofweek > 3  
group by Carrier,  
dayofweek\G  
...
```

```
    type: range  
possible_keys: covered  
    key: covered  
key_len: 2  
    ref: NULL  
rows: 2441781  
Extra: Using where; Using index; Using  
temporary; Using filesort
```

Range scan





ORDER BY and filesort



ORDER BY and filesort

- Find 10 cities in the US with the largest population

```
mysql> explain select district, name, population from  
      City where CountryCode = 'USA' order by population  
      desc limit 10\G
```

```
      table: City  
      type: ALL  
possible_keys: NULL  
      key: NULL  
key_len: NULL  
      ref: NULL  
      rows: 4079  
Extra: Using where; Using filesort
```


Fixing Filesort: Adding Index

```
mysql> alter table City add key my_sort2 (CountryCode,  
      population);
```

```
mysql> explain select district, name, population from  
      City where CountryCode = 'USA' order by population  
      desc limit 10\G
```

```
table: City  
type: ref  
key: my_sort2  
key_len: 3  
ref: const  
rows: 207  
Extra: Using where
```

No filesort

Sorting and Limit

```
mysql> alter table ontime_2010 add key (DepDelayMinutes);  
Query OK, 0 rows affected (38.68 sec)
```

```
mysql> explain select * from ontime_2010  
where dayofweek in (6,7) order by DepDelayMinutes desc  
limit 10\G
```

```
           type: index  
possible_keys: DayOfWeek,covered  
           key: DepDelayMinutes  
        key_len: 5  
           ref: NULL  
          rows: 24  
       Extra: Using where
```

```
10 rows in set (0.00 sec)
```

1. Index is sorted
2. Scan whole table in the order of the index
3. Filter results
4. Stop after finding 10 rows matching the “where” condition



Subqueries and Joins Optimizations



Subqueries, I

- Subquery inside select

```
SELECT (SELECT s1 FROM t2) FROM t1;
```

- Subquery inside where

```
SELECT * FROM t1 WHERE  
column1 in (SELECT column2 FROM t2);
```

- Subquery in FROM and joins

```
SELECT sb1,sb2,sb3 FROM (SELECT s1 AS sb1, s2  
    AS sb2, s3*2 AS sb3 FROM t1) AS sb  
WHERE sb1 > 1;
```

Subqueries, II

- Subquery inside select

```
SELECT (SELECT max(s1) FROM t2), ... FROM t1
where mydate>now();
10000 rows in set
```

- Will execute subquery SELECT max(s1) FROM t2 10,000 times
- Can rewrite it:

```
SELECT max(s1) into @m FROM t2
SELECT @m, ... FROM t1 where ...
```

Subqueries, III

- Subquery inside where
`SELECT * FROM t1 WHERE
column1 in (SELECT column2 FROM t2) ;`
- Will not use index on **column1**
 - <http://bugs.mysql.com/bug.php?id=8139>
- Can rewrite query as join

Subqueries, IV

- Subquery in FROM and joins

```
SELECT sb1,sb2,sb3 FROM (SELECT s1 AS sb1, s2  
    AS sb2, s3*2 AS sb3 FROM t1) AS sb  
WHERE sb1 > 1;
```

- MySQL will create a temporary table for (SELECT s1 AS sb1, s2 AS sb2, s3*2 AS sb3 FROM t1) with no indexes
- Rewrite as join

Joins, I

- Use INT to join tables rather than VARCHAR

```
select * from city ct
```

```
join country c on ct.CountryCode = c.Code;
```

```
CREATE TABLE `country` (  
  `Code` varchar(50) NOT NULL DEFAULT '',  
  ...  
  PRIMARY KEY (`Code`)  
) ENGINE=InnoDB DEFAULT CHARSET=latin1
```

```
CREATE TABLE `city` (  
  `ID` int(11) NOT NULL AUTO_INCREMENT,  
  ...  
  `Name` char(35) NOT NULL DEFAULT '',  
  `CountryCode` varchar(50) NOT NULL DEFAULT '',  
  PRIMARY KEY (`ID`),  
  KEY `CountryCode` (`CountryCode`),  
  CONSTRAINT `city_ibfk_1` FOREIGN KEY (`CountryCode`) REFERENCES `country` (`Code`)
```

Use code_id INT instead

Joins, II

- Use INT to join tables rather than VARCHAR

```
select * from city ct
```

```
join country c on ct.CountryCode = c.Code;
```

```
CREATE TABLE `country` (  
  `Code_ID` INT NOT NULL AUTO_INCREMENT DEFAULT '',  
  ...  
  PRIMARY KEY (`Code_ID`)  
) ENGINE=InnoDB DEFAULT CHARSET=latin1
```

Use code_id INT instead

```
CREATE TABLE `city` (  
  `ID` int(11) NOT NULL AUTO_INCREMENT,  
  ...  
  `Name` char(35) NOT NULL DEFAULT '',  
  `CountryCode_ID` INT NOT NULL DEFAULT '',  
  PRIMARY KEY (`ID`),  
  KEY `CountryCode` (`CountryCode`),  
  CONSTRAINT `city_ibfk_1` FOREIGN KEY (`CountryCode`) REFERENCES `country` (`Code`)
```



Reporting Queries



ORACLE®

The World's Most Popular Open Source Database

Reporting Example: Table

```
mysql> CREATE TABLE `lineitem` (  
  `l_shipdate` date NOT NULL,  
  `l_orderkey` int(11) NOT NULL,  
  `l_partkey` int(11) NOT NULL,  
  `l_suppkey` int(11) NOT NULL,  
  `l_linenumber` int(11) NOT NULL,  
  `l_quantity` decimal(15,2) NOT NULL,  
  `l_extendedprice` decimal(15,2) NOT NULL,  
  `l_discount` decimal(15,2) NOT NULL,  
  `l_tax` decimal(15,2) NOT NULL,  
  `l_returnflag` char(1) NOT NULL,  
  `l_linestatus` char(1) NOT NULL,  
  KEY `lineitem_fk2` (`l_suppkey`),  
  KEY `lineitem_fk3` (`l_partkey`, `l_suppkey`),  
  KEY `li_shp_dt_idx` (`l_shipdate`),  
  KEY `li_com_dt_idx` (`l_commitdate`),  
  KEY `li_rcpt_dt_idx` (`l_receiptdate`),  
);
```

Reporting Example: Query

- Group by year(date)

```
mysql> explain
select sum(l_extendedprice), year(l_shipdate) as yr
from lineitem group by yr limit 10\G
***** 1. row

      id: 1
  select_type: SIMPLE
        table: lineitem
         type: ALL
possible_keys: NULL
          key: NULL
       key_len: NULL
         ref: NULL
        rows: 116771866
   Extra: Using temporary; Using filesort
```

year(field) = calculated
MySQL can't use index

Adding Year/Month

```
mysql> alter table lineitem add yr year, add key(yr);  
mysql> update lineitem set yr = year(l_shipdate);
```

```
mysql> explain select sum(l_extendedprice), yr  
from lineitem group by yr desc limit 10\G  
***** 1. row *****  
      id: 1  
    select_type: SIMPLE  
      table: lineitem  
        type: index  
possible_keys: NULL  
         key: yr  
    key_len: 2  
         ref: NULL  
        rows: 116771866  
      Extra:  
1 row in set (0.00 sec)
```

No temporary, no filesort
Add covered index for
better performance

Dimension Table

```
CREATE TABLE dates (  
    date_id          INT UNSIGNED NOT NULL PRIMARY KEY,  
    date             DATE NOT NULL,  
    day_of_week      INT NOT NULL,  
    month            VARCHAR(10) NOT NULL,  
    month_day        INT NOT NULL,  
    year             INT NOT NULL,  
    UNIQUE KEY `date` (`date`)  
);
```

Flexible for different queries

```
mysql> alter table lineitem  
add date_id INT UNSIGNED NOT NULL,  
add key (date_id);
```

Dimension Table: Test

```
mysql> explain select d.year, d.month_day, sum(l_extendedprice)
      from dates d, lineitem1 l where l.date_id = d.date_id group
      by d.year, d.month_day\G
```

```
      table: l
      type: ALL
possible_keys: date_id
      key: NULL
     key_len: NULL
      ref: NULL
     rows: 116771866
```

```
      Extra: Using temporary; Using filesort
```

```
***** 2. row *****
```

```
      table: d
      type: eq_ref
possible_keys: PRIMARY,date_id
      key: PRIMARY
     key_len: 4
      ref: tpch.l.date_id
     rows: 1
     Extra:
```

Flexible, but slow!

No index, temporary table

Summary Tables, I

```
mysql> create table lineitem_summary as  
select year(l_shipdate) as yr,  
month(l_shipdate) as mon,  
sum(l_extendedprice) as revenue,  
count(*) as num_orders  
from lineitem group by yr, mon;
```

Query OK, 365 rows affected
Records: 365 Duplicates: 0

Data is already aggregated in summary table

Summary Tables, II

- Aggregate by Year, based on summary table

```
mysql> select yr,  
sum(l_extendedprice) as revenue,  
count(*) as num_orders  
from lineitem_summary group by yr;
```

Only 356 rows for 1 year or 3560 for 10 years

Data already aggregated
Small number of records
= Queries are much faster!



Summary Tables, III

- Advantages
 - Significantly faster for queries
 - Smaller number of rows
- Disadvantages
 - Needs to be updated: cron or manually
 - More data to store

Make sense for reporting
Use MySQL slave server

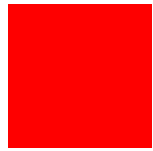


MySQL Full Text Index

- Available in MyISAM tables before MySQL 5.6
 - InnoDB supports full text index in MySQL 5.6
- Natural language search and Boolean search
- 4 char per word by default
- Stop word list by default
- Frequency-based ranking
 - Distance between words is not counted

MySQL Full Text: Creating Full Text Indexes

```
mysql> CREATE TABLE articles (  
-> id INT UNSIGNED AUTO_INCREMENT  
-> NOT NULL PRIMARY KEY,  
-> title VARCHAR(200),  
-> body TEXT,  
-> FULLTEXT (title,body)  
-> ) engine=MyISAM;
```



MySQL Full Text: Natural Language Mode

```
mysql> SELECT * FROM articles  
-> WHERE MATCH (title,body)  
-> AGAINST ('database' IN NATURAL  
LANGUAGE MODE);
```

```
+-----+-----+  
| title | body |  
+-----+-----+  
| MySQL vs. YourSQL | In the following database comparison ... |  
| MySQL Tutorial | DBMS stands for DataBase ... |
```

In natural language mode:
Default sorting by relevance!



MySQL Full Text: Boolean Mode

```
mysql> SELECT * FROM articles  
-> WHERE MATCH (title,body)  
-> AGAINST ('cat AND dog' IN  
BOOLEAN MODE) ;
```

No default sorting in Boolean mode!



New in MySQL 5.6: Features, Performance and Monitoring



ORACLE®

The World's Most Popular Open Source Database

Monitoring: Performance Schema, I

```
| Tables_in_performance_schema |  
+-----+  
| events_statements_current |  
| events_statements_history |  
| events_statements_history_long |  
| events_statements_summary_by_host_by_event_name |  
| events_statements_summary_by_thread_by_event_name |  
| events_statements_summary_by_user_by_event_name |  
| events_statements_summary_by_user_host_by_event_name |  
| events_statements_summary_global_by_event_name |
```

Statement list for all queries
Used for query profiling

Monitoring: Performance Schema, II

- How to find queries creating disk temporary tables

```
mysql> select * from events_statements_history_long  
where   CREATED_TMP_DISK_TABLES > 0 limit 10\G
```

...

```
      SQL_TEXT: SELECT DISTINCT c from sbtest where id  
between 847399 and 847499 order by c
```

```
      ROWS_SENT: 1
```

```
      ROWS_EXAMINED: 103
```

```
CREATED_TMP_DISK_TABLES: 1
```

```
      CREATED_TMP_TABLES: 1
```

```
      SELECT_FULL_JOIN: 0
```

```
      SELECT_FULL_RANGE_JOIN: 0
```

```
      SELECT_RANGE: 1
```

```
      SELECT_RANGE_CHECK: 0
```

```
      SELECT_SCAN: 0
```

```
      SORT_MERGE_PASSES: 0
```

```
      SORT_RANGE: 0
```

```
      SORT_ROWS: 1
```

```
      SORT_SCAN: 1
```

```
      NO_INDEX_USED: 0
```

```
      NO_GOOD_INDEX_USED: 0
```

MySQL 5.6 “labs release”
from labs.mysql.com/

Monitoring: Performance Schema, III

- List of queries creating disk temporary tables

```
mysql> select sql_text, count(*) as cnt from
       events_statements_history_long
where   CREATED_TMP_DISK_TABLES > 0
group by sql_text order by cnt desc limit 10;
```

sql_text	cnt
SELECT DISTINCT c from sbtest where id between 242012 and 242112 order by c	2
SELECT DISTINCT c from sbtest where id between 797388 and 797488 order by c	2
SELECT DISTINCT c from sbtest where id between 973150 and 973250 order by c	1
SELECT DISTINCT c from sbtest where id between 478783 and 478883 order by c	1
SELECT DISTINCT c from sbtest where id between 967035 and 967135 order by c	1
SELECT DISTINCT c from sbtest where id between 602102 and 602202 order by c	1
SELECT DISTINCT c from sbtest where id between 123827 and 123927 order by c	1
SELECT DISTINCT c from sbtest where id between 980527 and 980627 order by c	1
SELECT DISTINCT c from sbtest where id between 450354 and 450454 order by c	1
SELECT DISTINCT c from sbtest where id between 674804 and 674904 order by c	1

10 rows in set (0.04 sec)



Optimizer Improvements, I

- Index Condition Pushdown
 - Moves more of the processing for WHERE clauses to the storage engine
 - = less I/O overhead
 - = less internal communication overhead
 - InnoDB, MyISAM, and NDBCLUSTER
- Multi-Range Read (MRR)
 - Faster to read data sequentially than to do random accesses
 - MRR-
 - Scans one or more index ranges used in query
 - Sorts the associated disk blocks for the row data
 - Reads those disk blocks using larger sequential I/O requests.
 - Works for all storage engines




Optimizer Improvements, II

- File Sort Optimization
 - ORDER BY *non_indexed_column* LIMIT n
 - speeds up the sort when the contents of n rows can fit into the sort buffer
 - Works for all storage engines



MySQL 5.6 InnoDB Improvements

- Full Text Search Indexes for InnoDB tables
- Split Kernel Mutex
 - Remove bottlenecks in busy systems
- Multi-Threaded Purge and Separate Flush Thread
 - Increases performance and provides better scalability
- Persistent Optimizer Stats
 - Provides improved accuracy of InnoDB index statistics and consistency across MySQL restarts
- New InnoDB monitoring tables in INFORMATION_SCHEMA
- Much more!



The presentation is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, and timing of any features or functionality described for Oracle's products remains at the sole discretion of Oracle.