Oracle® Database Database Performance Tuning Guide





Oracle Database Database Performance Tuning Guide, 23ai

F47529-06

Copyright © 2007, 2024, Oracle and/or its affiliates.

Contributing Authors: Glenn Maxey, Manish Garodia, Rajesh Bhatiya, Immanuel Chan, Lance Ashdown

Contributors: Hermann Baer, Deba Chatterjee, Maria Colgan, Mikael Fries, Prabhaker Gongloor, Kevin Jernigan, Sue K. Lee, William Lee, David McDermid, Uri Shaft, Oscar Suro, Trung Tran, Sriram Vrinda, Yujun Wang

This software and related documentation are provided under a license agreement containing restrictions on use and disclosure and are protected by intellectual property laws. Except as expressly permitted in your license agreement or allowed by law, you may not use, copy, reproduce, translate, broadcast, modify, license, transmit, distribute, exhibit, perform, publish, or display any part, in any form, or by any means. Reverse engineering, disassembly, or decompilation of this software, unless required by law for interoperability, is prohibited.

The information contained herein is subject to change without notice and is not warranted to be error-free. If you find any errors, please report them to us in writing.

If this is software, software documentation, data (as defined in the Federal Acquisition Regulation), or related documentation that is delivered to the U.S. Government or anyone licensing it on behalf of the U.S. Government, then the following notice is applicable:

U.S. GOVERNMENT END USERS: Oracle programs (including any operating system, integrated software, any programs embedded, installed, or activated on delivered hardware, and modifications of such programs) and Oracle computer documentation or other Oracle data delivered to or accessed by U.S. Government end users are "commercial computer software," "commercial computer software documentation," or "limited rights data" pursuant to the applicable Federal Acquisition Regulation and agency-specific supplemental regulations. As such, the use, reproduction, duplication, release, display, disclosure, modification, preparation of derivative works, and/or adaptation of i) Oracle programs (including any operating system, integrated software, any programs embedded, installed, or activated on delivered hardware, and modifications of such programs), ii) Oracle computer documentation and/or iii) other Oracle data, is subject to the rights and limitations specified in the license contained in the applicable contract. The terms governing the U.S. Government's use of Oracle cloud services are defined by the applicable contract for such services. No other rights are granted to the U.S. Government.

This software or hardware is developed for general use in a variety of information management applications. It is not developed or intended for use in any inherently dangerous applications, including applications that may create a risk of personal injury. If you use this software or hardware in dangerous applications, then you shall be responsible to take all appropriate fail-safe, backup, redundancy, and other measures to ensure its safe use. Oracle Corporation and its affiliates disclaim any liability for any damages caused by use of this software or hardware in dangerous applications.

Oracle®, Java, MySQL, and NetSuite are registered trademarks of Oracle and/or its affiliates. Other names may be trademarks of their respective owners.

Intel and Intel Inside are trademarks or registered trademarks of Intel Corporation. All SPARC trademarks are used under license and are trademarks or registered trademarks of SPARC International, Inc. AMD, Epyc, and the AMD logo are trademarks or registered trademarks of Advanced Micro Devices. UNIX is a registered trademark of The Open Group.

This software or hardware and documentation may provide access to or information about content, products, and services from third parties. Oracle Corporation and its affiliates are not responsible for and expressly disclaim all warranties of any kind with respect to third-party content, products, and services unless otherwise set forth in an applicable agreement between you and Oracle. Oracle Corporation and its affiliates will not be responsible for any loss, costs, or damages incurred due to your access to or use of third-party content, products, or services, except as set forth in an applicable agreement between you and Oracle.

Contents

	Preface	
	Audience	XV
	Documentation Accessibility	XV
	Related Documents	XV
	Conventions	xvi
	Changes in This Release for Oracle Database Performance T	Tuning Guide
	Changes in Oracle Database Release 23ai	xvii
Part	Database Performance Fundamentals	
1	Performance Tuning Overview	
	Introduction to Performance Tuning	1-1
	Performance Planning	1-1
	Instance Tuning	1-1
	Performance Principles	1-2
	Baselines	1-2
	The Symptoms and the Problems	1-2
	When to Tune	1-3
	SQL Tuning	1-4
	Query Optimizer and Execution Plans	1-4
	Introduction to Performance Tuning Features and Tools	1-4
	Automatic Performance Tuning Features	1-5
	Additional Oracle Database Tools	1-5
	V\$ Performance Views	1-6
2	Designing and Developing for Performance	
	Oracle Methodology	2-1
	Understanding Investment Options	2-1
	Understanding Scalability	2-2
	Understanding Scalability	



	What is Scalability?	2-2
	System Scalability	2-3
	Factors Preventing Scalability	2-4
	System Architecture	2-5
	Hardware and Software Components	2-5
	Hardware Components	2-5
	Software Components	2-6
	Configuring the Right System Architecture for Your Requirements	2-7
	Application Design Principles	2-9
	Simplicity In Application Design	2-9
	Data Modeling	2-10
	Table and Index Design	2-10
	Appending Columns to an Index or Using Index-Organized Tables	2-11
	Using a Different Index Type	2-11
	Finding the Cost of an Index	2-12
	Serializing within Indexes	2-12
	Ordering Columns in an Index	2-12
	Using Views	2-13
	SQL Execution Efficiency	2-13
	Implementing the Application	2-14
	Trends in Application Development	2-16
	Workload Testing, Modeling, and Implementation	2-16
	Sizing Data	2-17
	Estimating Workloads	2-17
	Application Modeling	2-18
	Testing, Debugging, and Validating a Design	2-18
	Deploying New Applications	2-19
	Rollout Strategies	2-19
	Performance Checklist	2-19
3	Performance Improvement Methods	
	The Oracle Performance Improvement Method	3-1
	Steps in the Oracle Performance Improvement Method	3-2
	A Sample Decision Process for Performance Conceptual Modeling	3-3
	Top Ten Mistakes Found in Oracle Systems	3-4
	Emergency Performance Methods	3-5
	Steps in the Emergency Performance Method	3-6
4	Configuring a Database for Performance	
	Performance Considerations for Initial Instance Configuration	4-1



	Initialization Parameters	4-1
	Undo Space	4-3
	Redo Log Files	4-4
	Tablespaces	4-4
	Creating and Maintaining Tables for Optimal Performance	4-6
	Table Compression	4-6
	Reclaiming Unused Space	4-8
	Indexing Data	4-8
	Performance Considerations for Shared Servers	4-9
	Identifying and Reducing Contention Using the Dispatcher-Specific Views	4-9
	Identifying Contention for Shared Servers	4-11
	Improved Client Connection Performance Due to Prespawned Processes	4-12
Part	II Diagnosing and Tuning Database Performance	
5	Measuring Database Performance	
	About Database Statistics	5-1
	Time Model Statistics	5-1
	Active Session History Statistics	5-2
	Wait Events Statistics	5-3
	Session and System Statistics	5-4
	Interpreting Database Statistics	5-4
	Using Hit Ratios	5-5
	Using Wait Events with Timed Statistics	5-5
	Using Wait Events without Timed Statistics	5-5
	Using Idle Wait Events	5-6
	Comparing Database Statistics with Other Factors	5-6
	Using Computed Statistics	5-6
6	Gathering Database Statistics	
	About Gathering Database Statistics	6-1
	Automatic Workload Repository	6-2
	Snapshots	6-2
	Baselines	6-3
	Fixed Baselines	6-3
	Moving Window Baselines	6-3
	Baseline Templates	6-4
	Space Consumption	6-4
	Adaptive Thresholds	6-5
	Percentage of Maximum Thresholds	6-6



Significance Level Thresholds	6-6
Managing the Automatic Workload Repository	6-7
Enabling the Automatic Workload Repository	6-7
Managing Snapshots	6-8
User Interfaces for Managing Snapshots	6-8
Creating Snapshots	6-9
Dropping Snapshots	6-9
Modifying Snapshot Settings	6-10
Managing Baselines	6-11
User Interface for Managing Baselines	6-12
Creating a Baseline	6-12
Dropping a Baseline	6-13
Renaming a Baseline	6-14
Displaying Baseline Metrics	6-14
Resizing the Default Moving Window Baseline	6-14
Managing Baseline Templates	6-15
User Interfaces for Managing Baseline Templates	6-15
Creating a Single Baseline Template	6-16
Creating a Repeating Baseline Template	6-16
Dropping a Baseline Template	6-17
Transporting Automatic Workload Repository Data to Another System	6-18
Exporting AWR Data	6-18
Importing AWR Data	6-19
Using Automatic Workload Repository Views	6-20
Managing Automatic Workload Repository in a Multitenant Environment	6-21
Categorization of AWR Data in a Multitenant Environment	6-22
AWR Data Storage and Retrieval in a Multitenant Environment	6-22
Viewing AWR Data in a Multitenant Environment	6-25
Managing Automatic Workload Repository in Active Data Guard Standby Databases	6-26
Configuring the Remote Management Framework (RMF)	6-28
Managing Snapshots for Active Data Guard Standby Databases	6-32
Viewing AWR Data in Active Data Guard Standby Databases	6-35
Generating Automatic Workload Repository Reports	6-36
User Interface for Generating an AWR Report	6-36
Generating an AWR Report Using the Command-Line Interface	6-36
Generating an AWR Report for the Local Database	6-37
Generating an AWR Report for a Specific Database	6-38
Generating an AWR Report for the Local Database in Oracle RAC	6-39
Generating an AWR Report for a Specific Database in Oracle RAC	6-39
Generating an AWR Report for a SQL Statement on the Local Database	6-41
Generating an AWR Report for a SQL Statement on a Specific Database	6-41
Generating Performance Hub Active Report	6-43



About Performance Hub Active Report Tabs About Performance Hub Active Report Types Command-Line User Interface for Generating a Performance Hub Active Report	6-43
Command-Line User Interface for Generating a Performance Hub Active Report	6-44
	6-45
Generating a Performance Hub Active Report Using a SQL Script	6-45
Performance Workload Analysis	6-46
Automatic Performance Diagnostics	
Overview of the Automatic Database Diagnostic Monitor	7-1
ADDM Analysis	7-2
Using ADDM with Oracle Real Application Clusters	7-4
Using ADDM in a Multitenant Environment	7-4
Enabling ADDM in a Pluggable Database	7-6
Real-Time ADDM Analysis	7-7
Real-Time ADDM Connection Modes	7-7
Real-Time ADDM Triggers	7-7
Real-Time ADDM Trigger Controls	7-8
ADDM Analysis Results	7-9
Reviewing ADDM Analysis Results: Example	7-9
Setting Up ADDM	7-10
Diagnosing Database Performance Problems with ADDM	7-11
Running ADDM in Database Mode	7-11
Running ADDM in Instance Mode	7-12
Running ADDM in Partial Mode	7-13
Displaying an ADDM Report	7-13
Re-Running Failed ADDM Tasks	7-14
FAILED_AUTO_TASKS_REPORT Function	7-14
REEXECUTE_FAILED_AUTO_TASKS Procedure	7-15
ADDM Views	7-15
ADDM Spotlight	7-16
Monitoring Databases with Performance Hub	
About Performance Hub	8-1
What is Performance Hub?	8-1
How Performance Hub Works	8-1
Performance Hub Features	8-2
ASH Analytics	8-3
ASH Sample Resolution	8-4
Activity Tables	8-4
SQL Monitoring	8-4



Real-time SQL Monitoring	8-5
Graphical Explain Plan	8-5
Automatic Database Diagnostic Monitor (ADDM)	8-5
Workload	8-6
Monitored and Analyzed Time Indications	8-6
Regions	8-6
Blocking Sessions	8-6
Exadata	8-7
Exadata Details	8-7
Top Activity Lite	8-8
Simplified UI Options	8-9
AWR Snapshot Collection	8-10
Reports	8-10
Automatic Workload Repository (AWR)	8-10
Active Session History (ASH)	8-11
Performance Hub	8-11
Using Performance Hub to Analyze Database Performance	8-12
Access	8-12
For Autonomous Database	8-12
For Base Database Service and Exadata Databases	8-13
For External Database Service	8-14
View	8-15
Average Active Sessions Data	8-15
SQL Monitoring Report	8-17
Real-time SQL Activity	8-18
Graphical Explain Plan	8-20
Workload Metrics	8-21
Blocking and Waiting Sessions	8-22
ADDM Data	8-24
Exadata Statistics	8-28
Top Activity Lite Data	8-28
Manage	8-29
AWR Settings	8-29
Download	8-30
Automatic Workload Repository (AWR) Report	8-31
Active Sessions History (ASH) Report	8-32
Performance Hub Report	8-32

Part III Tuning Database Memory



9 Database Memory Allocation

	About Database Memory Caches and Other Memory Structures	9-1
	Database Memory Management Methods	9-2
	Automatic Memory Management	9-3
	Automatic Shared Memory Management	9-4
	Manual Shared Memory Management	9-5
	Automatic PGA Memory Management	9-5
	Manual PGA Memory Management	9-5
	Summary of Memory Management Methods	9-6
	Using Automatic Memory Management	9-8
	Monitoring Memory Management	9-8
10	Tuning the System Global Area	
	Using Unified Memory	10-1
	Using Automatic Shared Memory Management	10-1
	User Interfaces for Setting the SGA_TARGET Parameter	10-2
	Setting the SGA_TARGET Parameter in Oracle Enterprise Manager Cloud Control	10-2
	Setting the SGA_TARGET Parameter in the Command-Line Interface	10-2
	Setting the SGA_TARGET Parameter	10-3
	Enabling Automatic Shared Memory Management	10-3
	Disabling Automatic Shared Memory Management	10-3
	Unified Program Global Area	10-3
	Sizing the SGA Components Manually	10-4
	SGA Sizing Unit	10-5
	Maximum Size of the SGA	10-5
	Application Considerations	10-5
	Operating System Memory Use	10-6
	Reduce Paging	10-6
	Fit the SGA into Main Memory	10-6
	Allow Adequate Memory to Individual Users	10-7
	Iteration During Configuration	10-7
	Monitoring Shared Memory Management	10-7
	Improving Query Performance with the In-Memory Column Store	10-8
	Enabling High Performance Data Streaming with the Memoptimized Rowstore	10-9
	About the Memoptimized Rowstore	10-9
	Using Fast Ingest	10-9
	Prerequisites for Fast Ingest Table	10-12
	Enabling a Table for Fast Ingest	10-14
	Specifying a Hint for Using Fast Ingest for Data Inserts	10-14
	Disabling a Table for Fast Ingest	10-15



10-15
10-17
10-17
10-18
10-20
10-21
10-21
11-1
11-1
11-2
11-4
11-5
11-5
11-6
11-6
11-7
11-7
11-8
11-8
11-8
11-9
11-9
11-9
11-10
11-11
11-11
11-12
11-13
11-13
11-14
11-14
11-15
11-15
11-16
12-1
12-1



Shared Pool Concepts	12-1
Library Cache Concepts	12-2
Data Dictionary Cache Concepts	12-3
SQL Sharing Criteria	12-3
Using the Shared Pool	12-4
Use Shared Cursors	12-5
Use Single-User Logon and Qualified Table Reference	12-6
Use PL/SQL	12-6
Avoid Performing DDL Operations	12-6
Cache Sequence Numbers	12-6
Control Cursor Access	12-6
Controlling Cursor Access Using OCI	12-7
Controlling Cursor Access Using Oracle Precompilers	12-7
Controlling Cursor Access Using SQLJ	12-8
Controlling Cursor Access Using JDBC	12-8
Controlling Cursor Access Using Oracle Forms	12-8
Maintain Persistent Connections	12-8
Configuring the Shared Pool	12-8
Sizing the Shared Pool	12-9
Using Library Cache Statistics	12-9
Using Shared Pool Advisory Statistics	12-11
Using Dictionary Cache Statistics	12-13
Increasing Memory Allocated to the Shared Pool	12-15
Reducing Memory Allocated to the Shared Pool	12-15
Deallocating Cursors	12-15
Caching Session Cursors	12-16
About the Session Cursor Cache	12-16
Enabling the Session Cursor Cache	12-17
Sizing the Session Cursor Cache	12-17
Sharing Cursors	12-19
About Cursor Sharing	12-19
Forcing Cursor Sharing	12-19
Keeping Large Objects to Prevent Aging	12-20
Configuring the Reserved Pool	12-21
Sizing the Reserved Pool	12-22
Increasing Memory Allocated to the Reserved Pool	12-22
Reducing Memory Allocated to the Reserved Pool	12-23
Configuring the Large Pool	12-23
Configuring the Large Pool for Shared Server Architecture	12-24
Configuring the Large Pool for Parallel Query	12-25
Sizing the Large Pool	12-25
Limiting Memory Use for User Sessions	12-26



13 Tuning the Result Cache

About the Result Cache	13-1
Server Result Cache Concepts	13-1
Benefits of Using the Server Result Cache	13-1
Understanding How the Server Result Cache Works	13-2
Client Result Cache Concepts	13-3
Benefits of Using the Client Result Cache	13-4
Understanding How the Client Result Cache Works	13-4
Brownouts and Automatic Exclusion for the Result Cache	13-5
Configuring the Result Cache	13-5
Configuring the Server Result Cache	13-6
Sizing the Server Result Cache Using Initialization Parameters	13-6
Managing the Server Result Cache Using DBMS_RESULT_CACHE	13-7
Configuring the Client Result Cache	13-8
Setting the Result Cache Mode	13-9
Requirements for the Result Cache	13-11
Read Consistency Requirements	13-12
Query Parameter Requirements	13-12
Restrictions for the Result Cache	13-12
Specifying Queries for Result Caching	13-12
Using SQL Result Cache Hints	13-13
Using the RESULT_CACHE Hint	13-13
Using the NO_RESULT_CACHE Hint	13-13
Using the RESULT_CACHE Hint in Views	13-13
Using Result Cache Table Annotations	13-14
Using the DEFAULT Table Annotation	13-15
Using the FORCE Table Annotation	13-15
Monitoring the Result Cache	13-16
Tuning the Process Global Area	
About the Process Global Area	14-1
Work Area Sizes	14-1
Sizing the Program Global Area Using Automatic Memory Management	14-2
Configuring Automatic PGA Memory Management	14-3
Setting the Initial Value for PGA_AGGREGATE_TARGET	14-3
Monitoring Automatic PGA Memory Management	
Using the V\$PGASTAT View	14-5
Using the V\$PROCESS View	14-7



14

	Using the V\$PROCESS_MEMORY View	14-7
	Using the V\$SQL_WORKAREA_HISTOGRAM View	14-8
	Using the V\$WORKAREA_ACTIVE View	14-10
	Using the V\$SQL_WORKAREA View	14-10
	Tuning PGA_AGGREGATE_TARGET	14-12
	Enabling Automatic Generation of PGA Performance Advisory Views	14-12
	Using the V\$PGA_TARGET_ADVICE View	14-13
	Using the V\$PGA_TARGET_ADVICE_HISTOGRAM View	14-15
	Using the V\$SYSSTAT and V\$SESSTAT Views	14-16
	Tutorial: How to Tune PGA_AGGREGATE_TARGET	14-17
	Sizing the Program Global Area by Specifying an Absolute Limit	14-18
	Sizing the Program Global Area Using the PGA_AGGREGATE_LIMIT Parameter	14-18
	Sizing the Program Global Area Using the Resource Manager	14-19
Par	t IV Managing System Resources	
15	I/O Configuration and Design	
13	About I/O	15-1
	I/O Configuration	15-2
	Lay Out the Files Using Operating System or Hardware Striping	15-2
	Requested I/O Size	15-2
	Concurrency of I/O Requests	15-3
	Alignment of Physical Stripe Boundaries with Block Size Boundaries	15-4
	Manageability of the Proposed System	15-4
	Manually Distributing I/O	15-5
	When to Separate Files	15-5
	Tables, Indexes, and TEMP Tablespaces	15-6
	Redo Log Files	15-6
	Archived Redo Logs	15-6
	Three Sample Configurations	15-7
	Stripe Everything Across Every Disk	15-7
	Move Archive Logs to Different Disks	15-8
	Move Redo Logs to Separate Disks	15-8
	Oracle Managed Files	15-8
	Choosing Data Block Size	15-9
	Reads	15-9
	Writes	15-10
	Block Size Advantages and Disadvantages	15-10
	I/O Calibration Inside the Database	15-10
	Prerequisites for I/O Calibration	15-11



	Running I/O Calibration	15-11
	I/O Calibration with the Oracle Orion Calibration Tool	15-13
	Introduction to the Oracle Orion Calibration Tool	15-13
	Orion Test Targets	15-14
	Orion for Oracle Administrators	15-14
	Getting Started with Orion	15-14
	Orion Input Files	15-15
	Orion Parameters	15-16
	Orion Required Parameter	15-16
	Orion Optional Parameters	15-18
	Orion Command Line Samples	15-20
	Orion Configuration for Exadata Storage Cells	15-21
	Orion Output Files	15-22
	Orion Sample Output Files	15-22
	Orion Troubleshooting	15-25
16	Managing Operating System Resources	
	Understanding Operating System Performance Issues	16-1
	Using Operating System Caches	16-1
	Asynchronous I/O	16-2
	FILESYSTEMIO_OPTIONS Initialization Parameter	16-2
	Limiting Asynchronous I/O in NFS Server Environments	16-3
	Improving I/O Performance Using Direct NFS Client	16-3
	Memory Usage	16-4
	Buffer Cache Limits	16-4
	Parameters Affecting Memory Usage	16-4
	Using Operating System Resource Managers	16-4
	Resolving Operating System Issues	16-6
	Performance Hints on UNIX-Based Systems	16-6
	Performance Hints on Windows Systems	16-6
	Performance Hints on HP OpenVMS Systems	16-6
	Understanding CPU	16-7
	Resolving CPU Issues	16-8
	Finding and Tuning CPU Utilization	16-8
	Checking Memory Management	16-9
	Checking I/O Management	16-10
	Checking Network Management	16-10
	Checking Process Management	16-10
	Managing CPU Resources Using Oracle Database Resource Manager	16-11
	Managing CPU Resources Using Instance Caging	16-12



17 Working With Automatic Materialized Views

Overview of Automatic Materialized Views	17-1
Workload Information Provided by the Object Activity Tracking System	17-2
Data Dictionary Views That Provide Information About Automatic Materialized Views and OATS	17-2
The DBMS_AUTO_MV Package	17-4
The DBMS_ACTIVITY Package	17-12
Glossary	

Index



Preface

This preface contains these topics:

- Audience
- Documentation Accessibility
- Related Documents
- Conventions

Audience

This document is intended for database administrators (DBAs) who are responsible for the operation, maintenance, and performance of Oracle Database. This guide describes how to use Oracle Database performance tools to optimize database performance. This guide also describes performance best practices for creating an initial database and includes performance-related reference information.

See Also:

- Oracle Database SQL Tuning Guide for information about how to optimize and tune SQL performance
- Oracle Database 2 Day + Performance Tuning Guide to learn how to use Oracle Enterprise Manager Cloud Control (Cloud Control) to tune database performance

Documentation Accessibility

For information about Oracle's commitment to accessibility, visit the Oracle Accessibility Program website at http://www.oracle.com/pls/topic/lookup?ctx=acc&id=docacc.

Access to Oracle Support

Oracle customers that have purchased support have access to electronic support through My Oracle Support. For information, visit http://www.oracle.com/pls/topic/lookup?ctx=acc&id=info or visit http://www.oracle.com/pls/topic/lookup?ctx=acc&id=trs if you are hearing impaired.

Related Documents

Before reading this guide, you should be familiar with the following documents:

- Oracle Database Concepts
- Oracle Database Administrator's Guide



- Oracle Multitenant Administrator's Guide
- Oracle Database 2 Day DBA
- Oracle Database 2 Day + Performance Tuning Guide

To learn how to tune data warehouse environments, see *Oracle Database Data Warehousing Guide*.

Conventions

The following text conventions are used in this document:

Convention	Meaning
boldface	Boldface type indicates graphical user interface elements associated with an action, or terms defined in text or the glossary.
italic	Italic type indicates book titles, emphasis, or placeholder variables for which you supply particular values.
monospace	Monospace type indicates commands within a paragraph, URLs, code in examples, text that appears on the screen, or text that you enter.



Changes in This Release for Oracle Database Performance Tuning Guide

This preface contains:

Changes in Oracle Database Release 23ai

Changes in Oracle Database Release 23ai

The following are changes in *Oracle Database Performance Tuning Guide* for Oracle Database release 23ai.

New Features

The following features are new in this release:

- ADDM Spotlight provides a visual representation of critical performance categories, rates the overall impact on system performance, and provides recommendations.
 - Monitoring Databases with Performance Hub provides the ability to monitor database activity, diagnose issues, and tune queries to improve the performance of Oracle databases.
- Enhanced Result Cache provides better scalability, better control of what is being cached, and broadens the applicability to cache results beyond the limits of pure-memory storage.



See Setting the Result Cache Mode



Part I

Database Performance Fundamentals

This part contains the following chapters:

- Performance Tuning Overview
- Designing and Developing for Performance
- Performance Improvement Methods
- Configuring a Database for Performance



1

Performance Tuning Overview

This chapter provides an introduction to performance tuning and contains the following sections:

- Introduction to Performance Tuning
- Introduction to Performance Tuning Features and Tools

Introduction to Performance Tuning

This guide provides information about tuning Oracle Database for performance. Topics discussed in this guide include:

- · Performance Planning
- Instance Tuning
- SQL Tuning



Oracle Database 2 Day + Performance Tuning Guide to learn how to use Oracle Enterprise Manager Cloud Control (Cloud Control) to tune database performance

Performance Planning

Refer to the topic Database Performance Fundamentals before proceeding with the other parts of this documentation. Based on years of designing and performance experience, Oracle has designed a performance methodology. This topic describes activities that can dramatically improve system performance, such as:

- Understanding Investment Options
- Understanding Scalability
- System Architecture
- Application Design Principles
- Workload Testing, Modeling, and Implementation
- Deploying New Applications

Instance Tuning

Diagnosing and Tuning Database Performance discusses the factors involved in the tuning and optimizing of an Oracle database instance.

When considering instance tuning, take care in the initial design of the database to avoid bottlenecks that could lead to performance problems. In addition, you must consider:

- Allocating memory to database structures
- Determining I/O requirements of different parts of the database
- Tuning the operating system for optimal performance of the database

After the database instance has been installed and configured, you must monitor the database as it is running to check for performance-related problems.

Performance Principles

Performance tuning requires a different, although related, method to the initial configuration of a system. Configuring a system involves allocating resources in an ordered manner so that the initial system configuration is functional.

Tuning is driven by identifying the most significant bottleneck and making the appropriate changes to reduce or eliminate the effect of that bottleneck. Usually, tuning is performed reactively, either while the system is in preproduction or after it is live.

Baselines

The most effective way to tune is to have an established performance baseline that you can use for comparison if a performance issue arises. Most database administrators (DBAs) know their system well and can easily identify peak usage periods. For example, the peak periods could be between 10.00am and 12.00pm and also between 1.30pm and 3.00pm. This could include a batch window of 12.00am midnight to 6am.

It is important to identify these peak periods at the site and install a monitoring tool that gathers performance data for those high-load times. Optimally, data gathering should be configured from when the application is in its initial trial phase during the QA cycle. Otherwise, this should be configured when the system is first in production.

Ideally, baseline data gathered should include the following:

- Application statistics (transaction volumes, response time)
- Database statistics
- Operating system statistics
- Disk I/O statistics
- Network statistics

In the Automatic Workload Repository, baselines are identified by a range of snapshots that are preserved for future comparisons. See "Automatic Workload Repository".

The Symptoms and the Problems

A common pitfall in performance tuning is to mistake the symptoms of a problem for the actual problem itself. It is important to recognize that many performance statistics indicate the symptoms, and that identifying the symptom is not sufficient data to implement a remedy. For example:

Slow physical I/O

Generally, this is caused by poorly-configured disks. However, it could also be caused by a significant amount of unnecessary physical I/O on those disks issued by poorly-tuned SQL.

Latch contention



Rarely is latch contention tunable by reconfiguring the instance. Rather, latch contention usually is resolved through application changes.

Excessive CPU usage

Excessive CPU usage usually means that there is little idle CPU on the system. This could be caused by an inadequately-sized system, by untuned SQL statements, or by inefficient application programs.

When to Tune

There are two distinct types of tuning:

- Proactive Monitoring
- Bottleneck Elimination

Proactive Monitoring

Proactive monitoring usually occurs on a regularly scheduled interval, where several performance statistics are examined to identify whether the system behavior and resource usage has changed. Proactive monitoring can also be considered as proactive tuning.

Usually, monitoring does not result in configuration changes to the system, unless the monitoring exposes a serious problem that is developing. In some situations, experienced performance engineers can identify potential problems through statistics alone, although accompanying performance degradation is usual.

Experimenting with or tweaking a system when there is no apparent performance degradation as a proactive action can be a dangerous activity, resulting in unnecessary performance drops. Tweaking a system should be considered reactive tuning, and the steps for reactive tuning should be followed.

Monitoring is usually part of a larger capacity planning exercise, where resource consumption is examined to see changes in the way the application is being used, and the way the application is using the database and host resources.

Bottleneck Elimination

Tuning usually implies fixing a performance problem. However, tuning should be part of the life cycle of an application—through the analysis, design, coding, production, and maintenance stages. Often, the tuning phase is left until the database is in production. At this time, tuning becomes a reactive process, where the most important bottleneck is identified and fixed.

Usually, the purpose for tuning is to reduce resource consumption or to reduce the elapsed time for an operation to complete. Either way, the goal is to improve the effective use of a particular resource. In general, performance problems are caused by the overuse of a particular resource. The overused resource is the bottleneck in the system. There are several distinct phases in identifying the bottleneck and the potential fixes. These are discussed in the sections that follow.

Remember that the different forms of contention are symptoms that can be fixed by making changes in the following places:

- Changes in the application, or the way the application is used
- · Changes in Oracle
- Changes in the host hardware configuration

Often, the most effective way of resolving a bottleneck is to change the application.



SQL Tuning

Many application programmers consider SQL a messaging language, because queries are issued and data is returned. However, client tools often generate inefficient SQL statements. Therefore, a good understanding of the database SQL processing engine is necessary for writing optimal SQL. This is especially true for high transaction processing systems.

Typically, SQL statements issued by online transaction processing (OLTP) applications operate on relatively few rows at a time. If an index can point to the exact rows that are required, then Oracle Database can construct an accurate plan to access those rows efficiently through the shortest possible path. In decision support system (DSS) environments, selectivity is less important, because they often access most of a table's rows. In such situations, full table scans are common, and indexes are not even used. This book is primarily focussed on OLTP applications.

See Also:

- Oracle Database SQL Tuning Guide for detailed information on the process of tuning and optimizing SQL statements
- Oracle Database Data Warehousing Guide for detailed information on decision support systems (DSS) and mixed environments

Query Optimizer and Execution Plans

When a SQL statement is executed on an Oracle database, the query optimizer determines the most efficient execution plan after considering many factors related to the objects referenced and the conditions specified in the query. This determination is an important step in the processing of any SQL statement and can greatly affect execution time.

During the evaluation process, the query optimizer reviews statistics gathered on the system to determine the best data access path and other considerations. You can override the execution plan of the query optimizer with hints inserted in SQL statement.

Introduction to Performance Tuning Features and Tools

Effective data collection and analysis is essential for identifying and correcting performance problems. Oracle Database provides several tools that allow a performance engineer to gather information regarding database performance. In addition to gathering data, Oracle Database provides tools to monitor performance, diagnose problems, and tune applications.

The Oracle Database gathering and monitoring features are mainly automatic, managed by Oracle background processes. To enable automatic statistics collection and automatic performance features, the STATISTICS_LEVEL initialization parameter must be set to TYPICAL or ALL. You can administer and display the output of the gathering and tuning tools with Oracle Enterprise Manager Cloud Control (Cloud Control), or with APIs and views. For ease of use and to take advantage of its numerous automated monitoring and diagnostic tools, Cloud Control is recommended.



See Also:

- Oracle Database 2 Day DBA to learn how to use Cloud Control to manage Oracle Database
- Oracle Database 2 Day + Performance Tuning Guide for a quick lesson on how to use Cloud Control to tune database performance, and Enterprise Manager Cloud Control 13c for details on everything Cloud Control has to offer.
- Oracle Database PL/SQL Packages and Types Reference for detailed information on the DBMS_ADVISOR, DBMS_SQLTUNE, DBMS_AUTO_SQLTUNE, and DBMS_WORKLOAD_REPOSITORY packages
- Oracle Database Reference for information about the STATISTICS_LEVEL initialization parameter

Automatic Performance Tuning Features

The Oracle Database automatic performance tuning features include:

- Automatic Workload Repository (AWR) collects, processes, and maintains performance statistics for problem detection and self-tuning purposes. See "Automatic Workload Repository".
- Automatic Database Diagnostic Monitor (ADDM) analyzes the information collected by AWR for possible performance problems with the Oracle database. See "Overview of the Automatic Database Diagnostic Monitor".
- SQL Tuning Advisor allows a quick and efficient technique for optimizing SQL statements without modifying any statements. See Oracle Database SQL Tuning Guide.
- SQL Access Advisor provides advice on materialized views, indexes, and materialized view logs. See Oracle Database SQL Tuning Guide.
- End to End Application tracing identifies excessive workloads on the system by specific user, service, or application component. See *Oracle Database SQL Tuning Guide*.
- Server-generated alerts automatically provide notifications when impending problems are detected. See Oracle Database Administrator's Guide to learn how to monitor the operation of the database with server-generated alerts.
- Additional advisors that can be launched from Oracle Enterprise Manager Cloud Control (Cloud Control), such as memory advisors to optimize memory for an instance. The memory advisors are commonly used when automatic memory management is not set up for the database. Other advisors are used to optimize mean time to recovery (MTTR), shrinking of segments, and undo tablespace settings. To learn about the advisors available with Cloud Control, see Oracle Database 2 Day + Performance Tuning Guide.
- The Database Performance page in Cloud Control displays host, instance service time, and throughput information for real time monitoring and diagnosis. The page can be set to refresh automatically in selected intervals or manually. To learn about the Database Performance page, see *Oracle Database 2 Day + Performance Tuning Guide*.

Additional Oracle Database Tools

This section describes additional Oracle Database tools that you can use for determining performance problems.



V\$ Performance Views

The V\$ views are the performance information sources used by all Oracle Database performance tuning tools. The V\$ views are based on memory structures initialized at instance startup. The memory structures, and the views that represent them, are automatically maintained by Oracle Database for the life of the instance.

Note:

Oracle recommends using the Automatic Workload Repository to gather performance data. These tools have been designed to capture all of the data needed for performance analysis.

See Also:

- "Instance Tuning Using Performance Views" for more information about diagnosing database performance problems using the V\$ performance views
- Oracle Database Reference for more information about dynamic performance views



Designing and Developing for Performance

Optimal system performance begins with design and continues throughout the life of your system. Carefully consider performance issues during the initial design phase so that you can tune your system more easily during production.

This chapter contains the following sections:

- Oracle Methodology
- Understanding Investment Options
- Understanding Scalability
- System Architecture
- Application Design Principles
- Workload Testing, Modeling, and Implementation
- · Deploying New Applications

Oracle Methodology

System performance has become increasingly important as computer systems get larger and more complex as the Internet plays a bigger role in business applications. To accommodate this, Oracle has produced a performance methodology based on years of designing and performance experience. This methodology explains clear and simple activities that can dramatically improve system performance.

Performance strategies vary in their effectiveness, and systems with different purposes—such as operational systems and decision support systems—require different performance skills. This book examines the considerations that any database designer, administrator, or performance expert should focus their efforts on.

System performance is designed and built into a system. It does not just happen. Performance problems are usually the result of contention for, or exhaustion of, some system resource. When a system resource is exhausted, the system cannot scale to higher levels of performance. This new performance methodology is based on careful planning and design of the database, to prevent system resources from becoming exhausted and causing down-time. By eliminating resource conflicts, systems can be made scalable to the levels required by the business.

Understanding Investment Options

With the availability of relatively inexpensive, high-powered processors, memory, and disk drives, there is a temptation to buy more system resources to improve performance. In many situations, new CPUs, memory, or more disk drives can indeed provide an immediate performance improvement. However, any performance increases achieved by adding hardware should be considered a short-term relief to an immediate problem. If the demand and load rates on the application continue to grow, then the chance of the same problem occurring soon is likely.

In other situations, additional hardware does not improve the system's performance at all. Poorly designed systems perform poorly no matter how much extra hardware is allocated. Before purchasing additional hardware, ensure that serialization or single threading is not occurring within the application. Long-term, it is generally more valuable to increase the efficiency of your application in terms of the number of physical resources used for each business transaction.

Understanding Scalability

The word *scalability* is used in many contexts in development environments. The following section provides an explanation of scalability that is aimed at application designers and performance specialists.

This section covers the following topics:

- What is Scalability?
- System Scalability
- Factors Preventing Scalability

What is Scalability?

Scalability is a system's ability to process more workload, with a proportional increase in system resource usage.

In a scalable system, if you double the workload, then the system uses twice as many system resources. This sounds obvious, but due to conflicts within the system, the resource usage might exceed twice the original workload.

Examples of poor scalability due to resource conflicts include the following:

- Applications requiring significant concurrency management as user populations increase
- Increased locking activities
- Increased data consistency workload
- Increased operating system workload
- Transactions requiring increases in data access as data volumes increase
- Poor SQL and index design resulting in a higher number of logical I/Os for the same number of rows returned
- Reduced availability, because database objects take longer to maintain

An application is said to be unscalable if it exhausts a system resource to the point where no more throughput is possible when its workload is increased. Such applications result in fixed throughputs and poor response times.

Examples of resource exhaustion include the following:

- Hardware exhaustion
- Table scans in high-volume transactions causing inevitable disk I/O shortages
- Excessive network requests, resulting in network and scheduling bottlenecks
- Memory allocation causing paging and swapping
- Excessive process and thread allocation causing operating system thrashing



This means that application designers must create a design that uses the same resources, regardless of user populations and data volumes, and does not put loads on the system resources beyond their limits.

System Scalability

Applications that are accessible through the Internet have more complex performance and availability requirements.

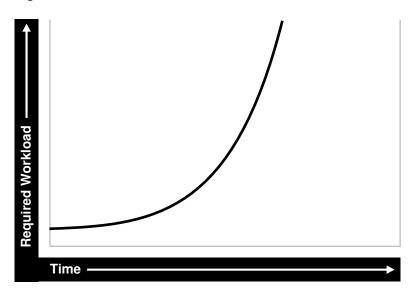
Some applications are designed and written only for Internet use, but even typical back-office applications—such as a general ledger application—might require some or all data to be available online.

Characteristics of Internet age applications include the following:

- Availability 24 hours a day, 365 days a year
- · Unpredictable and imprecise number of concurrent users
- Difficulty in capacity planning
- Availability for any type of query
- Multitier architectures
- Stateless middleware
- Rapid development timescale
- Minimal time for testing

The following figure illustrates the classic workload growth curve, with demand growing at an increasing rate. Applications must scale with the increase of workload and also when additional hardware is added to support increasing demand. Design errors can cause the implementation to reach its maximum, regardless of additional hardware resources or re-design efforts.

Figure 2-1 Workload Growth Curve



Applications are challenged by very short development timeframes with limited time for testing and evaluation. However, bad design typically means that you must later rearchitect and reimplement the system. If you deploy an application with known architectural and implementation limitations on the Internet, and if the workload exceeds the anticipated



demand, then failure is a real possibility. From a business perspective, poor performance can mean a loss of customers. If Web users do not get a response in seven seconds, then the user's attention could be lost forever.

In many cases, the cost of re-designing a system with the associated downtime costs in migrating to new implementations exceeds the costs of properly building the original system. The moral of the story is simple: design and implement with scalability in mind from the start.

Factors Preventing Scalability

When building applications, designers and architects should aim for as close to perfect scalability as possible. This is sometimes called *linear* scalability, where system throughput is directly proportional to the number of CPUs.

In the real world, linear scalability is impossible for reasons beyond a designer's control. However, making the application design and implementation as scalable as possible should ensure that current and future performance objectives can be achieved through expansion of hardware components and the evolution of CPU technology.

Factors that may prevent linear scalability include:

Poor application design, implementation, and configuration

The application has the biggest impact on scalability. For example:

- Poor schema design can cause expensive SQL that do not scale.
- Poor transaction design can cause locking and serialization problems.
- Poor connection management can cause poor response times and unreliable systems.

However, the design is not the only problem. The physical implementation of the application can be the weak link. For example:

- Systems can move to production environments with bad I/O strategies.
- The production environment could might different execution plans from those generated in testing.
- Memory-intensive applications that allocate a large amount of memory without much thought for freeing the memory at run time can cause excessive memory usage.
- Inefficient memory usage and memory leaks put a high stress on the operating virtual memory subsystem. This impacts performance and availability.
- Incorrect sizing of hardware components

Bad capacity planning of all hardware components is becoming less of a problem as relative hardware prices decrease. However, too much capacity can mask scalability problems as the workload is increased on a system.

Limitations of software components

All software components have scalability and resource usage limitations. This applies to application servers, database servers, and operating systems. Application design should not place demands on the software beyond what it can handle.

Limitations of hardware components

Hardware is not perfectly scalable. Most multiprocessor computers can get close to linear scaling with a finite number of CPUs, but after a certain point each additional CPU can increase performance overall, but not proportionately. There might come a time when an additional CPU offers no increase in performance, or even degrades performance. This behavior is very closely linked to the workload and the operating system setup.





These factors are based on Oracle Server Performance group's experience of tuning unscalable systems.

System Architecture

There are two main parts to a system's architecture:

- Hardware and Software Components
- Configuring the Right System Architecture for Your Requirements

Hardware and Software Components

A system architecture mainly contains hardware and software components.

- Hardware Components
- Software Components

Hardware Components

Today's designers and architects are responsible for sizing and capacity planning of hardware at each tier in a multitier environment. It is the architect's responsibility to achieve a balanced design. This is analogous to a bridge designer who must consider all the various payload and structural requirements for the bridge. A bridge is only as strong as its weakest component. As a result, a bridge is designed in balance, such that all components reach their design limits simultaneously.

The following are the main hardware components of a system.

CPU

There can be one or more CPUs, and they can vary in processing power from simple CPUs found in hand-held devices to high-powered server CPUs. Sizing of other hardware components is usually a multiple of the CPUs on the system.

Memory

Database and application servers require considerable amounts of memory to cache data and avoid time-consuming disk access.

I/O Subsystem

The I/O subsystem can vary between the hard disk on a client PC and high performance disk arrays. Disk arrays can perform thousands of I/Os each second and provide availability through redundancy in terms of multiple I/O paths and hot pluggable mirrored disks.

Network

All computers in a system are connected to a network, from a modem line to a high speed internal LAN. The primary concerns with network specifications are bandwidth (volume) and latency (speed).



Software Components

The same way computers have common hardware components, applications have common functional components. By dividing software development into functional components, it is possible to better comprehend the application design and architecture. Some components of the system are performed by existing software bought to accelerate application implementation, or to avoid re-development of common components.

The difference between software components and hardware components is that while hardware components only perform one task, a piece of software can perform the roles of various software components. For example, a disk drive only stores and retrieves data, but a client program can manage the user interface and perform business logic.

Most applications involve the following software components:

User Interface

This component is the most visible to application users, and includes the following functions:

- Displaying the screen to the user
- Collecting user data and transferring it to business logic
- Validating data entry
- Navigating through levels or states of the application

Business Logic

This component implements core business rules that are central to the application function. Errors made in this component can be very costly to repair. This component is implemented by a mixture of declarative and procedural approaches. An example of a declarative activity is defining unique and foreign keys. An example of procedure-based logic is implementing a discounting strategy.

Common functions of this component include:

- Moving a data model to a relational table structure
- Defining constraints in the relational table structure
- Coding procedural logic to implement business rules

Resources for Managing User Requests

This component is implemented in all pieces of software. However, there are some requests and resources that can be influenced by the application design and some that cannot.

In a multiuser application, most resource allocation by user requests are handled by the database server or the operating system. However, in a large application where the number of users and their usage pattern is unknown or growing rapidly, the system architect must be proactive to ensure that no single software component becomes overloaded and unstable.

Common functions of this component include:

- Connection management with the database
- Executing SQL efficiently (cursors and SQL sharing)
- · Managing client state information
- Balancing the load of user requests across hardware resources



- Setting operational targets for hardware and software components
- Persistent queuing for asynchronous execution of tasks

Data and Transactions

This component is largely the responsibility of the database server and the operating system.

Common functions of this component include:

- Providing concurrent access to data using locks and transactional semantics
- Providing optimized access to the data using indexes and memory cache
- Ensuring that data changes are logged in the event of a hardware failure
- Enforcing any rules defined for the data

Configuring the Right System Architecture for Your Requirements

Configuring the initial system architecture is a largely iterative process. System architects must satisfy the system requirements within budget and schedule constraints. If the system requires interactive users transacting business-making decisions based on the contents of a database, then user requirements drive the architecture. If there are few interactive users on the system, then the architecture is process-driven.

Examples of interactive user applications:

- Accounting and bookkeeping applications
- Order entry systems
- Email servers
- Web-based retail applications
- Trading systems

Examples of process-driven applications:

- Utility billing systems
- Fraud detection systems
- Direct mail

In many ways, process-driven applications are easier to design than multiuser applications because the user interface element is eliminated. However, because the objectives are process-oriented, system architects not accustomed to dealing with large data volumes and different success factors can become confused. Process-driven applications draw from the skills sets used in both user-based applications and data warehousing. Therefore, this book focuses on evolving system architectures for interactive users.



Generating a system architecture is not a deterministic process. It requires careful consideration of business requirements, technology choices, existing infrastructure and systems, and actual physical resources, such as budget and manpower.

The following questions should stimulate thought on system architecture, though they are not a definitive guide to system architecture. These questions demonstrate how business



requirements can influence the architecture, ease of implementation, and overall performance and availability of a system. For example:

How many users must the system support?

Most applications fall into one of the following categories:

Very few users on a lightly-used or exclusive computer

For this type of application, there is usually one user. The focus of the application design is to make the single user as productive as possible by providing good response time, yet make the application require minimal administration. Users of these applications rarely interfere with each other and have minimal resource conflicts.

A medium to large number of users in a corporation using shared applications

For this type of application, the users are limited by the number of employees in the corporation actually transacting business through the system. Therefore, the number of users is predictable. However, delivering a reliable service is crucial to the business. The users must share a resource, so design efforts must address response time under heavy system load, escalation of resource for each session usage, and room for future growth.

An infinite user population distributed on the Internet

For this type of application, extra engineering effort is required to ensure that no system component exceeds its design limits. This creates a bottleneck that halts or destabilizes the system. These applications require complex load balancing, stateless application servers, and efficient database connection management. In addition, use statistics and governors to ensure that the user receives feedback if the database cannot satisfy their requests because of system overload.

What will be the user interaction method?

The choices of user interface range from a simple Web browser to a custom client program.

Where are the users located?

The distance between users influences how the application is engineered to cope with network latencies. The location also affects which times of the day are busy, when it is impossible to perform batch or system maintenance functions.

What is the network speed?

Network speed affects the amount of data and the conversational nature of the user interface with the application and database servers. A highly conversational user interface can communicate with back-end servers on every key stroke or field level validation. A less conversational interface works on a screen-sent and a screen-received model. On a slow network, it is impossible to achieve high data entry speeds with a highly conversational user interface.

How much data will the user access, and how much of that data is largely read only?

The amount of data queried online influences all aspects of the design, from table and index design to the presentation layers. Design efforts must ensure that user response time is not a function of the size of the database. If the application is largely read only, then replication and data distribution to local caches in the application servers become a viable option. This also reduces workload on the core transactional server.

What is the user response time requirement?

Consideration of the user type is important. If the user is an executive who requires accurate information to make split second decisions, then user response time cannot be



compromised. Other types of users, such as users performing data entry activities, might not need such a high level of performance.

Do users expect 24 hour service?

This is mandatory for today's Internet applications where trade is conducted 24 hours a day. However, corporate systems that run in a single time zone might be able to tolerate after-hours downtime. You can use this after-hours downtime to run batch processes or to perform system administration. In this case, it might be more economic not to run a fully-available system.

Must all changes be made in real time?

It is important to determine whether transactions must be executed within the user response time, or if they can be queued for asynchronous execution.

The following are secondary questions, which can also influence the design, but really have more impact on budget and ease of implementation. For example:

How big will the database be?

This influences the sizing of the database server. On servers with a very large database, it might be necessary to have a bigger computer than dictated by the workload. This is because the administration overhead with large databases is largely a function of the database size. As tables and indexes grow, it takes proportionately more CPUs to allow table reorganizations and index builds to complete in an acceptable time limit.

- What is the required throughput of business transactions?
- What are the availability requirements?
- Do skills exist to build and administer this application?
- What compromises are forced by budget constraints?

Application Design Principles

This section describes the following design decisions that are involved in building applications:

- Simplicity In Application Design
- Data Modeling
- Table and Index Design
- Using Views
- SQL Execution Efficiency
- Implementing the Application
- Trends in Application Development

Simplicity In Application Design

Applications are no different than any other designed and engineered product. Well-designed structures, computers, and tools are usually reliable, easy to use and maintain, and simple in concept. In the most general terms, if the design looks correct, then it probably is. This principle should always be kept in mind when building applications.

Consider the following design issues:

 If the table design is so complicated that nobody can fully understand it, then the table is probably poorly designed.



- If SQL statements are so long and involved that it would be impossible for any optimizer to
 effectively optimize it in real time, then there is probably a bad statement, underlying
 transaction, or table design.
- If there are indexes on a table and the same columns are repeatedly indexed, then there is probably a poor index design.
- If queries are submitted without suitable qualification for rapid response for online users, then there is probably a poor user interface or transaction design.
- If the calls to the database are abstracted away from the application logic by many layers of software, then there is probably a bad software development method.

Data Modeling

Data modeling is important to successful relational application design. You must perform this modeling in a way that quickly represents the business practices. Heated debates may occur about the correct data model. The important thing is to apply greatest modeling efforts to those entities affected by the most frequent business transactions. In the modeling phase, there is a great temptation to spend too much time modeling the non-core data elements, which results in increased development lead times. Use of modeling tools can then rapidly generate schema definitions and can be useful when a fast prototype is required.

Table and Index Design

Table design is largely a compromise between flexibility and performance of core transactions. To keep the database flexible and able to accommodate unforeseen workloads, the table design should be very similar to the data model, and it should be normalized to at least 3rd normal form. However, certain core transactions required by users can require selective denormalization for performance purposes.

Examples of this technique include storing tables pre-joined, the addition of derived columns, and aggregate values. Oracle Database provides numerous options for storage of aggregates and pre-joined data by clustering and materialized view functions. These features allow a simpler table design to be adopted initially.

Again, focus and resources should be spent on the business critical tables, so that optimal performance can be achieved. For non-critical tables, shortcuts in design can be adopted to enable a more rapid application development. However, if prototyping and testing a non-core table becomes a performance problem, then remedial design effort should be applied immediately.

Index design is also a largely iterative process, based on the SQL generated by application designers. However, it is possible to make a sensible start by building indexes that enforce primary key constraints and indexes on known access patterns, such as a person's name. As the application evolves, and as you perform testing on realistic amounts of data, you may need to improve the performance of specific queries by building a better index. Consider the following list of indexing design ideas when building a new index:

- Appending Columns to an Index or Using Index-Organized Tables
- Using a Different Index Type
- Finding the Cost of an Index
- Serializing within Indexes
- · Ordering Columns in an Index



Appending Columns to an Index or Using Index-Organized Tables

One of the easiest ways to speed up a query is to reduce the number of logical I/Os by eliminating a table access from the execution plan. This can be done by appending to the index all columns referenced by the query. These columns are the select list columns, and any required join or sort columns. This technique is particularly useful in speeding up online applications response times when time-consuming I/Os are reduced. This is best applied when testing the application with properly sized data for the first time.

The most aggressive form of this technique is to build an index-organized table (IOT). However, you must be careful that the increased leaf size of an IOT does not undermine the efforts to reduce I/O.

Using a Different Index Type

There are several index types available, and each index has benefits for certain situations. The following list gives performance ideas associated with each index type.

B-Tree Indexes

These indexes are the standard index type, and they are excellent for primary key and highly-selective indexes. Used as concatenated indexes, the database can use B-tree indexes to retrieve data sorted by the index columns.

Bitmap Indexes

These indexes are suitable for columns that have a relatively low number of distinct values, where the benefit of adding a B-tree index is likely to be limited. These indexes are suitable for data warehousing applications where there is low DML activity and ad hoc filtering patterns. Combining bitmap indexes on columns allows efficient AND and OR operations with minimal I/O. Further, through compression techniques they can generate a large number of rowids with minimal I/Os. Bitmap indexes are particularly efficient in queries with COUNT(), because the query can be satisfied within the index.

Function-based Indexes

These indexes allow access through a B-tree on a value derived from a function on the base data. Function-based indexes have some limitations with regards to the use of nulls, and they require that you have the query optimizer enabled.

Function-based indexes are particularly useful when querying on composite columns to produce a derived result or to overcome limitations in the way data is stored in the database. An example is querying for line items in an order exceeding a certain value derived from (sales price - discount) x quantity, where these were columns in the table. Another example is to apply the UPPER function to the data to allow case-insensitive searches.

Partitioned Indexes

Partitioning a global index allows partition pruning to take place within an index access, which results in reduced I/Os. By definition of good range or list partitioning, fast index scans of the correct index partitions can result in very fast guery times.

Reverse Key Indexes

These indexes are designed to eliminate index hot spots on insert applications. These indexes are excellent for insert performance, but they are limited because the database cannot use them for index range scans.



Building and maintaining an index structure can be expensive, and it can consume resources such as disk space, CPU, and I/O capacity. Designers must ensure that the benefits of any index outweigh the negatives of index maintenance.

Use this simple estimation guide for the cost of index maintenance: each index maintained by an INSERT, DELETE, or UPDATE of the indexed keys requires about three times as much resource as the actual DML operation on the table. Thus, if you INSERT into a table with three indexes, then the insertion is approximately 10 times slower than an INSERT into a table with no indexes. For DML, and particularly for INSERT-heavy applications, the index design should be seriously reviewed, which might require a compromise between the query and INSERT performance.



Oracle Database Administrator's Guide to learn how to monitor index usage

Serializing within Indexes

Use of sequences or timestamps to generate key values that are indexed themselves can lead to database hotspot problems, which affect response time and throughput. This is usually the result of a monotonically growing key that results in a right-growing index. To avoid this problem, try to generate keys that insert over the full range of the index so as to make a workload more scalable. You can achieve this by using any of the following methods:

- using a reverse key index
- using a hash partitioned index
- using a cycling sequence to prefix sequence values
- using a scalable sequence

See Also:

Oracle Database Administrator's Guide for more information about the scalable sequences

Ordering Columns in an Index

Designers should be flexible in defining any rules for index building. Depending on your circumstances, use one of the following two ways to order the keys in an index:

- Order columns with most selectivity first. This method is the most commonly used because
 it provides the fastest access with minimal I/O to the actual rowids required. This technique
 is used mainly for primary keys and for very selective range scans.
- Order columns to reduce I/O by clustering or sorting data. In large range scans, I/Os can
 usually be reduced by ordering the columns in the least selective order, or in a manner that
 sorts the data in the way it should be retrieved.

Using Views

Views can speed up and simplify application design. A simple view definition can mask data model complexity from the programmers whose priorities are to retrieve, display, collect, and store data.

However, while views provide clean programming interfaces, they can cause sub-optimal, resource-intensive queries. The worst type of view use is when a view references other views, and when they are joined in queries. In many cases, developers can satisfy the query directly from the table without using a view. Usually, because of their inherent properties, views make it difficult for the optimizer to generate the optimal execution plan.

SQL Execution Efficiency

In the design and architecture phase of any system development, care should be taken to ensure that the application developers understand SQL execution efficiency. To achieve this goal, the development environment must support the following characteristics:

Good database connection management

Connecting to the database is an expensive operation that is highly unscalable. Therefore, the number of concurrent connections to the database should be minimized as much as possible. A simple system, where a user connects at application initialization, is ideal. However, in a Web-based or multitiered application, where application servers are used to multiplex database connections to users, this can be difficult. With these types of applications, design efforts should ensure that database connections are pooled and are not reestablished for each user request.

Good cursor usage and management

Maintaining user connections is equally important to minimizing the parsing activity on the system. Parsing is the process of interpreting a SQL statement and creating an execution plan for it. This process has many phases, including syntax checking, security checking, execution plan generation, and loading shared structures into the shared pool. There are two types of parse operations:

Hard parsing

A SQL statement is submitted for the first time, and no match is found in the shared pool. Hard parses are the most resource-intensive and unscalable, because they perform all the operations involved in a parse.

Soft parsing

A SQL statement is submitted for the first time, and a match *is* found in the shared pool. The match can be the result of previous execution by another user. The SQL statement is shared, which is good for performance. However, soft parses are not ideal, because they still require syntax and security checking, which consume system resources.

Because parsing should be minimized as much as possible, application developers should design their applications to parse SQL statements once and execute them many times. This is done through cursors. Experienced SQL programmers should be familiar with the concept of opening and re-executing cursors.

Application developers must also ensure that SQL statements are shared within the shared pool. To achieve this goal, use bind variables to represent the parts of the query that change from execution to execution. If this is not done, then the SQL statement is likely to

be parsed once and never re-used by other users. To ensure that SQL is shared, use bind variables and do not use string literals with SQL statements. For example:

Statement with string literals:

```
SELECT * FROM employees
  WHERE last name LIKE 'KING';
```

Statement with bind variables:

```
SELECT * FROM employees
WHERE last name LIKE :1;
```

The following example shows the results of some tests on a simple OLTP application:

```
Test #Users Supported
No Parsing all statements 270
Soft Parsing all statements 150
Hard Parsing all statements 60
Re-Connecting for each Transaction 30
```

These tests were performed on a four-CPU computer. The differences increase as the number of CPUs on the system increase.

Implementing the Application

The choice of development environment and programming language is largely a function of the skills available in the development team and architectural decisions made when specifying the application. There are, however, some simple performance management rules that can lead to scalable, high-performance applications.

- Choose a development environment suitable for software components, and do not let it limit your design for performance decisions. If it does, then you probably chose the wrong language or environment.
 - User interface

The programming model can vary between HTML generation and calling the windowing system directly. The development method should focus on response time of the user interface code. If HTML or Java is being sent over a network, then try to minimize network volume and interactions.

Business logic

Interpreted languages, such as Java and PL/SQL, are ideal to encode business logic. They are fully portable, which makes upgrading logic relatively easy. Both languages are syntactically rich to allow code that is easy to read and interpret. If business logic requires complex mathematical functions, then a compiled binary language might be needed. The business logic code can be on the client computer, the application server, and the database server. However, the application server is the most common location for business logic.

User requests and resource allocation

Most of this is not affected by the programming language, but tools and fourth generation languages that mask database connection and cursor management might use inefficient mechanisms. When evaluating these tools and environments, check their database connection model and their use of cursors and bind variables.

Data management and transactions

Most of this is not affected by the programming language.

- 2. When implementing a software component, implement its function and not the functionality associated with other components. Implementing another component's functionality results in sub-optimal designs and implementations. This applies to all components.
- 3. Do not leave gaps in functionality or have software components under-researched in design, implementation, or testing. In many cases, gaps are not discovered until the application is rolled out or tested at realistic volumes. This is usually a sign of poor architecture or initial system specification. Data archival and purge modules are most frequently neglected during initial system design, build, and implementation.
- 4. When implementing procedural logic, implement in a procedural language, such as C, Java, or PL/SQL. When implementing data access (queries) or data changes (DML), use SQL. This rule is specific to the business logic modules of code where procedural code is mixed with data access (nonprocedural SQL) code. There is great temptation to put procedural logic into the SQL access. This tends to result in poor SQL that is resource-intensive. SQL statements with DECODE case statements are very often candidates for optimization, as are statements with a large amount of OR predicates or set operators, such as UNION and MINUS.
- 5. Cache frequently accessed, rarely changing data that is expensive to retrieve on a repeated basis. However, make this cache mechanism easy to use, and ensure that it is indeed cheaper than accessing the data in the original method. This is applicable to all modules where frequently used data values should be cached or stored locally, rather than be repeatedly retrieved from a remote or expensive data store.

The most common examples of candidates for local caching include the following:

- Today's date. SELECT SYSDATE FROM DUAL can account for over 60% of the workload on a database.
- The current user name.
- Repeated application variables and constants, such as tax rates, discounting rates, or location information.
- Caching data locally can be further extended into building a local data cache into the
 application server middle tiers. This helps take load off the central database servers.
 However, care should be taken when constructing local caches so that they do not
 become so complex that they cease to give a performance gain.
- Local sequence generation.

The design implications of using a cache should be considered. For example, if a user is connected at midnight and the date is cached, then the user's date value becomes invalid.

- Optimize the interfaces between components, and ensure that all components are used in the most scalable configuration. This rule requires minimal explanation and applies to all modules and their interfaces.
- 7. Use foreign key references. Enforcing referential integrity through an application is expensive. You can maintain a foreign key reference by selecting the column value of the child from the parent and ensuring that it exists. The foreign key constraint enforcement supplied by Oracle—which does not use SQL—is fast, easy to declare, and does not create network traffic.
- **8.** Consider setting up action and module names in the application to use with End to End Application Tracing. This allows greater flexibility in tracing workload problems.



Trends in Application Development

The two biggest challenges in application development today are the increased use of Java to replace compiled C or C++ applications, and increased use of object-oriented techniques, influencing the schema design.

Java provides better portability of code and availability to programmers. However, there are several performance implications associated with Java. Because Java is an interpreted language, it is slower at executing similar logic than compiled languages, such as C. As a result, resource usage of client computers increases. This requires more powerful CPUs to be applied in the client or middle-tier computers and greater care from programmers to produce efficient code.

Because Java is an object-oriented language, it encourages insulation of data access into classes not performing the business logic. As a result, programmers might invoke methods without knowledge of the efficiency of the data access method being used. This tends to result in minimal database access and uses the simplest and crudest interfaces to the database.

With this type of software design, queries do not always include all the WHERE predicates to be efficient, and row filtering is performed in the Java program. This is very inefficient. In addition, for DML operations—and especially for INSERTS—single INSERTS are performed, making use of the array interface impossible. In some cases, this is made more inefficient by procedure calls. More resources are used moving the data to and from the database than in the actual database calls.

In general, it is best to place data access calls next to the business logic to achieve the best overall transaction design.

The acceptance of object-orientation at a programming level has led to the creation of object-oriented databases within the Oracle Server. This has manifested itself in many ways, from storing object structures within BLOBS and only using the database effectively as an indexed card file to the use of the Oracle Database object-relational features.

If you adopt an object-oriented approach to schema design, then ensure that you do not lose the flexibility of the relational storage model. In many cases, the object-oriented approach to schema design ends up in a heavily denormalized data structure that requires considerable maintenance and REF pointers associated with objects. Often, these designs represent a step backward to the hierarchical and network database designs that were replaced with the relational storage method.

In summary, if you are storing your data in your database for the long-term, and if you anticipate a degree of ad hoc queries or application development on the same schema, then the relational storage method probably gives the best performance and flexibility.

Workload Testing, Modeling, and Implementation

This section describes workload estimation, modeling, implementation, and testing. This section covers the following topics:

- Sizing Data
- Estimating Workloads
- Application Modeling
- Testing, Debugging, and Validating a Design



Sizing Data

You could experience errors in your sizing estimates when dealing with variable length data if you work with a poor sample set. As data volumes grow, your key lengths could grow considerably, altering your assumptions for column sizes.

When the system becomes operational, it becomes more difficult to predict database growth, especially for indexes. Tables grow over time, and indexes are subject to the individual behavior of the application in terms of key generation, insertion pattern, and deletion of rows. The worst case is where you insert using an ascending key, and then delete most rows from the left-hand side but not all the rows. This leaves gaps and wasted space. If you have index use like this, then ensure that you know how to use the online index rebuild facility.

DBAs should monitor space allocation for each object and look for objects that may grow out of control. A good understanding of the application can highlight objects that may grow rapidly or unpredictably. This is a crucial part of both performance and availability planning for any system. When implementing the production database, the design should attempt to ensure that minimal space management takes place when interactive users are using the application. This applies for all data, temp, and rollback segments.

Estimating Workloads

Considering the number of variables involved, estimation of workloads for capacity planning and testing purposes is extremely difficult. However, designers must specify computers with CPUs, memory, and disk drives, and eventually roll out an application. There are several techniques used for sizing, and each technique has merit. When sizing, it is best to use the following methods to validate your decision-making process and provide supporting documentation.

Extrapolating From a Similar System

This is an entirely empirical approach where an existing system of similar characteristics and known performance is used as a basis system. The specification of this system is then modified by the sizing specialist according to the known differences. This approach has merit in that it correlates with an existing system, but it provides little assistance when dealing with the differences.

This approach is used in nearly all large engineering disciplines when preparing the cost of an engineering project, such as a large building, a ship, a bridge, or an oil rig. If the reference system is an order of magnitude different in size from the anticipated system, then some components may have exceeded their design limits.

Benchmarking

The benchmarking process is both resource and time consuming, and it might not produce the correct results. By simulating an application in early development or prototype form, there is a danger of measuring something that has no resemblance to the actual production system. This sounds strange, but over the many years of benchmarking customer applications with the database development organization, Oracle has yet to see reliable correlation between the benchmark application and the actual production system. This is mainly due to the number of application inefficiencies introduced in the development process.

However, benchmarks have been used successfully to size systems to an acceptable level of accuracy. In particular, benchmarks are very good at determining the actual I/O requirements and testing recovery processes when a system is fully loaded.



Benchmarks by their nature stress all system components to their limits. As the benchmark stresses all components, be prepared to see all errors in application design and implementation manifest themselves while benchmarking. Benchmarks also test database, operating system, and hardware components. Because most benchmarks are performed in a rush, expect setbacks and problems when a system component fails. Benchmarking is a stressful activity, and it takes considerable experience to get the most out of a benchmarking exercise.

Application Modeling

Modeling the application can range from complex mathematical modeling exercises to the classic simple calculations performed on the back of an envelope. Both methods have merit, with one attempting to be very precise and the other making gross estimates. The downside of both methods is that they do not allow for implementation errors and inefficiencies.

The estimation and sizing process is an imprecise science. However, by investigating the process, some intelligent estimates can be made. The whole estimation process makes no allowances for application inefficiencies introduced by poor SQL, index design, or cursor management. A sizing engineer should build in margin for application inefficiencies. A performance engineer should discover the inefficiencies and make the estimates look realistic. The Oracle performance method describes how to discover the application inefficiencies.

Testing, Debugging, and Validating a Design

The testing process mainly consists of functional and stability testing. At some point in the process, performance testing is performed.

The following list describes some simple rules for performance testing an application. If correctly documented, then this list provides important information for the production application and the capacity planning process after the application has gone live.

- Use the Automatic Database Diagnostic Monitor (ADDM) and SQL Tuning Advisor for design validation
- Test with realistic data volumes and distributions

All testing must be done with fully populated tables. The test database should contain data representative of the production system in terms of data volume and cardinality between tables. All the production indexes should be built and the schema statistics should be populated correctly.

Use the correct optimizer mode

Perform all testing with the optimizer mode that you plan to use in production. All Oracle Database research and development effort is focused on the query optimizer. Therefore, the use of the query optimizer is recommended.

Test a single user performance

Test a single user on an idle or lightly-used database for acceptable performance. If a single user cannot achieve acceptable performance under ideal conditions, then multiple users cannot achieve acceptable performance under real conditions.

Obtain and document plans for all SQL statements

Obtain an execution plan for each SQL statement. Use this process to verify that the optimizer is obtaining an optimal execution plan, and that the relative cost of the SQL statement is understood in terms of CPU time and physical I/Os. This process assists in identifying the heavy use transactions that require the most tuning and performance work in the future.



Attempt multiuser testing

This process is difficult to perform accurately, because user workload and profiles might not be fully quantified. However, transactions performing DML statements should be tested to ensure that there are no locking conflicts or serialization problems.

Test with the correct hardware configuration

Test with a configuration as close to the production system as possible. Using a realistic system is particularly important for network latencies, I/O subsystem bandwidth, and processor type and speed. Failing to use this approach may result in an incorrect analysis of potential performance problems.

Measure steady state performance

When benchmarking, it is important to measure the performance under steady state conditions. Each benchmark run should have a ramp-up phase, where users are connected to the application and gradually start performing work on the application. This process allows for frequently cached data to be initialized into the cache and single execution operations—such as parsing—to be completed before the steady state condition. Likewise, at the end of a benchmark run, there should be a ramp-down period, where resources are freed from the system and users cease work and disconnect.

Deploying New Applications

The following are the key design decisions involved in deploying applications:

- Rollout Strategies
- Performance Checklist

Rollout Strategies

When new applications are rolled out, two strategies are commonly adopted:

- Big Bang approach all users migrate to the new system at once
- Trickle approach users slowly migrate from existing systems to the new one

Both approaches have merits and disadvantages. The Big Bang approach relies on reliable testing of the application at the required scale, but has the advantage of minimal data conversion and synchronization with the old system, because it is simply switched off. The Trickle approach allows debugging of scalability issues as the workload increases, but might mean that data must be migrated to and from legacy systems as the transition takes place.

It is difficult to recommend one approach over the other, because each method has associated risks that could lead to system outages as the transition takes place. Certainly, the Trickle approach allows profiling of real users as they are introduced to the new application, and allows the system to be reconfigured while only affecting the migrated users. This approach affects the work of the early adopters, but limits the load on support services. This means that unscheduled outages only affect a small percentage of the user population.

The decision on how to roll out a new application is specific to each business. Any adopted approach has its own unique pressures and stresses. The more testing and knowledge that you derive from the testing process, the more you realize what is best for the rollout.

Performance Checklist

To assist in the rollout, build a list of tasks that increase the chance of optimal performance in production and enable rapid debugging of the application. Do the following:

- When you create the control file for the production database, allow for growth by setting MAXINSTANCES, MAXDATAFILES, MAXLOGFILES, MAXLOGMEMBERS, and MAXLOGHISTORY to values higher than what you anticipate for the rollout. This technique results in more disk space usage and larger control files, but saves time later should these need extension in an emergency.
- Set block size to the value used to develop the application. Export the schema statistics from the development or test environment to the production database if the testing was done on representative data volumes and the current SQL execution plans are correct.
- Set the minimal number of initialization parameters. Ideally, most other parameters should be left at default. If there is more tuning to perform, then this appears when the system is under load.
- 4. Be prepared to manage block contention by setting storage options of database objects. Tables and indexes that experience high INSERT/UPDATE/DELETE rates should be created with automatic segment space management. To avoid contention of rollback segments, use automatic undo management.
- 5. All SQL statements should be verified to be optimal and their resource usage understood.
- 6. Validate that middleware and programs that connect to the database are efficient in their connection management and do not logon or logoff repeatedly.
- 7. Validate that the SQL statements use cursors efficiently. The database should parse each SQL statement once and then execute it multiple times. The most common reason this does not happen is because bind variables are not used properly and WHERE clause predicates are sent as string literals. If you use precompilers to develop the application, then make sure to reset the parameters MAXOPENCURSORS, HOLD_CURSOR, and RELEASE CURSOR from the default values before precompiling the application.
- 8. Validate that all schema objects have been correctly migrated from the development environment to the production database. This includes tables, indexes, sequences, triggers, packages, procedures, functions, Java objects, synonyms, grants, and views. Ensure that any modifications made in testing are made to the production system.
- 9. As soon as the system is rolled out, establish a baseline set of statistics from the database and operating system. This first set of statistics validates or corrects any assumptions made in the design and rollout process.
- **10.** Start anticipating the first bottleneck (which is inevitable) and follow the Oracle performance method to make performance improvement.



Performance Improvement Methods

This chapter discusses Oracle Database improvement methods and contains the following sections:

- The Oracle Performance Improvement Method
- Emergency Performance Methods

The Oracle Performance Improvement Method

Oracle performance methodology helps you to identify performance problems in an Oracle database. This involves identifying bottlenecks and fixing them. It is recommended that changes be made to a system only after you have confirmed that there is a bottleneck.

Performance improvement, by its nature, is iterative. For this reason, removing the first bottleneck might not lead to performance improvement immediately, because another bottleneck might be revealed. Also, in some cases, if serialization points move to a more inefficient sharing mechanism, then performance could degrade. With experience, and by following a rigorous method of bottleneck elimination, applications can be debugged and made scalable.

Performance problems generally result from either a lack of throughput, unacceptable user/job response time, or both. The problem might be localized between application modules, or it might be for the entire system.

Before looking at any database or operating system statistics, it is crucial to get feedback from the most important components of the system: the users of the system and the people ultimately paying for the application. Typical user feedback includes statements like the following:

- "The online performance is so bad that it prevents my staff from doing their jobs."
- "The billing run takes too long."
- "When I experience high amounts of Web traffic, the response time becomes unacceptable, and I am losing customers."
- "I am currently performing 5000 trades a day, and the system is maxed out. Next month, we roll out to all our users, and the number of trades is expected to quadruple."

From candid feedback, it is easy to set critical success factors for any performance work. Determining the performance targets and the performance engineer's exit criteria make managing the performance process much simpler and more successful at all levels. These critical success factors are better defined in terms of real business goals rather than system statistics.

Some real business goals for these typical user statements might be:

- "The billing run must process 1,000,000 accounts in a three-hour window."
- "At a peak period on a Web site, the response time must not exceed five seconds for a page refresh."
- "The system must be able to process 25,000 trades in an eight-hour window."

The ultimate measure of success is the user's perception of system performance. The performance engineer's role is to eliminate any bottlenecks that degrade performance. These bottlenecks could be caused by inefficient use of limited shared resources or by abuse of shared resources, causing serialization. Because all shared resources are limited, the goal of a performance engineer is to maximize the number of business operations with efficient use of shared resources. At a very high level, the entire database server can be seen as a shared resource. Conversely, at a low level, a single CPU or disk can be seen as shared resources.

You can apply the Oracle performance improvement method until performance goals are met or deemed impossible. This process is highly iterative. Inevitably, some investigations may have little or no impact on database performance. Time and experience are necessary to develop the skills to accurately and quickly pinpoint critical bottlenecks. However, prior experience can sometimes work against the experienced engineer who neglects to use the data and statistics available. This type of behavior encourages database tuning by myth and folklore. This is a very risky, expensive, and unlikely to succeed method of database tuning.

The Automatic Database Diagnostic Monitor (ADDM) implements parts of the performance improvement method and analyzes statistics to provide automatic diagnosis of major performance issues. Using ADDM can significantly shorten the time required to improve the performance of a system.

Systems are so different and complex that hard and fast rules for performance analysis are impossible. In essence, the Oracle performance improvement method defines a way of working, but not a definitive set of rules. With bottleneck detection, the only rule is that there are no rules! The best performance engineers use the data provided and think laterally to determine performance problems.

Steps in the Oracle Performance Improvement Method

- Perform the following initial standard checks:
 - a. Get candid feedback from users. Determine the performance project's scope and subsequent performance goals, and performance goals for the future. This process is key in future capacity planning.
 - b. Get a full set of operating system, database, and application statistics from the system when the performance is both good and bad. If these are not available, then get whatever is available. Missing statistics are analogous to missing evidence at a crime scene: They make detectives work harder and it is more time-consuming.
 - c. Sanity-check the operating systems of all computers involved with user performance. By sanity-checking the operating system, you look for hardware or operating system resources that are fully utilized. List any over-used resources as symptoms for analysis later. In addition, check that all hardware shows no errors or diagnostics.
- 2. Check for the top ten most common mistakes with Oracle Database, and determine if any of these are likely to be the problem. List these as symptoms for later analysis. These are included because they represent the most likely problems. ADDM automatically detects and reports nine of these top ten issues.
- Build a conceptual model of what is happening on the system using the symptoms as clues to understand what caused the performance problems. See "A Sample Decision Process for Performance Conceptual Modeling".
- 4. Propose a series of remedy actions and the anticipated behavior to the system, then apply them in the order that can benefit the application the most. ADDM produces recommendations each with an expected benefit. A golden rule in performance work is that you only change one thing at a time and then measure the differences. Unfortunately, system downtime requirements might prohibit such a rigorous investigation method. If



- multiple changes are applied at the same time, then try to ensure that they are isolated so that the effects of each change can be independently validated.
- 5. Validate that the changes made have had the desired effect, and see if the user's perception of performance has improved. Otherwise, look for more bottlenecks, and continue refining the conceptual model until your understanding of the application becomes more accurate.
- Repeat the last three steps until performance goals are met or become impossible due to other constraints.

This method identifies the biggest bottleneck and uses an objective approach to performance improvement. The focus is on making large performance improvements by increasing application efficiency and eliminating resource shortages and bottlenecks. In this process, it is anticipated that minimal (less than 10%) performance gains are made from instance tuning, and large gains (100% +) are made from isolating application inefficiencies.

A Sample Decision Process for Performance Conceptual Modeling

Conceptual modeling is almost deterministic. However, as you gain experience in performance tuning, you begin to appreciate that no real rules exist. A flexible heads-up approach is required to interpret statistics and make good decisions.

For a quick and easy approach to performance tuning, use ADDM. ADDM automatically monitors your Oracle system and provides recommendations for solving performance problems should problems occur. For example, suppose a DBA receives a call from a user complaining that the system is slow. The DBA simply examines the latest ADDM report to see which of the recommendations should be implemented to solve the problem.

The following steps illustrate how a performance engineer might look for bottlenecks without using automatic diagnostic features. These steps are only intended as a guideline for the manual process. With experience, performance engineers add to the steps involved. This analysis assumes that statistics for both the operating system and the database have been gathered.

- 1. Is the response time/batch run time acceptable for a single user on an empty or lightly loaded computer?
 - If it is not acceptable, then the application is probably not coded or designed optimally, and it will never be acceptable in a multiple user situation when system resources are shared. In this case, get application internal statistics, and get SQL Trace and SQL plan information. Work with developers to investigate problems in data, index, transaction SQL design, and potential deferral of work to batch and background processing.
- 2. Is all the CPU being utilized?

If the kernel utilization is over 40%, then investigate the operating system for network transfers, paging, swapping, or process thrashing. Continue to check CPU utilization in user space to verify if there are any non-database jobs consuming CPU on the system limiting the amount of shared CPU resources, such as backups, file transforms, print queues, and so on. After determining that the database is using most of the CPU, investigate the top SQL by CPU utilization. These statements form the basis of all future analysis. Check the SQL and the transactions submitting the SQL for optimal execution. Oracle Database provides CPU statistics in V\$SQL and V\$SQLSTATS.



Oracle Database Reference for more information about V\$SQL and V\$SQLSTATS

If the application is optimal and no inefficiencies exist in the SQL execution, then consider rescheduling some work to off-peak hours or using a bigger computer.

At this point, the system performance is unsatisfactory, yet the CPU resources are not fully utilized.

In this case, you have serialization and unscalable behavior within the server. Get the $\mathtt{WAIT_EVENTS}$ statistics from the server, and determine the biggest serialization point. If there are no serialization points, then the problem is most likely outside the database, and this should be the focus of investigation. Elimination of $\mathtt{WAIT_EVENTS}$ involves modifying application SQL and tuning database parameters. This process is very iterative and requires the ability to drill down on the $\mathtt{WAIT_EVENTS}$ systematically to eliminate serialization points.

Top Ten Mistakes Found in Oracle Systems

This section lists the most common mistakes found in Oracle databases. By following the Oracle performance improvement methodology, you should be able to avoid these mistakes altogether. If you find these mistakes in your system, then re-engineer the application where the performance effort is worthwhile.

1. Bad connection management

The application connects and disconnects for each database interaction. This problem is common with stateless middleware in application servers. It has over two orders of magnitude impact on performance, and is totally unscalable.

2. Bad use of cursors and the shared pool

Not using cursors results in repeated parses. If bind variables are not used, then there is hard parsing of all SQL statements. This has an order of magnitude impact in performance, and it is totally unscalable. Use cursors with bind variables that open the cursor and execute it many times. Be suspicious of applications generating dynamic SQL.

Bad SQL

Bad SQL is SQL that uses more resources than appropriate for the application requirement. This can be a decision support systems (DSS) query that runs for more than 24 hours, or a query from an online application that takes more than a minute. You should investigate SQL that consumes significant system resources for potential improvement. ADDM identifies high load SQL. SQL Tuning Advisor can provide recommendations for improvement.

4. Use of nonstandard initialization parameters

These might have been implemented based on poor advice or incorrect assumptions. Most databases provide acceptable performance using only the set of basic parameters. In particular, parameters associated with SPIN_COUNT on latches and undocumented optimizer features can cause a great deal of problems that can require considerable investigation.

Likewise, optimizer parameters set in the initialization parameter file can override proven optimal execution plans. For these reasons, schemas, schema statistics, and optimizer settings should be managed as a group to ensure consistency of performance.



- Oracle Database Administrator's Guide for information about initialization parameters and database creation
- Oracle Database Reference for details on initialization parameters

5. Getting database I/O wrong

Many sites lay out their databases poorly over the available disks. Other sites specify the number of disks incorrectly, because they configure disks by disk space and not I/O bandwidth.

6. Online redo log setup problems

Many sites run with too few online redo log files and files that are too small. Small redo log files cause system checkpoints to continuously put a high load on the buffer cache and I/O system. If too few redo log files exist, then the archive cannot keep up, and the database must wait for the archiver to catch up.

Serialization of data blocks in the buffer cache due to lack of free lists, free list groups, transaction slots (INITRANS), or shortage of rollback segments.

This is particularly common on INSERT-heavy applications, in applications that have raised the block size above 8K, or in applications with large numbers of active users and few rollback segments. Use automatic segment-space management (ASSM) and automatic undo management to solve this problem.

8. Long full table scans

Long full table scans for high-volume or interactive online operations could indicate poor transaction design, missing indexes, or poor SQL optimization. Long table scans, by nature, are I/O intensive and unscalable.

9. High amounts of recursive (SYS) SQL

Large amounts of recursive SQL executed by SYS could indicate space management activities, such as extent allocations, taking place. This is unscalable and impacts user response time. Use locally managed tablespaces to reduce recursive SQL due to extent allocation. Recursive SQL executed under another user ID is probably SQL and PL/SQL, and this is not a problem.

10. Deployment and migration errors

In many cases, an application uses too many resources because the schema owning the tables has not been successfully migrated from the development environment or from an older implementation. Examples of this are missing indexes or incorrect statistics. These errors can lead to sub-optimal execution plans and poor interactive user performance. When migrating applications of known performance, export the schema statistics to maintain plan stability using the DBMS STATS package.

Although these errors are not directly detected by ADDM, ADDM highlights the resulting high load SQL.

Emergency Performance Methods

This section provides techniques for dealing with performance emergencies. You presumably have a methodology for establishing and improving application performance. However, in an

emergency situation, a component of the system has changed to transform it from a reliable, predictable system to one that is unpredictable and not satisfying user requests.

In this case, the performance engineer must rapidly determine what has changed and take appropriate actions to resume normal service as quickly as possible. In many cases, it is necessary to take immediate action, and a rigorous performance improvement project is unrealistic.

After addressing the immediate performance problem, the performance engineer must collect sufficient debugging information either to get better clarity on the performance problem or to at least ensure that it does not happen again.

The method for debugging emergency performance problems is the same as the method described in the performance improvement method earlier in this book. However, shortcuts are taken in various stages because of the timely nature of the problem. Keeping detailed notes and records of facts found as the debugging process progresses is essential for later analysis and justification of any remedial actions. This is analogous to a doctor keeping good patient notes for future reference.

Steps in the Emergency Performance Method

The Emergency Performance Method is as follows:

- Survey the performance problem and collect the symptoms of the performance problem.
 This process should include the following:
 - User feedback on how the system is underperforming. Is the problem throughput or response time?
 - Ask the question, "What has changed since we last had good performance?" This
 answer can give clues to the problem. However, getting unbiased answers in an
 escalated situation can be difficult. Try to locate some reference points, such as
 collected statistics or log files, that were taken before and after the problem.
 - Use automatic tuning features to diagnose and monitor the problem. In addition, you
 can use Oracle Enterprise Manager Cloud Control (Cloud Control) performance
 features to identify top SQL and sessions.
- 2. Sanity-check the hardware utilization of all components of the application system. Check where the highest CPU utilization is, and check the disk, memory usage, and network performance on all the system components. This quick process identifies which tier is causing the problem. If the problem is in the application, then shift analysis to application debugging. Otherwise, move on to database server analysis.
- 3. Determine if the database server is constrained on CPU or if it is spending time waiting on wait events. If the database server is CPU-constrained, then investigate the following:
 - Sessions that are consuming large amounts of CPU at the operating system level and database; check V\$SESS TIME MODEL for database CPU usage
 - Sessions or statements that perform many buffer gets at the database level; check V\$SESSTAT and V\$SQLSTATS
 - Execution plan changes causing sub-optimal SQL execution; these can be difficult to locate
 - Incorrect setting of initialization parameters
 - Algorithmic issues caused by code changes or upgrades of all components

If the database sessions are waiting on events, then follow the wait events listed in $V\$SESSION_WAIT$ to determine what is causing serialization. The



V\$ACTIVE_SESSION_HISTORY view contains a sampled history of session activity which you can use to perform diagnosis even after an incident has ended and the system has returned to normal operation. In cases of massive contention for the library cache, it might not be possible to logon or submit SQL to the database. In this case, use historical data to determine why there is suddenly contention on this latch. If most waits are for I/O, then examine V\$ACTIVE_SESSION_HISTORY to determine the SQL being run by the sessions that are performing all of the inputs and outputs.

- 4. Apply emergency action to stabilize the system. This could involve actions that take parts of the application off-line or restrict the workload that can be applied to the system. It could also involve a system restart or the termination of job in process. These naturally have service level implications.
- 5. Validate that the system is stable. Having made changes and restrictions to the system, validate that the system is now stable, and collect a reference set of statistics for the database. Now follow the rigorous performance method described earlier in this book to bring back all functionality and users to the system. This process may require significant application re-engineering before it is complete.



4

Configuring a Database for Performance

This chapter contains an overview of the Oracle methodology for configuring a database for performance. Although performance modifications can be made to Oracle Database on an ongoing basis, significant benefits can be gained by proper initial configuration of the database.

This chapter contains the following sections:

- Performance Considerations for Initial Instance Configuration
- Creating and Maintaining Tables for Optimal Performance
- Performance Considerations for Shared Servers
- Improved Client Connection Performance Due to Prespawned Processes

Performance Considerations for Initial Instance Configuration

The initial database instance configuration options that have important performance impact on the database are:

- Initialization Parameters
- Undo Space
- Redo Log Files
- Tablespaces

Note:

If you use the Database Configuration Assistant (DBCA) to create a database, then the supplied seed database includes the necessary basic initialization parameters and meet the performance recommendations that are mentioned in this document.

See Also:

- Oracle Database Administrator's Guide to learn how to create a database with the Database Configuration Assistant
- Oracle Database Administrator's Guide to learn how to create a database with a SQL statement

Initialization Parameters

A running Oracle database instance is configured using initialization parameters, which are set in the initialization parameter file. These parameters influence the behavior of the running instance, including influencing performance. In general, a very simple initialization file with few

relevant settings covers most situations, and the initialization file should not be the first place you expect to do performance tuning, except for the few parameters described in the following table.

The following table describes the parameters necessary in a minimal initialization file. Although these parameters are necessary, they have no performance impact.

Table 4-1 Necessary Initialization Parameters Without Performance Impact

Parameter	Description
DB_NAME	Name of the database. This should match the <code>ORACLE_SID</code> environment variable.
DB_DOMAIN	Location of the database in Internet dot notation.
OPEN_CURSORS	Limit on the maximum number of cursors (active SQL statements) for each session. The setting is application-dependent; 500 is recommended.
CONTROL_FILES	Set to contain at least two files on different disk drives to prevent failures from control file loss.
DB_FILES	Set to the maximum number of files that can assigned to the database.

The following table includes the most important parameters to set with performance implications:

Table 4-2 Important Initialization Parameters With Performance Impact

Parameter	Description
COMPATIBLE	Specifies the release with which the Oracle database must maintain compatibility. It lets you take advantage of the maintenance improvements of a new release immediately in your production systems without testing the new functionality in your environment. If your application was designed for a specific release of Oracle Database, and you are actually installing a later release, then you might want to set this parameter to the version of the previous release.
DB_BLOCK_SIZE	Sets the size of the Oracle database blocks stored in the database files and cached in the SGA. The range of values depends on the operating system, but it is typically 8192 for transaction processing systems and higher values for database warehouse systems.
SGA_TARGET	Specifies the total size of all SGA components. If SGA_TARGET is specified, then the buffer cache (DB_CACHE_SIZE), Java pool (JAVA_POOL_SIZE), large pool (LARGE_POOL_SIZE), and shared pool (SHARED_POOL_SIZE) memory pools are automatically sized.
PGA_AGGREGATE_TARGET	Specifies the target aggregate PGA memory available to all server processes attached to the instance.
PROCESSES	Sets the maximum number of processes that can be started by that instance. This is the most important primary parameter to set, because many other parameter values are deduced from this.
SESSIONS	This is set by default from the value of processes. However, if you are using the shared server, then the deduced value is likely to be insufficient.
UNDO_MANAGEMENT	Specifies the undo space management mode used by the database. The default is ${\tt AUTO}.$ If unspecified, the database uses ${\tt AUTO}.$



Table 4-2 (Cont.) Important Initialization Parameters With Performance Impact

Parameter	Description
UNDO_TABLESPACE	Specifies the undo tablespace to be used when an instance starts.

The following guides for more information about these initialization parameters:

- Oracle Database Administrator's Guide
- Oracle Database Reference

Undo Space

The database uses undo space to store data used for read consistency, recovery, and rollback statements. This data exists in one or more undo tablespaces. If you use the Database Configuration Assistant (DBCA) to create a database, then the undo tablespace is created automatically. To manually create an undo tablespace, add the UNDO TABLESPACE clause to the CREATE DATABASE statement.

To automate the management of undo data, Oracle Database uses **automatic undo management**, which transparently creates and manages undo segments. To enable automatic undo management, set the <code>UNDO_MANAGEMENT</code> initialization parameter to <code>AUTO</code> (the default setting). If unspecified, then the <code>UNDO_MANAGEMENT</code> initialization parameter uses the <code>AUTO</code> setting. Oracle strongly recommends using automatic undo management because it significantly simplifies database management and eliminates the need for any manual tuning of undo (rollback) segments. Manual undo management using rollback segments is supported for backward compatibility.

The V\$UNDOSTAT view contains statistics for monitoring and tuning undo space. Using this view, you can better estimate the amount of undo space required for the current workload. Oracle Database also uses this information to help tune undo usage. The V\$ROLLSTAT view contains information about the behavior of the undo segments in the undo tablespace.

To facility the re-use of the undo space, the undo tablespace is implemented as multiple circular buffers. Its size can fixed or auto-extensible with MAXSIZE. The undo data is used by active transactions for rollback and features reliant on undo, like AS OF SCN, flashback queries, and flashback data archive, which can use the undo data until the undo retention time is reached. Expired undo data can be overwritten by new transactions.



- Oracle Database 2 Day DBA and Oracle Enterprise Manager Cloud Control (Cloud Control) online help to learn about the Undo Management Advisor
- Oracle Database Administrator's Guide for information about managing undo space using automatic undo management
- Oracle Database Reference for more information about the V\$ROLLSTAT view
- Oracle Database Reference for more information about the V\$UNDOSTAT view

Redo Log Files

The size of the redo log files can influence performance, because the behavior of the database writer and archiver processes depend on the redo log sizes. Generally, larger redo log files provide better performance. Undersized log files increase checkpoint activity and reduce performance.

Although the size of the redo log files does not affect LGWR performance, it can affect DBWR and checkpoint behavior. Checkpoint frequency is affected by several factors, including log file size and the setting of the FAST_START_MTTR_TARGET initialization parameter. If the FAST_START_MTTR_TARGET parameter is set to limit the instance recovery time, Oracle Database automatically tries to checkpoint as frequently as necessary. Under this condition, the size of the log files should be large enough to avoid additional checkpointing due to under sized log files. The optimal size can be obtained by querying the OPTIMAL_LOGFILE_SIZE column from the V\$INSTANCE_RECOVERY view. You can also obtain sizing advice on the **Redo Log Groups** page of Oracle Enterprise Manager Cloud Control (Cloud Control).

It may not always be possible to provide a specific size recommendation for redo log files, but redo log files in the range of 100 MB to a few gigabytes are considered reasonable. Size online redo log files according to the amount of redo your system generates. A rough guide is to switch log files at most once every 20 minutes.

See Also:

Oracle Database Administrator's Guide for information about managing the online redo log

Tablespaces

If you use the Database Configuration Assistant (DBCA) to create a database, then the seed database automatically includes the necessary tablespaces. If you choose not to use DBCA, then you must create extra tablespaces after creating the database.

All databases should have several tablespaces in addition to the SYSTEM and SYSAUX tablespaces. These additional tablespaces include:

- A temporary tablespace, which is used for operations such as sorting
- An undo tablespace to contain information for read consistency, recovery, and undo statements



 At least one tablespace for application use (in most cases, applications require several tablespaces)

For extremely large tablespaces with many data files, you can run multiple ALTER TABLESPACE ... ADD DATAFILE statements in parallel. During tablespace creation, the data files that make up the tablespace are initialized with special empty block images. Temporary files are not initialized.

Oracle Database does this to ensure that it can write all data files in their entirety, but this can obviously be a lengthy process if done serially. Therefore, run multiple CREATE TABLESPACE statements concurrently to speed up tablespace creation. For permanent tables, the choice between local and global extent management on tablespace creation can greatly affect performance. For any permanent tablespace that has moderate to large insert, modify, or delete operations compared to reads, choose local extent management.

Permanent Tablespaces - Automatic Segment-Space Management

For permanent tablespaces, Oracle recommends using automatic segment-space management. Such tablespaces, often referred to as bitmap tablespaces, are locally managed tablespaces with bitmap segment space management.

See Also:

- Oracle Database Concepts for a discussion of free space management
- Oracle Database Administrator's Guide for more information on creating and using automatic segment-space management for tablespaces

Temporary Tablespaces

Properly configuring the temporary tablespace helps optimize disk sort performance. Temporary tablespaces can be dictionary-managed or locally managed. Oracle recommends the use of locally managed temporary tablespaces with a UNIFORM extent size of 1 MB.

You should monitor temporary tablespace activity to check how many extents the database allocates for the temporary segment. If an application extensively uses temporary tables, as in a situation when many users are concurrently using temporary tables, then the extent size could be set smaller, such as 256K, because every usage requires at least one extent. The EXTENT MANAGEMENT LOCAL clause is optional for temporary tablespaces because all temporary tablespaces are created with locally managed extents of a uniform size. The default for SIZE is 1M.

A temporary tablespace can grow to a very large size due to spikes in temp usage by queries. Sorts, hash joins, and query transformations are examples that might cause high temp usage. Automatic Temp Tablespace Sizing is a feature which shrinks the temporary tablespace in the background when temp usage has subsided. In addition, if temp usage is increasing, the feature would pre-emptively grow the temporary tablespace to ensure query performance is not impacted.

This feature alleviates the need for a DBA to manually size the temporary tablespace. These are the two scenarios which would benefit:

- Shrinking the temporary tablespace to reclaim unsued space.
- Growing the temporary tablespace in anticipation of high temp usage.



- Oracle Database Administrator's Guide for more information on managing temporary tablespaces
- Oracle Database Concepts for more information on temporary tablespaces
- Oracle Database SQL Language Reference for more information on using the CREATE and ALTER TABLESPACE statements with the TEMPORARY clause

Creating and Maintaining Tables for Optimal Performance

When installing applications, an initial step is to create all necessary tables and indexes. When you create a segment, such as a table, the database allocates space for the data. If subsequent database operations cause the data volume to increase and exceed the space allocated, then Oracle Database extends the segment.

When creating tables and indexes, note the following:

- Specify automatic segment-space management for tablespaces
 In this way Oracle Database automatically manages segment space for best performance.
- Set storage options carefully

Applications should carefully set storage options for the intended use of the table or index. This includes setting the value for PCTFREE. Note that using automatic segment-space management eliminates the necessity of specifying PCTUSED.



Use of free lists is not recommended. To use automatic segment-space management, create locally managed tablespaces, with the segment space management clause set to ${\tt AUTO}.$

Table Compression

You can store heap-organized tables in a compressed format that is transparent for any kind of application. Compressed data in a database block is self-contained, which means that all information needed to re-create the uncompressed data in a block is available within the block. A block is also compressed in the buffer cache. Table compression not only reduces the disk storage but also the memory usage, specifically the buffer cache requirements. Performance improvements are accomplished by reducing the amount of necessary I/O operations for accessing a table and by increasing the probability of buffer cache hits.

Oracle Database has an advanced compression option that enables you to boost the performance of any type of application workload—including data warehousing and OLTP applications—while reducing the disk storage that is required by the database. You can use the advanced compression feature for all types of data, including structured data, unstructured data, backup data, and network data.



Oracle Database Administrator's Guide for information about using table compression

Estimating the Compression Factor

Table compression works by eliminating column value repetitions within individual blocks. Duplicate values in all the rows and columns in a block are stored once at the beginning of the block, in what is called a symbol table for that block. All occurrences of such values are replaced with a short reference to the symbol table. The compression is higher in blocks that have more repeated values.

Before compressing large tables you should estimate the expected compression factor. The compression factor is defined as the number of blocks necessary to store the information in an uncompressed form divided by the number of blocks necessary for a compressed storage. The compression factor can be estimated by sampling a small number of representative data blocks of the table to be compressed and comparing the average number of records for each block for the uncompressed and compressed case. Experience shows that approximately 1000 data blocks provides a very accurate estimation of the compression factor. Note that the more blocks you are sampling, the more accurate the results become.

Tuning to Achieve a Better Compression Ratio

Oracle Database achieves a good compression factor in many cases with no special tuning. As a DBA or application developer, you can try to tune the compression factor by reorganizing the records when the compression takes place. Tuning can improve the compression factor slightly in some cases and substantially in other cases.

To improve the compression factor you must increase the likelihood of value repetitions within a data block. The achievable compression factor depends on the cardinality of a specific column or column pairs (representing the likelihood of column value repetitions) and on the average row length of those columns. Table compression not only compresses duplicate values of a single column but tries to use multi-column value pairs whenever possible. Without a detailed understanding of the data distribution it is very difficult to predict the most optimal order.

Using Attribute-Clustered Tables

An attribute-clustered table is a heap-organized table that stores data in close proximity on disk based on user-specified clustering directives. The directives determine if the data stored in a table is ordered based on specified columns, or on a special algorithm that permits multicolumn I/O reduction. Attribute clustering is only available for bulk insert operations—such as the INSERT/*+APPEND*/ or ALTER TABLE ... MOVE PARTITION commands—and is ignored for conventional DML.

By reducing physical I/O in conjunction with zone maps, using attribute-clustered tables can significant reduce the I/O costs of table scans. Furthermore, it can also improve data compression because data can be more easily compressed when the same values are closer to each other on disk.



- Oracle Database Concepts for information about attribute-clustered tables
- Oracle Database Data Warehousing Guide for information about using attributeclustered tables

Reclaiming Unused Space

Over time, it is common for segment space to become fragmented or for a segment to acquire a lot of free space as the result of update and delete operations. The resulting sparsely populated objects can suffer performance degradation during queries and DML operations. If an object does have space available for reclamation, then you can compact and shrink segments or deallocate unused space at the end of a segment.

Oracle Database provides a Segment Advisor that provides advice on whether an object has space available for reclamation based on the level of space fragmentation within an object.



Oracle Database Administrator's Guide for a discussion on managing space for schema objects and the Segment Advisor

Indexing Data

The most efficient time to create indexes is after data has been loaded. In this way, space management becomes simpler, and no index maintenance takes place for each row inserted. SQL*Loader automatically uses this technique, but if you are using other methods to do initial data load, then you may need to create indexes manually. Additionally, you can perform index creation in parallel using the PARALLEL clause of the CREATE INDEX statement. However, SQL*Loader is not able to parallelize index creation, so you must manually create indexes in parallel after loading data.



Oracle Database Utilities for information about SQL*Loader

Specifying Memory for Sorting Data

During index creation on tables that contain data, the data must be sorted. This sorting is done in the fastest possible way, if all available memory is used for sorting. Oracle recommends that you enable automatic sizing of SQL working areas by setting the PGA_AGGREGATE_TARGET initialization parameter.



- "Tuning the Program Global Area" for information about PGA memory management
- Oracle Database Reference for information about the PGA_AGGREGATE_TARGET initialization parameter

Performance Considerations for Shared Servers

Using shared servers reduces the number of processes and the amount of memory consumed on the database host. Shared servers are beneficial for databases where there are many OLTP users performing intermittent transactions.

Using shared servers rather than dedicated servers is also generally better for systems that have a high connection rate to the database. With shared servers, when a connect request is received, a dispatcher is available to handle concurrent connection requests. With dedicated servers, however, a connection-specific dedicated server is sequentially initialized for each connection request.

Performance of certain database features can improve when a shared server architecture is used, and performance of certain database features can degrade slightly when a shared server architecture is used. For example, a session can be prevented from migrating to another shared server while parallel execution is active.

A session can remain nonmigratable even after a request from the client has been processed, because not all the user information has been stored in the UGA. If a server were to process the request from the client, then the part of the user state that was not stored in the UGA would be inaccessible. To avoid this situation, individual shared servers often need to remain bound to a user session.

See Also:

- Oracle Database Administrator's Guide to learn how to manage shared servers
- Oracle Database Net Services Administrator's Guide to learn how to configure dispatchers for shared servers

When using some features, you may need to configure more shared servers, because some servers might be bound to sessions for an excessive amount of time.

This section discusses how to reduce contention for processes used by Oracle Database architecture:

- Identifying and Reducing Contention Using the Dispatcher-Specific Views
- Identifying Contention for Shared Servers

Identifying and Reducing Contention Using the Dispatcher-Specific Views

The following views provide dispatcher performance statistics:



- V\$DISPATCHER: general information about dispatcher processes
- V\$DISPATCHER RATE: dispatcher processing statistics

The V\$DISPATCHER_RATE view contains current, average, and maximum dispatcher statistics for several categories. Statistics with the prefix CUR_ are statistics for the current sample. Statistics with the prefix AVG_ are the average values for the statistics after the collection period began. Statistics with the prefix MAX_ are the maximum values for these categories after statistics collection began.

To assess dispatcher performance, query the V\$DISPATCHER_RATE view and compare the current values with the maximums. If your present system throughput provides adequate response time and current values from this view are near the average and less than the maximum, then you likely have an optimally tuned shared server environment.

If the current and average rates are significantly less than the maximums, then consider reducing the number of dispatchers. Conversely, if current and average rates are close to the maximums, then you might need to add more dispatchers. A general rule is to examine V\$DISPATCHER_RATE statistics during both light and heavy system use periods. After identifying your shared server load patterns, adjust your parameters accordingly.

If necessary, you can also mimic processing loads by running system stress tests and periodically polling V\$DISPATCHER_RATE statistics. Proper interpretation of these statistics varies from platform to platform. Different types of applications also can cause significant variations on the statistical values recorded in V\$DISPATCHER_RATE.

See Also:

 Oracle Database Reference for detailed information about the V\$DISPATCHER and V\$DISPATCHER RATE views

Reducing Contention for Dispatcher Processes

To reduce contention, consider the following points:

- Adding dispatcher processes
 - The total number of dispatcher processes is limited by the value of the initialization parameter MAX_DISPATCHERS. You might need to increase this value before adding dispatcher processes.
- Enabling connection pooling
 - When system load increases and dispatcher throughput is maximized, it is not necessarily a good idea to immediately add more dispatchers. Instead, consider configuring the dispatcher to support more users with connection pooling.
- Enabling Session Multiplexing
 - Multiplexing is used by a connection manager process to establish and maintain network sessions from multiple users to individual dispatchers. For example, several user processes can connect to one dispatcher by way of a single connection from a connection manager process. Session multiplexing is beneficial because it maximizes use of the dispatcher process connections. Multiplexing is also useful for multiplexing database link sessions between dispatchers.



- Oracle Database Administrator's Guide to learn how to configure dispatcher processes
- Oracle Database Net Services Administrator's Guide to learn how to configure connection pooling
- Oracle Database Reference to learn about the DISPATCHERS and MAX DISPATCHERS initialization parameters

Identifying Contention for Shared Servers

Steadily increasing wait times in the requests queue indicate contention for shared servers. To examine wait time data, use the dynamic performance view V\$QUEUE. This view contains statistics showing request queue activity for shared servers. By default, this view is available only to the user SYS and to other users with SELECT ANY TABLE system privilege, such as SYSTEM. Table 4-3 lists the columns showing the wait times for requests and the number of requests in the queue.

Table 4-3 Wait Time and Request Columns in V\$QUEUE

Column	Description
WAIT	Displays the total waiting time, in hundredths of a second, for all requests that have ever been in the queue
TOTALQ	Displays the total number of requests that have ever been in the queue

Monitor these statistics occasionally while your application is running by issuing the following SQL statement:

```
SELECT DECODE(TOTALQ, 0, 'No Requests',

WAIT/TOTALQ || ' HUNDREDTHS OF SECONDS') "AVERAGE WAIT TIME PER REQUESTS"

FROM V$QUEUE

WHERE TYPE = 'COMMON';
```

This guery returns the results of a calculation that show the following:

```
AVERAGE WAIT TIME PER REQUEST
.090909 HUNDREDTHS OF SECONDS
```

From the result, you can tell that a request waits an average of 0.09 hundredths of a second in the queue before processing.

You can also determine how many shared servers are currently running by issuing the following query:

```
SELECT COUNT(*) "Shared Server Processes"
FROM V$SHARED_SERVER
WHERE STATUS != 'QUIT';
```

The result of this query could look like the following:

```
Shared Server Processes
```



If you detect resource contention with shared servers, then first ensure that this is not a memory contention issue by examining the shared pool and the large pool. If performance remains poor, then you might want to create more resources to reduce shared server process contention. You can do this by modifying the optional server process initialization parameters:

- MAX_DISPATCHERS
- MAX SHARED SERVERS
- DISPATCHERS
- SHARED SERVERS



Oracle Database Administrator's Guide to learn how to set the shared server process initialization parameters

Improved Client Connection Performance Due to Prespawned Processes

Oracle Database prespawns pools of server processes when dedicated broker connection mode is enabled or threaded execution mode is enabled. In this case, whenever a client requests for a database connection, it gets a dedicated connection to an existing server process from the process pools, thus improving the efficiency of client connections.

The $V\$PROCESS_POOL$ view shows information about these server process pools, and you can manage these pools using the <code>DBMS_PROCESS</code> package.

See Also:

- Oracle Database Administrator's Guide for more information about managing prespawned processes in Oracle Database
- Oracle Database PL/SQL Packages and Types Reference for information about the DBMS PROCESS package



Part II

Diagnosing and Tuning Database Performance

This part contains the following chapters:

- Measuring Database Performance
- Gathering Database Statistics
- Automatic Performance Diagnostics
- Comparing Database Performance Over Time
- Analyzing Sampled Data
- Instance Tuning Using Performance Views



5

Measuring Database Performance

This chapter describes how to measure the performance of Oracle Database using database statistics.

This chapter contains the following topics:

- About Database Statistics
- Interpreting Database Statistics

About Database Statistics

Database statistics provide information about the type of database load and the resources being used by the database. To effectively measure database performance, statistics must be available.

Oracle Database generates many types of cumulative statistics for the system, sessions, segments, services, and individual SQL statements. Cumulative values for statistics are generally accessible using dynamic performance views, or v views. When analyzing database performance in any of these scopes, look at the change in statistics (delta value) over the period you are interested in. Specifically, focus on the difference between the cumulative values of a statistic at the start and the end of the period.

This section describes some of the more important database statistics that are used to measure the performance of Oracle Database:

- Time Model Statistics
- Active Session History Statistics
- Wait Events Statistics
- Session and System Statistics



Oracle Database SQL Tuning Guide for information about optimizer statistics

Time Model Statistics

Time model statistics use time to identify quantitative effects about specific actions performed on the database, such as logon operations and parsing. The most important time model statistic is database time, or DB time. This statistic represents the total time spent in database calls for foreground sessions and is an indicator of the total instance workload. DB time is measured cumulatively from the time of instance startup and is calculated by aggregating the CPU and wait times of all foreground sessions not waiting on idle wait events (non-idle user sessions).

Note:

Because DB time is calculated by combining the times from all non-idle user foreground sessions, it is possible that the DB time can exceed the actual time elapsed after the instance started. For example, an instance that has been running for 30 minutes could have four active user sessions whose cumulative DB time is approximately 120 minutes.

When tuning an Oracle database, each component has its own set of statistics. To look at the system as a whole, it is necessary to have a common scale for comparisons. Many Oracle Database advisors and reports thus describe statistics in terms of time.

Ultimately, the objective in tuning an Oracle database is to reduce the time that users spend in performing an action on the database, or to simply reduce DB time. Time model statistics are accessible from the V\$SESS TIME MODEL and V\$SYS TIME MODEL views.

See Also:

Oracle Database Reference for information about the $v\$sess_time_model$ and $v\$sys_time_model$ views

Active Session History Statistics

An active session is any session that is either connected to the database or on a CPU, and is waiting for an event that does not belong to the Idle wait class. Oracle Database samples active sessions every second and stores the sampled data in a circular buffer in the shared global area (SGA).

The sampled session activity is accessible using the V\$ACTIVE_SESSION_HISTORY view. Each session sample contains a set of rows and the V\$ACTIVE_SESSION_HISTORY view returns one row for each active session per sample, starting with the latest session sample rows. Because the active session samples are stored in a circular buffer in the SGA, the greater the system activity, the smaller the number of seconds of session activity that can be stored. This means that the duration for which a session sample is displayed in the V\$ view is completely dependent on the level of database activity. Because the content of the V\$ view can become quite large during heavy system activity, only a portion of the session samples is written to disk.

By capturing only active sessions, a manageable set of data can be captured with its size being directly related to the work being performed, rather than the number of sessions allowed on the system. Active Session History (ASH) enables you to examine and perform detailed analysis on both current data in the V\$ACTIVE_SESSION_HISTORY view and historical data in the DBA_HIST_ACTIVE_SESS_HISTORY view, often avoiding the need to replay the workload to trace additional performance information. ASH also contains execution plan information for each captured SQL statement. You can use this information to identify which part of SQL execution contributed most to the SQL elapsed time. The data present in ASH can be rolled up in various dimensions that it captures, including:

- SQL identifier of SQL statement
- SQL plan identifier and hash value of the SQL plan used to execute the SQL statement
- SQL execution plan information



- Object number, file number, and block number
- Wait event identifier and parameters
- Session identifier and session serial number
- Module and action name
- Client identifier of the session
- Service hash identifier
- Consumer group identifier

You can gather this information over a specified duration into an ASH report.

Active session history sampling is also available for Active Data Guard physical standby instances and Oracle Automatic Storage Management (Oracle ASM) instances. On these instances, the current session activity is collected and displayed in the V\$ACTIVE SESSION HISTORY view, but not written to disk.

See Also:

- "Analyzing Sampled Data" for information about ASH reports
- Oracle Data Guard Concepts and Administration for information about Active Data Guard physical standby databases
- Oracle Automatic Storage Management Administrator's Guide for information about Oracle ASM instances

Wait Events Statistics

Wait events are statistics that are incremented by a server process or thread to indicate that it had to wait for an event to complete before processing could continue. Wait event data reveals various symptoms of problems that might be impacting performance, such as latch contention, buffer contention, and I/O contention.

To enable easier high-level analysis of wait events, Oracle Database groups events into the following classes:

- Administrative
- Application
- Cluster
- Commit
- Concurrency
- Configuration
- Idle
- Network
- Other
- Scheduler
- System I/O
- User I/O



The wait classes are based on a common solution that usually applies to fixing a problem with the particular wait event. For example, exclusive TX locks are generally an application-level issue and HW locks are generally a configuration issue. The following list includes common examples of wait events in some of the wait classes:

- Application: locks waits caused by row level locking or explicit lock commands
- Commit: waits for redo log write confirmation after a commit
- Idle: wait events that signify the session is inactive, such as SQL*Net message from client
- Network: waits for data to be sent over the network
- User I/O: wait for blocks to be read off a disk

Wait event statistics for a database instance include statistics for both background and foreground processes. Because tuning is typically focused in foreground activities, overall database instance activity is categorized into foreground and background statistics in the relevant V\$ views to facilitate tuning.

The V\$SYSTEM_EVENT view shows wait event statistics for the foreground activities of a database instance and the wait event statistics for the database instance. The V\$SYSTEM_WAIT_CLASS view shows these foreground and wait event statistics at the instance level after aggregating to wait classes. V\$SESSION_EVENT and V\$SESSION_WAIT_CLASS show wait event and wait class statistics at the session level.



Oracle Database Reference for information about wait events

Session and System Statistics

A large number of cumulative database statistics on a system and session level are accessible using the V\$SYSSTAT and V\$SESSTAT views.



Oracle Database Reference for information about the V\$SYSSTAT and V\$SESSTAT views

Interpreting Database Statistics

When initially examining performance data, you can formulate potential interpretations of the data by examining the database statistics. To ensure that your interpretation is accurate, crosscheck with other data to establish if a statistic or event is truly relevant. Because foreground activities are tunable, it is recommended to first analyze the statistics from foreground activities before analyzing the statistics from background activities.

The following sections provide tips for interpreting the various types of database statistics to measure database performance:

Using Hit Ratios

- Using Wait Events with Timed Statistics
- Using Wait Events without Timed Statistics
- Using Idle Wait Events
- Comparing Database Statistics with Other Factors
- Using Computed Statistics

Using Hit Ratios

When tuning, it is common to compute a ratio that helps determine if a problem exists. Such ratios may include the buffer cache hit ratio, the soft-parse ratio, and the latch hit ratio. Do not use these ratios as definitive identifiers of whether a performance bottleneck exists. Instead, use them as indicators. To identify whether a performance bottleneck exists, examine other related performance data. For information about how to calculate the buffer cache hit ratio, see "Calculating the Buffer Cache Hit Ratio".

Using Wait Events with Timed Statistics

Setting TIMED_STATISTICS to TRUE at the instance level directs the database to gather wait time for events, in addition to available wait counts. This data is useful for comparing the total wait time for an event to the total elapsed time between the data collections. For example, if the wait event accounts for only 30 seconds out of a 2-hour period, then very little performance improvement can be gained by investigating this event, even if it is the highest ranked wait event when ordered by time waited. However, if the event accounts for 30 minutes of a 45-minute period, then the event is worth investigating. For information about wait events, see "Wait Events Statistics".

Note:

Timed statistics are automatically collected for the database if the initialization parameter <code>STATISTICS_LEVEL</code> is set to <code>TYPICAL</code> or <code>ALL</code>. If <code>STATISTICS_LEVEL</code> is set to <code>BASIC</code>, then you must set <code>TIMED_STATISTICS</code> to <code>TRUE</code> to enable collection of timed statistics. Note that setting <code>STATISTICS_LEVEL</code> to <code>BASIC</code> disables many automatic features and is not recommended.

If you explicitly set <code>DB_CACHE_ADVICE</code>, <code>TIMED_STATISTICS</code>, or <code>TIMED_OS_STATISTICS</code>, either in the initialization parameter file or by using <code>ALTER_SYSTEM</code> or <code>ALTER SESSION</code>, then the explicitly set value overrides the value derived from <code>STATISTICS LEVEL</code>.

See Also:

Oracle Database Reference for information about the STATISTICS_LEVEL initialization parameter

Using Wait Events without Timed Statistics

If TIMED_STATISTICS is set to FALSE, then the amount of time spent waiting for an event is not available. Therefore, it is only possible to order wait events by the number of times each event

was waited for. Although the events with the largest number of waits might indicate a potential bottleneck, they might not be the main bottleneck. This situation can happen when an event is waited for a large number of times, but the total time waited for that event is small. Conversely, an event with fewer waits might be a bigger bottleneck if the wait time accounts for a significant proportion of the total wait time. Without the wait times to use for comparison, it is difficult to determine whether a wait event is worth investigating.

Using Idle Wait Events

Oracle Database uses some wait events to indicate whether the Oracle server process is idle. Typically, these events are of no value when investigating performance problems, and should be ignored when examining wait events.

Comparing Database Statistics with Other Factors

When evaluating statistics, it is important to consider other factors that may influence whether the statistic is of value. Such factors may include the user load and hardware capability. Even an event that had a wait of 30 minutes in a 45-minute period might not be indicative of a performance problem if you discover that there were 2000 users on the system, and the host hardware was a 64-node computer.

Using Computed Statistics

When interpreting computed statistics (such as rates, statistics normalized over transactions, or ratios), verify the computed statistic with the actual statistic counts. This comparison can confirm whether the derived rates are really of interest because small statistic counts usually can discount an unusual ratio. For example, on initial examination, a soft-parse ratio of 50% generally indicates a potential area for tuning. If, however, there was only one hard parse and one soft parse during the data collection interval, then the soft-parse ratio would be 50%, even though the statistic counts show this is not impacting performance. In this case, the ratio is not important due to the low raw statistic counts.



6

Gathering Database Statistics

This chapter describes how to gather database statistics for Oracle Database and contains the following topics:

- About Gathering Database Statistics
- Managing the Automatic Workload Repository
- Generating Automatic Workload Repository Reports
- Generating Performance Hub Active Report

About Gathering Database Statistics

Oracle Database automatically persists the cumulative and delta values for most of the statistics at all levels (except the session level) in the Automatic Workload Repository (AWR). This process is repeated on a regular time period and the results are captured in an AWR snapshot. The delta values captured by the snapshot represent the changes for each statistic over the time period.

A statistical baseline is a collection of statistic rates usually taken over a time period when the system is performing well at an optimal level. Use statistical baselines to diagnose performance problems by comparing statistics captured in a baseline to those captured during a period of poor performance. This enables you to identify specific statistics that may have increased significantly and could be the cause of the problem. AWR supports the capture of baseline data by enabling you to specify and preserve a pair or range of AWR snapshots as a baseline.

A metric is typically the rate of change in a cumulative statistic. You can measure this rate against a variety of units, including time, transactions, or database calls. For example, the number database calls per second is a metric. Metric values are exposed in some V\$ views, where the values are the averages over a fairly small time interval, typically 60 seconds. A history of recent metric values is available through V\$ views, and some data is also persisted by AWR snapshots.

The following sections describe various Oracle Database features that enable you to more effectively gather database statistics:

- Automatic Workload Repository
- Snapshots
- Baselines
- Space Consumption
- Adaptive Thresholds



- Data visibility and privilege requirements may differ when using AWR features with pluggable databases (PDBs). For information about how manageability features, including the AWR features, work in a multitenant container database (CDB), see Oracle Multitenant Administrator's Guide.
- License for Oracle Diagnostic Pack is required to use the AWR features described in this chapter.

Automatic Workload Repository

AWR collects, processes, and maintains performance statistics for problem detection and self-tuning purposes. This gathered data is stored both in memory and in the database, and is displayed in both reports and views.

The statistics collected and processed by AWR include:

- Object statistics that determine both access and usage statistics of database segments
- Time model statistics based on time usage for activities, displayed in the V\$SYS TIME MODEL and V\$SESS TIME MODEL views
- Some of the system and session statistics collected in the V\$SYSSTAT and V\$SESSTAT views
- SQL statements that are producing the highest load on the system, based on criteria such as elapsed time and CPU time
- Active Session History (ASH) statistics, representing the history of recent sessions activity

See Also:

- "About Database Statistics" for information about the various types of database statistics
- Oracle Database Reference for more information about the views
 V\$SYS_TIME_MODEL, V\$SESS_TIME_MODEL, V\$SYSSTAT, and V\$SESSTAT

Snapshots

Snapshots are sets of historical data for specific time periods that are used for performance comparisons by Automatic Database Diagnostic Monitor (ADDM). By default, Oracle Database automatically generates snapshots of the performance data once every hour and retains the statistics in AWR for 8 days. You can also manually create snapshots or change the snapshot retention period, but it is usually not necessary.

AWR compares the difference between snapshots to determine which SQL statements to capture based on the effect on the system load. This reduces the number of SQL statements that must be captured over time. After the snapshots are created, ADDM analyzes the data captured in the snapshots to perform its performance analysis.



See Also:

"Managing Snapshots" for information about managing snapshots

Baselines

A baseline is a set of snapshots from a specific time period that is preserved for comparison with other snapshots when a performance problem occurs. The snapshots contained in a baseline are excluded from the automatic AWR purging process and are retained indefinitely.

There are several types of available baselines:

- Fixed Baselines
- Moving Window Baselines
- Baseline Templates

Fixed Baselines

A fixed baseline corresponds to a fixed, contiguous time period in the past that you specify. Before creating a fixed baseline, carefully consider the time period you choose as a baseline, because the baseline should represent the system operating at an optimal level. In the future, you can compare the baseline with other baselines or snapshots captured during periods of poor performance to analyze performance degradation over time.

See Also:

"Managing Baselines" for information about managing fixed baselines

Moving Window Baselines

A moving window baseline corresponds to all AWR data that exists within the AWR retention period. This is useful when using adaptive thresholds because the database can use AWR data in the entire AWR retention period to compute metric threshold values.

Oracle Database automatically maintains a system-defined moving window baseline. The default window size for the system-defined moving window baseline is the current AWR retention period, which by default is 8 days. If you are planning to use adaptive thresholds, then consider using a larger moving window—such as 30 days—to accurately compute threshold values. You can resize the moving window baseline by changing the number of days in the moving window to a value that is equal to or less than the number of days in the AWR retention period. Therefore, to increase the size of a moving window, you must first increase the AWR retention period accordingly.

See Also:

"Resizing the Default Moving Window Baseline" for information about resizing a moving window baseline



Baseline Templates

Baseline templates enable you to create baselines for a contiguous time period in the future. There are two types of baseline templates:

- Single Baseline Templates
- Repeating Baseline Templates



"Managing Baseline Templates" for information about managing baseline templates

Single Baseline Templates

Use a single baseline template to create a baseline for a single contiguous time period in the future. This is useful if you know beforehand of a time period that you intend to capture in the future. For example, you may want to capture AWR data during a system test that is scheduled for the upcoming weekend. In this case, you can create a single baseline template to automatically capture the time period when the test occurs.

Repeating Baseline Templates

Use a repeating baseline template to create and drop baselines based on a repeating time schedule. This is useful if you want Oracle Database to automatically capture a contiguous time period on an ongoing basis. For example, you may want to capture AWR data during every Monday morning for a month. In this case, you can create a repeating baseline template to automatically create baselines on a repeating schedule for every Monday, and automatically remove older baselines after a specified expiration interval, such as one month.

Space Consumption

The space consumed by AWR is determined by several factors:

- Number of active sessions in the database at any given time
- Snapshot interval

The snapshot interval determines the frequency at which snapshots are captured. A smaller snapshot interval increases the frequency, which increases the volume of data collected by AWR.

Historical data retention period

The retention period determines how long this data is retained before being purged. A longer retention period increases the space consumed by AWR.

By default, Oracle Database captures snapshots once every hour and retains them in the database for 8 days. With these default settings, a typical system with an average of 10 concurrent active sessions can require approximately 200 to 300 MB of space for its AWR data.

To reduce AWR space consumption, increase the snapshot interval and reduce the retention period. When reducing the retention period, note that several Oracle Database self-managing features depend on AWR data for proper functioning. Not having enough data can affect the validity and accuracy of these components and features, including:



- Automatic Database Diagnostic Monitor (ADDM)
- SQL Tuning Advisor
- Undo Advisor
- Segment Advisor

If possible, Oracle recommends that you set the AWR retention period large enough to capture at least one complete workload cycle. If your system experiences weekly workload cycles, such as OLTP workload during weekdays and batch jobs during the weekend, then you do not need to change the default AWR retention period of 8 days. However, if your system is subjected to a monthly peak load during month-end book closing, then you may need to set the retention period to one month.

Under exceptional circumstances, you can disable automatic snapshot collection by setting the snapshot interval to 0. Under this condition, the automatic collection of the workload and statistical data is stopped, and most of the Oracle Database self-management functionality is not operational. In addition, you cannot manually create snapshots. For this reason, Oracle strongly recommends against disabling automatic snapshot collection.



Oracle Database uses the SYSAUX tablespace to store AWR data by default.

See Also:

"Modifying Snapshot Settings" for information about changing the default values for the snapshot interval and retention period

Adaptive Thresholds

Adaptive thresholds enable you to monitor and detect performance issues, while minimizing administrative overhead. Adaptive thresholds automatically set warning and critical alert thresholds for some system metrics using statistics derived from metric values captured in the moving window baseline. The statistics for these thresholds are recomputed weekly and might result in new thresholds as system performance evolves over time. Additionally, adaptive thresholds can compute different thresholds values for different times of the day or week based on periodic workload patterns.

For example, many databases support an online transaction processing (OLTP) workload during the day and batch processing at night. The performance metric for response time per transaction can be useful for detecting degradation in OLTP performance during the day. However, a useful OLTP threshold value is usually too low for batch workloads, where long-running transactions might be common. As a result, threshold values appropriate to OLTP might trigger frequent false performance alerts during batch processing. Adaptive thresholds can detect such a workload pattern and automatically set different threshold values for daytime and nighttime.

There are two types of adaptive thresholds:

- Percentage of Maximum Thresholds
- Significance Level Thresholds



Percentage of Maximum Thresholds

The threshold value for percentage of maximum thresholds is computed as a percentage multiple of the maximum value observed for the data in the moving window baseline.

Percentage of maximum thresholds are most useful when a system is sized for peak workloads, and you want to be alerted when the current workload volume approaches or exceeds previous high values. Metrics that have an unknown but definite limiting value are prime candidates for these settings. For example, the redo generated per second metric is typically a good candidate for a percentage of maximum threshold.

Significance Level Thresholds

The threshold value for significance level thresholds is set to a statistical percentile that represents how unusual it is to observe values above the threshold value based the data in the moving window baseline.

Significance level thresholds are most useful for metrics that exhibit statistically stable behavior when the system is operating normally, but might vary over a wide range when the system is performing poorly. For example, the response time per transaction metric should be stable for a well-tuned OLTP system, but may fluctuate widely when performance issues arise. Significance level thresholds are meant to generate alerts when conditions produce both unusual metric values and unusual system performance.

Significance level thresholds can be set to one of the following levels:

High (.95)

Only 5 in 100 observations are expected to exceed this value.

Very High (.99)

Only 1 in 100 observations are expected to exceed this value.

Severe (.999)

Only 1 in 1,000 observations are expected to exceed this value.

Extreme (.9999)

Only 1 in 10,000 observations are expected to exceed this value.

When you specify a significance level threshold, Oracle Database performs an internal calculation to set the threshold value. In some cases, Oracle Database cannot establish the threshold value at higher significance levels using the data in the baseline, and the significance level threshold is not set.

If you specified a Severe (.999) or Extreme (.9999) significance level threshold and are not receiving alerts as expected, try setting the significance level threshold to a lower value, such as Very High (.99) or High (.95). Alternatively, consider using a percentage of maximum threshold instead. If you change the threshold and find that you are receiving too many alerts, try increasing the number of occurrences to trigger an alert.



The primary interface for managing baseline metrics is Oracle Enterprise Manager Cloud Control (Cloud Control). To create an adaptive threshold for a baseline metric, use Cloud Control as described in *Oracle Database 2 Day + Performance Tuning Guide*.

✓ See Also:

- "Moving Window Baselines" for information about moving window baselines
- "Managing Baselines" for information about managing baseline metrics

Managing the Automatic Workload Repository

This section describes how to manage AWR features of Oracle Database and contains the following topics:

- Enabling the Automatic Workload Repository
- Managing Snapshots
- Managing Baselines
- · Managing Baseline Templates
- Transporting Automatic Workload Repository Data to Another System
- Using Automatic Workload Repository Views
- Managing Automatic Workload Repository in a Multitenant Environment
- Managing Automatic Workload Repository in Active Data Guard Standby Databases



"Automatic Workload Repository" for a description of AWR

Enabling the Automatic Workload Repository

Gathering database statistics using AWR is enabled by default and is controlled by the STATISTICS LEVEL initialization parameter.

The database can be configured so that the AWR reports use the database time zone instead of the operating system time zone. Refer to TIME_AT_DBTIMEZONE.

To enable statistics gathering by AWR:

• Set the STATISTICS LEVEL parameter to TYPICAL or ALL.

The default setting for this parameter is TYPICAL.



Setting STATISTICS_LEVEL to BASIC disables many Oracle Database features, including AWR, and is not recommended. If STATISTICS_LEVEL is set to BASIC, you can still manually capture AWR statistics using the DBMS_WORKLOAD_REPOSITORY package. However, because in-memory collection of many system statistics—such as segments statistics and memory advisor information—will be disabled, the statistics captured in these snapshots may not be complete.

See Also:

Oracle Database Reference for information about the STATISTICS_LEVEL initialization parameter

Managing Snapshots

By default, Oracle Database generates snapshots once every hour, and retains the statistics in the workload repository for 8 days. When necessary, you can manually create or drop snapshots and modify snapshot settings.

This section describes how to manage snapshots and contains the following topics:

- User Interfaces for Managing Snapshots
- Creating Snapshots
- Dropping Snapshots
- Modifying Snapshot Settings

See Also:

"Snapshots" for information about snapshots

User Interfaces for Managing Snapshots

The primary interface for managing snapshots is Oracle Enterprise Manager Cloud Control (Cloud Control). Whenever possible, you should manage snapshots using Cloud Control.

If Cloud Control is unavailable, then manage snapshots using the <code>DBMS_WORKLOAD_REPOSITORY</code> package in the command-line interface. The DBA role is required to invoke the <code>DBMS_WORKLOAD_REPOSITORY</code> procedures.

See Also:

Oracle Database PL/SQL Packages and Types Reference for information about the DBMS WORKLOAD REPOSITORY package



Creating Snapshots

By default, Oracle Database automatically generates snapshots once every hour. However, you may want to manually create snapshots to capture statistics at times different from those of the automatically generated snapshots.

Creating Snapshots Using the Command-Line Interface

To manually create snapshots, use the <code>CREATE_SNAPSHOT</code> procedure. The following example shows a <code>CREATE_SNAPSHOT</code> procedure call.

```
BEGIN
   DBMS_WORKLOAD_REPOSITORY.CREATE_SNAPSHOT ();
END;
/
```

In this example, a snapshot is created immediately on the local database instance. To view information about an existing snapshot, use the DBA HIST SNAPSHOT view.

Note:

You can specify value for the <code>flush_level</code> parameter of the <code>CREATE_SNAPSHOT</code> procedure to either TYPICAL or ALL. The default value for the flush level is TYPICAL.

The flush level signifies the breadth and depth of the AWR statistics to be captured. If you want to capture all the AWR statistics, then set the flush level to ALL. If you want to skip few AWR statistics, such as, SQL statistics, segment statistics, and files and tablespace statistics for performance reasons, then set the flush level to TYPICAL.

See Also:

- Oracle Database PL/SQL Packages and Types Reference for information about the DBMS WORKLOAD REPOSITORY package
- Oracle Database Reference for information about the DBA HIST SNAPSHOT view

Dropping Snapshots

By default, Oracle Database automatically purges snapshots that have been stored in AWR for over 8 days. However, you may want to manually drop a range of snapshots to free up space.

Dropping Snapshots Using the Command-Line Interface

To manually drop a range of snapshots, use the DROP_SNAPSHOT_RANGE procedure. The following example shows a DROP SNAPSHOT RANGE procedure call.



dbid => 3310949047); END;

In the example, snapshots with snapshot IDs ranging from 22 to 32 are dropped immediately from the database instance with the database identifier of 3310949047. Any ASH data that were captured during this snapshot range are also purged.



Tip:

To determine which snapshots to drop, use the $\mbox{DBA_HIST_SNAPSHOT}$ view to review the existing snapshots

See Also:

- Oracle Database PL/SQL Packages and Types Reference for information about the DBMS WORKLOAD REPOSITORY package
- Oracle Database Reference for information about the DBA HIST SNAPSHOT view

Modifying Snapshot Settings

You can adjust the interval, retention period, and number of top SQL to flush for snapshot generation, but note that this can affect the precision of the Oracle Database diagnostic tools.

Modifying Snapshot Settings Using the Command-Line Interface

You can modify snapshot settings using the following parameters of the DBMS_WORKLOAD_REPOSITORY.MODIFY_SNAPSHOT_SETTINGS procedure:

Parameter	Description
INTERVAL	This setting affects how often the database automatically generates snapshots.
RETENTION	This setting affects how long the database stores snapshots in AWR.
TOPNSQL	This setting affects the number of top SQL to flush for each SQL criteria (elapsed time, CPU time, parse calls, sharable memory, and version count).
	This setting is not affected by the statistics/flush level and overrides the system default behavior for AWR SQL collection. It is possible to set the value for this setting to MAXIMUM to capture the complete set of SQL in the shared SQL area, though doing so (or by setting the value to a very high number) may lead to possible space and performance issues because there will be more data to collect and store.

The following example shows how to modify snapshot settings using the DBMS_WORKLOAD_REPOSITORY.MODIFY_SNAPSHOT_SETTINGS procedure:

In this example, snapshot settings for the database with the database identifier of 3310949047 are modified as follows:

- The retention period is specified as 43200 minutes (30 days).
- The interval between each snapshot is specified as 30 minutes.
- The number of top SQL to flush for each SQL criteria is specified as 100.

To get information about the current snapshot settings for your database, use the DBA HIST WR CONTROL view as shown in the following example:

```
SQL> select snap_interval, retention from DBA_HIST_WR_CONTROL;

SNAP_INTERVAL RETENTION
+00000 01:00:00.0 +00008 00:00:00.0
```

The snap interval and retention values are displayed in the format:

```
+[days] [hours]:[minutes]:[seconds].[nanoseconds]
```

See Also:

- Oracle Database PL/SQL Packages and Types Reference for more information about the DBMS WORKLOAD REPOSITORY.MODIFY SNAPSHOT SETTINGS procedure
- Oracle Database Reference for more information about the DBA HIST WR CONTROL view

Managing Baselines

By default, Oracle Database automatically maintains a system-defined moving window baseline. When necessary, you can manually create, drop, or rename a baseline and view the baseline threshold settings. Additionally, you can manually resize the window size of the moving window baseline.

This section describes how to manage baselines and contains the following topics:

- User Interface for Managing Baselines
- Creating a Baseline
- Dropping a Baseline
- · Renaming a Baseline
- Displaying Baseline Metrics
- Resizing the Default Moving Window Baseline





"Baselines" for information about baselines

User Interface for Managing Baselines

The primary interface for managing baselines is Oracle Enterprise Manager Cloud Control (Cloud Control). Whenever possible, manage baselines using Cloud Control.

If Cloud Control is unavailable, then manage baselines using the DBMS_WORKLOAD_REPOSITORY package in the command-line interface. The DBA role is required to invoke the DBMS WORKLOAD REPOSITORY procedures.

See Also:

- Oracle Database 2 Day + Performance Tuning Guide for more information about managing baselines using Cloud Control
- Oracle Database PL/SQL Packages and Types Reference for information about the DBMS WORKLOAD REPOSITORY package

Creating a Baseline

By default, Oracle Database automatically maintains a system-defined moving window baseline. However, you may want to manually create a fixed baseline that represents the system operating at an optimal level, so you can compare it with other baselines or snapshots captured during periods of poor performance.

To create baselines using command-line interface, use the <code>CREATE_BASELINE</code> procedure as shown in the following example:

```
BEGIN

DBMS_WORKLOAD_REPOSITORY.CREATE_BASELINE (start_snap_id => 270,
end_snap_id => 280,
baseline_name => 'peak baseline',
dbid => 3310949047,
expiration => 30);

END;
/
```

In this example, a baseline is created on the database instance with the database identifier of 3310949047 with the following settings:

- The start snapshot sequence number is 270.
- The end snapshot sequence number is 280.
- The name of baseline is peak baseline.
- The expiration of the baseline is 30 days.

Oracle Database automatically assigns a unique ID to the new baseline when the baseline is created.



Tip:

To determine the range of snapshots to include in a baseline, use the $\verb|DBA_HIST_SNAPSHOT|$ view to review the existing snapshots

See Also:

- Oracle Database PL/SQL Packages and Types Reference for information about the DBMS WORKLOAD REPOSITORY package
- Oracle Database Reference for information about the DBA HIST SNAPSHOT view

Dropping a Baseline

To conserve disk space, consider periodically dropping a baseline that is no longer being used. The snapshots associated with a baseline are retained indefinitely until you explicitly drop the baseline or the baseline has expired.

To drop a baseline using command-line interface, use the <code>DROP_BASELINE</code> procedure as shown in the following example:

In the example, the baseline peak baseline is dropped from the database instance with the database identifier of 3310949047 and the associated snapshots are preserved.



Tip:

To determine which baseline to drop, use the $DBA_HIST_BASELINE$ view to review the existing baselines.



Tip:

To drop the associated snapshots along with the baseline, set the ${\tt cascade}$ parameter to ${\tt TRUE}.$



Oracle Database PL/SQL Packages and Types Reference for information about the DBMS WORKLOAD REPOSITORY package

Renaming a Baseline

To rename a baseline using command-line interface, use the RENAME_BASELINE procedure. The following example shows a RENAME BASELINE procedure call.

In this example, the name of the baseline on the database instance with the database identifier of 3310949047 is renamed from peak baseline to peak mondays.



Oracle Database PL/SQL Packages and Types Reference for information about the DBMS WORKLOAD REPOSITORY package

Displaying Baseline Metrics

When used with adaptive thresholds, a baseline contains AWR data that the database can use to compute metric threshold values.

To display the summary statistics for metric values in a baseline period using the command-line interface, use the <code>SELECT BASELINE METRICS</code> function:

```
DBMS_WORKLOAD_REPOSITORY.SELECT_BASELINE_METRICS (baseline_name IN VARCHAR2,
dbid IN NUMBER DEFAULT NULL,
instance_num IN NUMBER DEFAULT NULL)
RETURN awr baseline metric type table PIPELINED;
```

✓ See Also:

- "Adaptive Thresholds" for information about baseline metric thresholds
- Oracle Database PL/SQL Packages and Types Reference for information about the DBMS WORKLOAD REPOSITORY package

Resizing the Default Moving Window Baseline

By default, Oracle Database automatically maintains a system-defined moving window baseline. The default window size for the system-defined moving window baseline is the current AWR retention period, which by default is 8 days. In certain circumstances, you may want to modify the window size of the default moving window baseline, such as increasing its size to more accurately compute threshold values for adaptive thresholds.

To modify the window size of the default moving window baseline using the command-line interface, use the <code>MODIFY_BASELINE_WINDOW_SIZE</code> procedure as shown in the following example:

In this example, the default moving window is resized to 30 days on the database instance with the database identifier of 3310949047.



The window size must be set to a value that is equal to or less than the value of the AWR retention setting. To set a window size that is greater than the current AWR retention period, you must first increase the value of the retention parameter as described in "Modifying Snapshot Settings".

See Also:

- "Moving Window Baselines" for information about moving window baselines
- Oracle Database PL/SQL Packages and Types Reference for information about the DBMS WORKLOAD REPOSITORY package

Managing Baseline Templates

Baseline templates enable you to automatically create baselines to capture specified time periods in the future. This section describes how to manage baseline templates and contains the following topics:

- User Interfaces for Managing Baseline Templates
- Creating a Single Baseline Template
- Creating a Repeating Baseline Template
- Dropping a Baseline Template



"Baseline Templates" for information about baseline templates

User Interfaces for Managing Baseline Templates

The primary interface for managing baseline templates is Oracle Enterprise Manager Cloud Control (Cloud Control). Whenever possible, manage baseline templates using Cloud Control.

If Cloud Control is unavailable, then manage baseline templates using the DBMS_WORKLOAD_REPOSITORY package in the command-line interface. The DBA role is required to invoke the DBMS_WORKLOAD_REPOSITORY procedures.

See Also:

Oracle Database 2 Day + Performance Tuning Guide for more information about managing baseline templates using Cloud Control

Creating a Single Baseline Template

You can use a single baseline template to create a baseline during a single, fixed time interval in the future. For example, you can create a single baseline template to generate a baseline that is captured on April 2, 2012 from 5:00 p.m. to 8:00 p.m.

To create a single baseline template using command-line interface, use the CREATE BASELINE TEMPLATE procedure as shown in the following example:

In this example, a baseline template named <code>template_120402</code> is created that will generate a baseline named <code>baseline_120402</code> for the time period from 5:00 p.m. to 8:00 p.m. on April 2, 2012 on the database with a database ID of <code>3310949047</code>. The baseline will expire after 30 days.

See Also:

Oracle Database PL/SQL Packages and Types Reference for information about the DBMS WORKLOAD REPOSITORY package

Creating a Repeating Baseline Template

You can use a repeating baseline template to automatically create baselines that repeat during a particular time interval over a specific period in the future. For example, you can create a repeating baseline template to generate a baseline that repeats every Monday from 5:00 p.m. to 8:00 p.m. for the year 2012.

To create a repeating baseline template using command-line, use the CREATE_BASELINE_TEMPLATE procedure as shown in the following example:

```
17:00:00 PST',

end_time => '2012-12-31 20:00:00

PST',

baseline_name_prefix =>

'baseline_2012_mondays_',

template_2012_mondays',

dbid => 3310949047);

END;

/
```

In this example, a baseline template named <code>template_2012_mondays</code> is created that will generate a baseline on every Monday from 5:00 p.m. to 8:00 p.m. beginning on April 2, 2012 at 5:00 p.m. and ending on December 31, 2012 at 8:00 p.m. on the database with a database ID of <code>3310949047</code>. Each of the baselines will be created with a baseline name with the prefix <code>baseline_2012_mondays_</code> and will expire after 30 days.



Oracle Database PL/SQL Packages and Types Reference for information about the DBMS_WORKLOAD_REPOSITORY package

Dropping a Baseline Template

Periodically, you may want to remove baselines templates that are no longer used to conserve disk space.

To drop a baseline template using command-line, use the <code>DROP_BASELINE_TEMPLATE</code> procedure as shown in the following example:

In this example, the baseline template named template_2012_mondays is dropped from the database instance with the database identifier of 3310949047.



Tip:

To determine which baseline template to drop, use the DBA_HIST_BASELINE_TEMPLATE view to review the existing baseline templates.

See Also:

Oracle Database PL/SQL Packages and Types Reference for information about the DBMS WORKLOAD REPOSITORY package

Transporting Automatic Workload Repository Data to Another System

Oracle Database enables you to transport AWR data between systems. This is useful in cases where you want to use a separate system to perform analysis of AWR data, so as to reduce the overhead caused by performance analysis on a production system.

To transport AWR data from one system to another, first export the AWR data from the database on the source system, and then import it into the database on the target system.

This section contains the following topics:

- "Exporting AWR Data"
- "Importing AWR Data"

Exporting AWR Data

The awrextr.sql script exports AWR data for a range of snapshots from the database into a Data Pump export file. After it is created, you can transport this dump file to another database where you can import the exported AWR data. To run the awrextr.sql script, you must be connected to the database as the SYS user.

To export AWR data:

1. At the SQL prompt, enter:

```
@$ORACLE_HOME/rdbms/admin/awrextr.sql
```

A list of the databases in the AWR schema is displayed.

Specify the database from which AWR data needs to be exported:

```
Enter value for db id: 1377863381
```

In this example, the database with the database identifier of 1377863381 is specified.

3. Specify the number of days for which you want to view all the snapshot IDs:

```
Enter value for num days: 2
```

In this example, all the snapshots captured in the last 2 days are displayed.

4. Define the range of snapshots for which AWR data needs to be exported by specifying the beginning and the ending snapshot IDs:

```
Enter value for begin_snap: 30
Enter value for end snap: 40
```

In this example, the snapshot ID of 30 is specified as the beginning snapshot, and the snapshot ID of 40 is specified as the ending snapshot.

A list of directory objects is displayed.

5. Specify the directory object pointing to the directory where the export dump file needs to be stored:

```
Enter value for directory_name: DATA_PUMP_DIR
```

In this example, the directory object <code>DATA_PUMP_DIR</code> is specified that points to the directory <code>ORACLE_HOME/rdbms/log</code>, where <code>ORACLE_HOME</code> is <code>/u01/app/oracle/product/database_release_number/dbhome_1</code>.



6. Specify a name for the export dump file without the file extension. By default, the file extension of .dmp is used.

```
Enter value for file_name: awrdata_30_40
```

In this example, an export dump file named <code>awrdata_30_40.dmp</code> is created in the directory specified in the directory object <code>DATA_PUMP_DIR</code>:

```
Dump file set for SYS.SYS_EXPORT_TABLE_01 is: /u01/app/oracle/product/database_release_number/dbhome_1/rdbms/log/awrdata_30_40.dmp Job "SYS"."SYS_EXPORT_TABLE_01" successfully completed at 08:58:20
```

Depending on the amount of AWR data that must be exported, the AWR export operation may take a while to complete. After the dump file is created, you can use Data Pump to transport the file to another system.



Oracle Database Utilities for information about using Data Pump

Importing AWR Data

After the export dump file is transported to the target system, import the exported AWR data using the awrload.sql script. The awrload.sql script creates a staging schema where the snapshot data is transferred from the Data Pump file into the database. The data is then transferred from the staging schema into the appropriate AWR tables. To run the awrload.sql script, you must be connected to the database as the SYS user.

To import AWR data:

At the SQL prompt, enter:

```
@$ORACLE HOME/rdbms/admin/awrload.sql
```

A list of directory objects is displayed.

2. Specify the directory object pointing to the directory where the export dump file is located:

```
Enter value for directory name: DATA PUMP DIR
```

In this example, the directory object ${\tt DATA_PUMP_DIR}$ is specified that points to the directory where the export dump file is located.

3. Specify the name of the export dump file without the file extension. By default, the file extension of .dmp is used.

```
Enter value for file name: awrdata 30 40
```

In this example, the export dump file named awrdata 30 40.dmp is selected.

Specify the name of the staging schema where the AWR data needs to be imported:

```
Enter value for schema name: AWR STAGE
```

In this example, a staging schema named AWR STAGE is created.

5. Specify the default tablespace for the staging schema:

```
Enter value for default_tablespace: SYSAUX
```



In this example, the SYSAUX tablespace is specified.

Specify the temporary tablespace for the staging schema:

```
Enter value for temporary_tablespace: TEMP
```

In this example, the TEMP tablespace is specified.

7. First the AWR data is imported into the AWR_STAGE schema and then it is transferred to the AWR tables in the SYS schema:

```
Processing object type TABLE_EXPORT/TABLE/CONSTRAINT/CONSTRAINT Completed 113 CONSTRAINT objects in 11 seconds
Processing object type TABLE_EXPORT/TABLE/CONSTRAINT/REF_CONSTRAINT
Completed 1 REF_CONSTRAINT objects in 1 seconds
Job "SYS"."SYS_IMPORT_FULL_03" successfully completed at 09:29:30
... Dropping AWR_STAGE user
End of AWR Load
```

Depending on the amount of AWR data that must be imported, the AWR import operation may take a while to complete. After AWR data is imported, the staging schema will be dropped automatically.

Using Automatic Workload Repository Views

Typically, you would view AWR data using Oracle Enterprise Manager Cloud Control (Cloud Control) or AWR reports. However, you can also view historical data stored in the AWR using the following DBA_HIST views.



In a multitenant environment, these <code>DBA_HIST</code> views can also be interchanged with the <code>AWR_ROOT</code> views and <code>AWR_PDB</code> views at the CDB level and the PDB level respectively. For example, you can use the <code>AWR_PDB_ACTIVE_SESS_HISTORY</code> view for retrieving the AWR data about the active session history at the PDB level, which is equivalent to the <code>DBA_HIST_ACTIVE_SESS_HISTORY</code> view in an independent database in a non-multitenant environment. The <code>AWR_PDB</code> views will not show any AWR data, if the PDB level snapshots have not been collected.

Table 6-1 DBA_HIST Views

DBA_HIST View	Description
DBA_HIST_ACTIVE_SESS_HISTORY	Displays the history of the contents of the in-memory active session history for recent system activity.
DBA_HIST_BASELINE	Displays information about the baselines captured on the system, such as the time range of each baseline and the baseline type.
DBA_HIST_BASELINE_DETAILS	Displays details about a specific baseline.
DBA_HIST_BASELINE_TEMPLATE	Displays information about the baseline templates used by the system to generate baselines.
DBA_HIST_CON_SYS_TIME_MODEL	Displays historical system time model statistics, including OLAP timed statistics.
DBA HIST CON SYSMETRIC HIST	Displays the historical information about the system metric values.



Table 6-1 (Cont.) DBA_HIST Views

DBA_HIST View	Description
DBA_HIST_CON_SYSMETRIC_SUMM	Displays history of the statistical summary of all the metric values in the system metrics for the long duration (60 seconds) group.
DBA_HIST_CON_SYSSTAT	Displays historical system statistics, including OLAP kernel statistics.
DBA_HIST_CON_SYSTEM_EVENT	Displays historical information about the total waits for an event.
DBA_HIST_DATABASE_INSTANCE	Displays information about the database environment.
DBA_HIST_DB_CACHE_ADVICE	Displays historical predictions of the number of physical reads for the cache size corresponding to each row.
DBA_HIST_DISPATCHER	Displays historical information for each dispatcher process at the time of the snapshot.
DBA_HIST_DYN_REMASTER_STATS	Displays statistical information about the dynamic remastering process.
DBA_HIST_IOSTAT_DETAIL	Displays historical I/O statistics aggregated by file type and function.
DBA_HIST_RSRC_PDB_METRIC	Displays historical information about the Resource Manager metrics for pluggable databases (PDBs) for the past one hour.
DBA_HIST_RSRC_METRIC	Displays historical information about the Resource Manager metrics for consumer groups for the past one hour.
DBA_HIST_SHARED_SERVER_SUMMARY	Displays historical information for shared servers, such as shared server activity, common queues and dispatcher queues.
DBA_HIST_SNAPSHOT	Displays information on snapshots in the system.
DBA_HIST_SQL_PLAN	Displays the SQL execution plans.
DBA_HIST_WR_CONTROL	Displays the settings for controlling AWR.
DBA_HIST_WR_SETTINGS	Displays the settings and metadata of the AWR.
DBA_HIST_PROCESS_WAITTIME	Displays CPU and wait time for a process type.

See Also:

Oracle Database Reference for more information about the DBA HIST views

Managing Automatic Workload Repository in a Multitenant Environment

A centralized Automatic Workload Repository (AWR) stores the performance data related to CDB and PDBs in a multitenant environment.

CDBs and individual PDBs can store, view, and manage AWR data. You can take an AWR snapshot at the CDB level or at the PDB level.

A multitenant container database is the only supported architecture in Oracle Database 21c and later releases. While the documentation is being revised, legacy terminology may persist. In most cases, "database" and "non-CDB" refer to a CDB or PDB, depending on context. In some contexts, such as upgrades, "non-CDB" refers to a non-CDB from a previous release.

This section contains the following topics:

- Categorization of AWR Data in a Multitenant Environment
- AWR Data Storage and Retrieval in a Multitenant Environment
- Viewing AWR Data in a Multitenant Environment

Categorization of AWR Data in a Multitenant Environment

In a multitenant environment, AWR data falls into different categories.

The categories are as follows:

General AWR data

This data has no security implications. It is safe to be shared among all tenants in a CDB. This data is accessible by all PDBs and is captured in both CDB-level and PDB-level snapshots. Examples of general AWR data include the names of statistics, latches, and parameters.

AWR data for a CDB

This category aggregates data for all tenants in a CDB. This data contains the status of the database as a whole and is useful only for the CDB administrator. This data is captured only in the CDB-level snapshots.

AWR data for individual PDBs

This data describes the individual PDBs in a CDB. It shows container-specific data that represents the contribution of each individual PDB to the whole database instance. Therefore, this data is useful for both the CDB and the PDB administrators. This data is captured in both CDB-level and PDB-level snapshots.

AWR Data Storage and Retrieval in a Multitenant Environment

This section describes the process of managing snapshots, and exporting and importing AWR data in a multitenant environment.

Managing Snapshots

Starting with Oracle Database 12c Release 2 (12.2), you can take an AWR snapshot at a CDB-level, that is, on a CDB root, as well as at a PDB-level, that is, on an individual PDB. By default, the CDB-level snapshot data is stored in the SYSAUX tablespace of a CDB root and the PDB-level snapshot data is stored in the SYSAUX tablespace of a PDB.

A CDB-level snapshot contains information about the CDB statistics as well as all the PDB statistics, such as ASH, SQL statistics, and file statistics. The CDB administrator can perform CDB-specific management operations, such as setting AWR data retention period, setting snapshot schedule, taking manual snapshots, and purging snapshot data for a CDB root.



A PDB-level snapshot contains the PDB statistics and also some global statistics that can be useful for diagnosing the performance problems related to the PDB. The PDB administrator can perform PDB-specific management operations, such as setting AWR data retention period, setting snapshot schedule, taking manual snapshots, and purging snapshot data for a PDB.

The CDB-level and PDB-level snapshot operations, such as creating snapshots and purging snapshots, can be performed in either the *automatic mode* or the *manual mode*.

The automatic snapshot operations are scheduled, so that they get executed automatically at a particular time. The <code>AWR_PDB_AUTOFLUSH_ENABLED</code> initialization parameter enables you to specify whether to enable or disable automatic snapshots for all the PDBs in a CDB or for individual PDBs in a CDB. The automatic snapshot operations are enabled by default for a CDB and (in Oracle Database 23ai) for a PDB. To enable automatic snapshots for a PDB, the PDB administrator must connect to that PDB, set the value for the <code>AWR_PDB_AUTOFLUSH_ENABLED</code> parameter to <code>true</code>, and set the snapshot generation interval to a value greater than 0.

See Also:

Oracle Database Reference for more information about the AWR_PDB_AUTOFLUSH_ENABLED initialization parameter

The manual snapshot operations are explicitly initiated by users. The automatic snapshots and manual snapshots capture the same AWR information. Oracle recommends to generally use manual snapshots for a PDB. You should enable automatic snapshots only selectively for a PDB for performance reasons.

The primary interface for managing snapshots is Oracle Enterprise Manager Cloud Control (Cloud Control). If Cloud Control is not available, then you can use the procedures in the <code>DBMS_WORKLOAD_REPOSITORY</code> package to manage snapshots. The Oracle DBA role is required to use the procedures in the <code>DBMS_WORKLOAD_REPOSITORY</code> package. The SQL procedures to create, drop, and modify snapshots for a CDB root and a PDB are the same as that for a non-CDB. These SQL procedures perform their operations on the local database by default, if the target database information is not provided in their procedure call.



- The PDB-level snapshots have unique snapshot IDs and are not related to the CDB-level snapshots.
- The plugging and unplugging operations of a PDB in a CDB do not affect the AWR data stored on a PDB.
- The CDB administrator can use the PDB lockdown profiles to disable the AWR functionality for a PDB by executing the following SQL statement on that PDB:

```
SQL> alter lockdown profile profile_name disable
feature=('AWR ACCESS');
```

Once the AWR functionality is disabled on a PDB, snapshot operations cannot be performed on that PDB.

The AWR functionality can be enabled again for a PDB by executing the following SQL statement on that PDB:

```
SQL> alter lockdown profile profile_name enable
feature=('AWR ACCESS');
```

 Snapshot data is stored in the SYSAUX tablespace of a CDB and a PDB by default. Starting with Oracle Database 19c, you can specify any other tablespace to store snapshot data for a CDB and a PDB by modifying snapshot settings.

Note:

A multitenant container database is the only supported architecture in Oracle Database 21c and later releases. While the documentation is being revised, legacy terminology may persist. In most cases, "database" and "non-CDB" refer to a CDB or PDB, depending on context. In some contexts, such as upgrades, "non-CDB" refers to a non-CDB from a previous release.

See Also:

- "Creating Snapshots"
- "Dropping Snapshots"
- "Modifying Snapshot Settings"
- Oracle Database Security Guide for more information about the PDB lockdown profiles

Exporting and Importing AWR Data

The process of exporting and importing AWR data for a CDB root and a PDB in a multitenant environment is similar to the process of exporting and importing AWR data for a non-CDB.

See Also:

- "Exporting AWR Data" for information about exporting AWR data from an Oracle database
- "Importing AWR Data" for information about importing AWR data into an Oracle database

Viewing AWR Data in a Multitenant Environment

You can view the AWR data in a multitenant environment using various Oracle Database reports and views.

AWR Reports

The primary interface for generating AWR reports is Oracle Enterprise Manager Cloud Control (Cloud Control). Whenever possible, generate AWR reports using Cloud Control.

See Also:

Oracle Database 2 Day + Performance Tuning Guide for more information about generating AWR report using Cloud Control

If Cloud Control is unavailable, then you can generate the AWR reports by running SQL scripts as described below. The DBA role is required to run these scripts.

- You can generate a CDB-specific AWR report from a CDB root that shows the *global* system data statistics for the whole multitenant environment. You can generate this AWR report using the SQL scripts described in the section "Generating an AWR Report for the Local Database".
- You can generate a PDB-specific AWR report from a PDB that shows the statistics related to that PDB. You can generate this AWR report using the SQL scripts described in the section "Generating an AWR Report for the Local Database".
- You can generate a PDB-specific AWR report from a CDB root that shows the statistics related to a specific PDB. You can generate this AWR report using the SQL scripts described in the section "Generating an AWR Report for a Specific Database".

AWR Views

The following table lists the Oracle Database views for accessing the AWR data stored on the CDB root and the individual PDBs in a multitenant environment.

See Also:

"Using Automatic Workload Repository Views" for more information about these AWR views



Table 6-2 Views for Accessing AWR Data in a Multitenant Environment

Views	Description
DBA_HIST Views	The DBA_HIST views show the AWR data present only on the CDB root.
	 When the DBA_HIST views are accessed from a CDB root, they show all the AWR data stored on the CDB root.
	 When the DBA_HIST views are accessed from a PDB, they show the subset of the CDB root AWR data, which is specific to that PDB.
DBA_HIST_CON Views	 The DBA_HIST_CON views are similar to the DBA_HIST views, but they provide more fine grained information about each container, and thus, they have more data than the DBA_HIST views.
	• The DBA_HIST_CON views show the AWR data present only on the CDB root.
	 When the DBA_HIST_CON views are accessed from a CDB root, they show all the AWR data stored on the CDB root.
	 When the DBA_HIST_CON views are accessed from a PDB, they show the subset of the CDB root AWR data, which is specific to that PDB.
AWR_ROOT Views	 The AWR_ROOT views are available starting with Oracle Database 12c Release 2 (12.2) and are available only in the Multitenant environment.
	 The AWR_ROOT views are equivalent to the DBA_HIST views.
	 The AWR_ROOT views show the AWR data present only on the CDB root.
	 When the AWR_ROOT views are accessed from a CDB root, they show all the AWR data stored on the CDB root.
	 When the AWR_ROOT views are accessed from a PDB, they show the subset of the CDB root AWR data, which is specific to that PDB.
AWR_PDB Views	 The AWR_PDB views are available starting with Oracle Database 12c Release 2 (12.2).
	 The AWR_PDB views show the local AWR data present on a CDB root or a PDB.
	 When the AWR_PDB views are accessed from a CDB root, they show the AWR data stored on the CDB root.
	 When the AWR_PDB views are accessed from a PDB, they show the AWR data stored on that PDB.
CDB_HIST Views	The CDB_HIST views show the AWR data stored on the PDBs.
	 When the CDB_HIST views are accessed from a CDB root, they show the union of the AWR data stored on all the PDBs.
	 When the CDB_HIST views are accessed from a PDB, they show the AWR data stored on that PDB.

Managing Automatic Workload Repository in Active Data Guard Standby Databases

Automatic Workload Repository (AWR) data can be captured for Active Data Guard (ADG) standby databases. This feature enables analyzing any performance-related issues for ADG standby databases.

AWR snapshots for ADG standby databases are called *remote snapshots*. A database node, called *destination*, is responsible for storing snapshots that are collected from remote ADG standby database nodes, called *sources*.

A destination can be either an ADG primary database or a non-ADG database. If a destination is an ADG primary database, then it is also a source database, and its snapshots are *local* snapshots.

A source is identified by a unique name or source name by which it is known to a destination.

You can assign a name to a destination node or a source node during its configuration. Otherwise, the value of the initialization parameter DB_UNIQUE_NAME is assigned as a name for a node.

Each source must have two database links, a destination-to-source database link and a source-to-destination database link. These database links are configured for each source during the ADG deployment. You must manually reconfigure these database links after certain ADG events, such as failovers, switchovers, and addition and removal of hosts, so that the database applications continue functioning properly after these events.

You can take the remote snapshots either automatically at scheduled time intervals or manually. The remote snapshots are always started by the destination node. After the destination initiates the snapshot creation process, sources *push* their snapshot data to the destination using database links. The snapshot data or *AWR data* stored on the destination can be accessed using AWR reports, Oracle Database import and export functions, and user-defined queries. The Automatic Database Diagnostic Monitor (ADDM) application can use the AWR data for analyzing any database performance-related issues.

AWR snapshot at an ADG Standby is enabled by default (Oracle Database 23ai). There is no need to run <code>dbms_workload_repository.enable_snapshot_service()</code> to enable it. It enables automatic snapshots for CDB root and all PDBs at ADG.

The snapshot retention is 8 days by default. To change the retention time, use dbms workload repository.modify snapshot settings (retention).

After a switch-over, snapshots work correctly for both Primary and Standby.

Destination Database Responsibilities

A destination database manages the following tasks:

- Registering sources
- Assigning unique identifier for each source
- Creating database links between destination and sources
- Scheduling and initiating automatic snapshots for sources
- Managing destination workload by coordinating snapshots among sources
- Managing snapshot settings for each source
- Assigning identifiers to newly generated snapshots
- Partitioning the AWR tables
- Storing the performance data in the local AWR
- Purging the AWR data of destination and sources

Source Database Responsibilities

A source database manages the following tasks:

- Storing its performance data in the local AWR
- Sending its AWR data to the destination
- Responding to service requests from the destination
- Extracting the AWR data from the destination



Major Steps for Managing AWR in ADG Standby Databases

The following are the major steps for managing AWR in ADG standby databases:

- 1. Configuring the Remote Management Framework (RMF)
- 2. Managing Snapshots for Active Data Guard Standby Databases
- 3. Viewing AWR Data in Active Data Guard Standby Databases



Before you start configuring AWR for ADG environment, make sure that the database links for all the ADG standby databases are already configured during the ADG deployment.

Configuring the Remote Management Framework (RMF)

The Remote Management Framework (RMF) is an architecture for capturing performance statistics (AWR data) in an Oracle database.

Note:

RMF can be used only for ADG standby databases and standalone databases.

The RMF *topology* is a centralized architecture that consists of all the participating database nodes along with their metadata and connection information. The RMF topology has one database node, called *destination*, which is responsible for storing and managing performance data (AWR data) that is collected from the database nodes, called *sources*. A *candidate destination* is a source that can be configured in such way that it can replace the original destination, when the original destination is unavailable or is downgraded. A topology can have only one destination, and one or more candidate destinations.

Each database node in a topology must be assigned a unique name. This can be done using the procedure <code>DBMS_UMF.configure_node()</code> during configuring a node. If the name for a node is not provided in this procedure, then the value of the initialization parameter <code>DB_UNIQUE_NAME</code> is used as the name for a node.

The database nodes in a topology communicate with each other using database links. The database links between destination to source and source to destination must be created for each ADG standby database during the ADG deployment.

A *service* is an application running on a topology. For example, an AWR service running on a topology enables remote AWR snapshots for all the database nodes in that topology.

The RMF APIs are the PL/SQL procedures and functions that can be used to configure the RMF topology. The RMF APIs are declared in the PL/SQL package DBMS UMF.



- The SYS\$UMF user is the default database user that has all the privileges to
 access the system-level RMF views and tables. All the AWR related operations in
 RMF can be performed only by the SYS\$UMF user. The SYS\$UMF user is locked by
 default and it must be unlocked before deploying the RMF topology.
- You need to provide password for the SYS\$UMF user when creating database links in the RMF topology. If the password for the SYS\$UMF user is changed, all the database links in the RMF topology must be recreated.

See Also:

Oracle Database PL/SQL Packages and Types Reference for more information about the DBMS UMF package

Setting Up the RMF Topology

You need to set up the RMF topology for collecting performance statistics for an Oracle database.

The following are the prerequisites for setting up the RMF topology:

 You must create destination to source and source to destination database links for all the database nodes to be registered in the RMF topology. This setup should be done during the ADG deployment.

The following are the steps for setting up the RMF topology:

- Configure database nodes to add to the topology.
- Create the topology.
- 3. Register database nodes with the topology.
- 4. (Optional) Register database links between the nodes in the topology. This configuration is required when a destination becomes unavailable and a candidate destination needs to connect to the remaining nodes in the topology using database links.

Example for Setting Up the RMF Topology

In this example, the three database nodes T, S0, and S1 are added to the topology $Topology_1$. Node T is the destination node and nodes S0 and S1 are the source nodes. Node S1 is a candidate destination, that is, when the original destination T is not available, node S1 becomes the new destination. The AWR service is enabled for all the sources in the topology.

Assume that the following database links are already created during the ADG deployment:

- DBLINK T to S0: Database link from T to S0.
- DBLINK T to S1: Database link from T to S1.
- DBLINK SO to T: Database link from SO to T.
- DBLINK_S0_to_S1: Database link from S0 to S1.
- DBLINK S1 to T: Database link from S1 to T.



DBLINK S1 to S0: Database link from S1 to S0.

The following is a sample code for setting up the RMF topology:

```
/* Configure the nodes T, SO, and S1 by executing these procedures */
/* Execute this procedure on node T */
SQL> exec DBMS UMF.configure node ('T');
/* Execute this procedure on node S0 */
SQL> exec DBMS UMF.configure node ('S0', 'DBLINK S0 to T');
/* Execute this procedure on node S1 */
SQL> exec DBMS UMF.configure node ('S1', 'DBLINK S1 to T');
/* Execute all the following procedures on the destination node T */
/* Create the topology 'Topology 1' */
SQL> exec DBMS UMF.create topology ('Topology 1');
/* Register the node SO with the topology 'Topology 1' */
SQL> exec DBMS UMF.register node ('Topology 1',
                                   'SO',
                                   'DBLINK T to SO',
                                   'DBLINK SO to T',
                                   'TRUE' /* Set it as a source */,
                                   'FALSE' /* Set it as not a candidate destination */);
/* Register the node S1 with the topology 'Topology 1' */
SQL> exec DBMS UMF.register node ('Topology 1',
                                   'S1',
                                   'DBLINK T to S1',
                                   'DBLINK S1 to T',
                                   'TRUE' /* Set it as a source */,
                                   'TRUE' /* Set it as a candidate destination */);
/* Register the database links between the nodes S0 and S1 in the topology 'Topology 1'.
 * When destination T is unavailable at the time of failover, the source SO can connect
 * to the candidate destination S1 using this database link.
SQL> exec DBMS UMF.create link ('Topology 1',
                                 'S0',
                                 'S1',
                                 'DBLINK S0 to S1',
                                 'DBLINK S1 to S0');
^{\prime \star} Enable the AWR service on the node S0 in the topology 'Topology 1' ^{\star}
SQL> exec DBMS WORKLOAD REPOSITORY.register remote database(node name=>'S0');
/* Enable the AWR service on the node S1 in the topology 'Topology 1' */
SQL> exec DBMS WORKLOAD REPOSITORY.register remote database(node name=>'S1');
```

The AWR service can be disabled for a node using the procedure:

```
SQL> exec
DBMS_WORKLOAD_REPOSITORY.unregister_remote_database(node_name)
```

Managing ADG Role Transition

An ADG role transition occurs when the ADG Primary or original destination fails (*failover* event) or when an ADG standby database or candidate destination takes over the role of the ADG Primary during the maintenance phase (*switchover* event).

Oracle recommends that you perform the following configuration steps before making the role change, that is, before making the candidate destination as the new destination due to the failover or switchover event:

Create database links between the sources and the candidate destination. This
configuration must be done for all the sources by executing the following procedure on
each source:

Note:

Oracle recommends that you create database links among all the nodes in a topology to avoid any unanticipated issues that may arise at the time of role change.

Take an AWR snapshot on the candidate destination.

Note:

To generate an AWR report for the candidate destination after the role change, take at least one snapshot for the candidate destination before the role change.

3. Restart the candidate destination as well as all the sources.

After completing the preceding configuration steps, you can make the role change by executing the following procedure on the candidate destination:

```
SQL> EXEC DBMS UMF.SWITCH DESTINATION(topology name, force switch=>FALSE);
```

Oracle recommends that you do not take any snapshots for the sources during the role transition period. After the role change process is complete by executing the <code>DBMS_UMF.SWITCH_DESTINATION</code> procedure, you can take snapshots for the sources. If you want to generate AWR reports for the sources after the role change, then you must choose only those snapshots that were taken after the role change.

Getting the Details of Registered RMF Topologies

The RMF views described below show the configuration information about all the registered RMF topologies in a multi-database environment.

Table 6-3 RMF Views

RMF View	Description
DBA_UMF_TOPOLOGY	Shows all the registered topologies in a multi-database environment. Each topology has a topology name, a destination ID, and topology state. To enable RMF, the topology state of at least one topology should be ACTIVE.
DBA_UMF_REGISTRATION	Shows all the registered nodes in all the topologies in a multi- database environment.
DBA_UMF_LINK	Shows all the registered database links in all the topologies in a multi-database environment.
DBA_UMF_SERVICE	Shows all the registered services in all the topologies in a multi- database environment.

See Also:

Oracle Database Reference for more information about these RMF views

Managing Snapshots for Active Data Guard Standby Databases

The AWR snapshots for ADG standby databases are called *remote* snapshots. Similar to local AWR snapshots, remote AWR snapshots can be generated automatically at scheduled time intervals or can be generated manually. The *Push-on-Demand* mechanism is used for generating remote snapshots, where the snapshots generation process is initiated by the destination, which then instructs the sources to start *pushing* the snapshot data to the destination over database links. The destination periodically initiates automatic snapshots based on the snapshot time interval configured for each of the sources.



The destination is responsible for purging the expired remote snapshots based on the snapshot data or *AWR data* retention settings for individual sources. Purging of locally generated snapshots occurs as part of the regularly scheduled purging process. By default, Oracle Database automatically purges snapshots that have been stored in AWR for over 8 days. The partitioning of AWR table for remote snapshots is done in the same way as that of the local snapshots.

Creating, Modifying, and Deleting Remote Snapshots

The APIs for creating, modifying, and deleting *remote* snapshots are same as that for the *local* snapshots.

Note:

For creating remote snapshots, you can also use the <code>DBMS_WORKLOAD_REPOSITORY.CREATE_REMOTE_SNAPSHOT</code> API. This API works similar to the local snapshot creation API <code>DBMS_WORKLOAD_REPOSITORY.CREATE_SNAPSHOT</code>, but it takes the additional parameter of RMF topology name.

See Also:

- "Creating Snapshots"
- "Modifying Snapshot Settings"
- "Dropping Snapshots"
- Oracle Database PL/SQL Packages and Types Reference for the syntax of the DBMS_WORKLOAD_REPOSITORY.CREATE_REMOTE_SNAPSHOT API.

Managing Baselines for Remote Snapshots

The APIs for managing baselines for *remote* snapshots are same as that for the *local* snapshots.



"Managing Baselines"



Exporting and Importing Remote Snapshots



You cannot execute the AWR export and import scripts related to remote snapshots on an ADG standby database, that is, on a source database. Always execute these scripts on a destination database.

The process of exporting and importing AWR data for *remote* snapshots is same as that for the *local* snapshots. Starting with Oracle Database 12c Release 2 (12.2), the AWR data export and import scripts <code>awrextr.sql</code> and <code>awrload.sql</code> use the *source name* identifier to distinguish snapshots originating from a particular source. A source name is stored in a dump file during an export operation and is used as a default source name during an import operation.



"Transporting Automatic Workload Repository Data to Another System" for information about exporting and importing AWR data for local snapshots.

Exporting Remote Snapshots Using the awrextr.sql Script

The process of exporting *remote* snapshots is similar to exporting *local* snapshots using the awrextr.sql script described in the section "Exporting AWR Data" with the following differences:

- The default export log file directory is the same as that of the dump file, but you can also specify any other directory for an export log file.
- The .dmp suffix can be specified to the name of the dump file to export.
- The export script displays the values of SOURCE_DBID and SOURCE_NAME columns of AWR tables before prompting for the Mapped Database ID value to export.

Importing Remote Snapshots Using the awrload.sql Script

The process of importing *remote* snapshots is similar to importing *local* snapshots using the awrload.sql script described in the section "Importing AWR Data" with the following differences:

- The default import log file directory is the same as that of the dump file, but you can also specify any other directory for an import log file. This is particularly useful when the dump file resides in a read-only directory.
- The .dmp suffix can be specified to the name of the dump file to import.
- The import script uses the values of SOURCE_DBID and SOURCE_NAME columns present in the dump file to determine the appropriate Mapped Database ID to use for storing the snapshot data in AWR.



The snapshot import operation is not affected by the version of the Oracle database from which the snapshot dump was generated.

Viewing AWR Data in Active Data Guard Standby Databases

You can view the AWR data stored in the ADG standby databases using Oracle supplied AWR views and AWR reports.

Viewing AWR Data Using AWR Views

You can view the historical data stored in AWR using the <code>DBA_HIST</code> views described in the section "Using Automatic Workload Repository Views".

Note:

Starting with Oracle Database 12c Release 2 (12.2), the view DBA_HIST_DATABASE_INSTANCE contains the column DB_UNIQUE_NAME to support AWR for ADG standby databases. The column DB_UNIQUE_NAME stores the unique identifier of a source by which it is known to the destination.

Viewing AWR Data Using AWR Reports

You can view the performance statistics related to ADG standby databases using AWR reports. The primary interface for generating AWR reports is Oracle Enterprise Manager Cloud Control (Cloud Control). Whenever possible, generate AWR reports using Cloud Control. If Cloud Control is unavailable, then generate AWR reports using the Oracle supplied SQL scripts. The DBA role is required to run these scripts.

The AWR data can be queried for a particular source using the source name-mapped database ID pair. The mapped database ID is similar to the database identifier (DBID) that is used by AWR to identify a database instance and is stored in the DBID column in the AWR tables. The AWR DBID value is derived as follows for the ADG standby databases:

- For a destination, the AWR DBID value is the value of V\$DATABASE.CON DBID.
- For a source, the AWR DBID value is the value of <code>DBMS_UMF.GET_NODE_ID_LOCAL()</code> or the value of the column <code>NODE_ID</code> in the <code>DBA_UMF_REGISTRATION</code> view.

As snapshot IDs are not unique across sources, the pair of snapshot ID-mapped database ID identifies a snapshot for a particular source.

See Also:

"Generating an AWR Report for a Specific Database" for information about generating AWR reports using Oracle supplied SQL scripts.



Generating Automatic Workload Repository Reports

An AWR report shows data captured between two snapshots (or two points in time). AWR reports are divided into multiple sections. The content of the report contains the workload profile of the system for the selected range of snapshots. The HTML report includes links that can be used to navigate quickly between sections.



If you run a report on a database that does not have any workload activity during the specified range of snapshots, then calculated percentages for some report statistics can be less than 0 or greater than 100. This result means that there is no meaningful value for the statistic.

This section describes how to generate AWR reports and contains the following topics:

- · User Interface for Generating an AWR Report
- Generating an AWR Report Using the Command-Line Interface

User Interface for Generating an AWR Report

The primary interface for generating AWR reports is Oracle Enterprise Manager Cloud Control (Cloud Control). Whenever possible, generate AWR reports using Cloud Control.

If Cloud Control is unavailable, then generate AWR reports by running SQL scripts. The DBA role is required to run these scripts.

See Also:

Oracle Database 2 Day + Performance Tuning Guide for more information about generating AWR reports using Cloud Control

Generating an AWR Report Using the Command-Line Interface

This section describes how to generate AWR reports by running SQL scripts in the commandline interface. The DBA role is required to run these scripts. Click on an appropriate task link in the following table for the detailed steps to generate the required AWR report.

Table 6-4 SQL Scripts for Generating AWR Reports

Task	SQL Script	Description
Generating an AWR Report for the Local Database	awrrpt.sql	Generates an AWR report in HTML or text format that displays statistics from a range of snapshot IDs in the local database instance.
Generating an AWR Report for a Specific Database	awrrpti.sql	Generates an AWR report in HTML or text format that displays statistics from a range of snapshot IDs in a specific database instance.



Table 6-4 (Cont.) SQL Scripts for Generating AWR Reports

Task	SQL Script	Description
Generating an AWR Report for the Local Database in Oracle RAC	awrgrpt.sql	Generates an AWR report in HTML or text format that displays statistics from a range of snapshot IDs in the local database instance in an Oracle RAC environment.
Generating an AWR Report for a Specific Database in Oracle RAC	awrgrpti.sql	Generates an AWR report in HTML or text format that displays statistics from a range of snapshot IDs in a specific database instance in an Oracle RAC environment.
Generating an AWR Report for a SQL Statement on the Local Database	awrsqrpt.sql	Generates an AWR report in HTML or text format that displays statistics for a particular SQL statement from a range of snapshot IDs in the local database instance.
Generating an AWR Report for a SQL Statement on a Specific Database	awrsqrpi.sql	Generates an AWR report in HTML or text format that displays statistics for a particular SQL statement from a range of snapshot IDs in a specific database instance.

Generating an AWR Report for the Local Database

The awrrpt.sql SQL script generates an HTML or text report that displays statistics from a range of snapshot IDs.

To generate an AWR report on the local database instance using the command-line interface:

1. At the SQL prompt, enter:

```
@$ORACLE HOME/rdbms/admin/awrrpt.sql
```

2. Specify whether you want an HTML or a text report:

```
Enter value for report type: text
```

In this example, a text report is chosen.

3. Specify the number of days for which you want to list snapshot IDs.

```
Enter value for num days: 2
```

A list of existing snapshots for the specified time range is displayed. In this example, snapshots captured in the last 2 days are displayed.

4. Specify a beginning and ending snapshot ID for the workload repository report:

```
Enter value for begin_snap: 150
Enter value for end_snap: 160
```

In this example, the snapshot with a snapshot ID of 150 is selected as the beginning snapshot, and the snapshot with a snapshot ID of 160 is selected as the ending snapshot.

5. Enter a report name, or accept the default report name:

```
Enter value for report_name:
Using the report name awrrpt 1 150 160
```



In this example, the default name is accepted and an AWR report named ${\tt awrrpt}\ 1\ 150\ 160$ is generated.

Generating an AWR Report for a Specific Database

The awrrpti.sql SQL script generates an HTML or text report that displays statistics from a range of snapshot IDs using a specific database instance. This script enables you to specify a database identifier and instance for which the AWR report will be generated.

To generate an AWR report on a specific database instance using the command-line interface:

1. At the SQL prompt, enter:

```
@$ORACLE HOME/rdbms/admin/awrrpti.sql
```

2. Specify whether you want an HTML or a text report:

```
Enter value for report type: text
```

In this example, a text report is chosen.

A list of available database identifiers and instance numbers are displayed:

Instances in this Workload Repository schema

DB Id	Inst Num	DB Name	Instance	Host
3309173529	_	MAIN	main	examp1690
3309173529		TINT251	tint251	samp251

3. Enter the values for the database identifier (dbid) and instance number (inst num):

```
Enter value for dbid: 3309173529
Using 3309173529 for database Id
Enter value for inst num: 1
```

Note:

For an ADG standby database, the value for dbid can be determined as follows:

- For a Destination node, use the value of v\$database.con dbid.
- For a Source node, use the value of dbms_umf.get_node_id_local().
- 4. Specify the number of days for which you want to list snapshot IDs.

```
Enter value for num days: 2
```

A list of existing snapshots for the specified time range is displayed. In this example, snapshots captured in the last 2 days are displayed.

5. Specify a beginning and ending snapshot ID for the workload repository report:

```
Enter value for begin_snap: 150
Enter value for end snap: 160
```

In this example, the snapshot with a snapshot ID of 150 is selected as the beginning snapshot, and the snapshot with a snapshot ID of 160 is selected as the ending snapshot.

6. Enter a report name, or accept the default report name:

```
Enter value for report_name:
Using the report name awrrpt_1_150_160
```

In this example, the default name is accepted and an AWR report named awrrpt_1_150_160 is generated on the database instance with a database ID value of 3309173529.

Generating an AWR Report for the Local Database in Oracle RAC

The awrgrpt.sql SQL script generates an HTML or text report that displays statistics from a range of snapshot IDs using the current database instance in an Oracle Real Application Clusters (Oracle RAC) environment.



In an Oracle RAC environment, Oracle recommends generating an HTML report (instead of a text report) because it is much easier to read.

To generate an AWR report for Oracle RAC on the local database instance using the command-line interface:

1. At the SQL prompt, enter:

```
@$ORACLE HOME/rdbms/admin/awrgrpt.sql
```

Specify whether you want an HTML or a text report:

```
Enter value for report type: html
```

In this example, an HTML report is chosen.

3. Specify the number of days for which you want to list snapshot IDs.

```
Enter value for num days: 2
```

A list of existing snapshots for the specified time range is displayed. In this example, snapshots captured in the last day are displayed.

4. Specify a beginning and ending snapshot ID for the workload repository report:

```
Enter value for begin_snap: 150
Enter value for end_snap: 160
```

In this example, the snapshot with a snapshot ID of 150 is selected as the beginning snapshot, and the snapshot with a snapshot ID of 160 is selected as the ending snapshot.

5. Enter a report name, or accept the default report name:

```
Enter value for report_name:
Using the report name awrrpt_rac_150_160.html
```

In this example, the default name is accepted and an AWR report named $awrrpt_rac_150_160.html$ is generated.

Generating an AWR Report for a Specific Database in Oracle RAC

The awrgrpti.sql SQL script generates an HTML or text report that displays statistics from a range of snapshot IDs using specific databases instances running in an Oracle RAC

environment. This script enables you to specify database identifiers and a comma-delimited list of database instances for which the AWR report will be generated.



In an Oracle RAC environment, Oracle recommends generating an HTML report (instead of a text report) because it is much easier to read.

To generate an AWR report for Oracle RAC on a specific database instance using the command-line interface:

1. At the SQL prompt, enter:

@\$ORACLE HOME/rdbms/admin/awrgrpti.sql

2. Specify whether you want an HTML or a text report:

```
Enter value for report type: html
```

In this example, an HTML report is chosen.

A list of available database identifiers and instance numbers are displayed:

Instances in this Workload Repository schema

DB Id	Inst Num	DB Name	Instance	Host
3309173529 3309173529	_	MAIN TINT251	main tint251	examp1690 samp251
3309173529	2	TINT251	tint252	samp252

3. Enter the value for the database identifier (dbid):

```
Enter value for dbid: 3309173529
Using 3309173529 for database Id
```

4. Enter the value for the instance numbers (instance_numbers_or_all) of the Oracle RAC instances you want to include in the report:

```
Enter value for instance_numbers_or_all: 1,2
```

Specify the number of days for which you want to list snapshot IDs.

```
Enter value for num days: 2
```

A list of existing snapshots for the specified time range is displayed. In this example, snapshots captured in the last 2 days are displayed.

Specify a beginning and ending snapshot ID for the workload repository report:

```
Enter value for begin_snap: 150
Enter value for end_snap: 160
```

In this example, the snapshot with a snapshot ID of 150 is selected as the beginning snapshot, and the snapshot with a snapshot ID of 160 is selected as the ending snapshot.

7. Enter a report name, or accept the default report name:

```
Enter value for report_name:
Using the report name awrrpt rac 150 160.html
```



In this example, the default name is accepted and an AWR report named awrrpt_rac_150_160.html is generated on the database instance with a database ID value of 3309173529.

Generating an AWR Report for a SQL Statement on the Local Database

The awrsqrpt.sql SQL script generates an HTML or text report that displays statistics of a particular SQL statement from a range of snapshot IDs. Run this report to inspect or debug the performance of a SQL statement.

To generate an AWR report for a SQL statement on the local database instance using the command-line interface:

1. At the SQL prompt, enter:

```
@$ORACLE HOME/rdbms/admin/awrsqrpt.sql
```

2. Specify whether you want an HTML or a text report:

```
Enter value for report type: html
```

In this example, an HTML report is chosen.

Specify the number of days for which you want to list snapshot IDs.

```
Enter value for num days: 1
```

A list of existing snapshots for the specified time range is displayed. In this example, snapshots captured in the previous day are displayed.

4. Specify a beginning and ending snapshot ID for the workload repository report:

```
Enter value for begin_snap: 146
Enter value for end snap: 147
```

In this example, the snapshot with a snapshot ID of 146 is selected as the beginning snapshot, and the snapshot with a snapshot ID of 147 is selected as the ending snapshot.

Specify the SQL ID of a particular SQL statement to display statistics:

```
Enter value for sql_id: 2b064ybzkwf1y
```

In this example, the SQL statement with a SQL ID of 2b064ybzkwf1y is selected.

6. Enter a report name, or accept the default report name:

```
Enter value for report_name:
Using the report name awrrpt_1_146_147.html
```

In this example, the default name is accepted and an AWR report named ${\tt awrrpt}\ 1\ 146\ 147$ is generated.

Generating an AWR Report for a SQL Statement on a Specific Database

The awrsqrpi.sql SQL script generates an HTML or text report that displays statistics of a particular SQL statement from a range of snapshot IDs using a specific database instance. This script enables you to specify a database identifier and instance for which the AWR report will be generated. Run this report to inspect or debug the performance of a SQL statement on a specific database and instance.



To generate an AWR report for a SQL statement on a specific database instance using the command-line interface:

1. At the SQL prompt, enter:

```
@$ORACLE HOME/rdbms/admin/awrsqrpi.sql
```

2. Specify whether you want an HTML or a text report:

```
Enter value for report_type: html
```

In this example, an HTML report is chosen.

A list of available database identifiers and instance numbers are displayed:

3. Enter the values for the database identifier (dbid) and instance number (inst num):

```
Enter value for dbid: 3309173529
Using 3309173529 for database Id
Enter value for inst_num: 1
Using 1 for instance number
```

4. Specify the number of days for which you want to list snapshot IDs.

```
Enter value for num_days: 1
```

A list of existing snapshots for the specified time range is displayed. In this example, snapshots captured in the previous day are displayed.

Specify a beginning and ending snapshot ID for the workload repository report:

```
Enter value for begin_snap: 146
Enter value for end snap: 147
```

In this example, the snapshot with a snapshot ID of 146 is selected as the beginning snapshot, and the snapshot with a snapshot ID of 147 is selected as the ending snapshot.

Specify the SQL ID of a particular SQL statement to display statistics:

```
Enter value for sql_id: 2b064ybzkwf1y
```

In this example, the SQL statement with a SQL ID of 2b064ybzkwf1y is selected.

7. Enter a report name, or accept the default report name:

```
Enter value for report_name:
Using the report name awrrpt_1_146_147.html
```

In this example, the default name is accepted and an AWR report named <code>awrrpt_1_146_147</code> is generated on the database instance with a database ID value of <code>3309173529</code>.



Generating Performance Hub Active Report

Performance Hub feature of EM provides an active report with a consolidated view of all performance data for a specified time period. The report is fully interactive; its contents are saved in a HTML file, which you can access offline using a web browser.



Oracle Database 2 Day DBA for more information about Performance Hub feature of FM.

This section describes how to generate Performance Hub active report and contains the following topics:

- Overview of Performance Hub Active Report
- · Command-Line User Interface for Generating a Performance Hub Active Report
- Generating a Performance Hub Active Report Using a SQL Script
- Performance Workload Analysis

Overview of Performance Hub Active Report

Performance Hub active report enables you to view all performance data available for a time period that you specify. Different tabs are available in the Performance Hub, depending on whether real-time or historical data is selected for the time period. When real-time data is selected, more granular data is presented, because real-time data for the last hour is displayed. When historical data is selected, more detailed data is presented, but the data points are averaged out to the Automatic Workload Repository (AWR) interval for the selected time period.

This section describes Performance Hub active report and contains the following topics:

- About Performance Hub Active Report Tabs
- About Performance Hub Active Report Types

About Performance Hub Active Report Tabs

Performance Hub active report contains interactive tabs that enable you to view and navigate through performance data categorized into various performance areas.

The tabs contained in a Performance Hub active report include the following:

Summary

The Summary tab provides an overview of system performance, including resource consumption, average active sessions, and load profile information. This tab is available for real-time data as well as historical data.

Activity

The Activity tab displays ASH analytics. This tab is available for real-time data as well as historical data.

Workload



The Workload tab displays metric information about the workload profile, such as call rates, logon rate, and top SQL. This tab is available for real-time data as well as historical data.

RAC

The RAC tab displays metrics specific to Oracle RAC, such as the number of global cache blocks received and the average block latency. This tab is only available in Oracle RAC environments. This tab is available for real-time data as well as historical data.

Monitored SQL

The Monitored SQL tab displays information about monitored SQL statements. This tab is available for real-time data as well as historical data.

ADDM

The ADDM tab displays information for ADDM analysis tasks and Real-Time ADDM analysis reports. This tab is available for real-time data as well as historical data.

Current ADDM Findings

The Current ADDM Findings tab displays a real-time analysis of system performance for the past 5 minutes. This tab is only available if the specified time period for the Performance Hub active report is within the past hour. This tab is available only for real-time data.

Database time

The Database Time tab displays wait events by category for various metrics. This tab is available only for historical data.

Resources

The Resources tab displays operating system and I/O usage statistics. This tab is available only for historical data.

System Statistics

The System Statistics tab displays database and system statistics. This tab is available only for historical data.

About Performance Hub Active Report Types

You can choose the level of details displayed within each tab of the Performance Hub active report by selecting the report type.

The available report types for the Performance Hub active report include the following:

Basic

Only the basic information for all the tabs is saved to the report.

Typical

In addition to the information saved in the basic report type, the SQL Monitor information for the top SQL statements contained in the Monitored SQL tab and the ADDM reports are saved to the report.

All

In addition to the information saved in the typical report type, the SQL Monitor information for all SQL statements contained in the Monitored SQL tab and all detailed reports for all tabs are saved to the report.



Command-Line User Interface for Generating a Performance Hub Active Report

You can generate a Performance Hub active report using the command-line interface in one of two ways:

- Using a SQL script, as described in "Generating a Performance Hub Active Report Using a SQL Script".
- Using the DBMS_PERF package, as described in Oracle Database PL/SQL Packages and Types Reference.

Generating a Performance Hub Active Report Using a SQL Script

This section describes how to generate Performance Hub active report by running the perfhubrpt.sql SQL script in the command-line interface. The DBA role is required to run this script.

To generate a Performance Hub active report:

1. At the SQL prompt, enter:

```
@$ORACLE HOME/rdbms/admin/perfhubrpt.sql
```

2. Specify the desired report type:

```
Please enter report type: typical
```

For information about the available report types, see "About Performance Hub Active Report Types".

3. Enter the value for the database identifier of the database you want to use:

```
Please enter database ID: 3309173529
```

To use the local database, enter a null value (the default value). If you specify a database identifier for a database other than the local database, then the Performance Hub active report is generated from imported AWR data.

4. Enter the value for the instance number of the database instance you want to use:

```
Please enter instance number: all instances
```

To specify all instances, enter all instances (the default value).

5. Enter the desired time period by specifying an end time and a start time in the format of dd:mm:yyyy hh:mi:ss:

```
Please enter end time in format of dd:mm:yyyy hh24:mi:ss: 03:04:2014 17:00:00 Please enter start time in format of dd:mm:yyyy hh24:mi:ss: 03:04:2014 16:00:00
```

6. Enter a report name, or accept the default report name:

```
Enter value for report_name: my_perfhub_report.html
```

In this example, a Performance Hub active report named my_perfhub_report is generated on all database instances with a database ID value of 3309173529 for the specified time period from 4:00 p.m. to 5:00 p.m on April 3, 2014.



Performance Workload Analysis

Workload Analysis

The **Workload Analysis** is a tool to quickly identify issues with the workload by comparing a given workload with another reference, which could be fixed or rolling. The analysis itself can be a one time analysis, for instance comparing it two periods, or the analysis can be a reoccurring task that is scheduled with various flexible parameters. The time period for the analysis is selected from the **View Data** drop-down with the default ranges of **Last 24 Hours**, **Last 7 Days**, **Last 30 Days**, and **All**.

Scheduled Analysis

The **Scheduled Analysis** tab of the **Workload Analysis** page displays five boxes with counters to alert the number of differences to:

- Changed Workloads
- Monitored Workloads
- Regressed Metrics
- Improved Metrics
- Unchanged Metrics

The counter overview provides insight into how changes to workloads improved or regressed for the time period.

The **Create Analysis Task** for the **Scheduled Analysis** requires that you to provide additional information about the re-occurring task. Here are some settings, some of them optional, needed for establishing a re-occurring analysis in the **Workload Capture** form.

- SQL Tuning Set Name Prefix
- Load SQL Statements using AWR Snapshot,
 - Specify Custom Time Range with a Start Time and End Time
 - Quick Select from Snapshots created in the past with specified units of hours or days, and only up to 7 days

Total number of SQL Statements Captured

- Capture All
- Capture Top N
- Filter Option
- Workload Comparison
 - Subsequent Comparisons: Compare using a rolling reference, or compare using a fixed reference
 - Optional Initial reference workload
 - Comparison Metric such as CPU Time, Elapsed Time, Buffer Gets, Disk Reads,
 Direct Writes, Physical I/Os, and Optimizer Cost.
 - Change Threshold (%)
 - Consumer Resource Group
 - Days to retain purge results



Schedule

Time zone: Local (America/Denver), Database (GMT), or UTC

Start Date: Immediate, or Later specified with date field

End Date: None, or Specified Date **Repeat Every**: Select units first.

(1-24) Hours

(1-7) Days

* (1-4) Weeks

* (1-12) Months

Repeat On: Specified Date, or Custom Time

When an analysis task is created, it appears under the **Scheduled Analysis**. The arrow next to the Analysis name in the table expands to a list of various instances when that scheduled analysis was run and links to the resulting Comparison Report (or **Workload Analysis Report**)

One-Time Analysis

The **One-Time Analysis** tab of the **Workload Analysis** page displays five boxes with counters to alert the number of differences to:

- Analysis With Difference
- Total Analyses
- Regressed Metrics
- Improved Metrics
- Unchanged Metrics

The counter overview provides insight into how changes to workloads improved or regressed for the time period.

Create Analysis Task is similar for One-Time Analysis. The task requires

- A unique name for the task
- Select the **Reference Workload** by name or searching for them.
- Select the Compared Workload by name or searching for them.
- The task needs one or more comparison metrics, such as CPU Time, Elapsed Time,
 Buffer Gets, Disk Reads, Direct Writes, Physical I/Os, and Optimizer Cost.

When an analysis task is created, it appears under the **One-Time Analysis**. The arrow next to the Analysis name in the table expands to a list of various instances when that scheduled analysis was run and links to the resulting Comparison Report (or **Workload Analysis Report**)

Workload Analysis Report

The **Workload Analysis Report** is available after an analysis task has been run, whether regularly or one-time. Each analysis has its own report.

The information in the report varies according to the criteria established from **Create Analysis Task**. The reports contain data about the reference and/or comparison followed by the summary using a comparison metric, such as **CPU Time**, **Elapsed Time**, **Buffer Gets**, **Disk Reads**, **Direct Writes**, **Physical I/Os**, **Optimizer Cost**.



The Summary area lists, for instance, the Workload Impact by Overall, Common SQL, Improved SQL, Regressed SQL, Missing SQL, and New SQL:

The **Breakdown** area provides graphs for various performance topics, such as:

- **CPU Time** is a graph with the Before and After to the workload.
- **SQL Statement by Performance** is a graph with Improved, Regressed, Unchanged, Missing, and New categories and a bar indicating their number.
- SQL Statement by Plan Change shows the New Plan and Same Plan impacts side by side.

Top SQL Statements By Workload Impacts area is a table for each impact by:

- Impact
- SQL ID
- SQL Test
- Plan Change
- Net Absolute Impact on Workload (%)
- CPU Time Change (ms)



7

Automatic Performance Diagnostics

This chapter describes Oracle Database automatic features for performance diagnosing and tuning.

This chapter contains the following topics:

- Overview of the Automatic Database Diagnostic Monitor
- Setting Up ADDM
- Diagnosing Database Performance Problems with ADDM
- ADDM Views
- ADDM Spotlight



Oracle Database 2 Day + Performance Tuning Guide for information about using Oracle Enterprise Manager Cloud Control (Cloud Control) to diagnose and tune the database with the Automatic Database Diagnostic Monitor

Overview of the Automatic Database Diagnostic Monitor

The Automatic Workload Repository (AWR) stores performance related statics for an Oracle database. The Automatic Database Diagnostic Monitor (ADDM) is a diagnostic tool that analyzes the AWR data on a regular basis, locates root causes of any performance problems, provides recommendations for correcting the problems, and identifies non-problem areas of the system. Because AWR is a repository of historical performance data, ADDM can analyze performance issues after the event, often saving time and resources in reproducing a problem.

In most cases, ADDM output should be the first place that a DBA looks when notified of a performance problem. ADDM provides the following benefits:

- Automatic performance diagnostic report every hour by default
- Problem diagnosis based on decades of tuning expertise
- Time-based quantification of problem impacts and recommendation benefits
- Identification of root cause, not symptoms
- Recommendations for treating the root causes of problems
- Identification of non-problem areas of the system
- Minimal overhead to the system during the diagnostic process
- Accurate diagnosis of performance issues due to misconfiguration, incorrect usage, etc.
- Identification of objects that are candidates for in-memory placement as well as objects that are in-memory but not utilized for in-memory scans.

- Detection of hard disks in use instead of flash, low cell offloading, I/O workload throttling or imbalance across disks or cells in the system, which improve performance diagnosis capabilities particularly for Exadata.
- Diagnosability into failed tasks and associated error messages through DBA accessible views

Tuning is an iterative process, and fixing one problem can cause the bottleneck to shift to another part of the system. Even with the benefit of ADDM analysis, it can take multiple tuning cycles to reach acceptable system performance. ADDM benefits apply beyond production systems; on development and test systems, ADDM can provide an early warning of performance issues.

Note:

Data visibility and privilege requirements may differ when using ADDM features with pluggable databases (PDBs). For information about how manageability features, including ADDM features, work in a multitenant container database (CDB), see *Oracle Multitenant Administrator's Guide*.

ADDM Analysis

An ADDM analysis can be performed on a pair of AWR snapshots and a set of instances from the same database. The pair of AWR snapshots define the time period for analysis, and the set of instances define the target for analysis.

If you are using Oracle Real Application Clusters (Oracle RAC), then ADDM has three analysis modes:

Database

In Database mode, ADDM analyzes all instances of the database.

Instance

In Instance mode, ADDM analyzes a particular instance of the database.

Partial

In Partial mode, ADDM analyzes a subset of all database instances.

If you are not using Oracle RAC, then ADDM can only function in Instance mode because only one instance of the database exists.

An ADDM analysis is performed each time an AWR snapshot is taken and the results are saved in the database. The time period analyzed by ADDM is defined by the last two snapshots (the last hour by default). ADDM will always analyze the specified instance in Instance mode. For non-Oracle RAC or single instance environments, the analysis performed in the Instance mode is the same as a database-wide analysis. If you are using Oracle RAC, then ADDM also analyzes the entire database in Database mode, as described in "Using ADDM with Oracle Real Application Clusters".

After ADDM completes its analysis, you can view the ADDM results using either Cloud Control, or DBMS_ADDM package subprograms, or DBA_ADDM_* and DBA_ADVISOR_* views.

ADDM analysis is performed top down, first identifying symptoms, and then refining them to reach the root causes of performance problems. The goal of the analysis is to reduce a single throughput metric called DB time. DB time is the fundamental measure of database

performance, and is the cumulative time spent by the database in processing user requests. It includes wait time and CPU time of all non-idle user foreground sessions. DB time is displayed in the V\$SESS TIME MODEL and V\$SYS TIME MODEL views.

By reducing DB time, the database is able to support more user requests using the same resources, which increases throughput. The problems reported by ADDM are sorted by the amount of DB time they are responsible for. System areas that are not responsible for a significant portion of DB time are reported as non-problem areas.

The types of problems that ADDM considers include the following:

- CPU bottlenecks Is the system CPU bound by Oracle Database or some other application?
- Undersized Memory Structures Are the Oracle Database memory structures, such as the SGA, PGA, and buffer cache, adequately sized?
- I/O capacity issues Is the I/O subsystem performing as expected?
- High-load SQL statements Are there any SQL statements which are consuming excessive system resources?
- High-load PL/SQL execution and compilation, and high-load Java usage
- Oracle RAC specific issues What are the global cache hot blocks and objects; are there any interconnect latency issues?
- Sub-optimal use of Oracle Database by the application Are there problems with poor connection management, excessive parsing, or application level lock contention?
- Database configuration issues Is there evidence of incorrect sizing of log files, archiving issues, excessive checkpoints, or sub-optimal parameter settings?
- Concurrency issues Are there buffer busy problems?
- Hot objects and top SQL for various problem areas

Note:

This is not a comprehensive list of all problem types that ADDM considers in its analysis.

ADDM also documents the non-problem areas of the system. For example, wait event classes that are not significantly impacting the performance of the system are identified and removed from the tuning consideration at an early stage, saving time and effort that would be spent on items that do not impact overall system performance.

See Also:

- Oracle Database Reference for information about the V\$SESS_TIME_MODEL and V\$SYS TIME MODEL views
- "Time Model Statistics" for a discussion of time model statistics and DB time
- Oracle Database Concepts for information about server processes



Using ADDM with Oracle Real Application Clusters

If you are using Oracle RAC, then run ADDM in Database analysis mode to analyze the throughput performance of all instances of the database. In Database mode, ADDM considers DB time as the sum of the database time for all database instances. Using the Database analysis mode enables you to view all findings that are significant to the entire database in a single report, instead of reviewing a separate report for each instance.

The Database mode report includes findings about database resources (such as I/O and interconnect). The report also aggregates findings from the various instances if they are significant to the entire database. For example, if the CPU load on a single instance is high enough to affect the entire database, then the finding appears in the Database mode analysis, which points to the particular instance responsible for the problem.



Oracle Real Application Clusters Administration and Deployment Guide for information about using ADDM with Oracle RAC

Using ADDM in a Multitenant Environment

Starting with Oracle Database 12c, ADDM is enabled by default in the root container of a multitenant container database (CDB). Starting with Oracle Database 19c, you can also use ADDM in a pluggable database (PDB).

Note:

A multitenant container database is the only supported architecture in Oracle Database 21c and later releases. While the documentation is being revised, legacy terminology may persist. In most cases, "database" and "non-CDB" refer to a CDB or PDB, depending on context. In some contexts, such as upgrades, "non-CDB" refers to a non-CDB from a previous release.

In a CDB, ADDM works in the same way as it works in a non-CDB, that is, the ADDM analysis is performed each time an AWR snapshot is taken on a CDB root or a PDB, and the ADDM results are stored on the same database system where the snapshot is taken. The time period analyzed by ADDM is defined by the last two snapshots (the last hour by default).

After ADDM completes its analysis, you can view the ADDM results using any of the following methods:

- Using Enterprise Manager Cloud Control (Cloud Control)
- Using the DBA ADDM * and DBA ADVISOR * views



Note:

- ADDM is enabled by default in a CDB root.
- ADDM does not work in a PDB by default, because automatic AWR snapshots are disabled by default in a PDB. To use ADDM in a PDB, you must enable automatic AWR snapshots in the PDB.
- A user whose current container is the CDB root can view ADDM results for the entire CDB. The ADDM results can include information about multiple PDBs. ADDM results related to a PDB are not included if the PDB is unplugged. The ADDM results stored on the CDB root cannot be viewed when the current container is a PDB.
- ADDM results on a PDB provide only PDB-specific findings and recommendations. A user whose current container is a PDB can view ADDM results for the current PDB only. The ADDM results exclude findings that apply to the CDB as a whole, for example, I/O problems relating to the buffer cache size.
- Enabling AWR snapshots on a PDB does not change the ADDM report on the CDB root.
- AWR data on a PDB cannot be accessed from the CDB root.

PDB-Level ADDM Restrictions

Unlike in a non-CDB, ADDM does not report the following issues in a PDB, because these issues apply to a CDB as a whole and do not apply to an individual PDB:

- I/O problems due to:
 - undersized buffer cache
 - undersized streams pool
 - excessive temporary writes
 - excessive checkpoint writes
 - excessive undo writes
 - excessive PQ checkpoint writes
 - excessive truncate writes
 - excessive tablespace DDL checkpoint
 - I/O capacity limit
- SQL hard parsing issues due to:
 - cursor aging
 - out-of-memory failed parse
- SGA sizing issues

ADDM also does not report the following issues in a PDB, because these issues cannot be resolved at a PDB level:

- Cluster messaging related issues, such as network latency, congestion, contention, and lost blocks
- Log file switch waits on archiving and on checkpoint incomplete



- Too many free-buffer waits
- Contention on log buffer waits
- Waits due to CPU bottleneck
- Operating system VM paging
- Session slot wait event
- CPU quantum wait event
- RMAN related wait events, such as PQ queued wait event, PGA limit wait event, and I/O queue wait event

See Also:

- "Enabling ADDM in a Pluggable Database" for information about how to enable ADDM in a PDB
- "ADDM Views" for more information about the ADDM views $\tt DBA_ADDM_*$ and $\tt DBA_ADVISOR~*$
- "Diagnosing Database Performance Problems with ADDM" for information about how to run ADDM in an Oracle database using the DBMS_ADDM package subprograms
- Oracle Database 2 Day + Performance Tuning Guide for information about how to run ADDM in an Oracle database using Cloud Control

Enabling ADDM in a Pluggable Database

ADDM does not work in a pluggable database (PDB) by default, because automatic AWR snapshots are disabled by default in a PDB. To use ADDM in a PDB, you must enable automatic AWR snapshots in the PDB by setting the AWR_PDB_AUTOFLUSH_ENABLED initialization parameter to TRUE and AWR snapshot interval greater than 0.

To enable ADDM in a PDB:

1. The AWR_PDB_AUTOFLUSH_ENABLED initialization parameter has a default value of TRUE. If it is not set to TRUE, use the following command to change it:

```
SQL> ALTER SYSTEM SET AWR PDB AUTOFLUSH ENABLED=TRUE;
```

Set the AWR snapshot interval greater than 0 in the PDB using the command as shown in the following example:

```
SQL> EXEC dbms workload repository.modify snapshot settings(interval=>60);
```



See Also:

- Oracle Database Reference for more information about the AWR PDB AUTOFLUSH ENABLED initialization parameter
- Oracle Database PL/SQL Packages and Types Reference for more information about the DBMS_WORKLOAD_REPOSITORY.MODIFY_SNAPSHOT_SETTINGS procedure

Real-Time ADDM Analysis

Introduced in Oracle Enterprise Manager Cloud Control (Cloud Control) 12c, Real-Time ADDM helps you to analyze and resolve problems in unresponsive or hung databases that traditionally require you to restart the database. Real-Time ADDM runs through a set of predefined criteria to analyze the current performance of the database. After analyzing the problem, Real-Time ADDM helps you to resolve the identified issues—such as deadlocks, no response, shared pool contention, and other exception situations—without having to restart the database.

This section describes Real-Time ADDM and contains the following topics:

- Real-Time ADDM Connection Modes
- Real-Time ADDM Triggers
- Real-Time ADDM Trigger Controls



Oracle Database 2 Day + Performance Tuning Guide for information about using Real-Time ADDM with Cloud Control

Real-Time ADDM Connection Modes

Depending on the database state, Real-Time ADDM uses two different types of connection modes when connecting to the database using Cloud Control:

Normal connection

In this mode, Real-Time ADDM performs a normal JDBC connection to the database. This mode is intended to perform extensive performance analysis of the database when some connectivity is available.

Diagnostic connection

In this mode, Real-Time ADDM performs a latch-less connection to the database. This mode is intended for extreme non-responsive situations when a normal JDBC connection is not possible.

Real-Time ADDM Triggers

Real-Time ADDM proactively detects transient database performance issues. To do this, Real-Time ADDM runs automatically every 3 seconds and uses in-memory data to diagnose any performance spikes in the database.



Real-Time ADDM triggers an analysis automatically when a performance problem is detected, as described in the following steps:

- Every 3 seconds, the manageability monitor process (MMON) performs an action to obtain performance statistics without lock or latch.
- The MMON process checks these statistics and triggers a Real-Time ADDM analysis if any of the issues listed in Table 7-1 are detected.
- 3. The MMON child process creates the report and stores it in the AWR.

To view metadata for the report, use the DBA HIST REPORTS view.

Table 7-1 lists the issues and conditions that trigger a Real-Time ADDM analysis.

Table 7-1 Triggering Issues and Conditions for Real-Time ADDM

Issue	Condition
High load	Average active sessions are greater than 3 times the number of CPU cores
I/O bound	I/O impact on active sessions based on single block read performance
CPU bound	Active sessions are greater than 10% of total load and CPU utilization is greater than 50%
Over-allocated memory	Memory allocations are over 95% of physical memory
Interconnect bound	Based on single block interconnect transfer time
Session limit	Session limit is close to 100%
Process limit	Process limit is close to 100%
Hung session	Hung sessions are greater than 10% of total sessions
Deadlock detected	Any deadlock is detected

See Also

Oracle Database Reference for information about the DBA_HIST_REPORTS view

Real-Time ADDM Trigger Controls

To ensure that the automatic triggers do not consume too many system resources and overwhelm the system, Real-Time ADDM employs the following controls:

Duration between reports

If a Real-Time ADDM report was created in the past 5 minutes by the automatic trigger, then no new reports will be generated.

Oracle RAC control

Automatic triggers are local to the database instance. For Oracle RAC, only one database instance can create a Real-Time ADDM report at a given time because a lock is required and a query is performed by the MMON child process before the report is actually generated.

Repeated triggers

An automatic trigger for any issue must have an impact of 100% or higher than the previous report with the same triggering issue within the past 45 minutes. For example, if a



report is triggered for active sessions with an impact of 8 sessions, then in order for another report to trigger within the next 45 minutes, there must be at least 16 active sessions. In this case, the reported problem with the database is becoming more severe over time. On the other hand, if the same report is being generated once every 45 minutes, then the database is experiencing a persistent problem that has a consistent impact.

Newly identified issues

If a new issue is detected (that was not previously detected within the past 45 minutes), then a new report is generated. For example, if a report is triggered for 8 active sessions and a new deadlock issue is detected, then a new report is generated regardless of the new active sessions load.

ADDM Analysis Results

In addition to problem diagnostics, ADDM recommends possible solutions. ADDM analysis results are represented as a set of findings. See Example 7-1 for an example of an ADDM analysis result. Each ADDM finding can belong to one of the following types:

- Problem findings describe the root cause of a database performance problem.
- Symptom findings contain information that often lead to one or more problem findings.
- Information findings are used for reporting information that are relevant to understanding
 the performance of the database, but do not constitute a performance problem (such as
 non-problem areas of the database and the activity of automatic database maintenance).
- Warning findings contain information about problems that may affect the completeness or accuracy of the ADDM analysis (such as missing data in AWR).

Each problem finding is quantified by an impact that is an estimate of the portion of DB time caused by the finding's performance issue. A problem finding can be associated with a list of recommendations for reducing the impact of the performance problem. The types of recommendations include:

- · Hardware changes: adding CPUs or changing the I/O subsystem configuration
- Database configuration: changing initialization parameter settings
- Schema changes: hash partitioning a table or index, or using automatic segment-space management (ASSM)
- Application changes: using the cache option for sequences or using bind variables
- Using other advisors: running SQL Tuning Advisor on high-load SQL or running Segment Advisor on hot objects

A list of recommendations can contain various alternatives for solving the same problem; you do not have to apply all the recommendations to solve a specific problem. Each recommendation has a benefit which is an estimate of the portion of DB time that can be saved if the recommendation is implemented. Recommendations are composed of actions and rationales. You must apply all the actions of a recommendation to gain the estimated benefit. The rationales are used for explaining why the set of actions were recommended and to provide additional information to implement the suggested recommendation.

Reviewing ADDM Analysis Results: Example

Consider the following section of an ADDM report in Example 7-1.



Example 7-1 Example ADDM Report

In Example 7-1, the finding points to a particular root cause, the usage of literals in SQL statements, which is estimated to have an impact of about 31% of total DB time in the analysis period.

The finding has a recommendation associated with it, composed of one action and one rationale. The action specifies a solution to the problem found and is estimated to have a maximum benefit of up to 31% DB time in the analysis period. Note that the benefit is given as a portion of the total DB time and not as a portion of the finding's impact. The rationale provides additional information on tracking potential SQL statements that were using literals and causing this performance issue. Using the specified plan hash value of SQL statements that could be a problem, a DBA could quickly examine a few sample statements.

When a specific problem has multiple causes, ADDM may report multiple problem and symptom findings. In this case, the impacts of these multiple findings can contain the same portion of DB time. Because the performance issues of findings can overlap, the sum of the impacts of the findings can exceed 100% of DB time. For example, if a system performs many reads, then ADDM might report a SQL statement responsible for 50% of DB time due to I/O activity as one finding, and an undersized buffer cache responsible for 75% of DB time as another finding.

When multiple recommendations are associated with a problem finding, the recommendations may contain alternatives for solving the problem. In this case, the sum of the recommendations' benefits may be higher than the finding's impact.

When appropriate, an ADDM action may have multiple solutions for you to choose from. In the example, the most effective solution is to use bind variables. However, it is often difficult to modify the application. Changing the value of the <code>CURSOR_SHARING</code> initialization parameter is much easier to implement and can provide significant improvement.

Setting Up ADDM

Automatic database diagnostic monitoring is enabled by default and is controlled by the CONTROL_MANAGEMENT_PACK_ACCESS and the STATISTICS_LEVEL initialization parameters.

The <code>CONTROL_MANAGEMENT_PACK_ACCESS</code> parameter should be set to <code>DIAGNOSTIC</code> or <code>DIAGNOSTIC+TUNING</code> to enable automatic database diagnostic monitoring. The default setting is <code>DIAGNOSTIC+TUNING</code>. Setting <code>CONTROL MANAGEMENT PACK ACCESS</code> to <code>NONE</code> disables ADDM.

The STATISTICS_LEVEL parameter should be set to the TYPICAL or ALL to enable automatic database diagnostic monitoring. The default setting is TYPICAL. Setting STATISTICS_LEVEL to BASIC disables many Oracle Database features, including ADDM, and is strongly discouraged.



See Also:

microseconds.

Oracle Database Reference for information about the CONTROL_MANAGEMENT_PACK_ACCESS and STATISTICS_LEVEL initialization parameters

ADDM analysis of I/O performance partially depends on a single argument, <code>DBIO_EXPECTED</code>, that describes the expected performance of the I/O subsystem. The value of <code>DBIO_EXPECTED</code> is the average time it takes to read a single database block in microseconds. Oracle Database uses the default value of 10 milliseconds, which is an appropriate value for most modern hard drives. If your hardware is significantly different—such as very old hardware or very fast RAM disks—then consider using a different value.

To determine the correct setting for the DBIO_EXPECTED parameter:

- Measure the average read time of a single database block read for your hardware.
 Note that this measurement is for random I/O, which includes seek time if you use standard hard drives. Typical values for hard drives are between 5000 and 20000
- 2. Set the value one time for all subsequent ADDM executions.

For example, if the measured value if 8000 microseconds, you should execute the following command as SYS user:

Diagnosing Database Performance Problems with ADDM

To diagnose database performance problems, first review the ADDM analysis results that are automatically created each time an AWR snapshot is taken. If a different analysis is required (such as a longer analysis period, using a different DBIO_EXPECTED setting, or changing the analysis mode), you can run ADDM manually as described in this section.

ADDM can analyze any two AWR snapshots (on the same database), as long as both snapshots are still stored in AWR (have not been purged). ADDM can only analyze instances that are started before the beginning snapshot and remain running until the ending snapshot. Additionally, ADDM will not analyze instances that experience significant errors when generating AWR snapshots. In such cases, ADDM will analyze the largest subset of instances that did not experience these problems.

The primary interface for diagnostic monitoring is Cloud Control. Whenever possible, run ADDM using Cloud Control, as described in *Oracle Database 2 Day + Performance Tuning Guide*. If Cloud Control is unavailable, then run ADDM using the DBMS_ADDM package. To run the DBMS_ADDM APIs, the user must be granted the ADVISOR privilege.

Running ADDM in Database Mode

For Oracle RAC configurations, you can run ADDM in Database mode to analyze all instances of the databases. For single-instance configurations, you can still run ADDM in Database mode; ADDM will behave as if running in Instance mode.

To run ADDM in Database mode, use the DBMS ADDM.ANALYZE DB procedure:

The <code>task_name</code> parameter specifies the name of the analysis task that will be created. The <code>begin_snapshot</code> parameter specifies the snapshot number of the beginning snapshot in the analysis period. The <code>end_snapshot</code> parameter specifies the snapshot number of the ending snapshot in the analysis period. The <code>db_id</code> parameter specifies the database identifier of the database that will be analyzed. If unspecified, this parameter defaults to the database identifier of the database to which you are currently connected.

The following example creates an ADDM task in database analysis mode, and executes it to diagnose the performance of the entire database during the time period defined by snapshots 137 and 145:

```
VAR tname VARCHAR2(30);
BEGIN
  :tname := 'ADDM for 7PM to 9PM';
  DBMS_ADDM.ANALYZE_DB(:tname, 137, 145);
END;
//
```

Running ADDM in Instance Mode

To analyze a particular instance of the database, you can run ADDM in Instance mode. To run ADDM in Instance mode, use the DBMS ADDM.ANALYZE INST procedure:

The <code>task_name</code> parameter specifies the name of the analysis task that will be created. The <code>begin_snapshot</code> parameter specifies the snapshot number of the beginning snapshot in the analysis period. The <code>end_snapshot</code> parameter specifies the snapshot number of the ending snapshot in the analysis period. The <code>instance_number</code> parameter specifies the instance number of the instance that will be analyzed. If unspecified, this parameter defaults to the instance number of the instance to which you are currently connected. The <code>db_id</code> parameter specifies the database identifier of the database that will be analyzed. If unspecified, this parameter defaults to the database identifier of the database to which you are currently connected.

The following example creates an ADDM task in instance analysis mode, and executes it to diagnose the performance of instance number 1 during the time period defined by snapshots 137 and 145:

```
VAR tname VARCHAR2(30);
BEGIN
  :tname := 'my ADDM for 7PM to 9PM';
  DBMS_ADDM.ANALYZE_INST(:tname, 137, 145, 1);
```



```
END;
```

Running ADDM in Partial Mode

To analyze a subset of all database instances, you can run ADDM in Partial mode. To run ADDM in Partial mode, use the DBMS ADDM.ANALYZE PARTIAL procedure:

The task_name parameter specifies the name of the analysis task that will be created. The instance_numbers parameter specifies a comma-delimited list of instance numbers of instances that will be analyzed. The begin_snapshot parameter specifies the snapshot number of the beginning snapshot in the analysis period. The end_snapshot parameter specifies the snapshot number of the ending snapshot in the analysis period. The db_id parameter specifies the database identifier of the database that will be analyzed. If unspecified, this parameter defaults to the database identifier of the database to which you are currently connected.

The following example creates an ADDM task in partial analysis mode, and executes it to diagnose the performance of instance numbers 1, 2, and 4, during the time period defined by snapshots 137 and 145:

```
VAR tname VARCHAR2(30);
BEGIN
  :tname := 'my ADDM for 7PM to 9PM';
  DBMS_ADDM.ANALYZE_PARTIAL(:tname, '1,2,4', 137, 145);
END;
//
```

Displaying an ADDM Report

To display a text report of an executed ADDM task, use the DBMS ADDM.GET REPORT function:

The following example displays a text report of the ADDM task specified by its task name using the tname variable:

```
SET LONG 1000000 PAGESIZE 0;
SELECT DBMS ADDM.GET REPORT(:tname) FROM DUAL;
```

Note that the return type of a report is a CLOB, formatted to fit line size of 80. For information about reviewing the ADDM analysis results in an ADDM report, see "ADDM Analysis Results".



Re-Running Failed ADDM Tasks

The main purpose of Automatic Database Diagnostic Monitoring (ADDM) is to run automatically at the end of each AWR snapshot (default is every hour.) If an ADDM task fails, the failure is written to the database trace logs.

In addition, the DBA_ADDM_PENDING_AUTOTASKS view tracks the ADDM tasks being executed. The views are accessible at both CDB and PDB levels and contain information about the tasks executed at each level. ADB and Cloud users can see information about the tasks that were executed in their PDB, obtain more diagnostic information on failed tasks, and manually reexecute those tasks. Once a task is run successfully, the entry is removed, otherwise failed tasks remain in the view. When the corresponding AWR snapshots age out, any associated failed tasks age out.

The <code>FAILED_AUTO_TASKS_REPORT</code> function returns a report about the ADDM tasks that had failed and are visible in the <code>DBA_ADDM_PENDING_AUTOTASKS</code> views for the corresponding PDB level. The report can be generated to visualize either all failed tasks or only those that meet certain criteria, such as a time range, or identifier. The information in the report helps identify and diagnose the existence of failed ADDM tasks and their possible root causes, and thus useful to a DBA.

The REEXECUTE_FAILED_AUTO_TASKS procedure re-executes the tasks registered as failed in the DBA_ADDM_PENDING_AUTOTASKS view. To re-run a task, the user specifies the range for the snapshots or a time interval. The duration of the re-execution is time constrained by a user-specified parameter. Once a task has executed successfully, its corresponding row is removed from the DBA_ADDM_PENDING_AUTOTASKS view.

FAILED AUTO TASKS_REPORT Function

This function creates a plain text, user-readable report for the failed tasks registered in the \mbox{DBA} ADDM PENDING AUTOTASKS views.

Syntax

Parameters

Table 7-2 FAILED_AUTO_TASKS_REPORT Function Parameters

Parameter	Description
instance_number	Instance number for the tasks to be reported.
begin_snapshot	Earliest begin snapshot ID for the tasks to be reported.
end_snapshot	Latest end snapshot ID for the tasks to be reported.
dbid	Database ID for the tasks to be reported.



REEXECUTE_FAILED_AUTO_TASKS Procedure

This procedure re-executes the tasks registered in the DBA ADDM PENDING AUTOTASKS views.

This API can be called from both the CDB and PDB level. The user specifies the range for the snapshots or a time interval. The duration of the re-execution is time constrained by a user-specified parameter or a default value. Once a task has executed successfully, its corresponding row is removed from the DBA ADDM PENDING AUTOTASKS view.

Syntax

Parameters

Table 7-3 REEXECUTE FAILED AUTO TASKS Function Parameters

Parameter	Description
instance_number	Instance number for the tasks to be re executed.
begin_snapshot	Earliest begin snapshot ID for the tasks to be re executed.
end_snapshot	Latest end snapshot ID for the tasks to be re executed.
dbid	Database ID for the tasks to be re executed.
time_budget_in_sec	Forced time out to run the procedure after exiting.
max_attempts	Maximum number of attempts to re execute a task.

ADDM Views

Typically, you should view ADDM analysis using Cloud Control or DBMS_ADDM package subprograms.

However, you can also get ADDM information using the DBA_ADDM_* and DBA_ADVISOR_* views. This group of views includes:

• DBA ADVISOR FINDINGS

This view displays all the findings discovered by all advisors. Each finding is displayed with an associated finding ID, name, and type. For tasks with multiple executions, the name of each task execution associated with each finding is also listed.

DBA ADDM FINDINGS

This view contains a subset of the findings displayed in the related <code>DBA_ADVISOR_FINDINGS</code> view. This view only displays the ADDM findings discovered by all advisors. Each ADDM finding is displayed with an associated finding ID, name, and type.

DBA ADDM PENDING AUTOTASKS



This view tracks the automatic ADDM tasks that are being executed. The views are accessible at both CDB and PDB levels and contain information about the tasks executed at each level. ADB and Cloud users can see information about the tasks that were executed in their PDB, obtain more diagnostic information on failed tasks, and manually reexecute those tasks. Once a task is executed successfully, the entry is removed from the view. Failed tasks automatically age out and are removed when the corresponding AWR snapshots age out.

DBA_ADVISOR_FINDING_NAMES

This view lists all finding names registered with the advisor framework.

• DBA ADVISOR RECOMMENDATIONS

This view displays the results of completed diagnostic tasks with recommendations for the problems identified in each execution. The recommendations should be reviewed in the order of the RANK column, as this relays the magnitude of the problem for the recommendation. The BENEFIT column displays the benefit to the system you can expect after the recommendation is performed. For tasks with multiple executions, the name of each task execution associated with each advisor task is also listed.

DBA ADVISOR TASKS

This view provides basic information about existing tasks, such as the task ID, task name, and when the task was created. For tasks with multiple executions, the name and type of the last or current execution associated with each advisor task is also listed.

See Also:

- Oracle Database Reference for more information about DBA_ADDM_* and DBA ADVISOR * views
- "Displaying an ADDM Report"

ADDM Spotlight

Once the Automatic Database Diagnostic Monitor (ADDM) has been configured, you can use the ADDM Spotlight to consolidate the recommendations from ADDM for your Oracle database.

ADDM Spotlight provides a visual representation of critical performance categories, such as CPU Usage, Top SQL Statements, Resource Manager CPU Throttling, and "User I/O" wait Class. It rates the overall impact on system performance and provides recommendations. It helps locate root causes of any performance problems, provides recommendations for correcting the problems, and identifies problem areas of the system.

While in Enterprise Manager for your database, you select from the Performance drop-down menu the ADDM Spotlight item. The default time range is the last 24 hours when first entering the ADDM Spotlight. You can use the Quick Select pull-down menu to specify other common ranges, such as the last 7 days. Or you can specify a specific time range with time zone.

The Findings tab of ADDM Spotlight displays a graphic summary of impacts identified by category for that time range. Each point in time in the graph has a bar whose height represents the number of findings from all categories. Each bar is further subdivided by color-coded impact categories whose size on each individual bar is its number of findings. When you mouse-over a colored area of a specific bar, the tooltip's detailed information about that



category's impact at the database time is displayed, as shown in Figure 7-1. Within the Findings tab is a summary table showing each finding that was reported during the period. For each finding, it shows the category, frequency, impact, and number of recommendations.

Figure 7-1 The Findings tab of ADDM Spotlight

The Recommendations tab of ADDM Spotlight has radio buttons along the top for specifying the category: Database Parameters, SQL, and Schema Objects. Each point in time in the graphic summary for the category has a bar whose height represents the count of performance impacts. Further, each bar is subdivided by color-coded impact of specific items in that category.

For example, if the SQL category is selected as shown in Figure 7-2, the color-coded areas of each bar represents the count of individual SQL commands. The table below the summary graph provides details by SQL ID, the recommendation for improving its performance, the frequency of its occurrence, and the estimated benefit of implementing the recommendation for that SQL ID. Similarly, the Database Parameters and the Scheme Objects categories display recommendations for those respective items. The information provided for the SQL category is useful for subsequent actions, such as those performed with the SQL Tuning Advisor.



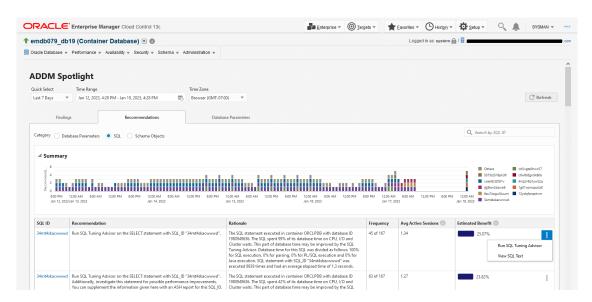
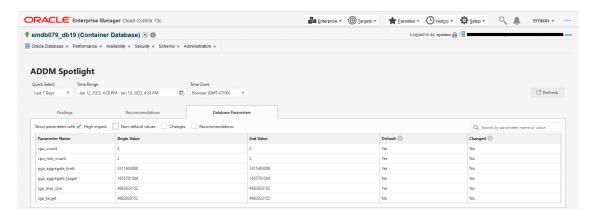


Figure 7-2 The Recommendations Tab of ADDM Spolight

If the Database Parameters tab of ADDM Spotlight is selected as shown in Figure 7-3, the table shows information abot the values of those database parameters, such as Parameter Name, Begin Value, End Value, Default, and Changed. Filter checkboxes for High Impact, Non-Default Values, Changes, and Recommendations help focus the database parameter table on areas of interest. The Changes column has a pop-up that shows the change history for that database parameter.

Figure 7-3 The Database Parameters Tab of ADDM Spotlight





8

Monitoring Databases with Performance Hub

Use Performance Hub to monitor database activity, diagnose issues, and tune queries to improve the performance of Oracle databases.

You can use Performance Hub to view real-time and historical performance data for a database. Performance Hub shows Active Session History (ASH) analytics, SQL monitoring and workload information.

About Performance Hub

With Performance Hub, you can monitor your databases for defined time periods and download statistical reports.

What is Performance Hub?

Performance Hub displays information about the performance of your database for the time period you specify.

With this tool, you can view real-time and historical performance data.

Performance Hub supports:

- Autonomous Database on shared and dedicated Exadata infrastructure
- Databases running on Base Database Service and Exadata Database Service on Dedicated Infrastructure systems
- Databases managed with OCI's External Database

How Performance Hub Works

For Autonomous Database, you can access the Performance Hub feature from the Autonomous Database Details page.

To use Performance Hub for databases running on Base Database Service, Exadata Cloud Service, or databases managed with the External Database Service:

Enable Database Management

While enabling Database Management for Base Database Service or Exadata Cloud service, the database administrator can select from two options, Basic Management or Full Management. For more information, see Enable Database Management.



If you enable Database Management for Autonomous Database, you can access Performance Hub from the Database Management service also.

In Database Management, you can set preferred credentials for users so they can automatically connect to the Managed Database and use Performance Hub.

For information about preferred credentials in Database Management, see Set Preferred Credentials.



Using Oracle Identity Cloud Service (IAM), you can create a policy that grants users access to Performance Hub while limiting actions they can take on the Oracle Database being analyzed. For more information, about policies and how to use them, see How Policies Work.

The following example shows a policy that grants access only to performance data without allowing general use access on Autonomous Databases.

```
Allow group <groupname>
   to inspect autonomous-database-family in compartment <name>
Allow group <groupname>
   to use autonomous-database-family in compartment <name>
   where request.operation = 'RetrieveAutonomousDatabasePerformanceBulkData'
```

Performance Hub Features

Performance Hub consists of a number of tabs and options, and provides reports so you can monitor your databases.



These tabs displayed in Performance Hub depends on the Database Management option that is enabled for the database you are monitoring, if the database is running in a Base Database Service or Exadata Database Service on Dedicated Infrastructure system, or is managed with External Database.

When a database is set to:

- Basic Management Performance Hub displays only the ASH Analytics and SQL Monitoring tabs.
- Full Management Performance Hub displays all the tabs described in the subsequent sections.

For more information, see About Management Options in Enable Database Management for Oracle Cloud Databases.



Performance Hub currently does not support Pluggable Databases in Base Database Service systems.

The Performance Hub page includes the following tabs:

- ASH Analytics
- SQL Monitoring
- Automatic Database Diagnostic Monitor (ADDM)
- Workload
- Blocking Sessions
- Exadata

This tab is available for databases running in Exadata Database Service on Dedicated Infrastructure and for Exadata Databases managed by External Database.

Additionally, Performance Hub provides options for:

AWR Snapshot Collection

Performance Hub is also available in another mode, namely Top Activity Lite. This mode is a simplified view of Performance Hub and focuses on fast response time to monitor the performance of database systems in real time.

You can explore these tabs and options to monitor a selected database and analyze its performance, which includes but is not limited to:

- How much is a database waiting for a resource, such as CPU or disk I/O
- Whether database performance degraded over a given time period, and what is the primary reason
- Which specific modules are causing a load on the system, and where most of the time a database spends on the modules
- Which SQL statements are the key contributors to changes in database performance, and which executions are causing them
- Which user sessions are causing performance bottlenecks
- Which sessions are currently blocking, and if outstanding requests exist for a lock
- Whether any databases are interfering with the current database
- Which databases, disks, and cells are causing high I/O load, and if the load is for Backup, Rebalance, or User I/O
- What cells are CPU bound, and how to determine the cell offload
- · What a database's configuration is, including software version and disk status

ASH Analytics

The ASH Analytics tab shows Active Session History (ASH) analytics charts that you can use to explore ASH data. This tab is displayed by default.

You can use this tab to drill down into database performance across multiple dimensions, such as **Consumer Group**, **Wait Class**, **SQL ID**, **User Name**, and so on.

In the ASH Analytics tab, you can select an Average Active Sessions dimension and view the top activity for that dimension for the selected time period. The Average Active Sessions chart has a control at the right end of the chart to select the displayed resolution of ASH data.

For more information about ASH, see:

Active Session History (ASH) in Oracle Database Concepts



ASH Sample Resolution

The ASH Sample Resolution menu gives users the ability to control the sampling of ASH data displayed in the Average Active Sessions chart.

Data resolution means displaying more or fewer data points in the sample data in given time period. Lower resolution displays coarser data with better performance and less effect on the database. Higher resolution aggregates more data to display finer detail, but can have a corresponding cost in latency and effect on the database.

The Sample Resolution menu is displayed at the right-side of the chart. The data resolution options are:

- Low: The chart displays the fewest data points in the selected data sample.
- **Medium**: The chart displays more data points in the selected data sample.
- High: The chart displays more data points in the selected data sample.
- Maximum: The chart displays the most data points available in the selected data sample.

For information about how to use this feature, see:

Average Active Sessions Data

Activity Tables

By default, two tables under the **Average Active Sessions** graph display the top SQLs and user sessions for the time period covered by the Average Active Sessions graph. To view activities by other dimensions, use the menus at the top-left of each of the two tables.

You can use the ASH Analytics tab to monitor real-time SQL activity.

For more information, see:

Real-time SQL Activity

SQL Monitoring

The SQL Monitoring tab contains a table that displays monitored SQL statements. To view this tab, click **SQL Monitoring** on the Performance Hub page.

The table in the SQL Monitoring tab displays SQL statement executions by dimensions, such as Last Active Time, CPU Time, Database Time, and so on. It not only provides information about SQL statements that are currently running but also SQL statements that completed, failed, or were terminated. The table consists of columns that include Status, Duration, SQL ID, and so on for monitored SQL statements.



SQL statements are monitored only if they have been running for at least five seconds or if they are run in parallel.

The **Status** column includes the following icons:

A spinning icon indicates that the SQL statement is running.



- A green check mark icon indicates that the SQL statement completed its execution during the specified time period.
- A red cross icon indicates that the SQL statement did not complete. The icon displays when an error occurs because the session was terminated.
- A clock icon indicates that the SQL statement is queued.

You can end session to terminate a running or gueued SQL statement.

For more information, see:

SQL Monitoring Report

Real-time SQL Monitoring

You can also click an **SQL ID** to go to the corresponding Real-time SQL Monitoring page. This page provides additional details that helps you tune the selected SQL statement.

For more information, see:

Real-time SQL Activity

Graphical Explain Plan

The Real-time SQL Monitoring page provides different options for viewing plan statistics for SQL execution, namely Tabular Execution Plan and Graphical Explain Plan.

- The default Tabular Execution Plan presents a tabular view of the execution plan.
- Graphical Explain Plan presents a graphical tree representation of the explain plan. In
 this plan, you can display as much or as little data as needed, and can rotate the tree to
 display the information either horizontally or vertically.

For more information, see:

Graphical Explain Plan

Automatic Database Diagnostic Monitor (ADDM)

Automatic Database Diagnostic Monitor (ADDM) analyzes the data in Automatic Workload Repository (AWR) and provides options to resolve performance problems.

As AWR is a repository of historical performance data, ADDM analyzes the performance problems after the diagnosis, thus saving time and resources required to reproduce the problem.

Benefits of ADDM

- Performs time-based quantification of application problems
- Identifies the root causes of performance problems
- Provides recommendations for fixing the problems
- Identifies the nonproblem areas of the application
- Helps you create and run ADDM tasks to analyze the current or past performance

Besides production systems, you can also use ADDM on development and test systems to find early warnings of application performance problems.



For more information about how to use ADDM, see:

ADDM Data

Workload

The Workload tab graphically displays four sets of statistics that you can use to monitor the database workload and identify spikes and bottlenecks.

Each set of statistics is displayed in a separate region, as described in the following sections.

Monitored and Analyzed Time Indications

The time slider has more functionality in the Workload tab than it does in the Active Session History and SQL Monitoring tabs. Note the following **Quick Select** time range options:

- Last Hour, Last 8 Hours, and Last 24 Hours The charts in the Workload tab display
 data for the entire time period of specified time range. A shadowed area is displayed in
 each chart that corresponds to the position of the time slider in the time range.
- Last Week The charts in the Workload tab display data for the selected time period of the time slider in the time range. A shadowed area is not displayed in this case.
- Custom The shadowed area display depends on whether the time period is up to and including 24 hours, or greater than 24 hours.

Regions

The tab contains four regions, namely CPU Statistics, Wait Time Statistics, Workload Profile, and Sessions.

Each region contains one or more charts that indicate the characteristics of the workload and the distribution of the resources. The data displayed on all the charts is for the same time period, as selected by the Time Range and time slider at the top of the window.

For more information about the Workload tab, its regions and the associated charts, see:

Workload Metrics

Blocking Sessions

The Blocking Sessions tab displays the current blocking and waiting sessions in a hierarchical display.

In this tab, you can:

- View detailed information about each blocking session, and view sessions blocked by each blocking session.
- Inspect or drill down into the SQL involved, to determine the cause of the blocking.
- Perform several troubleshooting operations, including terminating one or more listed sessions to resolve a waiting session problem.

The hierarchical display nests waiting sessions underneath the session that is blocking them in an easily viewable parent-child relationship. The hierarchy can contain any number of levels to correctly represent the structure of the sessions involved.

The sessions listed include sessions that are waiting for a resource, and sessions that hold a resource that is being waited on that creates a blocking condition.



For more information, see:

Blocking and Waiting Sessions

Exadata

The Exadata tab provides a unified view of Oracle Exadata hard disk and flash performance statistics.

It gives a deep insight into the health and performance of system components including the databases, Exadata storage cells, and Automatic Storage Management (ASM). You can use the Exadata tab for enhanced performance diagnostics for Exadata databases.

This tab is available for Exadata Databases managed through External Database and databases in Exadata Database Service on Dedicated Infrastructure deployments.



The Exadata tab is available only in the historical mode and supports only container databases. It does not provide real-time views, or any information about pluggable databases.

In the Exadata tab, you can:

- Analyze outliers that affect database performance, including finding a slow disk component that is affecting the system
- Analyze performance characteristics of multiple databases deployed in an Exadata system
- Identify a high I/O load and classify the load as Backup, Rebalance, User I/O, and others
- Identify CPU bound cells and determine cell offload
- Identify configurations such as versions and disk status

It provides various statistics about databases running on Exadata systems, including:

- Operating system statistics
- Storage server software statistics
- Smart scan statistics
- Database, disk, and cell statistics
- Exadata-specific ADDM recommendations
- Exadata system statistics

Exadata Details

The Exadata Details page contains nine tabs:

- Performance shows database performance, including a summary, latency statistics, flash I/O statistics, and Hard Disk I/O statistics
- **Health** shows the number and type of open alerts (Information, Warning, and Critical), the number of offline disks, open alert details, and offline disk details



 ADDM - displays Automatic Database Diagnostic Monitor (ADDM) performance findings specific to Exadata for the selected ADDM task, and makes specific recommendations to correct performance problems.

For more information, see:

- Automatic Database Diagnostic Monitor (ADDM)
- ADDM Data
- Top Consumers identifies workload distribution across all databases including the
 databases that consume a significant amount of I/O bandwidth. I/O requests and I/O
 throughput aggregates the data to report the top databases captured. The tab shows the
 number of requests, I/O speed, and usage of the top databases per cell in the Exadata
 system. It also displays details of IORM Queue Time and the top databases per cell.
- Cells and Disks compares the operational statistics of the flash drives and hard disks in the system.

It provides I/O related information, such as:

- how I/O is distributed between flash drives and hard disks
- realize the benefit the system is getting from Flash Cache
- how much disk I/O is caused by Flash Cache
- how I/O correlates to database single block reads

The tab displays statistics by cell, including the Top Cells, Outlier Cells, OS I/O Throughput, Cell Server I/O Requests, Cell Server I/O Throughput, % Disk Utilization, I/O Latency, and Cell Latency. You can select the default column group by which to display the data. Hover the cursor over a statistic to view details about it.

- Smart I/O displays summaries and details of Flash log and Flash Cache statistics by cell.
- **I/O Reasons** displays the number of requests and bytes for the set of selected cells to show the reasons for a high I/O load by category. Categories include backup, rebalance, and user I/O. The Requests and Bytes statistics bars for each cell are broken into sections that show the number of database control file reads per second, voting files I/O operations per second, and xrov I/O operations per second, and others.
- Configuration displays the configuration of the storage server model, the storage server version, and the objective of each cell.

For more information, see:

Exadata Statistics

Top Activity Lite

Top Activity Lite is a simplified yet highly efficient mode of Performance Hub in Database Management Service. In this mode, database administrators can monitor real-time performance of managed databases in a single view for quick visualization of issues.

The primary advantages of Top Activity Lite include streamlined UI options and quick response time for real-time monitoring while providing critical information required for performance diagnosis.



Simplified UI Options

In comparison to Performance Hub, the Top Activity Lite mode consists of only two tabs, **ASH Analytics** and **SQL Monitoring**. These tabs provide similar functionality as in Performance Hub. For more information, see ASH Analytics, SQL Monitoring.

Though the Top Activity Lite page appears similar to Performance Hub, the following distinguishes this mode as a separate view:

- Time Range The Time Range selector is not enabled, and the time range used for data analysis is fixed at the last hour.
- Auto Refresh The time interval to automatically retrieve new data from the database and update the current tab content.



This option is **Off** by default.

Average Active Sessions - A single chart combined with a time slider, which controls the
time period of the data shown in the activity tables. The chart always shows an hour of
ASH data. Every time the page refreshes, the chart is updated to show data for the last
hour.

As compared to Performance Hub, Top Activity Lite has fewer ASH dimensions available. However, you can use other functions, such as drill down a dimension, control sample resolution, change dimensions, and so on. For more information, see ASH Analytics.

Time slider - A time period selector on the Average Active Sessions chart with a duration
of five minutes. By default, the time slider is pinned to the right-side of the chart. When the
time slider is pinned, the activity tables are updated on each refresh to show data for the
last five minutes.

You can explicitly move the time slider to select a historical period. This unpins the time slider and updates the activity tables to show data for the selected historical period.

After it is unpinned, the time slider shifts left on the chart (covering the same historical period) until it reaches the left-side. Then the time slider is automatically pinned again to the right-side of the chart.

Activity tables - By default, these tables display the top SQL and user sessions for the
selected time period in the Average Active Sessions chart. You can open the details page
and view more information for each SQL or user session. You can also use the dimension
selector for an activity table to view top activity along a different ASH dimension.

You can monitor the Average Active Sessions chart to identify any variations or unusual data points, then use the activity tables to perform detailed analysis on a specific time range. You can use the time slider at any time to adjust the time period covered by the activity tables data.

- Data Source This option, by default, enables fetching data from Memory and AWR. To view data only from memory, you can change the data source to Memory.
- View Option The Top Activity Lite mode provides two options, the default bar chart and an area chart view.

For more information about how to use this mode, see:

Top Activity Lite Data



AWR Snapshot Collection

Oracle Database periodically collects statistical data or snapshots of its state, the workload, and the resources that is uses over a selected time.

It stores these snapshots in the Automatic Workload Repository (AWR) for problem detection and self-tuning purposes. You can use these snapshots to monitor and diagnose the database performance for a specific time period.

AWR compares the difference between snapshots to determine which SQL statements to capture based on the effect on the system load. This reduces the number of SQL statements that must be captured over time. After the snapshots are created, Automatic Database Diagnostic Monitor (ADDM) processes the data in the snapshots for performance analysis and resolution.

By default, Oracle Database automatically collects snapshots every hour and retains the statistics in AWR for eight days. However, you can use the **AWR Settings** option on the Performance Hub page and manually configure some settings for snapshot collection, such as:

- Enable and disable snapshot collection for the database.
- Set the interval for snapshots.

A snapshot interval determines the frequency at which snapshots are captured. A smaller interval means higher frequency, which increases the volume of data collected by AWR.

Specify the number of days to store the snapshots in AWR.

A retention period determines how long the statistical data is maintained before it gets deleted permanently. A longer retention period takes more space in AWR.

For more information about how to configure these settings, see:

AWR Settings

Reports

Performance Hub provides the following reports for your databases:

- Automatic Workload Repository (AWR)
- Active Session History (ASH)
- Performance Hub

Automatic Workload Repository (AWR)

Automatic Workload Repository (AWR) collects, processes, and maintains performance statistics for problem detection and self-tuning purposes. This data is both in memory and stored in the database.

From Performance Hub, you can generate and download a report of the gathered data.



AWR Report

An AWR report shows data captured between two points in time (or snapshots). Not only can you view diagnostic data of your database but you can also get instance level reports for Oracle Real Application Clusters (RAC) databases to view details about a specific instance.

AWR reports are divided into multiple sections. The content of the report contains the workload profile of the system for the selected range of snapshots. The HTML report includes links that you can use to navigate quickly between sections.

The statistics collected and processed by AWR include:

- Object statistics that determine both access and usage statistics of database segments
- Time model statistics based on time usage for activities, displayed in the V\$SYS TIME MODEL and V\$SESS TIME MODEL views
- Some of the system and session statistics collected in the V\$SYSSTAT and V\$SESSTAT views
- SQL statements that are producing the highest load on the system, based on criteria such as elapsed time and CPU time
- ASH statistics, representing the history of recent sessions activity

To generate and download an AWR report, see:

Automatic Workload Repository (AWR) Report

Active Session History (ASH)

Performance Hub includes the ability to generate, view, and download Active Session History (ASH) reports for your managed databases.

An ASH report contains data that helps you to:

- Triage transient performance problems that typically last for a few minutes
- Perform scoped or targeted performance analysis by various dimensions or their combinations, such as time, session, module, action, or SQL identifier.

Note:

When the Database Management set to is Full Management, the ASH report is available for:

- Databases running in Base Database Service and Exadata Database Service on Dedicated Infrastructure systems
- Oracle Databases managed with External Database

To generate and download an ASH report, see:

Active Sessions History (ASH) Report

Performance Hub

Performance Hub includes the ability to generate, view, and download active reports that contains information about the performance of your managed databases.



The report contains the Performance Hub UI backed by data collected and bundled with the report at the time you created it.

You can set the report to contain one of three levels of data:

- Basic includes only the tab contents
- Typical includes the tab contents with details for the top SQL statements
- · All includes the tab contents with details for all SQL statements

Note:

When the Database Management set to is Full Management, the Performance Hub report is available for:

- Databases running in Base Database Service and Exadata Database Service on Dedicated Infrastructure systems
- Oracle Databases managed with External Database

To generate and download a Performance Hub report, see:

Performance Hub Report

Using Performance Hub to Analyze Database Performance

Sections:

- Access
- View
- Manage
- Download

Access

You can navigate to the Performance Hub page in the Oracle Cloud Infrastructure Console for your managed databases.

To navigate to these database systems, see:

- For Autonomous Database
- For Base Database Service and Exadata Databases
- For External Database Service

For Autonomous Database

To access Performance Hub for Autonomous Database:

- Open the navigation menu and click Oracle Database.
- 2. Under Autonomous Database, select the type of database that you are managing:
 - Autonomous Data Warehouse



- Autonomous JSON Database
- Autonomous Transaction Processing
- 3. Choose your Compartment.
- 4. In the list of databases, click the display name of the database that you want to analyze using Performance Hub.
- 5. On the details page of the database, click **Performance Hub** to open the Performance Hub page.



To exit the Performance Hub page and to go back to the database details page, click **Close** at the bottom-left of the page.

To learn more about Autonomous Database, see:

Autonomous Database

For Base Database Service and Exadata Databases

To access Performance Hub for Base Database Service and Exadata databases:

- 1. Open the navigation menu and click **Oracle Database**.
- Select the type of database that you are managing:
 - Base Database Service (VM, BM)
 - Exadata
- Choose your Compartment.
- 4. Access the database details page for Base Database Service or Exadata databases.

For Base Database Service:

- a. Click a **Database System name**. The Database System page is displayed.
- b. In the list of databases, click a database name. The database details page is displayed.

For Exadata:

- a. Click the VM Cluster.
- **b.** Click the database name. The database details page is displayed.

Note:

The **Associated Services** section on the database details page shows the status of the Database Management service, whether the service is enabled for the database. It also shows which database management option (Basic or Full) is enabled.

- If Database Management is **Enabled**, you can click **Disable** to disable it.
- If Database Management is **Disabled**, you can click **Enable** to enable it.



Click **Edit** to display the Edit Database Management page. For information about changing database management settings, see About Management Options in Enable Database Management.

On the details page of the database, click **Performance Hub** to open the Performance Hub page.

Note:

Performance Hub is enabled only under the following conditions.

- The Database Management service must be enabled.
- The database must be an Enterprise Edition, version 12.1.0.0.0 or higher.

Note:

To exit the Performance Hub page and to go back to the database details page, click **Close** at the bottom-left of the page.

To learn more about Base Database Service, see:

Base Database Service

For External Database Service

To access Performance Hub for External Database Service:

- Open the navigation menu and click Oracle Database.
- 2. Click External Database.
- 3. Choose your Compartment.
- 4. Under **External databases**, select the type of external database you are managing:
 - Pluggable databases
 - Container databases
 - Non-container databases

Note:

The **Associated Services** section on the database details page shows the status of the Database Management service, whether the service is enabled for the database. It also shows which database management option (Basic or Full) is enabled.

- If Database Management is Enabled, you can click Disable to disable it.
- If Database Management is **Disabled**, you can click **Enable** to enable it.
- 5. On the details page of the database, click **Performance Hub** to open the Performance Hub page.



Performance Hub is enabled only under the following conditions.

- The Database Management service must be enabled.
- The database must be an Enterprise Edition, version 12.1.0.0.0 or higher.

Note:

To exit the Performance Hub page and to go back to the database details page, click **Close** at the bottom-left of the page.

To learn more about External Database, see:

External Database

View

You can view the performance of the databases you are managing from the Performance Hub page in the Oracle Cloud Infrastructure Console.

Performance Hub displays performance data through various tabs and UI options.

- Average Active Sessions Data
- SQL Monitoring Report
- Real-time SQL Activity
- Graphical Explain Plan
- Workload Metrics
- Blocking and Waiting Sessions
- ADDM Data
- Exadata Statistics
- Top Activity Lite Data

Average Active Sessions Data

To view Average Active Sessions Data by a selected dimension:

1. Open the Performance Hub page from the database details page, if not already open.



For more information on how to navigate to the Performance Hub page, see:

- For Autonomous Database
- For Base Database Service and Exadata Databases
- For External Database Service
- 2. You can adjust the time slider to select the time period for which performance data is displayed on the current tab.

The **ASH Analytics** tab displays the top activity for a selected dimension in the selected time period.

The **Quick Select** field allows you to quickly choose from a set of predetermined time ranges (for example, Last Hour, Last 8 Hours, and so on) or set a different custom time range.

Select a dimension in the ASH Dimension menu of the Average Active Sessions chart to display ASH analytics by that dimension.

Optionally, you can:

- Click the Maximum Threads check box to view the number of Max CPU Threads. The
 red line on the chart shows this limit.
- Click the Total Activity check box to view a black border that denotes total activity of all the components of the selected dimension on the chart. This option is selected by default when you use the filtering capabilities to only view the data for a particular component within a dimension.

For information about filtering capabilities, see Filter Average Active Sessions Data.

4. Use the Sample Resolution menu to select the sampling of ASH data displayed in the Average Active Sessions chart. The data resolution selections are Low, Medium, High, and Maximum.

For more information about sample resolutions, see ASH Analytics.

For the dimension selected in the ASH Dimension menu of the Average Active Sessions chart, you can further drill down into session details by selecting dimensions in the two tables at the bottom of the ASH Analytics tab.

By default, the following dimensions are selected for the two tables:

- SQL ID displays the SQL statements with the top average active sessions activity broken down by the main ASH dimension for the selected time period. You can rightclick the bar charts to sort the SQL statements in ascending or descending order or click the SQL ID to go the SQL Details page.
- User Session displays the user sessions with the top average active sessions
 activity broken down by the main ASH dimension for the selected time period. You can
 right-click the bar charts to sort the user sessions in ascending or descending order or
 click the user session to go to the User Session page.

Filter Average Active Sessions Data

To filter Average Active Sessions Data:

1. Open the Performance Hub page from the database details page, if not already open.

For more information on how to navigate to the Performance Hub page, see:

- For Autonomous Database
- For Base Database Service and Exadata Databases
- For External Database Service
- In the ASH Analytics tab, select a dimension in the ASH Dimension menu of the Average Active Sessions chart.



By default, **Consumer Group** is selected for Autonomous Databases. For all other databases, **Wait Class** is selected by default.

Each color in the chart denotes a component of the selected dimension. For example, the Consumer Group dimension has **High**, **Medium**, and **Low**, which are predefined service names assigned to your database to provide different levels of concurrency and performance.

3. Click a component in the legend.

The selected component is displayed in the **Applied Filters** field and the chart is updated to only display data pertaining to that component.

Note:

When ASH data is down sampled (maybe due to the selected data resolution), a new data request call is automatically triggered using the selected filters to ensure that precise data is displayed. Otherwise, filtering is performed on the UI layer only.

The total activity, which includes all the components of the dimension, is defined by a black outline and is displayed by default when you filter data.

SQL Monitoring Report

To view the SQL Monitoring report:

Open the Performance Hub page from the database details page, if not already open.



For more information on how to navigate to the Performance Hub page, see:

- For Autonomous Database
- For Base Database Service and Exadata Databases
- For External Database Service
- 2. Click **SQL Monitoring** to display the SQL monitoring tab.

Optionally, you can get detailed information about a specific SQL statement execution by clicking an ID number in the **SQL ID** column. When you click an ID number, it displays the Real-time SQL Monitoring page.

Click Save Report on the top-right corner to download the report for your selected SQL statement.

Real-time SQL Activity

To view Real-time SQL Activity:

1. Open the Performance Hub page from the database details page, if not already open.

Note:

For more information on how to navigate to the Performance Hub page, see:

- For Autonomous Database
- For Base Database Service and Exadata Databases
- For External Database Service
- Use Quick Select, Time Range, Time Zone, and time slider to set the time period of the SQL activity that you want to monitor.

You can view and monitor real-time SQL activity for databases in Oracle Cloud Infrastructure using the ASH Analytics tab or the SQL Monitoring tab.

View Real-time SQL Data from ASH Analytics

- On the Performance Hub page, click ASH Analytics to display the ASH Analytics tab, if not already open.
- 2. Scroll down to the table containing the list of SQL queries.
- 3. Click an ID number in the **SQL ID** column to display detailed information about a specific SQL statement. It opens the Performance Hub SQL Details page.

The Performance Hub SQL Details page contains the following tabs:

- Summary shows the overall information about the SQL statement and the SQL plans.
- ASH Analytics shows the active sessions and two tables. The default information shown for the SQL query is Wait Class and User Session. You can click the menu

containing the Wait Class option and select a specific parameter and subparameter to display the associated SQL data. Click a session in the **User Session** column to display the SQL activity during the selected user session.

• **Execution Statistics** - shows the SQL statement execution in Tabular Execution Plan or Graphical Explain Plan.

For more information, see Graphical Explain Plan.

SQL Monitoring - displays SQL data for the selected SQL statement. Click an SQL ID
to display detailed data about the SQL statement. Depending on the database
selected, the Details section includes Plan Statistics, Parallel, SQL Text, and Activity
tabs.

In the **Details** section of the page:

 Click Plan Statistics to view the Tabular Execution Plan and Graphical Explain Plan.

For more information, see Graphical Explain Plan.

- Click Parallel to view the Parallel Server and the Instances and Parallel Groups it contains.
- Click SQL Text to view the text of the active SQL query.
- Click Activity to view the active sessions of the database that is selected. Use the menu on the Details page to view either the Active Sessions Resource Type or Plan Line of the selected SQL operation.
- 4. Click the **Back** arrow at the top-left corner of the page to go back to the ASH Analytics tab.

View Real-time SQL Data from SQL Monitoring

- On the Performance Hub page, click SQL Monitoring to display the SQL monitoring tab.
- In the SQL ID column, click an ID to display detailed information about the specific SQL statement execution.

This action displays the Performance Hub Real-time SQL Monitoring page.

3. Scroll down to display the SQL activities.

Depending on the database selected, the Details section includes Plan Statistics, Parallel, SQL Text, and Activity tabs.

In the **Details** section of the page:

- Click Plan Statistics to view the Tabular Execution Plan and Graphical Explain Plan.
- Click Parallel to view the Parallel Server and the Instances and Parallel Groups it contains.
- Click SQL Text to view the text of the active SQL query.
- Click Activity to view the active sessions of the database that is selected.

Use the menu on the details page to view either the Active Sessions Resource Type or Plan Line of the selected SQL operation.

Click the Back arrow at the top-left corner of the page to go back to the SQL Monitoring tab.



Graphical Explain Plan

To view the Graphical Explain Plan:

1. Open the Performance Hub page from the database details page, if not already open.



For more information on how to navigate to the Performance Hub page, see:

- For Autonomous Database
- For Base Database Service and Exadata Databases
- For External Database Service
- 2. Click **SQL Monitoring** to display the SQL monitoring tab.
- In the SQL Monitoring tab, click an ID number in the SQL ID column to display detailed information on a specific SQL statement execution. It opens the Real-time SQL Monitoring page.

The Real-time SQL Monitoring page displays the **Plan Statistics** tab by default. The plan statistics for the SQL execution is displayed in the default **Tabular Execution Plan** view.

4. Click View Option and change it to Graphical Explain Plan.

The SQL explain plan is presented as a collapsible tree of nodes where each node represents an operation.

Using Graphical Explain Plan

By default, the Graphical Explain Plan displays the operation nodes horizontally, with the root node of the tree on the right-side of the screen. Sections of the tree within a node can be collapsed or expanded by clicking the minus or plus sign for the node.

- 1. You can zoom in or zoom out on the Graphical Explain Plan using the mouse or the keyboard +/- keys. You can also click and drag the plan within the viewport.
 - An overview section is located in the bottom-right corner of the content area. It provides a visual representation of the positioning of the viewport with respect to the overall content of the plan.
- 2. Click a **plus sign** to expand a node and display information for the operation subtree. Click a **minus sign** to collapse the operation subtree beneath a node.
- Hover over a node to view the operation name and estimated row count. Click a node to display additional information about the operation in a message box. Click Close to close the message box.
- 4. Click **Rotate** (to the right of the Graphical Explain menu) to rotate the entire explain plan 90 degrees counterclockwise.
 - Click **Rotate** subsequently to move the explain plan another 90 degrees counterclockwise. In any position of the tree, click the plus or minus sign to expand or contract the subtree of an operation node.
- Click the Back arrow at the top-left corner of the page to go back to the SQL Monitoring tab.



Workload Metrics

To view Workload metrics:

Open the Performance Hub page from the database details page, if not already open.



For more information on how to navigate to the Performance Hub page, see:

- For Autonomous Database
- For Base Database Service and Exadata Databases
- For External Database Service
- The Quick Select field allows you to quickly choose from a set of predetermined time ranges (for example, Last Hour, Last 8 Hours, and so on) or set a different custom time range.

You can adjust the time slider to select the time period for which performance data is displayed on the current tab.

- Click Workload to view the Workload tab. It displays four regions and their associated charts. All charts show data for the entire specified time range if within 24 hours.
 - The CPU Statistics region contains two charts, CPU Time and CPU Utilization (%).
 - Select CPU Time to display how much CPU time is being used by the foreground sessions per second. It identifies where the CPU time is mostly spent in the workload and pinpoints any unusual CPU spikes. When CPU time is selected optionally, click the Maximum Threads check box to show the maximum CPU time available. This shows the CPU time component of Average Active Sessions.
 - Select CPU Utilization (%) to display the CPU Utilization (%) chart. This chart
 indicates the percentage of CPU time aggregated by consumer group as
 calculated by the resource manager.
 - The Wait Time Statistics region contains a chart that displays the time used in different wait classes. To see the total average active sessions, select the DB Time check box. The activities are broken down by 13 wait classes.
 - The Workload Profile region contains a group of charts that indicate patterns of user calls, executions, transactions, and parses, as well as the number of running statements and queued statements. This region includes a menu that you can use to select the data to display.

Click the menu and select the metrics that you want to view in the Workload Profile.

- Select **User Calls** to display the combined number of logons, parses, and calls executed per second.
- Select Executions to display the combined number of user and recursive calls that executed SQL statements per second.
- Select **Transactions** to display the combined number of user commits and user rollbacks per second.
- Select Parses to display the combined number of hard and soft parses per second.



- Select Running Statements to display the number of running SQL statements, aggregated by consumer group.
- Select Queued Statements to display the number of queued parallel SQL statements, aggregated by consumer group.
- The Sessions region contains charts that show the number of current logons and sessions.

Click the menu and select the metric that you want to view in the Sessions region.

- Select Current Logons to display the number of current successful logons.
- Select **Sessions** to display the number of sessions.

For more information about the Workload tab, its regions and the associated charts, see:

Workload

Blocking and Waiting Sessions

To view blocking and waiting sessions:

1. Open the Performance Hub page from the database details page, if not already open.



For more information on how to navigate to the Performance Hub page, see:

- For Autonomous Database
- For Base Database Service and Exadata Databases
- For External Database Service
- The Time Range field displays a time period, which is the total amount of time available for analysis. You can adjust the time slider to select the time period in the Activity Summary.

For more information about the time selector fields, see Time Range.

Click Blocking Sessions to display details about current blocking and waiting sessions.



The Blocking Sessions option does not support analysis of historical sessions.

4. Click a link in each column to view details of the listed blocking and waiting sessions, as shown in the following table.

Tab Column	Description
User Name	The name of the user.
Status	Indicates whether the session is active, inactive, or expired.
Lock	The lock type for the session. Click lock type to display a table with more information about the session lock. It lists the Lock Type, Lock Mode, Lock Request, Object Type, Subobject Type, Time, ID1, ID2, Lock Object Address, and Lock Address of the selected session.



Tab Column	Description
User Session	Lists the Instance, SID, and Serial number.
SQL ID	The ID of the SQL associated with the session.
Wait Event	The wait event for the session. Click wait event to show additional wait event details.
Object Name	The name of the locked database object.
Blocking Time	The time that a blocking session has been blocking a session.
Wait Time	The time that a session is waiting.

Set the Minimum Wait Time

The minimum wait time works as a filter for the Blocking Sessions information. It sets the minimum time that a session must wait before it is displayed in the tab.

For example, if the minimum wait time is set to three seconds, and a session has waited only two seconds, then it is not displayed in the table. If you change the minimum wait time to one second, then the session that waited only two seconds is added to the display.



The minimum wait time default setting is three seconds.

Terminate a Session

- Click the check box at the left of the session User Name to select a session. The Kill Session button is enabled.
- 2. Click the **Kill Session** button. A confirmation dialog box is displayed for terminating the session.
- 3. Click the Kill Session button in the dialog box to end the session.

Display Lock Details

- In the session Lock column, click the name of the lock type (Lock or Exclusive Lock) for the session. The Wait Event Details message box is displayed.
- 2. You can note the information and take action, as needed.

Display Wait Event Information

- 1. In the session **Wait Event** column, click the name of the wait event for the selected session. The **Session Lock Information** table is displayed.
- 2. You can note the information and take action, as needed.

Display Session Details

 In the session User Session column, click Session Identifier for the session. The Performance Hub Session Details page is displayed.



Optionally move the time slider to display a specific time range of the session.

2. Use the Session Details page to explore additional details about the session.

Display SQL Details

1. In the session **SQL ID** column, click the SQL ID associated with the session. The Performance Hub SQL Details page is displayed.

Optionally move the time slider to display a specific time range of the session.

- Select one or more of the following tabs. You can note the information and take action, as needed.
 - Summary displays the SQL Overview and Source details
 - ASH Analytics displays the SQL average active sessions
 - Execution Statistics displays the SQL plans and plan details
 - SQL Monitoring displays information about monitored SQL executions
 - SQL Text displays the SQL

ADDM Data

To view the Automatic Database Diagnostic Monitor (ADDM) information:

1. Open the Performance Hub page from the database details page, if not already open.

Note:

For more information on how to navigate to the Performance Hub page, see:

- For Autonomous Database
- For Base Database Service and Exadata Databases
- For External Database Service
- 2. Click **ADDM** to open the ADDM tab.
- 3. Select a time range from the **Quick Select** field. The page displays data for the selected time period.
- **4.** In the **Activity Summary** area, click one of the gray AWR Snapshot icons to display findings for the associated ADDM task.

A white check mark in the gray icon indicates that there are problem findings available. When selected, the gray icon changes to blue.

For a wider time period, for example **Last Week**, there may be multiple ADDM tasks. The **Activity Summary** area displays icons indicating groups of tasks. Click an icon to open the list where you can select a particular ADDM task.



Alternatively, you can select a **Task Name** in the ADDM tab or by moving the time slider close to a snapshot icon.

Note:

When you manually change the ADDM task selection, either by clicking the gray icon for an associated AWR Snapshot, or by selecting an option from the ADDM task menu, the time slider position and size are adjusted to cover the analysis period for the ADDM task.

5. Hover over the icon to display a message about the AWR Snapshot and ADDM task, including the number of findings for the ADDM task.

The findings are displayed in two tables:

- Findings If there are findings, then this table shows the Name of the finding, Impact, Number of Recommendations, and Average Active Sessions for that finding. If there are no findings available, then the table displays a message that says there are no ADDM analysis available for the selected time period.
- Warnings and Information This table is displayed next to the Findings table.

It lists messages related to the findings.

- Warning messages identify issues such as missing data in the AWR that may affect the completeness or accuracy of the ADDM analysis.
- Information messages provide information that is relevant to understanding the
 performance of the database but does not represent a performance problem. This
 may include identification of nonproblem areas of the database and automatic
 database maintenance activity.

Note:

Both the Findings table and the Warnings and Information table are collapsible to save space when many findings are found. Click the minus icon (-) to collapse a table. Similarly, click the plus icon (+) to expand the table again.

- 6. If a finding has ADDM recommendations available, then the name of the finding is displayed as a link. Click the name of the finding to display more information about the finding, including a table of recommendations for corrective actions.
 - Each recommendation includes the problem area, the suggested action to take to solve it, and the estimated benefit that will result when the action is taken.
- Click the expand icon at the end of a row in the recommendations table to view a rationale for the recommendation.



Run an ADDM Task



This feature is supported for databases running on Base Database Service, External Databases, and Exadata Database Service on Dedicated Infrastructure, but not for Autonomous Database.

To set up and run an Automatic Database Diagnostic Monitor (ADDM) task:

1. Open the Performance Hub page from the database details page, if not already open.



For more information on how to navigate to the Performance Hub page, see:

- For Autonomous Database
- For Base Database Service and Exadata Databases
- For External Database Service
- Click the ADDM tab.
- Select a time range from the Quick Select field. The page displays data for the selected time period.
- 4. In the ADDM tab, click Run ADDM Task.
- 5. In the Run ADDM Task dialog box, select from the following options:
 - Capture and analyze current performance to create a current snapshot. This
 option is selected by default.
 - Analyze past performance to create a historical snapshot.

To create a current snapshot

- 1. Select Capture and analyze current performance.
- 2. Click Run ADDM.

The status of the task is shown in the top-right corner of the window.

See the section on how to view a completed ADDM task.

To create a historical snapshot

- 1. Click Analyze past performance.
- Select a Start Snapshot date and time and an End Snapshot date and time from the drop-down.
- 3. Click Run ADDM.

The status of the task is shown in the upper right corner of the window.

See the section on how to view a completed ADDM task.

Running multiple tasks at the same time is possible, but Oracle recommends that you run only one task at a time. This helps speed the processing of each task and presents the results sooner.

View a Completed ADDM Task

Note:

The time for an ADDM task to complete depends on the current system workload, the time length of the snapshot and the amount of data it contains. Performance Hub displays a message to check later if the job is taking a while to run.

When an ADDM task completes, the snapshot it created is displayed in the **Activity Summary** area. The time slider changes size to match the start and end times of the analyzed snapshot range and moves over the AWR snapshot icon in the graphical display.

 In the Activity Summary area, click the gray AWR snapshot icon to display the AWR snapshot. Hover over the icon to view the statistics of the snapshot, including the snap ID, snapshot times, task name, task owner, and the number of findings of the associated ADDM task.

A white check mark in the gray icon indicates that findings are available. When selected, the gray icon changes to blue.

Note:

If you have run multiple tasks concurrently, then an AWR icon for each snapshot appears under the activity summary graph and the time slider appears over the first task that you ran. To display a different snapshot, click the AWR snapshot icon for that task. The time slider moves over the AWR snapshot icon and expands or contracts to match the snapshot time.

- The task name, analysis period, and type are generated and displayed in the ADDM task list under the Activity Summary area.
- 3. The findings for the ADDM task are displayed in the Findings tables. To view a short description of the reason for the finding, scroll down to the Findings table. Then click the **information icon** at the end of the name of the finding you are interested in.

The message box explains the reason for the finding.

4. To view the task finding details, scroll down to the Findings table and click the name of the finding. It opens the **ADDM Performance Finding Details** page.

The page shows the name of the finding, the impact of the finding, the analysis period, the type of recommendation, the recommended action for the finding, and the impact of each recommendation.



Exadata Statistics

To view statistics of databases running on an Exadata system:

1. Open the Performance Hub page from the database details page, if not already opened.



For more information about how to navigate to the Performance Hub page, see For Base Database Service and Exadata Databases.

2. Click the **Exadata** tab. The Exadata summary page is displayed.

It includes the following sections:

- **Summary** displays open alerts and type, I/O requests per disk, cell server I/O throughput per disk, and disk utilization for the database you have selected.
- Latency displays the latency of OS and cell I/O, small reads and writes, and large reads and writes.
- Flash I/O System Total displays I/O requests, database throughput, and the maximum capacity of the flash drives used by the database.
- Hard Disk I/O System Total displays I/O requests, database throughput, and the maximum capacity of the hard disks used by the database.
- 3. Click **Exadata Details** to display the Exadata Details page.

The Exadata Details page contains the following tabs:

- Performance (displayed by default)
- Health
- ADDM
- Top Consumers
- Cells
- Disks
- Smart I/O
- I/O Reasons
- Configuration
- 4. Click a tab to view the statistics information in it.

For more information about each tab, see:

Exadata under Performance Hub Features

Top Activity Lite Data

To view the activities of your database in the Top Activity Lite mode of Performance Hub:

Open the Performance Hub page from the database details page, if not already open.



For more information on how to navigate to the Performance Hub page, see:

- For Autonomous Database
- For Base Database Service and Exadata Databases
- For External Database Service
- Click the Top Activity Lite button on the top-right corner to access the Top Activity Lite mode of Performance Hub.



The **Time Range** is fixed to the last hour.

- Select Auto Refresh and set the time interval to 15 Seconds. This option updates data in the current tab at the set interval.
 - Similarly, you can select the refresh rate as 30 Seconds or 1 Minute, or turn it off completely.
- In the ASH Analytics tab, move the time slider in the Average Active Sessions chart to a specific position of your choice.
 - The tables at the bottom of the page refreshes and displays the activities of your database for the selected time.
- In the activity tables, click an ID in the SQL ID column to view details for that particular SQL.
 - Click **Back** on the top-left corner to go back to the Top Activity Lite page.
- Click a session in the User Session column to view details for that session.
 - Click **Back** on the top-left corner to go back to the Top Activity Lite page.
- 7. Go to the **SQL Monitoring** tab to view information about SQL executions.

For more information about these tabs, see:

- Average Active Sessions Data
- SQL Monitoring Report

Manage

Performance Hub provides options to run administrative tasks for monitoring and managing your databases.

AWR Settings

AWR Settings

To view and modify AWR Settings for snapshot collection from Performance Hub:

- 1. Open the Performance Hub page from the database details page, if not already open.
- 2. Click the AWR Settings button in the top-right corner.



It opens the AWR Settings dialog box and displays the current settings for snapshot collection.



- Click the switch next to AWR Snapshot Collection to enable or disable snapshot collection for the database. When enabled, you can specify an interval for snapshot collection.
- 4. Set the **Snapshot Interval** frequency in minutes.



The default value for snapshot interval is 60 minutes.

- 5. You can use the Retain Forever check box to manage the retention period for snapshots as follows:
 - Select this check box to retain snapshots indefinitely.
 - Deselect this check box to specify a fixed retention period for snapshots.

Note:

The default value for snapshot retention period is 8 days.

Enter the number of days to maintain the snapshots in AWR before it gets deleted permanently.

6. Click **Save** to update the AWR settings, or click **Cancel** to discard any changes. It closes the AWR Settings dialog box.

The Performance Hub page displays a confirmation message at the top-right corner indicating that AWR snapshot settings are saved successfully. You can click **AWR Settings** again to open the AWR Settings dialog box for viewing or modifying the current settings for snapshot collection.

To reduce AWR space consumption, increase the snapshot interval and reduce the retention period.

Note:

Oracle recommends that you set the AWR retention period large enough to capture at least one complete workload cycle.

Download

Performance Hub provides options to download reports for your managed databases.

For more information, see:



- Automatic Workload Repository (AWR) Report
- Active Sessions History (ASH) Report
- Performance Hub Report

Automatic Workload Repository (AWR) Report

To download an AWR report, identify the database version and follow the respective steps.

Open the Performance Hub page from the database details page, if not already open.

Note:

For more information on how to navigate to the Performance Hub page, see:

- For Autonomous Database
- For Base Database Service and Exadata Databases
- For External Database Service

For Oracle Database 18c and older

1. Click the **Reports** button in the upper-right corner and select **Automatic Workload Repository**. It opens the Generate Automatic Workload Repository dialog box.

You can generate a report either from two snapshots close to the current date and time or from a custom time range of your choice.

To generate a report from a custom time range, select **Custom** and then select the start and end time for the range.

Click **Download** to download the AWR report.

Oracle Database generates a report in the format AWRReport_date_range.html for the selected time range.

For Oracle Database 19c and later

- 1. In the Quick Select menu, select a time period for the AWR report.
- 2. Click the **Reports** button in the upper-right corner and select **Automatic Workload Repository**. It opens the Generate Automatic Workload Repository dialog box.
- Select the begin time of the snapshot from Start Snapshot and the end time from End Snapshot.

You can also get information about a specific instance and monitor its performance. This option is not enabled by default. To include details of a particular instance in the AWR report, click the check box **Instance Level Report** and select the instance from the dropdown list.





4. Click **Download** to generate the report and download it.

Oracle Database generates the report in the format <code>AWRReport_date_range.html</code> for the selected time range. When the report is complete, the report name is displayed at the top of the screen, and the report is downloaded automatically.

Active Sessions History (ASH) Report

To download an ASH report:

Open the Performance Hub page from the database details page, if not already open.



For more information on how to navigate to the Performance Hub page, see:

- For Autonomous Database
- For Base Database Service and Exadata Databases
- For External Database Service
- From the menu in the upper-right corner, click Reports and select Active Session History. The Generate ASH Report dialog box is displayed.
- Select the date of the Start Time. In the calendar and time dialog boxes that are displayed, select the start date and time for the report.
 - Repeat this step for the **End Time** date and time.
- 4. Click **Download** to generate the report for the selected time range and download it.

While the report is being generated, a message Report generation is in progress appears at the upper-right corner of the screen. When the report is complete and downloaded, a confirmation message with the name of the report file is displayed.

Performance Hub Report

To download a Performance Hub report:

1. Open the Performance Hub page from the database details page, if not already open.



For more information on how to navigate to the Performance Hub page, see:

- For Autonomous Database
- For Base Database Service and Exadata Databases
- For External Database Service
- 2. From the menu in the upper right corner, click Reports and select Performance Hub.
- 3. In the Generate Performance Hub Report dialog box dialog box, select the type of report:
 - Basic saves tab contents but no further details



- Typical saves tab contents with details for top SQL statements
- All saves tab contents with details for all SQL statements
- 4. Click **Download** to generate the report for the selected time range and download it.

While the report is being generated, a message Report generation is in progress appears at the upper-right corner of the screen. When the report is complete and downloaded, a confirmation message with the name of the report file is displayed.



Part III

Tuning Database Memory

This part contains the following chapters:

- Database Memory Allocation
- Tuning the System Global Area
- Tuning the Database Buffer Cache
- Tuning the Shared Pool and the Large Pool
- Tuning the Result Cache
- Tuning the Program Global Area



9

Database Memory Allocation

This chapter describes memory allocation in Oracle Database and the various methods for managing memory.

This chapter contains the following topics:

- About Database Memory Caches and Other Memory Structures
- Database Memory Management Methods
- Using Automatic Memory Management
- Monitoring Memory Management

About Database Memory Caches and Other Memory Structures

Oracle Database stores information in memory caches and on disk. Memory access is much faster than disk access. Disk access (physical I/O) takes a significant amount of time, compared to memory access, typically in the order of 10 milliseconds. Physical I/O also increases the CPU resources required due to the path length in device drivers and operating system event schedulers. For this reason, it is more efficient for data requests of frequently accessed objects to be perform by memory, rather than also requiring disk access. Proper sizing and effective use of Oracle Database memory caches greatly improves database performance.

The main Oracle Database memory caches that affect performance include:

Database buffer cache

The database buffer cache stores data blocks read from disk.

Redo log buffer

The redo log buffer stores redo entries of changes made to data blocks in the buffer cache.

Shared pool

The shared pool caches many different types of data and is mainly comprised of the following components:

- Library cache
- Data dictionary cache
- Server result cache
- Large pool

The large pool provides large memory allocations for the following Oracle Database features:

- Shared server architecture
- Parallel query
- Recovery Manager (RMAN)
- Java pool

The Java pool stores session-specific Java code and Java Virtual Machine (JVM) data.

Streams pool

The Streams pool provides memory for Oracle Advanced Queuing (AQ) and replication processes.

Process-private memory

Process-private memory includes memory used for operations such as sorting and hash joins.

In-Memory Column Store (IM column store)

The IM column store is an optional, static SGA pool that stores copies of tables and partitions. In the IM column store, data is stored in a special columnar format, which improves performance of operations such as scans, joins, and aggregations.



The IM column store does not replace the buffer cache, but acts as a supplement so that both memory areas can store the same data in different formats.

• Unified memory allows better usage of memory. Use memory_size rather than sga_target and pga aggregate target/limit.

See Also:

Oracle Database Concepts for information about the Oracle Database memory architecture

Database Memory Management Methods

The goal of memory management is to reduce the physical I/O overhead as much as possible, either by making it more likely that the required data is in memory, or by making the process of retrieving the required data more efficient. To achieve this goal, proper sizing and effective use of Oracle Database memory caches is essential.

Figure 9-1shows a decision tree for memory management options. The following sections explain the options in detail.

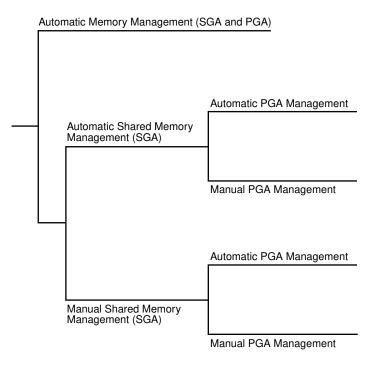


Figure 9-1 Memory Management Methods

- Automatic Memory Management
- Automatic Shared Memory Management
- Manual Shared Memory Management
- Automatic PGA Memory Management
- Manual PGA Memory Management
- Summary of Memory Management Methods

Automatic Memory Management

In **automatic memory management**, Oracle Database manages the shared global area (SGA) and program global area (instance PGA) memory completely. This method is the simplest and is strongly recommended by Oracle.

Before setting any memory pool sizes manually, strongly consider using automatic memory management.

The only user-specified controls are the target memory size initialization parameter (MEMORY_TARGET) and optional maximum memory size initialization parameter (MEMORY_MAX_TARGET). Oracle Database tunes to the target memory size, redistributing memory as needed between the SGA and the instance PGA.

The SGA can contain the In-Memory Column Store (IM column store) or memoptimize pool, both of which are optional. No matter which memory management method you use, size the IM column store with the INMEMORY_SIZE initialization parameter, and the memoptimize pool with the MEMOPTIMIZE_POOL_SIZE initialization parameter. The IM column store size are memoptimize pool are included in the memory target, but are not managed by the automatic resize algorithm. For example, if you set MEMORY_TARGET to 5 GB and INMEMORY_SIZE to 1 GB, then the overall memory target is 5 GB (not 6 GB), and the INMEMORY_SIZE is always 1 GB.



The following graphic shows a database that sometimes processes jobs submitted by online users and sometimes batch jobs. Using automatic memory management, the database automatically adjusts the size of the large pool and database buffer cache depending on which type of jobs are running. Note that the graphic does not depict the IM column store or memoptimize area.

Buffer Cache

Large Pool
Shared Pool
Java Pool
Streams Pool

Online Users

Buffer Cache

Large Pool
Shared Pool
Streams Pool

Batch Jobs

Figure 9-2 Automatic Memory Management

If you create a database with DBCA and choose the basic installation option, then Oracle Database enables automatic memory management by default.

See Also:

- Using Automatic Memory Management
- Oracle Database Administrator's Guide to learn about automatic memory management

Automatic Shared Memory Management

Oracle Database uses automatic shared memory management of the SGA memory when automatic memory management is disabled.

In this mode, Oracle Database automatically distributes memory to individual SGA components based on a target size that you set for the total SGA memory.

The database tunes the total SGA to the target size and dynamically tunes the sizes of SGA components. If you are using a server parameter file, then Oracle Database remembers the sizes of the automatically tuned components across database instance shutdowns.

For information about using automatic shared memory management, see "Using Automatic Shared Memory Management".

Manual Shared Memory Management

If both automatic memory management and automatic shared memory management are disabled, then you must manage SGA memory manually by sizing the individual memory pools in the SGA. Although this mode enables you to exercise complete control over how SGA memory is distributed, it requires the most effort because the SGA components must be manually tuned on an ongoing basis.



When automatic memory management is disabled, then in some cases the database can automatically adjust the relative sizes of the shared pool and buffer cache, based on user workload.

For information about using manual shared memory management, see "Sizing the SGA Components Manually".

Automatic PGA Memory Management

When automatic memory management (MEMORY_TARGET) is disabled and PGA_AGGREGATE_TARGET is set to a nonzero value, then Oracle Database uses automatic PGA memory management. In this mode, PGA_AGGREGATE_TARGET specifies a "soft" target size for the instance PGA. The target is soft because it only applies to specific types of memory allocations that can choose to use temporary space rather than the PGA. The database tunes the size of the instance PGA to this target and dynamically tunes the sizes of individual PGAs. If you do not explicitly set a target size, then the database automatically configures a reasonable default.

The PGA_AGGREGATE_LIMIT initialization parameter dynamically sets an instance-wide hard limit for PGA memory. Because the parameter responds to changing memory conditions, you do not need to set the parameter value explicitly. By default, PGA_AGGREGATE_LIMIT is set to the greater of the following:

- 2 GB
- 200% of PGA AGGREGATE TARGET initialization parameter setting
- (Value of PROCESSES initialization parameter setting) * 3 MB

A background process periodically compares the PGA size to the limit set by PGA_AGGREGATE_LIMIT. If the limit is reached or exceeded, then this process terminates calls for the sessions using the most untunable PGA memory. If these sessions still do not release enough memory, then they are also terminated.

For information about automatic PGA memory management, see Tuning the Program Global Area .

Manual PGA Memory Management

When automatic memory management is disabled and PGA_AGGREGATE_TARGET is set to 0 (automatic PGA memory management is disabled), the database defaults to manual PGA management. You must manage PGA memory manually by adjusting the portion of PGA memory dedicated to each work area.

This method can be very difficult because the workload is always changing and is not recommended by Oracle. Although manual PGA memory management is supported by Oracle Database, Oracle strongly recommends using automatic memory management or automatic PGA memory management instead.

Summary of Memory Management Methods

Memory management is either automatic or manual.

If you do not enable automatic memory management, then you must separately configure one memory management method for the SGA and one for the PGA.



When automatic memory management is disabled for the database instance as a whole, Oracle Database enables automatic PGA memory management by default.

The following table includes the <code>INMEMORY_SIZE</code> initialization parameter and the <code>MEMOPTIMIZE_POOL_SIZE</code> initialization parameter. Both control memory areas that are optional.

Table 9-1 Memory Management Methods

Instance	SGA	PGA	Description	Initialization Parameters
Auto	n/a	n/a	The database tunes the size of the instance based on a single instance target size.	You set: Total memory target size for the database instance (MEMORY_TARGET) Optional maximum memory size for the database instance (MEMORY_MAX_TARGET) Optional size for the IM column store (INMEMORY_SIZE) in the SGA
n/a	Auto	Auto	The database automatically tunes the SGA based on an SGA target. The database automatically tunes the PGA based on a PGA target.	You set: SGA target size (SGA_TARGET) Optional SGA maximum size (SGA_MAX_SIZE) Optional size for the IM column store (INMEMORY_SIZE) in the SGA Optional size for the memoptimize pool in the SGA (MEMOPTIMIZE_POOL_SIZE) PGA aggregate target size (PGA_AGGREGATE_TARGET)¹ The database automatically configures the PGA_AGGREGATE_LIMIT initialization parameter. You may set this parameter manually.



Table 9-1 (Cont.) Memory Management Methods

Instance	SGA	PGA	Description	Initialization Parameters
n/a	Auto	Manual	The database automatically tunes the SGA based on an SGA target. You control the PGA manually, setting the maximum work area size for each type of SQL operator.	You set: SGA target size (SGA_TARGET) Optional SGA maximum size (SGA_MAX_SIZE) Optional size for the IM column store in the SGA (INMEMORY_SIZE) Optional size for the memoptimize pool in the SGA (MEMOPTIMIZE_POOL_SIZE) PGA work area parameters such as SORT_AREA_SIZE, HASH_AREA_SIZE, and BITMAP MERGE AREA SIZE
n/a	Manual	Auto	You control the SGA manually by setting individual component sizes. The database automatically tunes the PGA based on a PGA target.	You set: Shared pool size (SHARED_POOL_SIZE) Buffer cache size (DB_CACHE_SIZE) Large pool size (LARGE_POOL_SIZE) Java pool size (JAVA_POOL_SIZE) Optional size for the IM column store (INMEMORY_SIZE) in the SGA Optional size for the memoptimize pool in the SGA (MEMOPTIMIZE_POOL_SIZE) PGA aggregate target size (PGA_AGGREGATE_TARGET) ² The database automatically configures the PGA_AGGREGATE_LIMIT initialization parameter. You may set this parameter manually.
n/a	Manual	Manual	You must manually configure SGA component sizes. You control the PGA manually, setting the maximum work area size for each type of SQL operator.	You must manually configure SGA component sizes. You set: Shared pool size (SHARED_POOL_SIZE) Buffer cache size (DB_CACHE_SIZE) Large pool size (LARGE_POOL_SIZE) Java pool size (JAVA_POOL_SIZE) Optional size for the IM column store (INMEMORY_SIZE) in the SGA Optional size for the memoptimize pool in the SGA (MEMOPTIMIZE_POOL_SIZE) PGA work area parameters such as SORT_AREA_SIZE, HASH_AREA_SIZE, and BITMAP_MERGE_AREA_SIZE

¹ The database automatically configures the PGA_AGGREGATE_LIMIT initialization parameter. You also choose to set this parameter manually.

✓ See Also:

Oracle Database Administrator's Guide because automatic memory management is not available on all platforms

² The database automatically configures the PGA_AGGREGATE_LIMIT initialization parameter. You also choose to set this parameter manually.

Using Automatic Memory Management

To use automatic memory management, set the following initialization parameters:

MEMORY TARGET

The MEMORY_TARGET initialization parameter specifies the target memory size. The database tunes to the value specified for this parameter, redistributing memory as needed between the SGA and the instance PGA. This parameter is dynamic, so its value can be changed at any time without restarting the database.

MEMORY MAX TARGET

The MEMORY_MAX_TARGET initialization parameter specifies the maximum memory size. The value specified for this parameter serves as the limit to which the MEMORY_TARGET initialization parameter can be set. This parameter is static, so its value cannot be changed after instance startup.

If you need tuning advice for the MEMORY_TARGET parameter, then use the V\$MEMORY TARGET ADVICE view.



Oracle Database Administrator's Guide for information about using automatic memory management

Monitoring Memory Management

Table 9-2 lists the views that provide information about memory resize operations.

Table 9-2 Memory Management Views

View	Description
V\$MEMORY_CURRENT_RESIZE_OPS	Displays information about memory resize operations (both automatic and manual) that are currently in progress.
V\$MEMORY_DYNAMIC_COMPONENTS	Displays information about the current sizes of all dynamically-tuned memory components, including the total sizes of the SGA and instance PGA.
V\$MEMORY_RESIZE_OPS	Displays information about the last 800 completed memory resize operations (both automatic and manual). This does not include operations that are currently in progress.
V\$MEMORY_TARGET_ADVICE	Displays tuning advice for the MEMORY_TARGET initialization parameter.

See Also

Oracle Database Reference for more information about these views



10

Tuning the System Global Area

This chapter describes how to tune the System Global Area (SGA). If you are using automatic memory management to manage the database memory on your system, then there is no need to tune the SGA as described in this chapter.

This chapter contains the following topics:

- Using Unified Memory
- Using Automatic Shared Memory Management
- Sizing the SGA Components Manually
- Monitoring Shared Memory Management
- Improving Query Performance with the In-Memory Column Store
- Enabling High Performance Data Streaming with the Memoptimized Rowstore

Using Unified Memory

Unified Memory is a flexible and simple memory configuration for Oracle Database 23ai that uses a single parameter to control database memory allocations. It reduces or eliminates the need for system restart to change memory configurations. Unified Memory is especially useful in multiple workload high density database consolidation environments.

Unified Memory simplifies memory management to run multiple workloads in a highly consolidated environment with minimum disruption. It is easier to set the single parameter MEMORY_SIZE for configuring the database instance memory instead of using separate parameters like SGA_TARGET and PGA_AGGREGATE_LIMIT.

Using Automatic Shared Memory Management

Automatic shared memory management simplifies the configuration of the SGA by automatically distributing the memory in the SGA for the following memory pools:

- Database buffer cache (default pool)
- Shared pool
- Large pool
- Java pool
- Streams pool

Automatic shared memory management is controlled by the SGA_TARGET parameter. Changes in the value of the SGA_TARGET parameter automatically resize these memory pools. If these memory pools are set to nonzero values, then automatic shared memory management uses these values as minimum levels. Oracle recommends that you set the minimum values based on the minimum amount of memory an application component requires to function properly.

The following memory caches are manually-sized components and are not controlled by automatic shared memory management:

Redo log buffer

The redo log buffer is sized using the LOG_BUFFER initialization parameter, as described in "Configuring the Redo Log Buffer".

Other buffer caches (such as KEEP, RECYCLE, and other nondefault block size)

The KEEP pool is sized using the DB_KEEP_CACHE_SIZE initialization parameter, as described in "Configuring the KEEP Pool".

The RECYCLE pool is sized using the DB_RECYCLE_CACHE_SIZE initialization parameter, as described in "Configuring the RECYCLE Pool".

Fixed SGA and other internal allocations

Fixed SGA and other internal allocations are sized using the DB_nK_CACHE_SIZE initialization parameter.

The memory allocated to these memory caches is deducted from the value of the SGA_TARGET parameter when automatic shared memory management computes the values of the automatically-tuned memory pools.

The following sections describe how to access and set the value of the SGA TARGET parameter:

- User Interfaces for Setting the SGA_TARGET Parameter
- Setting the SGA_TARGET Parameter

See Also:

- Oracle Database Concepts for information about the SGA
- Oracle Database Administrator's Guide for information about managing the SGA
- Oracle Database Administrator's Guide for information about using initialization parameters

User Interfaces for Setting the SGA TARGET Parameter

This section describes the user interfaces for setting the value of the SGA TARGET parameter.

This section contains the following topics:

- Setting the SGA_TARGET Parameter in Oracle Enterprise Manager Cloud Control
- Setting the SGA TARGET Parameter in the Command-Line Interface

Setting the SGA_TARGET Parameter in Oracle Enterprise Manager Cloud Control

You can change the value of the SGA_TARGET parameter in Oracle Enterprise Manager Cloud Control (Cloud Control) by accessing the SGA Size Advisor from the Memory Parameters SGA page.

Setting the SGA_TARGET Parameter in the Command-Line Interface

You can change the value of the SGA_TARGET parameter in the command-line interface by querying the V\$SGA_TARGET_ADVICE view and using the ALTER SYSTEM command.

Setting the SGA_TARGET Parameter

This section describes how to enable and disable automatic shared memory management by setting the value of the SGA TARGET parameter.

This section contains the following topics:

- Enabling Automatic Shared Memory Management
- Disabling Automatic Shared Memory Management

Enabling Automatic Shared Memory Management

To enable automatic shared memory management, set the following initialization parameters:

- STATISTICS LEVEL to TYPICAL or ALL
- SGA TARGET to a nonzero value

The SGA_TARGET parameter can be set to a value that is less than or equal to the value of the SGA_MAX_SIZE initialization parameter. Set the value of the SGA_TARGET parameter to the amount of memory that you intend to dedicate to the SGA.

Disabling Automatic Shared Memory Management

To disable automatic shared memory management, set the value of the SGA_TARGET parameter dynamically to 0 at instance startup.

This disables automatic shared memory management and the current auto-tuned sizes will be used for each memory pool. If necessary, you can manually resize each memory pool, as described in "Sizing the SGA Components Manually".

Unified Program Global Area

The unified program global area (PGA) pool is a shared global area (SGA) component that is used for PGA work areas by certain pluggable databases (PDBs) that have smaller SGA target values per-CPU compared to the multi-tenant container database (CDB).

The PGA allows the replacement of an Autonomous Transaction Processing (ATP) pluggable database (PDB) with an Autonomous Data Warehouse (ADW) PDB. This is achieved by transparently allocating some PGA for the ADW PDB from the SGA.

In an ATP-D environment, the memory is split between SGA and PGA usage. The SGA is supported by the large pages that are typically reserved during the virtual machine or host start-up time or in the early stage before the memory is fragmented. The SGA portion of system memory is configured based on transaction process requirements, typically 65% for SGA and 35% for PGA. However, ADW pluggable databases require more private memory and are typically in the range of 35% for SGA and 65% for PGA. Unfortunately, because the system is configured for ATP, the memory from huge pages (configured for SGA) cannot be released for PGA needs. It is difficult to reclaim the large pages in the future once it is fragmented.

The unified PGA pool is a new construct that allows the shared global area (SGA) to be used for pluggable databases' (PDB) program global area (PGA). The unified PGA only requests full granules from the buffer cache when it needs to grow, and it gives up full granules when no longer needed. At instance start-up, the PGA can have a minimum size of 0 or the size defined by an underscore size parameter.



Note:

Because only full granules can be consumed by the unified PGA pool, care is taken as to how often the granules are requested for growth and how soon they are given back to the system list.

Sizing the SGA Components Manually

If the system is not using automatic memory management or automatic shared memory management, then you must manually configure the sizes of the following SGA components:

Database buffer cache

The database buffer cache is sized using the <code>DB_CACHE_SIZE</code> initialization parameter, as described in "Configuring the Database Buffer Cache".

Shared pool

The shared pool is sized using the SHARED_POOL_SIZE initialization parameter, as described in "Configuring the Shared Pool".

Large pool

The large pool is sized using the LARGE_POOL_SIZE initialization parameter, as described in "Configuring the Large Pool".

Java pool

The Java pool is sized using the JAVA POOL SIZE initialization parameter.

Streams pool

The Streams pool is sized using the STREAMS POOL SIZE initialization parameter.

IM column store

The IM column store is sized using the INMEMORY SIZE initialization parameter.

The values for these parameters are also dynamically configurable using the ALTER SYSTEM statement.

Before configuring the sizes of these SGA components, take the following considerations into account:

- SGA Sizing Unit
- Maximum Size of the SGA
- Application Considerations
- Operating System Memory Use
- Iteration During Configuration



See Also:

- Oracle Database Java Developer's Guide for information about Java memory usage and the JAVA POOL SIZE initialization parameter
- Oracle Database In-Memory Guide for information about the INMEMORY_SIZE initialization parameter

SGA Sizing Unit

Memory for the buffer cache, shared pool, large pool, and Java pool is allocated in units of granules. If the SGA size is less than 1 GB, then the granule size is 4MB. If the SGA size is greater than 1 GB, the granule size changes to 16MB. The granule size is calculated and fixed when the database instance starts up. The size does not change during the lifetime of the instance.

To view the granule size that is currently being used for the SGA, use the V\$SGA_DYNAMIC_COMPONENTS view. The same granule size is used for all dynamic components in the SGA.

Maximum Size of the SGA

The maximum amount of memory usable by the database instance is determined at instance startup by the value of the SGA_MAX_SIZE initialization parameter. You can expand the total SGA size to a value equal to the SGA_MAX_SIZE parameter. The value of the SGA_MAX_SIZE parameter defaults to the aggregate setting of all the SGA components.

If the value of the SGA_MAX_SIZE parameter is not set, then decrease the size of one cache and reallocate that memory to another cache if necessary. Alternatively, you can set the value of the SGA_MAX_SIZE parameter to be larger than the sum of all of the SGA components, such as the buffer cache and the shared pool. Doing so enables you to dynamically increase a cache size without having to decrease the size of another cache.



The value of the SGA MAX SIZE parameter cannot be dynamically resized.

Application Considerations

When configuring memory, size the memory caches appropriately based on the application's needs. Conversely, tuning the application's use of the memory caches can greatly reduce resource requirements. Efficient use of the memory caches also reduces the load on related resources, such as latches, CPU, and the I/O system.

For optimal performance, consider the following:

- Design the cache to use the operating system and database resources in the most efficient manner.
- Allocate memory to Oracle Database memory structures to best reflect the needs of the application.

- If changes or additions are made to an existing application, resize Oracle Database memory structures to meet the needs of the modified application.
- If the application uses Java, investigate whether the default configuration for the Java pool needs to be modified.



Oracle Database Java Developer's Guide for information about Java memory usage

Operating System Memory Use

For most operating systems, it is important to consider the following when configuring memory:

- Reduce Paging
- Fit the SGA into Main Memory
- Allow Adequate Memory to Individual Users

See Also:

Your operating system hardware and software documentation, and the Oracle documentation specific to your operating system, for more information on tuning operating system memory usage

Reduce Paging

Paging occurs when an operating system transfers memory-resident pages to disk solely to load new pages into memory. Many operating systems page to accommodate large amounts of information that do not fit into real memory. On most operating systems, paging reduces performance.

To determine whether significant paging is occurring on the host system, use operating system utilities to examine the operating system. If significant paging is occurring, then the total system memory may not be large enough to hold the memory caches for which memory is allocated. Consider either increasing the total memory on the system, or decreasing the amount of memory allocated.

Fit the SGA into Main Memory

Because the purpose of the SGA is to store data in memory for fast access, the SGA should reside in the main memory. If pages of the SGA are swapped to disk, then the data is no longer quickly accessible. On most operating systems, the disadvantage of paging significantly outweighs the advantage of a large SGA.

This section contains the following topics:

- Viewing SGA Memory Allocation
- Locking the SGA into Physical Memory



Viewing SGA Memory Allocation

To view how much memory is allocated to the SGA and each of its internal structures, use the SHOW SGA statement in SQL*Plus as shown in the following example:

SOL> SHOW SGA

The output of this statement might look like the following:

Total System Global Area 840205000 bytes Fixed Size 279240 bytes Variable Size 520093696 bytes Database Buffers 318767104 bytes Redo Buffers 1064960 bytes

Locking the SGA into Physical Memory

To prevent the SGA from being paged out, consider locking the SGA into physical memory by enabling the LOCK_SGA parameter. The database does not use the MEMORY_TARGET and MEMORY_MAX_TARGET parameters when the LOCK_SGA parameter is enabled.

Allow Adequate Memory to Individual Users

When sizing the SGA, ensure that you allow enough memory for the individual server processes and any other programs running on the system.

Iteration During Configuration

Configuring memory allocation involves distributing available memory to Oracle Database memory structures, depending on the needs of the application. The distribution of memory to Oracle Database structures can affect the amount of physical I/O necessary for Oracle Database to operate properly. Having a proper initial memory configuration provides an indication of whether the I/O system is effectively configured.

After the initial pass through the memory configuration process, it may be necessary to repeat the steps of memory allocation. Subsequent passes enable you to make adjustments to earlier steps, based on changes in subsequent steps. For example, decreasing the size of the buffer cache enables you to increase the size of another memory structure, such as the shared pool.

Monitoring Shared Memory Management

Table 10-1 lists the views that provide information about SGA resize operations.

Table 10-1 Shared Memory Management Views

View	Description
V\$SGA_CURRENT_RESIZE_OPS	Displays information about SGA resize operations that are currently in progress.
V\$SGA_RESIZE_OPS	Displays information about the last 800 completed SGA resize operations. This does not include operations that are currently in progress.



Table 10-1 (Cont.) Shared Memory Management Views

View	Description
V\$SGA_DYNAMIC_COMPONENTS	Displays information about the dynamic components in the SGA. This view summarizes information of all completed SGA resize operations that occurred after instance startup.
V\$SGA_DYNAMIC_FREE_MEMORY	Displays information about the amount of SGA memory available for future dynamic SGA resize operations.

See Also:

Oracle Database Reference for information about these views

Improving Query Performance with the In-Memory Column Store

The In-Memory Column Store (IM column store) is an optional portion of the system global area (SGA) that stores copies of tables, partitions, and other database objects in columnar format, and this columnar data is optimized for rapid scans. Because the IM column store puts database objects in memory, Oracle Database can perform scans, queries, joins, and aggregates on that data much faster than on data stored in row format.

Note:

 The IM column store and database buffer cache store the same data, but in different formats. The IM column store does not replace the row-based storage in the database buffer cache, but supplements it for achieving better query performance.

In-memory scans are also permitted when not all columns in a table have been populated into the IM column store. This situation can occur when columns have been specified as ${\tt NO}$ INMEMORY to save space.

In-memory hybrid scans can access some data from the IM column store, and some data from the row store, improving performance by orders of magnitude over pure row store queries. The query is divided into two parts, with one part scanning the IM column store to perform filters, and the other part scanning the row store to project the filtered query results.

See Also:

Oracle Database In-Memory Guide for more information about the IM column store



Enabling High Performance Data Streaming with the Memoptimized Rowstore

The Memoptimized Rowstore enables high performance data streaming for applications, such as Internet of Things (IoT).

This section contains the following topics:

- About the Memoptimized Rowstore
- Using Fast Ingest
- Using Fast Lookup

About the Memoptimized Rowstore

The Memoptimized Rowstore enables high performance data streaming for applications, such as Internet of Things (IoT) applications, which typically stream small amounts of data in single-row inserts from a large number of clients simultaneously and also query data for clients at a very high frequency.

The Memoptimized Rowstore provides the following functionality:

Fast ingest

Fast ingest optimizes the processing of high-frequency, single-row data inserts into a database. Fast ingest uses the large pool for buffering the inserts before writing them to disk, so as to improve data insert performance.

Fast lookup

Fast lookup enables fast retrieval of data from a database for high-frequency queries. Fast lookup uses a separate memory area in the SGA called the *memoptimize pool* for buffering the data queried from tables, so as to improve query performance.



For using fast lookup, you must allocate appropriate memory size to the memoptimize pool using the MEMOPTIMIZE POOL SIZE initialization parameter.

Using Fast Ingest

Fast ingest optimizes the processing of high-frequency, single-row data inserts into database from applications, such as Internet of Things (IoT) applications.

Fast ingest uses the MEMOPTIMIZE_WRITE hint to insert data into tables specified as MEMOPTIMIZE FOR WRITE. The database temporarily buffers these inserts in the large pool and automatically commits the changes at the time of writing these buffered inserts to disk. The changes cannot be rolled back.

The inserts using fast ingest are also known as *deferred inserts*, because they are initially buffered in the large pool and later written to disk asynchronously by background processes.



Steps for using fast ingest for inserting data into a table

The following are the steps for using fast ingest for inserting data into a table:

1. Enable a table for fast ingest by specifying the MEMOPTIMIZE FOR WRITE clause in the CREATE TABLE Or ALTER TABLE statement.

```
SQL> create table test_fast_ingest (
id number primary key,
test_col varchar2(15))
segment creation immediate
memoptimize for write;
Table created.
```

See "Enabling a Table for Fast Ingest" for more information.

Enable fast ingest for inserts by specifying the MEMOPTIMIZE_WRITE hint in the INSERT statement.

The following is not how fast ingest is meant to be used, but demonstrates the mechanism.

```
SQL> insert /*+ memoptimize_write */ into test_fast_ingest values (1,
'test');

1 row created

SQL> insert /*+ memotimize_write */ into test_fast_ingest values (2,
'test');

1 row created
```

See "Specifying a Hint for Using Fast Ingest for Data Inserts" for more information.

The result of the two inserts above is to write data to the ingest buffer in the large pool of the SGA. At some point, that data is flushed to the <code>TEST_FAST_INGEST</code> table. Until that happens, the data is not durable.

Because the purpose of fast-ingest is to support high performance data streaming, a more realistic architecture would involve having one or more application or ingest servers collecting data and batching inserts to the database.

The first time an insert is run, the fast ingest area is allocated from the large pool. The amount of memory allocated is written to the alert.log.

```
Memoptimize Write allocated 2055M from large pool
```

If the request fails to allocate even the minimal memory requirement, then an error message is written to the alert.log.

Memoptimize Write disabled. Unable to allocate sufficient memory from large pool.



Details about Fast Ingest

The intent of fast-ingest is to support applications that generate lots of informational data that has important value in the aggregate but that doesn't necessarily require full ACID guarantees. Many applications in the Internet of Things (IoT) have a rapid "fire and forget" type workload, such as sensor data, smart meter data or even traffic cameras. For these applications, data might be collected and written to the database in high volumes for later analysis.

The following diagram shows how this might work with the Memoptimized Rowstore – Fast Ingest feature.

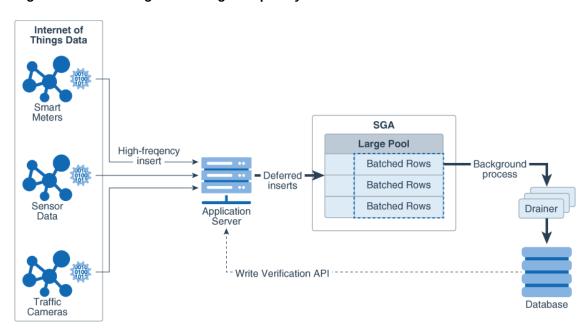


Figure 10-1 Fast-Ingest with high-frequency inserts.

The ingested data is batched in the large pool and is not immediately written to the database. Thus, the ingest process is very fast. Very large volumes of data can be ingested efficiently without having to process individual rows. However, if the database goes down before the ingested data is written out to the database files, it is possible to lose data.

Fast ingest is very different from normal Oracle Database transaction processing where data is logged and never lost once "written" to the database (i.e. committed). In order to achieve the maximum ingest throughput, the normal Oracle transaction mechanisms are bypassed, and it is the responsibility of the application to check to see that all data was indeed written to the database. Special APIs have been added that can be called to check if the data has been written to the database.

The commit operation has no meaning in the context of fast ingest, because it is not a transaction in the traditional Oracle sense. There is no ability to rollback the inserts. You also cannot query the data until it has been flushed from the fast ingest buffers to disk. You can see some administrative information about the fast ingest buffers by querying the view $V\$MEMOPTIMIZE \ WRITE \ AREA.$

You can also use the packages <code>DBMS_MEMOPTIMIZE</code> and <code>DBMS_MEMOPTIMIZE_ADMIN</code> to perform functions like flushing fast ingest data from the large pool and determining the sequence id of data that has been flushed.

Index operations and constraint checking is done only when the data is written from the fast ingest area in the large pool to disk. If primary key violations occur when the background processes write data to disk, then the database will not write those rows to the database.

Assuming (for most applications but not all) that all inserted data needs to be written to the database, it is critical that the application insert process checks to see that the inserted data has actually been written to the database before destroying that data. Only when that confirmation has occurred can the data be deleted from the inserting process.

See Also:

- Prerequisites for Fast Ingest Table
- Enabling a Table for Fast Ingest
- Specifying a Hint for Using Fast Ingest for Data Inserts
- Managing Fast Ingest Data in the Large Pool
- Disabling a Table for Fast Ingest
- Oracle Database Concepts for more information about the deferred insert mechanism

Prerequisites for Fast Ingest Table

Fast Ingest is not supported for tables with certain characteristics, objects, or partitioning.

If the table contains at least one column of BFILE datatype, or if the table contains more than 255 columns, fast-ingest is not supported.

When the Autonomous Database supports an item and does not have that limitation, it is so noted.

- Tables with the following characteristics cannot use fast ingest:
 - disk compression
 - in-memory compression
 - function-based indexes
 - domain indexes
 - bitmap indexes
 - bitmap join indexes
 - column default values [Autonomous Database supports column default values]
 - encryption [Autonomous Database supports encryption]
 - ref types
 - varray types
 - OID\$ types
 - unused columns
 - LOBs [Autonomous Database supports inline LOBs.]
 - triggers



- binary columns
- foreign keys
- row archival
- invisible columns
- virtual columns [Autonomous Database supports virtual columns.]
- The following objects cannot use fast ingest.
 - Temporary tables
 - Nested tables
 - Index organized tables
 - External tables
 - Materialized views with on-demand refresh
 - Sub-partitioning is not supported. [Autonomous Database supports sub-partitioning.]
- The following partitioning types are not supported.
 - REFERENCE
 - SYSTEM
 - INTERVAL [Autonomous Database supports this.]
 - AUTOLIST [Autonomous Database supports this.]

Unique Constraints using Non-Unique Indexes

Fast ingest does *not* support unique constraints that are enforced using non-unique indexes. The following is an invalid SQL code snippet that attempts this.

```
drop table t1;
create table t1 (c1 number, c2 number) memoptimize for write;
insert /*+ memoptimize_write */ into t1 values (20, 20);
create index ti on t1 (c1,c2);
alter table t1 add primary key (c1, c2) using index ti;
exec dbms_memoptimize.write_end;

Creating the index - "create index ti on t1 (c1,c2);
[...]
Error ORA-600
```

In the above example, Error ORA-600 is thrown during the flush.

Additional Considerations for Fast Ingest

The following are some additional considerations for fast ingest:

- Assuming (for most applications but not all) that all inserted data needs to be written to the
 database, it is critical that the application implement process checks to see that the
 inserted data has actually been written to the database before destroying that data. Only
 when that confirmation has occurred can the data be deleted from the inserting process.
- Because fast ingest buffers data in the large pool, there is a possibility of data loss in the event of a system failure. To avoid data loss, a client must keep a local copy of the data after performing inserts, so that it can replay the inserts in the event of a system failure before the data is written to disk. A client can use the DBMS MEMOPTIMIZE package



subprograms to track the durability of the inserts. After inserts are written to disk, a client can destroy its local copy of the inserted data.

- Queries do not read data from the large pool, hence data inserted using fast ingest cannot be queried until it is written to disk.
- Index operations are supported by fast ingest similar to the regular inserts. However, for fast ingest, database performs index operations while writing data to disk, and not while writing data into the large pool.
- The size allocated to the fast ingest buffers in the Large pool is fixed once created. If the buffer fills, further ingest waits until the background processes drain the buffer.
- Flush order is maintained within an object. If a table is (sub)partitioned, row order is maintained at the (sub)partition level and not at the table level.



A table can be configured for using both fast ingest and fast lookup.

Enabling a Table for Fast Ingest

You can enable a table for fast ingest by specifying the MEMOPTIMIZE FOR WRITE clause in the CREATE TABLE or ALTER TABLE statement.

To enable a table for fast ingest:

- 1. In SQL*Plus, log in to the database as a user with ALTER TABLE privileges.
- 2. Run the CREATE TABLE or ALTER TABLE statement with the MEMOPTIMIZE FOR WRITE clause.

The following example creates a new table <code>test_fast_ingest</code> and enables it for fast ingest:

The following example enables the existing table hr.employees for fast ingest:

```
ALTER TABLE hr.employees MEMOPTIMIZE FOR WRITE;
```

Specifying a Hint for Using Fast Ingest for Data Inserts

You can use fast ingest for data inserts by specifying the MEMOPTIMIZE_WRITE hint in INSERT statements.

Prerequisites

This task assumes:

- The table is already enabled for fast ingest.
- The optimizer is allowed to use hints, meaning optimizer_ignore_hints=FALSE.

To use fast ingest for data inserts:

- 1. In SQL*Plus, log in to the database as a user with the privileges to insert data into tables.
- 2. Run the INSERT statement with the MEMOPTIMIZE_WRITE hint for a table that is already enabled for fast ingest.

For example:

```
INSERT /*+ MEMOPTIMIZE WRITE */ INTO test fast ingest VALUES (1,'test');
```

See Also:

"Enabling a Table for Fast Ingest"

Disabling a Table for Fast Ingest

You can disable a table for fast ingest by specifying the NO MEMOPTIMIZE FOR WRITE clause in the ALTER TABLE statement.

To disable a table for fast ingest:

- 1. In SQL*Plus, log in to the database as a user with the ALTER TABLE privileges.
- 2. Run the ALTER TABLE statement with the NO MEMOPTIMIZE FOR WRITE clause.

The following example disables the table hr.employees for fast ingest:

ALTER TABLE hr.employees NO MEMOPTIMIZE FOR WRITE;

Managing Fast Ingest Data in the Large Pool

You can view the fast ingest data in the large pool using the V\$MEMOPTIMIZE_WRITE_AREA view. You can also view and control the fast ingest data in the large pool using the subprograms of the packages DBMS_MEMOPTIMIZE and DBMS_MEMOPTIMIZE_ADMIN.

Overview of the V\$MEMOPTIMIZE_WRITE_AREA view

The V\$MEMOPTIMIZE_WRITE_AREA view provides the following information about the memory usage and data inserts in the large pool by fast ingest:

- Total amount of memory allocated for fast ingest data in the large pool
- Total amount of memory currently used by fast ingest data in the large pool
- Total amount of memory currently free for storing fast ingest data in the large pool
- Number of fast ingest insert operations for which data is still in the large pool and is yet to be written to disk
- Number of clients currently using fast ingest for inserting data into the database



See Also:

Oracle Database Reference for information about the V\$MEMOPTIMIZE_WRITE_AREA view

Overview of the DBMS_MEMOPTIMIZE package subprograms

You can use the following subprograms of the DBMS_MEMOPTIMIZE package to view and control the fast ingest data in the large pool:

Subprogram	Description
GET_APPLY_HWM_SEQID	Returns the low high- water mark (low HWM) of sequence numbers of data records that are successfully written to disk by all the sessions.
GET_WRITE_HWM_SEQID	Returns the high-water mark (HWM) sequence number of the data record that is written to the large pool for the current session.
WRITE_END	Flushes all the fast ingest data from the large pool to disk for the current session.

See Also:

Oracle Database PL/SQL Packages and Types Reference for information about the $\tt DBMS \ MEMOPTIMIZE \ package$

Overview of the DBMS_MEMOPTIMIZE_ADMIN package subprograms

You can use the following subprograms of the <code>DBMS_MEMOPTIMIZE_ADMIN</code> package to control the fast ingest data in the large pool:

Subprogram	Description	
WRITES_FLUSH	Flushes all the fast ingest data from the large pool to disk for all the sessions.	

See Also:

Oracle Database PL/SQL Packages and Types Reference for information about the DBMS MEMOPTIMIZE ADMIN package



Using Fast Lookup

Fast lookup enables fast data retrieval from database tables for applications, such as Internet of Things (IoT) applications.

Fast lookup uses a hash index that is stored in the SGA buffer area called *memoptimize pool* to provide fast access to blocks of tables permanently pinned in the buffer cache, thus avoiding disk I/O and improving query performance.

Steps for using fast lookup for a table

The following are the steps for using fast lookup for a table:

1. Enable the memoptimize pool

This task is achieved by setting the MEMOPTIMIZE_POOL_SIZE initialization parameter to a non-zero value.

See "Enabling the Memoptimize Pool" for more information.

2. Enable a table for fast lookup

This task is achieved by specifying the MEMOPTIMIZE FOR READ clause in the CREATE TABLE or ALTER TABLE statement.

See "Enabling a Table for Fast Lookup" for more information.



A table can be configured for using both fast ingest and fast lookup.

See Also:

- Enabling the Memoptimize Pool
- Enabling a Table for Fast Lookup
- Disabling a Table for Fast Lookup
- Managing Fast Lookup Data in the Memoptimize Pool
- Oracle Database Concepts for information about the memoptimize pool memory architecture
- Oracle Database Reference for information about the MEMOPTIMIZE_POOL_SIZE initialization parameter

Prerequisites for Fast Lookup

Fast Lookup is not supported for tables with certain characteristics, objects, or partitioning.

The following are the limitations for using fast lookup:

- Tables enabled for fast lookup cannot be compressed.
- Tables enabled for fast lookup must have a primary key constraint enabled.



- Tables with the following characteristics cannot use fast READ:
 - disk compression
 - in-memory compression
 - function-based indexes
 - domain indexes
 - bitmap indexes
 - bitmap join indexes
 - ref types
 - varray types
 - OID\$ types
 - unused columns
 - LOBs
 - LONG columns
- The following objects cannot use fast READ.
 - Temporary tables
 - Nested tables
 - Index organized tables
 - External tables
 - Materialized views with on-demand refresh
 - Sub-partitioning is not supported.
 - AQ tables

Enabling the Memoptimize Pool

You must enable the memoptimize pool before using fast lookup. The memoptimize pool resides in the SGA, and stores the data and hash index for the tables that are enabled for fast lookup.

Prerequisites

This task assumes that the COMPATIBLE initialization parameter is set to 18.0.0 or higher.

To enable the memoptimize pool:

- 1. In SQL*Plus, log in to the database as a user with administrative privileges.
- 2. Set the MEMOPTIMIZE_POOL_SIZE initialization parameter to a non-zero value. The minimum setting is 100 MB. When you set this initialization parameter in a server parameter file (SPFILE) using the ALTER SYSTEM statement, you must specify SCOPE=SPFILE.

For example, the following statement sets the memoptimize pool size to 10 GB:

```
ALTER SYSTEM SET MEMOPTIMIZE POOL SIZE = 10G SCOPE=SPFILE;
```

3. Restart the database for the change to take effect.



Example: Enabling the Memoptimize Pool

Assume that the MEMOPTIMIZE_POOL_SIZE initialization parameter is initially set to 0. The following example enables the memoptimize pool by setting the MEMOPTIMIZE_POOL_SIZE to 10 GB:

SQL> SHOW PARAMETER MEMOPTIMIZE POOL SIZE

NAME TYPE VALUE
----memoptimize pool size big integer 0

SQL> ALTER SYSTEM SET MEMOPTIMIZE_POOL_SIZE=10G SCOPE=SPFILE;

System altered.

SQL> SHUTDOWN IMMEDIATE Database closed. Database dismounted. ORACLE instance shut down.

SQL> STARTUP ORACLE instance started.

Total System Global Area 1.1832E+10 bytes
Fixed Size 9010864 bytes
Variable Size 1.1799E+10 bytes
Database Buffers 16777216 bytes
Redo Buffers 7766016 bytes

Database mounted. Database opened.

SQL> SHOW PARAMETER MEMOPTIMIZE POOL SIZE

NAME TYPE VALUE
----memoptimize_pool_size big integer 10G

Note:

The MEMOPTIMIZE_POOL_SIZE value does count toward SGA_TARGET, but the database does not grow and shrink the memoptimize pool automatically. For example, if SGA_TARGET is 10 GB, and if MEMOPTIMIZE_POOL_SIZE is 1 GB, then a total of 9 GB is available for SGA memory other than the memoptimize pool.



See Also:

- *Oracle Database Concepts* for information about the memoptimize pool memory architecture
- Oracle Database Reference for information about the MEMOPTIMIZE_POOL_SIZE initialization parameter

Enabling a Table for Fast Lookup

You can enable a table for fast lookup by specifying the MEMOPTIMIZE FOR READ clause in the CREATE TABLE or ALTER TABLE statement.

Prerequisites

This task assumes that the memoptimize pool is enabled.

To enable a table for fast lookup:

- 1. In SQL*Plus, log in to the database as a user with ALTER TABLE privileges.
- 2. Run the CREATE TABLE or ALTER TABLE statement with the MEMOPTIMIZE FOR READ clause for the table that needs to be enabled for fast lookup.

The following example creates a new table $test_fast_lookup$ and enables it for fast lookup:

The following example enables the existing table hr.employees for fast lookup:

```
ALTER TABLE hr.employees MEMOPTIMIZE FOR READ;
```

See Also:

- Enabling the Memoptimize Pool
- Disabling a Table for Fast Lookup
- Managing Fast Lookup Data in the Memoptimize Pool



Disabling a Table for Fast Lookup

You can disable a table for fast lookup by specifying the NO MEMOPTIMIZE FOR READ clause in the ALTER TABLE statement.

Prerequisites

This task assumes that a table is already enabled for fast lookup.

To disable a table for fast lookup:

- 1. In SQL*Plus, log in to the database as a user with the ALTER TABLE privileges.
- 2. Run the ALTER TABLE statement with the NO MEMOPTIMIZE FOR READ clause for the table that needs to be disabled for fast lookup.

The following example disables the hr.employees table for fast lookup:

ALTER TABLE hr.employees NO MEMOPTIMIZE FOR READ;



"Enabling a Table for Fast Lookup"

Managing Fast Lookup Data in the Memoptimize Pool

The memoptimize pool stores the data (*fast lookup data*) of all the tables that are enabled for fast lookup. You can explicitly delete or populate fast lookup data for a table in the memoptimize pool using the DBMS MEMOPTIMIZE package subprograms.

Overview of the DBMS_MEMOPTIMIZE package subprograms

The following are the <code>DBMS_MEMOPTIMIZE</code> package subprograms that can be used to delete or populate fast lookup data for a table in the memoptimize pool:

Subprogram	Description
DROP_OBJECT	Removes fast lookup data for a table from the memoptimize pool.
POPULATE	Populates fast lookup data for a table in the memoptimize pool.



See Also:

- "Enabling a Table for Fast Lookup"
- "Enabling the Memoptimize Pool"
- Oracle Database PL/SQL Packages and Types Reference for information about the DBMS_MEMOPTIMIZE package
- Oracle Database Concepts for information about the memoptimize pool memory architecture



11

Tuning the Database Buffer Cache

This chapter describes how to tune the database buffer cache. If you are using automatic memory management to manage the database memory on your system, there is no need to manually tune the memory caches described in this chapter.

This chapter contains the following topics:

- About the Database Buffer Cache
- Configuring the Database Buffer Cache
- · Configuring Multiple Buffer Pools
- Configuring the Redo Log Buffer
- Configuring the Database Caching Mode

About the Database Buffer Cache

For many types of operations, Oracle Database uses the buffer cache to store data blocks read from disk. Oracle Database bypasses the buffer cache for particular operations, such as sorting and parallel reads.

To use the database buffer cache effectively, tune SQL statements for the application to avoid unnecessary resource consumption. To meet this goal, verify that frequently executed SQL statements and SQL statements that perform many buffer gets are well-tuned.

When using parallel query, consider configuring the database to use the database buffer cache instead of performing direct reads into the Program Global Area (PGA). This configuration may be appropriate when the system has a large amount of memory.

See Also:

- Oracle Database SQL Tuning Guide for information about tuning SQL statements
- Oracle Database VLDB and Partitioning Guide for information about parallel execution

Configuring the Database Buffer Cache

When configuring a new database instance, it is impossible to know the correct size for the buffer cache. Typically, a database administrator makes a first estimate for the cache size, then runs a representative workload on the instance and examines the relevant statistics to see whether the cache is under-configured or over-configured.

This section describes how to configure the database buffer cache. If you are using automatic shared memory management to configure the Shared Global Area (SGA), there is no need to manually tune the database buffer cache as described in this section.

This section contains the following topics:

- Using the V\$DB_CACHE_ADVICE View
- Calculating the Buffer Cache Hit Ratio
- Interpreting the Buffer Cache Hit Ratio
- Increasing Memory Allocated to the Database Buffer Cache
- Reducing Memory Allocated to the Database Buffer Cache

Using the V\$DB CACHE ADVICE View

The V\$DB_CACHE_ADVICE view shows the simulated miss rates for a range of potential buffer cache sizes. This view assists in cache sizing by providing information that predicts the number of physical reads for each potential cache size. The data also includes a physical read factor, which is a factor by which the current number of physical reads is estimated to change if the buffer cache is resized to a given value.

However, physical reads do not necessarily indicate disk reads in Oracle Database, because physical reads may be accomplished by reading from the file system cache. Hence, the relationship between successfully finding a block in the cache and the size of the cache is not always a smooth distribution. When sizing the buffer pool, avoid using additional buffers that do not contribute (or contribute very little) to the cache hit ratio.

The following figure illustrates the relationship between physical I/O ratio and buffer cache size.

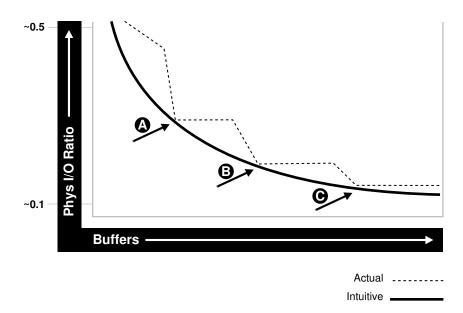


Figure 11-1 Physical I/O Ratio and Buffer Cache Size

Examining the example illustrated in the above figure leads to the following observations:

- As the number of buffers increases, the physical I/O ratio decreases.
- The decrease in the physical I/O between points A and B and points B and C is not smooth, as indicated by the dotted line in the graph.
- The benefit from increasing buffers from point A to point B is considerably higher than from point B to point C.



The benefit from increasing buffers decreases as the number of buffers increases.

There is some overhead associated with using this advisory view. When the advisory is enabled, there is a small increase in CPU usage, because additional bookkeeping is required. To reduce both the CPU and memory overhead associated with bookkeeping, Oracle Database uses sampling to gather cache advisory statistics. Sampling is not used if the number of buffers in a buffer pool is small to begin with.

To use the V\$DB CACHE ADVICE view:

Set the value of the DB CACHE ADVICE initialization parameter to ON.

This enables the advisory view. The DB_CACHE_ADVICE parameter is dynamic, so the advisory can be enabled and disabled dynamically to enable you to collect advisory data for a specific workload.

2. Run a representative workload on the database instance.

Allow the workload to stabilize before querying the V\$DB CACHE ADVICE view.

3. Query the V\$DB CACHE ADVICE view.

The following example shows a query of this view that returns the predicted I/O requirement for the default buffer pool for various cache sizes.

The output of this query might look like the following:

		-	Estd Phys	
Cache Size (MB)	Buffers	Read Factor	Reads	
30	3,802	18.70	192,317,943	10% of Current Size
60	7,604		131,949,536	
91	11,406			
121	15,208		51,111,658	
152	19,010	3.64	37,460,786	
182	22,812	2.50	25,668,196	
212	26,614	1.74	17,850,847	
243	30,416	1.33	13,720,149	
273	34,218	1.13	11,583,180	
304	38,020	1.00	10,282,475	Current Size
334	41,822	.93	9,515,878	
364	45,624	.87	8,909,026	
395	49,426	.83	8,495,039	
424	53,228	.79	8,116,496	
456	57,030	.76	7,824,764	
486	60,832	.74	7,563,180	
517	64,634	.71	7,311,729	
547	68,436	.69	7,104,280	
577	72,238	.67	6,895,122	
608	76,040	.66	6,739,731	200% of Current Size



In this example, the output shows that if the cache was 212 MB instead of the current size of 304 MB, the estimated number of physical reads would increase by a factor of 1.74, or 74%. Hence, it is not advisable to decrease the cache size to 212MB.

However, increasing the cache size to 334MB may potentially decrease reads by a factor of .93, or 7%. If an additional 30MB memory is available on the system and the value of the SGA_MAX_SIZE parameter allows for the increment, it is advisable to increase the default buffer cache pool size to 334MB.

Calculating the Buffer Cache Hit Ratio

The buffer cache hit ratio calculates how often a requested block has been found in the buffer cache without requiring disk access. This ratio is computed using data selected from the V\$SYSSTAT performance view. Use the buffer cache hit ratio to verify the physical I/O as predicted by the V\$DB CACHE ADVICE view.

Table 11-1 lists the statistics from the V\$SYSSTAT view used to calculate the buffer cache hit ratio.

Table 11-1 Statistics for Calculating the Buffer Cache Hit Ratio

Statistic	Description
consistent gets from cache	Number of times a consistent read was requested for a block from the buffer cache.
db block gets from cache	Number of times a CURRENT block was requested from the buffer cache.
physical reads cache	Total number of data blocks read from disk into buffer cache.

Example 11-1 shows a query of this view.

Example 11-1 Querying the V\$SYSSTAT View

```
SELECT name, value

FROM V$SYSSTAT

WHERE name IN ('db block gets from cache', 'consistent gets from cache',

'physical reads cache');
```

In this example, the query is simplified by using values selected directly from the V\$SYSSTAT view, rather than over an interval. It is recommended to calculate the delta of these statistics over an interval while the application is running, then use these delta values to determine the buffer cache hit ratio. For information about collecting statistics over an interval, see Automatic Performance Diagnostics .

Using the values from the output of this query, calculate the hit ratio for the buffer cache using the following formula:

```
1 - (('physical reads cache') / ('consistent gets from cache' +
'db block gets from cache'))
```



Oracle Database Reference for information about the V\$SYSSTAT view

Interpreting the Buffer Cache Hit Ratio

Before deciding whether to increase or decrease the buffer cache size, you should first examine the buffer cache hit ratio.

A low cache hit ratio does not necessarily imply that increasing the size of the buffer cache will benefit performance. Moreover, a high cache hit ratio may wrongly indicate that the buffer cache is adequately sized for the workload.

To interpret the buffer cache hit ratio, consider the following factors:

- Avoid repeated scanning of frequently accessed data by performing the processing in a single pass or by optimizing the SQL statement.
 - Repeated scanning of the same large table or index can artificially inflate a low cache hit ratio. Examine frequently executed SQL statements with a large number of buffer gets, to ensure that the execution plans for these SQL statements are optimal.
- Avoid requerying the same data by caching frequently accessed data in the client program or middle tier.
- In large databases running OLTP applications, many rows are accessed only once (or never). Hence, there is no purpose in keeping the block in memory following its use.
- Do not continuously increase the buffer cache size.
 - Continuous increases of the buffer cache size have no effect if the database is performing full table scans or operations that do not use the buffer cache.
- Consider poor hit ratios when large full table scans are occurring.

Database blocks accessed during a long full table scan are placed on the tail end of the Least Recently Used (LRU) list and not on the head of the list. Therefore, the blocks age out faster than blocks read when performing indexed lookups or small table scans.



Short table scans are scans performed on tables under a certain size threshold. The definition of a small table is the maximum of 2% of the buffer cache or 20, whichever is bigger.

Increasing Memory Allocated to the Database Buffer Cache

If the cache hit ratio is low and your application is tuned to avoid performing full table scans, consider increasing the size of the buffer cache. If possible, resize the buffer pools dynamically, rather than shutting down the instance to perform this change.

To increase the size of the database buffer cache:

- 1. Set the value of the DB CACHE ADVICE initialization parameter to ON.
- 2. Allow the buffer cache statistics to stabilize.
- 3. Examine the advisory data in the V\$DB_CACHE_ADVICE view to determine the next increment required to significantly decrease the amount of physical I/O performed, as described in "Using the V\$DB_CACHE_ADVICE View".



- If it is possible to allocate the extra memory required to the buffer cache without causing the system to page, then allocate this memory.
- 5. To increase the amount of memory allocated to the buffer cache, increase the value of the DB CACHE SIZE initialization parameter.

The DB_CACHE_SIZE parameter specifies the size of the default cache for the database's standard block size. To create and use tablespaces with block sizes other than the database's standard block sizes (such as for transportable tablespaces), configure a separate cache for each block size used. Use the DB_nK_CACHE_SIZE parameter to configure the nonstandard block size needed (where n is 2, 4, 8, 16 or 32 and not the standard block size).

Note:

- The process of choosing a cache size is the same, regardless of whether the cache is the default standard block size cache, the KEEP or RECYCLE cache, or a nonstandard block size cache.
- When the cache is resized significantly (greater than 20%), the old cache
 advisory value is discarded and the cache advisory is set to the new size.
 Otherwise, the old cache advisory value is adjusted to the new size by the
 interpolation of existing values.

See Also:

For more information about the DB $\it nK$ CACHE SIZE parameter, see:

- Oracle Database Administrator's Guide
- Oracle Database Reference

Reducing Memory Allocated to the Database Buffer Cache

If the cache hit ratio is high, then the buffer cache is likely large enough to store the most frequently accessed data. If this is the case and memory is required for another memory structure, consider reducing the size of the buffer cache.

To reduce the size of the database buffer cache:

- Examine the advisory data in the V\$DB_CACHE_ADVICE view to determine if decreasing the size of the buffer cache will significantly increase the number of physical I/Os, as described in "Using the V\$DB_CACHE_ADVICE View".
- To reduce the amount of memory allocated to the buffer cache, decrease the value of the DB_CACHE_SIZE initialization parameter.

Configuring Multiple Buffer Pools

For most systems, a single default buffer pool is generally adequate. However, database administrators with detailed knowledge of an application's buffer pool may benefit from configuring multiple buffer pools.

For segments that have atypical access patterns, consider storing blocks from these segments in two separate buffer pools: the KEEP pool and the RECYCLE pool. A segment's access pattern may be atypical if it is constantly accessed (sometimes referred to as hot) or infrequently accessed (such as a large segment that is accessed by a batch job only once a day).

Using multiple buffer pools enables you to address these irregularities. You can use the KEEP pool to maintain frequently accessed segments in the buffer cache, and the RECYCLE pool to prevent objects from consuming unnecessary space in the buffer cache. When an object is associated with a buffer cache, all blocks from that object are placed in that cache. Oracle Database maintains a DEFAULT buffer pool for objects that are not assigned to a specific buffer pool. The default buffer pool size is determined by the DB_CACHE_SIZE initialization parameter. Each buffer pool uses the same LRU replacement policy. For example, if the KEEP pool is not large enough to store all of the segments allocated to it, then the oldest blocks age out of the cache.

By allocating objects to appropriate buffer pools, you can:

- Reduce or eliminate I/Os
- Isolate or limit an object to a separate cache

This section describes how to configure multiple buffer pools and contains the following topics:

- Considerations for Using Multiple Buffer Pools
- Using Multiple Buffer Pools
- Using the V\$DB CACHE ADVICE View for Individual Buffer Pools
- Calculating the Buffer Pool Hit Ratio for Individual Buffer Pools
- Examining the Buffer Cache Usage Pattern
- Configuring the KEEP Pool
- Configuring the RECYCLE Pool

Considerations for Using Multiple Buffer Pools

When using multiple buffer pools, take the following considerations into account:

- Random Access to Large Segments
- Oracle Real Application Cluster Instances

Random Access to Large Segments

A problem may occur with an LRU aging method when a very large segment (compared to the size of the buffer cache) is accessed with a large or unbounded index range scan. Any single segment that accounts for a substantial portion (more than 10%) of nonsequential physical reads can be considered very large. Random reads to a large segment may cause buffers that contain data for other segments to be aged out of the cache. The large segment ends up consuming a large percentage of the buffer cache, but it does not benefit from the cache.

Very frequently accessed segments are not affected by large segment reads because their buffers are warmed frequently enough that they do not age out of the buffer cache. However, the problem affects warm segments that are not accessed frequently enough to survive the buffer aging caused by the large segment reads. There are three options for solving this problem:

• If the object accessed is an index, determine whether the index is selective. If not, tune the SQL statement to use a more selective index.

- If the SQL statement is tuned, move the large segment into a separate RECYCLE cache so it does not affect the other segments. The RECYCLE cache should be smaller than the DEFAULT buffer pool, and it should reuse buffers more quickly.
- Alternatively, consider moving the small, warm segments into a separate KEEP cache that is
 not used for large segments. Size the KEEP cache to minimize misses in the cache. You
 can make the response times for specific queries more predictable by storing the
 segments accessed by the queries in the KEEP cache to ensure that they do not age out.

Oracle Real Application Cluster Instances

In an Oracle Real Application Cluster (Oracle RAC) environment, consider creating multiple buffer pools for each database instance. It is not necessary to define the same set of buffer pools for each instance of the database. Among instances, the buffer pools can be different sizes or undefined. Tune each instance according to the application requirements for that instance.

Using Multiple Buffer Pools

To define a default buffer pool for an object, use the <code>BUFFER_POOL</code> keyword of the <code>STORAGE</code> clause. This clause is valid for the following SQL statements:

- CREATE TABLE
- CREATE CLUSTER
- CREATE INDEX
- ALTER TABLE
- ALTER CLUSTER
- ALTER INDEX

After a buffer pool is defined, all subsequent blocks read for the object are placed in that pool. If a buffer pool is defined for a partitioned table or index, then each partition of the object inherits the buffer pool from the table or index definition, unless if it is overridden by a specific buffer pool.

When the buffer pool of an object is changed using the ALTER statement, all buffers currently containing blocks of the altered segment remain in the buffer pool they were in before the ALTER statement. Newly loaded blocks and any blocks that age out and are reloaded are placed into the new buffer pool.



Oracle Database SQL Language Reference for information about specifying ${\tt BUFFER\ POOL\ in\ the\ STORAGE\ clause}$

Using the V\$DB_CACHE_ADVICE View for Individual Buffer Pools

As with the default buffer pool, you can use V\$DB_CACHE_ADVICE view to assist in cache sizing of other pools. After estimating the initial cache size and running a representative workload, query the V\$DB_CACHE_ADVICE view for the pool you want to use.



For more information about using the V\$DB_CACHE_ADVICE view, see "Using the V\$DB_CACHE_ADVICE View".

Example 11-2 shows a query of this view that queries data from the KEEP pool:

Example 11-2 Querying the V\$DB_CACHE_ADVICE View for the KEEP Pool

Calculating the Buffer Pool Hit Ratio for Individual Buffer Pools

The data in the V\$SYSSTAT view reflects the logical and physical reads for all buffer pools within one set of statistics. To determine the hit ratio for the buffer pools individually, query the V\$BUFFER_POOL_STATISTICS view. This view maintains statistics on the number of logical reads and writes for each pool.

See Also:

- "Calculating the Buffer Cache Hit Ratio" for more information about calculating hit ratios
- Oracle Database Reference for more information about V\$BUFFER_POOL_STATISTICS view

The following query calculates the hit ratio using the V\$BUFFER POOL STATISTICS view.

Example 11-3 Querying the V\$BUFFER_POOL_STATISTICS View

Examining the Buffer Cache Usage Pattern

The V\$BH view shows the data object ID of all blocks that currently reside in the SGA. To determine which segments have many buffers in the pool, use this view to examine the buffer cache usage pattern. You can either examine the buffer cache usage pattern for all segments or a specific segment, as described in the following sections:

- Examining the Buffer Cache Usage Pattern for All Segments
- Examining the Buffer Cache Usage Pattern for a Specific Segment

Examining the Buffer Cache Usage Pattern for All Segments

One method to determine which segments have many buffers in the pool is to query the number of blocks for all segments that reside in the buffer cache at a given time. Depending on buffer cache size, this might require a lot of sort space.

Example 11-4 shows a query that counts the number of blocks for all segments.

Example 11-4 Querying the Number of Blocks for All Segments

```
COLUMN object_name FORMAT A40
COLUMN number_of_blocks FORMAT 999,999,999,999

SELECT o.object_name, COUNT(*) number_of_blocks
FROM DBA_OBJECTS o, V$BH bh
WHERE o.data_object_id = bh.OBJD
AND o.owner != 'SYS'
GROUP BY o.object_Name
ORDER BY COUNT(*);
```

The output of this guery might look like the following:

OBJECT_NAME	NUMBER_OF_BLOCKS
OA PREF UNIQ KEY	1
SYS_C002651	1
••	
DS PERSON	78
OM_EXT_HEADER	701
OM SHELL	1,765
OM HEADER	5,826
OM_INSTANCE	12,644

Examining the Buffer Cache Usage Pattern for a Specific Segment

Another method to determine which segments have many buffers in the pool is to calculate the percentage of the buffer cache used by an individual object at a given time.

To calculate the percentage of the buffer cache used by an individual object:

1. Find the Oracle Database internal object number of the segment by querying the DBA OBJECTS view:

```
SELECT data_object_id, object_type
FROM DBA_OBJECTS
WHERE object name = UPPER('segment name');
```

Because two objects can have the same name (if they are different types of objects), use the <code>OBJECT_TYPE</code> column to identify the object of interest.

2. Find the number of buffers in the buffer cache for SEGMENT NAME:

```
SELECT COUNT(*) buffers
FROM V$BH
WHERE objd = data_object_id_value;
```

For data object id value, use the value of DATA OBJECT ID from the previous step.

Find the number of buffers in the database instance:

```
SELECT name, block_size, SUM(buffers)
FROM V$BUFFER_POOL
GROUP BY name, block_size
HAVING SUM(buffers) > 0;
```

4. Calculate the ratio of buffers to total buffers to obtain the percentage of the cache currently used by SEGMENT NAME:

```
% cache used by segment name = [buffers(Step2)/total buffers(Step3)]
```



This method works only for a single segment. For a partitioned object, run the query for each partition.

Configuring the KEEP Pool

The purpose of the KEEP buffer pool is to retain objects in memory, thus avoiding I/O operations. Each object kept in memory results in a trade-off. It is more beneficial to keep frequently-accessed blocks in the cache. Avoid retaining infrequently-used blocks in the cache, as this results in less space for other, more active blocks

If there are certain segments in your application that are referenced frequently, then consider storing the blocks from those segments in the KEEP buffer pool. Typical segments that are kept in the KEEP pool are small, frequently-used reference tables. To determine which tables are candidates, check the number of blocks from candidate tables by querying the V\$BH view, as described in "Examining the Buffer Cache Usage Pattern".

To configure the KEEP pool:

1. Compute an approximate size for the KEEP buffer pool.

The size of the KEEP buffer pool depends on the objects to be kept in the buffer cache. To estimate its size, add the blocks used by all objects assigned to this pool.

If you gathered statistics on the segments, query DBA_TABLES.BLOCKS and DBA TABLES.EMPTY BLOCKS to determine the number of blocks used.

2. Taking two snapshots of system performance at different times.

Query data from the KEEP pool for each snapshot using the V\$DB_CACHE_ADVICE view, as described in "Using the V\$DB_CACHE_ADVICE View for Individual Buffer Pools".

3. Subtract the more recent values for physical reads, block gets, and consistent gets from the older values, and use the results to calculate the hit ratio.

A buffer pool hit ratio of 100% may not be optimal. Oftentimes, you can decrease the size of the KEEP buffer pool and still maintain a sufficiently high hit ratio. Allocate blocks removed from the KEEP buffer pool to other buffer pools.

4. Allocate memory to the KEEP buffer pool by setting the value of the DB_KEEP_CACHE_SIZE parameter to the required size.

The memory for the KEEP pool is not a subset of the default pool.



If an object grows in size, then it might no longer fit in the KEEP buffer pool. You will begin to lose blocks out of the cache.

Configuring the RECYCLE Pool

You can configure a RECYCLE buffer pool for blocks belonging to those segments that you do not want to keep in memory. The purpose of the RECYCLE pool is to retain segments that are

scanned rarely or are not referenced frequently. If an application randomly accesses the blocks of a very large object, then it is unlikely for a block stored in the buffer pool to be reused before it is aged out. This is true regardless of the size of the buffer pool (given the constraint of the amount of available physical memory). Consequently, the object's blocks do not need to be cached; the cache buffers can be allocated to other objects.

Do not discard blocks from memory too quickly. If the buffer pool is too small, then blocks can age out of the cache before the transaction or SQL statement completes its execution. For example, an application might select a value from a table, use the value to process some data, and then update the record. If the block is removed from the cache after the SELECT statement, then it must be read from disk again to perform the update. The block should be retained for the duration of the user transaction.

To configure the RECYCLE POOL:

 Allocate memory to the RECYCLE buffer pool by setting the value of the DB_RECYCLE_CACHE_SIZE parameter to the required size.

The memory for the RECYCLE pool is not a subset of the default pool.

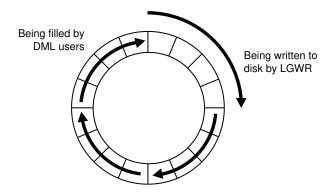
Configuring the Redo Log Buffer

Server processes making changes to data blocks in the buffer cache generate redo data into the log buffer. The log writer process (LGWR) begins writing to copy entries from the redo log buffer to the online redo log if any of the following conditions are true:

- · The redo log buffer becomes at least one-third full
- LGWR is posted by a server process performing a COMMIT or ROLLBACK
- A database writer process (DBWR) posts LGWR to do so

When LGWR writes redo entries from the redo log buffer to a redo log file or disk, user processes can copy new entries over the entries in memory that are written to disk, as illustrated in the following figure.

Figure 11-2 Redo Log Buffer



LGWR attempts to write fast enough to ensure that space is available in the redo log buffer for new entries, even if it is frequently accessed. Having a larger redo log buffer makes it more likely that there is space for new entries, and also enables LGWR to efficiently process redo records. On a system with large updates, if the redo log buffer is too small, LGWR will continuously flush redo to disk so that it remains two-thirds empty.

On systems with fast processors and relatively slow disks, the processors might be filling the rest of the redo log buffer in the time it takes the redo log writer to move a portion of the redo log buffer to disk. In this situation, a larger redo log buffer can temporarily mask the effects of slower disks. Alternatively, consider either improving:

- The checkpointing or archiving process
- The performance of LGWR by moving all online logs to fast raw devices

To improve the performance of the redo log buffer, ensure that you are:

- Batching commit operations for batch jobs, so that LGWR is able to write redo log entries
 efficiently
- Using NOLOGGING operations when loading large quantities of data

This section describes how to configure the redo log buffer and contains the following topics:

- Sizing the Redo Log Buffer
- Using Redo Log Buffer Statistics

Sizing the Redo Log Buffer

The default size of the redo log buffer is calculated as follows:

```
MAX(0.5M, (128K * number of cpus))
```

Applications that insert, modify, or delete large volumes of data may require changing the default size of the redo log buffer. Oracle recommends setting the redo log buffer size to minimum of 8 MB. Set it to a minimum of 64 MB for databases using flashback functionality and having 4GB or higher SGAs. Set it to a minimum of 256 MB if you are using Oracle Data Guard with asynchronous redo transport and have a high redo generation rate.

To determine if the size of the redo log buffer is too small, monitor the redo log buffer statistics, as described in "Using Redo Log Buffer Statistics". You can also check if the log buffer space wait event is a significant factor in the wait time for the database instance. If it is not, then the log buffer size is most likely adequately-sized.

To size the redo log buffer:

• Set the size of the redo log buffer by setting the value of the LOG_BUFFER initialization parameter to the required size.

The value of this parameter is expressed in bytes.



The size of the redo log buffer cannot be modified after instance startup.

Using Redo Log Buffer Statistics

The REDO BUFFER ALLOCATION RETRIES statistic reflects the number of times a user process waits for space in the redo log buffer. This statistic can be queried using the V\$SYSSTAT performance view.

You should monitor the redo buffer allocation retries statistic over a period while the application is running. The value of this statistic should be near zero over an interval. If this

value increases consistently, then it means user processes had to wait for space in the redo log buffer to become available. The wait can be caused by the redo log buffer being too small or by checkpointing. In this case, consider one of the following options:

- Increase the size of the redo log buffer, as described in "Sizing the Redo Log Buffer"
- Improve the checkpointing or archiving process

Example 11-5 shows a query of the V\$SYSSTAT view for this statistic.

Example 11-5 Querying the V\$SYSSTAT View

```
SELECT name, value
  FROM V$SYSSTAT
WHERE name = 'redo buffer allocation retries';
```

Configuring the Database Caching Mode

Starting with Oracle Database 12c Release 1 (12.1.0.2), there are two database caching modes: the default database caching mode used in previous versions of Oracle Database, and the force full database caching mode that is new to this release. In default caching mode, Oracle Database does not always cache the underlying data when a user queries a large table. In force full database caching mode, Oracle Database assumes that the buffer cache is large enough to cache the full database and tries to cache all the blocks that are accessed by queries.

This section contains the following topics:

- Default Database Caching Mode
- Force Full Database Caching Mode
- Determining When to Use Force Full Database Caching Mode
- Verifying the Database Caching Mode



Force full database caching mode is available starting with Oracle Database 12c Release 1 (12.1.0.2).

Default Database Caching Mode

By default, Oracle Database uses the default database caching mode when performing full table scans. In default caching mode, Oracle Database does not always cache the underlying data when a user queries a large table, because doing so might remove more useful data from the buffer cache.

If the Oracle Database instance determines that there is enough space to cache the full database in the buffer cache and that it would be beneficial to do so, then the instance automatically caches the full database in the buffer cache.

If the Oracle Database instance determines that there is not enough space to cache the full database in the buffer cache, then:

 Smaller tables are loaded into memory only when the table size is less than 2 percent of the buffer cache size.

- For medium tables, Oracle Database analyzes the interval between the last table scan and
 the aging timestamp of the buffer cache. If the size of the table reused in the last table
 scan is greater than the remaining buffer cache size, then the table is cached.
- Large tables are typically not loaded into memory, unless if you explicitly declare the table for the KEEP buffer pool.



In default caching mode, Oracle Database instance does not cache NOCACHE LOBs in the buffer cache.

See Also:

Oracle Database Concepts for information about the default database caching mode

Force Full Database Caching Mode

As more memory is added to a database, buffer cache sizes may continually grow. In some cases, the size of the buffer cache may become so large that the entire database can fit into memory. The ability to cache an entire database in memory can drastically improve database performance when performing full table scans or accessing LOBs.

In force full database caching mode, Oracle Database caches the entire database in memory when the size of the database is smaller than the database buffer cache size. All data files, including NOCACHE LOBS and LOBS that use SecureFiles, are loaded into the buffer cache as they are being accessed.

See Also:

- Oracle Database Concepts
- Oracle Database Administrator's Guide

Determining When to Use Force Full Database Caching Mode

To improve database performance for table scans and LOB data access, especially for workloads that are limited by I/O throughput or response time, consider using force full database caching mode whenever the size of the database buffer cache is greater than the size of the database.

Consider using force full database caching mode in the following situations:

The logical database size (or actual used space) is smaller than the individual buffer cache
of each database instance in an Oracle RAC environment. This is applicable for nonOracle RAC database as well.

- The logical database size is smaller than 80% of the combined buffer cache sizes of all the database instances for well-partitioned workloads (by instance access) in an Oracle RAC environment.
- The database uses SGA_TARGET or MEMORY_TARGET.
- The NOCACHE LOBs need to be cached. The NOCACHE LOBs are never cached unless force full database caching is used.

For the first three situations, you should monitor the system performance periodically to verify that the performance figures are according to your expectations.

When one Oracle RAC database instance uses force full database caching mode, then all the other database instances in the Oracle RAC environment will also use force full database caching mode.

In a multitenant environment, force full database caching mode applies to the entire container database (CDB), including all of its pluggable databases (PDBs).

Verifying the Database Caching Mode

By default, Oracle Database runs in the default database caching mode.

To verify if force full database caching mode is enabled:

Query the V\$DATABASE view as shown:

```
SELECT FORCE FULL DB CACHING FROM V$DATABASE;
```

If the query returns a value of YES, then force full database caching mode is enabled on the database. If the query returns a value of NO, then force full database caching mode is disabled and the database is in default database caching mode.

Note:

To enable force full database caching mode, use the following ALTER DATABASE command:

ALTER DATABASE FORCE FULL DATABASE CACHING;

See Also:

- Oracle Database Administrator's Guide for more information about enabling and disabling force full database caching mode
- Oracle Database Reference for more information about the V\$DATABASE view



Tuning the Shared Pool and the Large Pool

This chapter describes how to tune the shared pool and the large pool. If you are using automatic memory management to manage the database memory on your system, or automatic shared memory management to configure the Shared Global Area (SGA), there is no need to manually tune the shared pool and the large pool as described in this chapter.

This chapter contains the following topics:

- About the Shared Pool
- Using the Shared Pool
- Configuring the Shared Pool
- Configuring the Large Pool

About the Shared Pool

Oracle Database uses the shared pool to cache many different types of data. Cached data includes the textual and executable forms of PL/SQL blocks and SQL statements, dictionary cache data, result cache data, and other data.

This section describes the shared pool and contains the following topics:

- Benefits of Using the Shared Pool
- Shared Pool Concepts

Benefits of Using the Shared Pool

Proper use and sizing of the shared pool can reduce resource consumption in at least four ways:

- If the SQL statement is in the shared pool, parse overhead is avoided, resulting in reduced CPU resources on the system and elapsed time for the end user.
- Latching resource usage is significantly reduced, resulting in greater scalability.
- Shared pool memory requirements are reduced, because all applications use the same pool of SQL statements and dictionary resources.
- I/O is reduced, because dictionary elements that are in the shared pool do not require disk access.

Shared Pool Concepts

The main components of the shared pool include:

Library cache

The library cache stores the executable (parsed or compiled) form of recently referenced SQL and PL/SQL code.

Data dictionary cache

The data dictionary cache stores data referenced from the data dictionary.

Server result cache (depending on the configuration)

The server result cache is an optional cache that stores query and PL/SQL function results within the shared pool. For information about the server result cache, see "About the Result Cache".

Many of the caches in the shared pool—including the library cache and the dictionary cache—automatically increase or decrease in size, as needed. Old entries are aged out to accommodate new entries when the shared pool runs out of space.

A cache miss on the library cache or data dictionary cache is more expensive than a miss on the buffer cache. For this reason, the shared pool should be sized to ensure that frequently-used data is cached.

Several features require large memory allocations in the shared pool, such as shared server, parallel query, or Recovery Manager. Oracle recommends using a separate memory area—the large pool—to segregate the System Global Area (SGA) memory used by these features.

Allocation of memory from the shared pool is performed in chunks. This chunking enables large objects (over 5 KB) to be loaded into the cache without requiring a single contiguous area. In this way, the database reduces the possibility of running out of contiguous memory due to fragmentation.

Java, PL/SQL, or SQL cursors may sometimes make allocations out of the shared pool that are larger than 5 KB. To enable these allocations to occur more efficiently, Oracle Database segregates a small amount of the shared pool. The segregated memory, called the reserved pool, is used if the shared pool runs out of space.

The following sections provide more details about the main components of the shared pool:

- Library Cache Concepts
- Data Dictionary Cache Concepts
- SQL Sharing Criteria

Library Cache Concepts

The library cache stores executable forms of SQL cursors, PL/SQL programs, and Java classes, which are collectively referred to as the application code. This section focuses on tuning as it relates to the application code.

When the application code is executed, Oracle Database attempts to reuse existing code if it has been executed previously and can be shared. If the parsed representation of the SQL statement exists in the library cache and it can be shared, then the database reuses the existing code. This is known as a soft parse, or a library cache hit. If Oracle Database cannot use the existing code, then the database must build a new executable version of the application code. This is known as a hard parse, or a library cache miss. For information about when SQL and PL/SQL statements can be shared, see "SQL Sharing Criteria".

In order to perform a hard parse, Oracle Database uses more resources than during a soft parse. Resources used for a soft parse include CPU and library cache latch gets. Resources required for a hard parse include additional CPU, library cache latch gets, and shared pool latch gets. A hard parse may occur on either the parse step or the execute step when processing a SQL statement.

When an application makes a parse call for a SQL statement, if the parsed representation of the statement does not exist in the library cache, then Oracle Database parses the statement



and stores the parsed form in the shared pool. To reduce library cache misses on parse calls, ensure that all sharable SQL statements are stored in the shared pool whenever possible.

When an application makes an execute call for a SQL statement, if the executable portion of the SQL statement is aged out (or deallocated) from the library cache to make room for another statement, then Oracle Database implicitly reparses the statement to create a new shared SQL area for it, and executes the statement. This also results in a hard parse. To reduce library cache misses on execution calls, allocate more memory to the library cache.

For more information about hard and soft parsing, see "SQL Execution Efficiency".

Data Dictionary Cache Concepts

Information stored in the data dictionary cache includes:

- Usernames
- Segment information
- Profile data
- Tablespace information
- Sequence numbers

The data dictionary cache also stores descriptive information, or metadata, about schema objects. Oracle Database uses this metadata when parsing SQL cursors or during the compilation of PL/SQL programs.

SQL Sharing Criteria

Oracle Database automatically determines whether a SQL statement or PL/SQL block being issued is identical to another statement currently in the shared pool.

To compare the text of the SQL statement to the existing SQL statements in the shared pool, Oracle Database performs the following steps:

- 1. The text of the SQL statement is hashed.
 - If there is no matching hash value, then the SQL statement does not currently exist in the shared pool, and a hard parse is performed.
- If there is a matching hash value for an existing SQL statement in the shared pool, then the text of the matched statement is compared to the text of the hashed statement to verify if they are identical.

The text of the SQL statements or PL/SQL blocks must be identical, character for character, including spaces, case, and comments. For example, the following statements cannot use the same shared SQL area:

```
SELECT * FROM employees;
SELECT * FROM Employees;
SELECT * FROM employees;
```

Also, SQL statements that differ only in literals cannot use the same shared SQL area. For example, the following statements do not resolve to the same SQL area:

```
SELECT count(1) FROM employees WHERE manager_id = 121;
SELECT count(1) FROM employees WHERE manager id = 247;
```

The only exception to this rule is when the <code>CURSOR_SHARING</code> parameter is set to <code>FORCE</code>, in which case similar statements can share SQL areas. For information about the costs and benefits involved in cursor sharing, see "Sharing Cursors".

3. The objects referenced in the issued statement are compared to the referenced objects of all existing statements in the shared pool to ensure that they are identical.

References to schema objects in the SQL statements or PL/SQL blocks must resolve to the same object in the same schema. For example, if two users each issue the following SQL statement but they each have their own <code>employees</code> table, then this statement is not considered identical, because the statement references different tables for each user:

```
SELECT * FROM employees;
```

4. Bind variables in the SQL statements must match in name, data type, and length.

For example, the following statements cannot use the same shared SQL area, because the bind variable names are different:

```
SELECT * FROM employees WHERE department_id = :department_id;
SELECT * FROM employees WHERE department id = :dept id;
```

Many Oracle products, such as Oracle Forms and the precompilers, convert the SQL before passing statements to the database. Characters are uniformly changed to uppercase, white space is compressed, and bind variables are renamed so that a consistent set of SQL statements is produced.

The session's environment must be identical.

For example, SQL statements must be optimized using the same optimization goal.

See Also:

Oracle Database Reference for information about the CURSOR_SHARING initialization parameter

Using the Shared Pool

An important purpose of the shared pool is to cache the executable versions of SQL and PL/SQL statements. This enables multiple executions of the same SQL or PL/SQL code to be performed without the resources required for a hard parse, which results in significant reductions in CPU, memory, and latch usage.

The shared pool is also able to support unshared SQL in data warehousing applications, which execute low-concurrency, high-resource SQL statements. In this situation, using unshared SQL with literal values is recommended. Using literal values rather than bind variables enables the optimizer to make good column selectivity estimates, thus providing an optimal data access plan.

In a high-currency online transaction processing (OLTP) system, efficient use of the shared pool significantly reduces the probability of parse-related application scalability issues. There are several ways to ensure efficient use of the shared pool and related resources in an OLTP system:

- Use Shared Cursors
- Use Single-User Logon and Qualified Table Reference
- Use PL/SQL
- Avoid Performing DDL Operations
- Cache Sequence Numbers

- Control Cursor Access
- Maintain Persistent Connections



Oracle Database VLDB and Partitioning Guide for information about impact of parallel query execution on the shared pool

Use Shared Cursors

Reuse of shared SQL for multiple users running the same application avoids hard parsing. Soft parses provide a significant reduction in the use of resources, such as the shared pool and library cache latches.

To use shared cursors:

Use bind variables instead of literals in SQL statements whenever possible.

For example, the following two SQL statements cannot use the same shared area because they do not match character for character:

```
SELECT employee_id FROM employees WHERE department_id = 10; SELECT employee id FROM employees WHERE department id = 20;
```

Replacing the literals with a bind variable results in only one SQL statement which can be executed twice:

```
SELECT employee id FROM employees WHERE department id = :dept id;
```

For existing applications where rewriting the code to use bind variables is not possible, use the <code>CURSOR_SHARING</code> initialization parameter to avoid some of the hard parse overhead, as described in "Sharing Cursors".

 Avoid application designs that result in large numbers of users issuing dynamic, unshared SQL statements.

Typically, the majority of data required by most users can be satisfied using preset queries. Use dynamic SQL where such functionality is required.

- Ensure that users of the application do not change the optimization approach and goal for their individual sessions.
- Establish the following policies for application developers:
 - Standardize naming conventions for bind variables and spacing conventions for SQL statements and PL/SQL blocks.
 - Consider using stored procedures whenever possible.

Multiple users issuing the same stored procedure use the same shared PL/SQL area automatically. Because stored procedures are stored in a parsed form, their use reduces run-time parsing.

For SQL statements which are identical but are not being shared, query the V\$SQL SHARED CURSOR view to determine why the cursors are not shared.

This includes optimizer settings and bind variable mismatches.



See Also:

Oracle Database SQL Tuning Guide for more information about cursor sharing

Use Single-User Logon and Qualified Table Reference

In large OLTP systems where users log in to the database with their own user logon, qualifying the segment owner explicitly instead of using public synonyms significantly reduces the number of entries in the dictionary cache.

An alternative to qualifying table names is to connect to the database through a single user logon, rather than individual user logons. User-level validation can take place locally on the middle tier.

Use PL/SQL

Using stored PL/SQL packages can overcome many of the scalability issues for systems with thousands of users, each with individual user logon and public synonyms. This is because a package is executed as the owner, rather than the caller, which reduces the dictionary cache load considerably.



Oracle encourages the use of definer's rights packages to overcome scalability issues. The benefits of reduced dictionary cache load are not as great with invoker's rights packages.

Avoid Performing DDL Operations

Avoid performing DDL operations on high-usage segments during peak hours. Performing DDL operations on these segments often results in the dependent SQL being invalidated and reparsed in a later execution.

Cache Sequence Numbers

Allocating sufficient cache space for frequently updated sequence numbers significantly reduces the frequency of dictionary cache locks, which improves scalability.

To configure the number of cache entries for each sequence:

• Use the Cache keyword in the Create sequence or alter sequence statement.

Control Cursor Access

Depending on your application tool, you can control how frequently the application performs parse calls by controlling cursor access.

The frequency with which the application either closes cursors or reuses existing cursors for new SQL statements affects the amount of memory used by a session, and often the amount of parsing performed by that session. An application that closes cursors or reuses cursors (for a different SQL statement) does not require as much session memory as an application that keeps cursors open. Conversely, that same application may need to perform more parse calls, using more CPU and database resources

Cursors associated with SQL statements that are not executed frequently can be closed or reused for other statements, because the likelihood of re-executing (and reparsing) that statement is low. Extra parse calls are required when a cursor containing a SQL statement that will be re-executed is closed or reused for another statement. Had the cursor remained open, it may have been reused without the overhead of issuing a parse call.

The ways in which you control cursor access depends on your application development tool. This section describes the methods used for Oracle Database tools:

- Controlling Cursor Access Using OCI
- Controlling Cursor Access Using Oracle Precompilers
- Controlling Cursor Access Using SQLJ
- Controlling Cursor Access Using JDBC
- Controlling Cursor Access Using Oracle Forms



The tool-specific documentation for information about each tool

Controlling Cursor Access Using OCI

When using Oracle Call Interface (OCI), do not close and reopen cursors that you will be reexecuting. Instead, leave the cursors open, and change the literal values in the bind variables before execution.

Avoid reusing statement handles for new SQL statements when the existing SQL statement will be re-executed in the future.

Controlling Cursor Access Using Oracle Precompilers

When using the Oracle precompilers, you can control when cursors are closed by setting precompiler clauses. In Oracle mode, the clauses are as follow:

- HOLD CURSOR = YES
- RELEASE_CURSOR = NO
- MAXOPENCURSORS = desired_value

The precompiler clauses can be specified on the precompiler command line or within the precompiler program. Oracle recommends that you not use ANSI mode, in which the values of HOLD CURSOR and RELEASE CURSOR are switched.

See Also:

Your language's precompiler manual for information about these clauses



Controlling Cursor Access Using SQLJ

Prepare the SQL statement, then re-execute the statement with the new values for the bind variables. The cursor stays open for the duration of the session.



Starting with Oracle Database 12c Release 2 (12.2), server-side SQLJ code is not supported, that is, you cannot use SQLJ code inside stored procedures, functions, and triggers.

Controlling Cursor Access Using JDBC

Avoid closing cursors if they will be re-executed, because the new literal values can be bound to the cursor for re-execution. Alternatively, JDBC provides a SQL statement cache within the JDBC client using the <code>setStmtCacheSize()</code> method. Using this method, JDBC creates a SQL statement cache that is local to the JDBC program.

See Also

Oracle Database JDBC Developer's Guide for information about using the JDBC SQL statement cache

Controlling Cursor Access Using Oracle Forms

With Oracle Forms, it is possible to control some aspects of cursor access at run time, the trigger level, or the form level.

Maintain Persistent Connections

Large OLTP applications with middle tiers should maintain connections, instead of connecting and disconnecting for each database request. Maintaining persistent connections saves CPU resources and database resources, such as latches.

Configuring the Shared Pool

This section describes how to configure the shared pool and contains the following topics:

- Sizing the Shared Pool
- Deallocating Cursors
- Caching Session Cursors
- Sharing Cursors
- Keeping Large Objects to Prevent Aging
- Configuring the Reserved Pool



When configuring a new database instance, it is difficult to know the correct size for the shared pool cache. Typically, a DBA makes a first estimate for the cache size, then runs a representative workload on the instance, and examines the relevant statistics to see whether the cache is under-configured or over-configured.

For most OLTP applications, shared pool size is an important factor in application performance. Shared pool size is less important for applications that issue a very limited number of discrete SQL statements, such as decision support systems (DSS).

If the shared pool is too small, then extra resources are used to manage the limited amount of available space. This consumes CPU and latching resources, and causes contention. Ideally, the shared pool should be just large enough to cache frequently-accessed objects. Having a significant amount of free memory in the shared pool is a waste of memory. When examining the statistics after the database has been running, ensure that none of these mistakes are present in the workload.

This section describes how to size the shared pool and contains the following topics:

- Using Library Cache Statistics
- Using Shared Pool Advisory Statistics
- Using Dictionary Cache Statistics
- Increasing Memory Allocated to the Shared Pool
- Reducing Memory Allocated to the Shared Pool

Using Library Cache Statistics

When sizing the shared pool, the goal is to cache SQL statements that are executed multiple times in the library cache without over-allocating memory. To accomplish this goal, examine the following library cache statistics:

RELOADS

The RELOADS column in the V\$LIBRARYCACHE view shows the amount of reloading (or reparsing) of a previously-cached SQL statement that aged out of the cache. If the application reuses SQL effectively and runs on a system with an optimal shared pool size, this statistic should have a value near zero.

INVALIDATIONS

The INVALIDATIONS column in V\$LIBRARYCACHE view shows the number of times library cache data was invalidated and had to be reparsed. This statistic should have a value near zero, especially on OLTP systems during peak loads. This means SQL statements that can be shared were invalidated by some operation (such as a DDL).

Library cache hit ratio

The library cache hit ratio is a broad indicator of the library cache health. This value should be considered along with the other statistics, such as the rate of hard parsing and if there is any shared pool or library cache latch contention.

Amount of free memory in the shared pool

To view the amount of free memory in the shared pool, query the V\$SGASTAT performance view. Ideally, free memory should be as low as possible, without causing any reparsing on the system.



The following sections describe how to view and examine these library cache statistics:

- Using the V\$LIBRARYCACHE View
- Calculating the Library Cache Hit Ratio
- Viewing the Amount of Free Memory in the Shared Pool

Using the V\$LIBRARYCACHE View

Use the V\$LIBRARYCACHE view to monitor statistics that reflect library cache activity. These statistics reflect all library cache activity after the most recent database instance startup.

Each row in this view contains statistics for one type of item kept in the library cache. The item described by each row is identified by the value of the NAMESPACE column. Rows with the following NAMESPACE values reflect library cache activity for SQL statements and PL/SQL blocks:

- SQL AREA
- TABLE/PROCEDURE
- BODY
- TRIGGER

Rows with other NAMESPACE values reflect library cache activity for object definitions that Oracle Database uses for dependency maintenance.

Example 12-1 shows a query of this view to examine each namespace individually.

Example 12-1 Querying the V\$LIBRARYCACHE View

```
SELECT namespace, pins, pinhits, reloads, invalidations FROM V$LIBRARYCACHE ORDER BY namespace;
```

The output of this query might look like the following:

NAMESPACE	PINS	PINHITS	RELOADS	INVALIDATIONS
BODY	8870	8819	0	0
CLUSTER	393	380	0	0
INDEX	29	0	0	0
OBJECT	0	0	0	0
PIPE	55265	55263	0	0
SQL AREA	21536413	21520516	11204	2
TABLE/PROCEDURE	10775684	10774401	0	0
TRIGGER	1852	1844	0	0

In this example, the output shows that:

- For the SQL AREA namespace, there are 21,536,413 executions.
- 11,204 of these executions resulted in a library cache miss, requiring the database to
 implicitly reparse a statement or block, or reload an object definition because it aged out of
 the library cache.
- SQL statements are invalidated twice, again causing library cache misses.





This query returns data from instance startup. Using statistics gathered over an interval instead may better identify the problem. For information about gathering information over an interval, see Automatic Performance Diagnostics .

See Also:

Oracle Database Reference for information about the V\$LIBRARYCACHE view

Calculating the Library Cache Hit Ratio

To calculate the library cache hit ratio, use the following formula:

```
Library Cache Hit Ratio = sum(pinhits) / sum(pins)
```

Applying the library cache hit ratio formula to Example 12-1 results in the following library cache hit ratio:

In this example, the hit percentage is about 99.94%, which means that only .06% of executions resulted in reparsing.

Viewing the Amount of Free Memory in the Shared Pool

The amount of free memory in the shared pool is reported in the V\$SGASTAT view.

Example 12-2 shows a query of this view.

Example 12-2 Querying the V\$SGASTAT View

```
SELECT *
  FROM V$SGASTAT
WHERE name = 'free memory'
  AND pool = 'shared pool';
```

The output of this query might look like the following:

```
POOL NAME BYTES
-----
shared pool free memory 4928280
```

If free memory is always available in the shared pool, then increasing its size offers little or no benefit. Yet, just because the shared pool is full does not necessarily mean there is a problem. It may be indicative of a well-configured system.

Using Shared Pool Advisory Statistics

The amount of memory available for the library cache can drastically affect the parse rate of Oracle Database. To help you correctly size the library cache, Oracle Database provides the following shared pool advisory views:

- V\$SHARED POOL ADVICE
- V\$LIBRARY CACHE MEMORY
- V\$JAVA POOL ADVICE
- V\$JAVA_LIBRARY_CACHE_MEMORY

These shared pool advisory views provide information about library cache memory, enabling you to predict how changing the size of the shared pool can affect aging out of objects in the shared pool. The shared pool advisory statistics in these views track the library cache's use of shared pool memory and predict how the library cache will behave in shared pools of different sizes. Using these views enable you to determine:

- How much memory the library cache is using
- How much memory is currently pinned
- How much memory is on the shared pool's Least Recently Used (LRU) list
- How much time might be lost or gained by changing the size of the shared pool

These views display shared pool advisory statistics when the shared pool advisory is enabled. The statistics reset when the advisory is disabled.

The following sections describe these views in more detail:

- About the V\$SHARED POOL ADVICE View
- About the V\$LIBRARY CACHE MEMORY View
- About V\$JAVA POOL ADVICE and V\$JAVA LIBRARY CACHE MEMORY Views

About the V\$SHARED POOL ADVICE View

The V\$SHARED_POOL_ADVICE view displays information about estimated parse time in the shared pool for different pool sizes. The sizes range from 10% of the current shared pool size or the amount of pinned library cache memory—whichever is higher—to 200% of the current shared pool size, in equal intervals. The value of the interval depends on the current size of the shared pool.

See Also:

Oracle Database Reference for more information about the V\$SHARED_POOL_ADVICE view

About the V\$LIBRARY_CACHE_MEMORY View

The V\$LIBRARY_CACHE_MEMORY view displays information about memory allocated to library cache memory objects in different namespaces. A memory object is an internal grouping of memory for efficient management. A library cache object may consist of one or more memory objects.



See Also:

Oracle Database Reference for more information about the V\$LIBRARY_CACHE_MEMORY view

About V\$JAVA POOL ADVICE and V\$JAVA_LIBRARY_CACHE_MEMORY Views

The V\$JAVA_POOL_ADVICE and V\$JAVA_LIBRARY_CACHE_MEMORY views contain Java pool advisory statistics that track information about library cache memory used for Java and predict how changing the size of the Java pool affects the parse rate.

The V\$JAVA_POOL_ADVICE view displays information about estimated parse time in the Java pool for different pool sizes. The sizes range from 10% of the current Java pool size or the amount of pinned Java library cache memory—whichever is higher—to 200% of the current Java pool size, in equal intervals. The value of the interval depends on the current size of the Java pool.

See Also:

Oracle Database Reference for more information about the "V\$JAVA_POOL_ADVICE" and "V\$JAVA_LIBRARY_CACHE_MEMORY" views

Using Dictionary Cache Statistics

Typically, if the shared pool is adequately sized for the library cache, it will also be adequate sized for the dictionary cache data.

Misses on the data dictionary cache are to be expected in some cases. When the database instance starts up, the data dictionary cache does not contain any data. Therefore, any SQL statement issued is likely to result in cache misses. As more data is read into the cache, the likelihood of cache misses decreases. Eventually, the database reaches a steady state, in which the most frequently-used dictionary data is in the cache. At this point, very few cache misses occur.

Each row in the V\$ROWCACHE view contains statistics for a single type of data dictionary item. These statistics reflect all data dictionary activity since the most recent instance startup.

Table 12-1 lists the columns in the V\$ROWCACHE view that reflect the use and effectiveness of the data dictionary cache.

Table 12-1 Data Dictionary Columns in the V\$ROWCACHE View

Column	Description
PARAMETER	Identifies a particular data dictionary item. For each row, the value in this column is the item prefixed by dc For example, in the row that contains statistics for file descriptions, this column contains the value dc_files.



Table 12-1 (Cont.) Data Dictionary Columns in the V\$ROWCACHE View

Column	Description
GETS	Shows the total number of requests for information about the corresponding item. For example, in the row that contains statistics for file descriptions, this column contains the total number of requests for file description data.
GETMISSES	Shows the number of data requests that are not satisfied by the cache and required an I/O.
MODIFICATIONS	Shows the number of times data in the dictionary cache was updated.

Example 12-3 shows a query of this view to monitor the statistics over a period while the application is running. The derived column PCT_SUCC_GETS can be considered as the itemspecific hit ratio.

Example 12-3 Querying the V\$ROWCACHE View

```
column parameter format a21
column pct_succ_gets format 999.9
column updates format 999,999,999

SELECT parameter,
        sum(gets),
        sum(getmisses),
        100*sum(gets - getmisses) / sum(gets) pct_succ_gets,
        sum(modifications) updates

FROM V$ROWCACHE
WHERE gets > 0
GROUP BY parameter;
```

The output of this query might look like the following:

PARAMETER	SUM(GETS)	SUM (GETMISSES)	PCT_SUCC_GETS	UPDATES
dc database links	81	1	98.8	0
dc free extents	44876	20301	54.8	40,453
dc global oids	42	9	78.6	0
dc_histogram_defs	9419	651	93.1	0
dc_object_ids	29854	239	99.2	52
dc_objects	33600	590	98.2	53
dc_profiles	19001	1	100.0	0
dc_rollback_segments	47244	16	100.0	19
dc_segments	100467	19042	81.0	40,272
dc_sequence_grants	119	16	86.6	0
dc_sequences	26973	16	99.9	26,811
dc_synonyms	6617	168	97.5	0
dc_tablespace_quotas	120	7	94.2	51
dc_tablespaces	581248	10	100.0	0
dc_used_extents	51418	20249	60.6	42,811
dc_user_grants	76082	18	100.0	0
dc_usernames	216860	12	100.0	0
dc_users	376895	22	100.0	0

In this example, the output shows the following:

 There are large numbers of misses and updates for used extents, free extents, and segments. This implies that the database instance had a significant amount of dynamic space extension. Comparing the percentage of successful gets with the actual number of gets indicates the shared pool is large enough to adequately store dictionary cache data.

You can also calculate the overall dictionary cache hit ratio using the following query; however, summing up the data over all the caches will lose the finer granularity of data:

```
SELECT (SUM(gets - getmisses - fixed)) / SUM(gets) "row cache"
FROM V$ROWCACHE;
```

Increasing Memory Allocated to the Shared Pool

Increasing the amount of memory for the shared pool increases the amount of memory available to the library cache, the dictionary cache, and the result cache. Before doing so, review the shared pool statistics and examine:

- If the value of the V\$LIBRARYCACHE.RELOADS column is near zero
- If the ratio of total V\$ROWCACHE.GETMISSES column to total V\$ROWCACHE.GETS is less than 10% or 15% for frequently-accessed dictionary caches, depending on the application

If both of these conditions are met, then the shared pool is adequately sized and increasing its memory will likely not improve performance. On the other hand, if either of these conditions is not met, and the application is using the shared pool effectively, as described in "Using the Shared Pool", then consider increasing the memory of the shared pool.

To increase the size of the shared pool:

Increase the value of the SHARED_POOL_SIZE initialization parameter until the conditions are
met.

The maximum value for this parameter depends on your operating system. This measure reduces implicit reparsing of SQL statements and PL/SQL blocks on execution.

IC - Need link to "Managing Server Result Cache Memory with Init Parameters"

Reducing Memory Allocated to the Shared Pool

If the value of the V\$LIBRARYCACHE.RELOADS column is near zero, and there is a small amount of free memory in the shared pool, then the shared pool is adequately sized to store the most frequently-accessed data. If there are always significant amounts of free memory in the shared pool and you want to allocate this memory elsewhere, then consider reducing the shared pool size.

To decrease the size of the shared pool

• Reduce the value of the SHARED_POOL_SIZE initialization parameter, while ensuring that good performance is maintained.

Deallocating Cursors

If there are no library cache misses, then consider setting the value of the <code>CURSOR_SPACE_FOR_TIME</code> initialization parameter to <code>TRUE</code> to accelerate execution calls. This parameter specifies whether a cursor can be deallocated from the library cache to make room for a new SQL statement.

If the CURSOR SPACE FOR TIME parameter is set to:

 FALSE (the default), then a cursor can be deallocated from the library cache regardless of whether application cursors associated with its SQL statement are open. In this case, Oracle Database must verify that the cursor containing the SQL statement is in the library cache.

 TRUE, then a cursor can be deallocated only when all application cursors associated with its statement are closed.

In this case, Oracle Database does not need to verify that a cursor is in the library cache because it cannot be deallocated while an application cursor associated with it is open.

Setting the value of the parameter to TRUE saves Oracle Database a small amount of time and may slightly improve the performance of execution calls. This value also prevents the deallocation of cursors until associated application cursors are closed.

Do not set the value of the CURSOR SPACE FOR TIME parameter to TRUE if:

- Library cache misses are found in execution calls
 - Library cache misses indicate that the shared pool is not large enough to hold the shared SQL areas of all concurrently open cursors. If the shared pool does not have enough space for a new SQL statement and the value for this parameter is set to TRUE, then the statement cannot be parsed and Oracle Database returns an error indicating that there is not enough shared memory.
- The amount of memory available to each user for private SQL areas is scarce
 - This value also prevents the deallocation of private SQL areas associated with open cursors. If the private SQL areas for all concurrently open cursors fills the available memory so that there is no space for a new SQL statement, then the statement cannot be parsed and Oracle Database returns an error indicating that there is not enough memory.

If the shared pool does not have enough space for a new SQL statement and the value of this parameter is set to FALSE, then Oracle Database deallocates an existing cursor. Although deallocating a cursor may result in a subsequent library cache miss (if the cursor is reexecuted), this is preferable to an error halting the application because a SQL statement cannot be parsed.

Caching Session Cursors

The session cursor cache contains closed session cursors for SQL and PL/SQL, including recursive SQL. This cache can be useful to applications that use Oracle Forms because switching from one form to another closes all session cursors associated with the first form. If an application repeatedly issues parse calls on the same set of SQL statements, then reopening session cursors can degrade performance. By reusing cursors, the database reduces parse times, leading to faster overall execution times.

This section describes the session cursor cache and contains the following topics:

- About the Session Cursor Cache
- Enabling the Session Cursor Cache
- Sizing the Session Cursor Cache

About the Session Cursor Cache

A session cursor represents an instantiation of a shared child cursor, which is stored in the shared pool, for a specific session. Each session cursor stores a reference to a child cursor that it has instantiated.

Oracle Database checks the library cache to determine whether more than three parse requests have been issued on a given statement. If a cursor has been closed three times, then

Oracle Database assumes that the session cursor associated with the statement should be cached and moves the cursor into the session cursor cache.

Subsequent requests to parse a SQL statement by the same session search an array for pointers to the shared cursor. If the pointer is found, then the database dereferences the pointer to determine whether the shared cursor exists. To reuse a cursor from the cache, the cache manager checks whether the cached states of the cursor match the current session and system environment.



Reuse of a cached cursor still registers as a parse, even though it is not a hard parse.

An LRU algorithm removes entries in the session cursor cache to make room for new entries when needed. The cache also uses an internal time-based algorithm to age out cursors that have been idle for an certain amount of time.

Enabling the Session Cursor Cache

The following initialization parameters pertain to the session cursor cache:

SESSION CACHED CURSORS

This parameter sets the maximum number of cached closed cursors for each session. The default value is 50. Use this parameter to reuse cursors from the cache for the statements that get executed repeatedly in the same session.

OPEN CURSORS

This parameter specifies the maximum number of cursors a session can have open simultaneously. For example, if its value is set to 1000, then each session can have up to 1000 cursors open at one time.

These parameters are independent. For example, you can set the value of the SESSION_CACHED_CURSORS parameter higher than the value of the OPEN_CURSORS parameter because session cursors are not cached in an open state.

To enable the session cursor cache:

- **1.** Determine the maximum number of session cursors to keep in the cache.
- 2. Do one of the following:
 - To enable static caching, set the value of the SESSION_CACHED_CURSORS parameter to the number determined in the previous step.
 - To enable dynamic caching, execute the following statement:

```
ALTER SESSION SET SESSION CACHED CURSORS = value;
```

Sizing the Session Cursor Cache

Use the V\$SESSTAT view to determine if the session cursor cache is adequately sized for the database instance.

To size the session cursor cache:

- 1. Query the V\$SESSTAT view to determine how many cursors are currently cached in a particular session.
- 2. Query the V\$SESSTAT view to find the percentage of parse calls that found a cursor in the session cursor cache.
- Consider increasing the value of the SESSION_CACHED_CURSORS parameter if the following conditions are true:
 - The session cursor cache count is close to the maximum
 - The percentage of session cursor cache hits is low relative to the total parses
 - The application repeatedly performs parse calls for the same queries

Example 12-4 shows two queries of this view.

Example 12-4 Querying the V\$SESSTAT View

The following query finds how many cursors are currently cached in a particular session:

The output of this query might look like the following:

This output shows that the number of cursors currently cached for session 35 is close to the maximum.

The following query finds the percentage of parse calls that found a cursor in the session cursor cache:

The output of this query might look like the following:

This output shows that the number of hits in the session cursor cache for session 35 is low compared to the total number of parses.

In this example, setting the value of the SESSION_CACHED_CURSORS parameter to 100 may help boost performance.



Sharing Cursors

In the context of SQL parsing, an identical statement is a SQL statement whose text is identical to another statement, character for character, including spaces, case, and comments. A similar statement is identical except for the values of some literals.

The parse phase compares the statement text with statements in the shared pool to determine if the statement can be shared. If the value of the <code>CURSOR_SHARING</code> initialization parameter is set to <code>EXACT</code> (the default value), and if a statement in the shared pool is not identical, then the database does not share the SQL area. Instead, each SQL statement has its own parent cursor and its own execution plan based on the literal in the statement.

This section describes how cursors can be shared and contains the following topics:

- About Cursor Sharing
- Forcing Cursor Sharing

About Cursor Sharing

When SQL statements use literals rather than bind variables, setting the value of the CURSOR_SHARING initialization parameter to FORCE enables the database to replace literals with system-generated bind variables. Using this technique, the database may reduce the number of parent cursors in the shared SQL area.

When the value of the <code>CURSOR_SHARING</code> parameter is set to <code>FORCE</code>, the database performs the following steps during the parse phase:

- 1. Searches for an identical statement in the shared pool.
 - If an identical statement is found, then the database skips the next step and proceeds to step 3. Otherwise, the database proceeds to the next step.
- Searches for a similar statement in the shared pool.
 - If a similar statement is *not* found, then the database performs a hard parse. If a similar statement *is* found, then the database proceeds to the next step.
- Proceeds through the remaining steps of the parse phase to ensure that the execution plan of the existing statement is applicable to the new statement.
 - If the plan is not applicable, then the database performs a hard parse. If the plan is applicable, then the database proceeds to the next step.
- 4. Shares the SQL area of the statement.

For details about the various checks performed by the database, see "SQL Sharing Criteria".

Forcing Cursor Sharing

The best practice is to write sharable SQL and use the default value of EXACT for the CURSOR_SHARING initialization parameter. By default, Oracle Database uses adaptive cursor sharing to enable a single SQL statement that contains bind variables to use multiple execution plans. However, for applications with many similar statements that use literals instead of bind variables, setting the value of the CURSOR_SHARING parameter to FORCE may improve cursor sharing, resulting in reduced memory usage, faster parses, and reduced latch contention. Consider this approach when statements in the shared pool differ only in the values of literals, and when response time is poor because of a high number of library cache misses. In this case, setting the value of the CURSOR_SHARING parameter to FORCE maximizes cursor sharing



and leverages adaptive cursor sharing to generate multiple execution plans based on different literal value ranges.

If stored outlines are generated with the value of the <code>CURSOR_SHARING</code> parameter set to <code>EXACT</code>, then the database does not use stored outlines generated with literals. To avoid this problem, generate outlines with <code>CURSOR_SHARING</code> set to <code>FORCE</code> and use the <code>CREATE_STORED_OUTLINES</code> parameter.

Setting the value of the CURSOR SHARING parameter to FORCE has the following drawbacks:

- The database must perform extra work during the soft parse to find a similar statement in the shared pool.
- There is an increase in the maximum lengths (as returned by DESCRIBE) of any selected
 expressions that contain literals in a SELECT statement. However, the actual length of the
 data returned does not change.
- Star transformation is not supported.

When the value of the <code>CURSOR_SHARING</code> parameter is set to <code>FORCE</code>, the database uses one parent cursor and one child cursor for each distinct SQL statement. The same plan is used for each execution of the same statement. For example, consider the following SQL statement:

```
SELECT *
  FROM hr.employees
WHERE employee id = 101;
```

If the value of the CURSOR_SHARING parameter is set to FORCE, then the database optimizes this statement as if it contained a bind variable and uses bind peeking to estimate cardinality.

Note:

Starting with Oracle Database 11*g* Release 2, setting the value of the CURSOR_SHARING parameter to SIMILAR is obsolete. Consider using adaptive cursor sharing instead.

See Also:

- Oracle Database Reference for information about the CURSOR_SHARING initialization parameter
- Oracle Database SQL Tuning Guide for information about adaptive cursor sharing

Keeping Large Objects to Prevent Aging

After an entry is loaded into the shared pool, it cannot be moved. Sometimes, as entries are loaded and aged out, the free memory may become fragmented. Shared SQL and PL/SQL areas age out of the shared pool according to an LRU algorithm that is similar to database buffers. To improve performance and avoid reparsing, prevent large SQL or PL/SQL areas from aging out of the shared pool.

The DBMS_SHARED_POOL package enables you to keep objects in shared memory, so that they do not age out with the normal LRU mechanism. By using the DBMS_SHARED_POOL package to

load the SQL and PL/SQL areas before memory fragmentation occurs, the objects can be kept in memory. This ensures that memory is available and prevents the sudden slowdowns in user response times that occur when SQL and PL/SQL areas are accessed after being aged out.

Consider using the DBMS SHARED POOL package:

- When loading large PL/SQL objects, such as the STANDARD and DIUTIL packages.
 - When large PL/SQL objects are loaded, user response time may be affected if smaller objects must age out of the shared pool to make room for the larger objects. In some cases, there may be insufficient memory to load the large objects.
- To keep compiled triggers on frequently used tables in the shared pool.
- Support sequences.

Sequence numbers are lost when a sequence ages out of the shared pool. The <code>DBMS_SHARED_POOL</code> package keeps sequences in the shared pool, thus preventing the loss of sequence numbers.

To keep a SQL or PL/SQL area in shared memory:

- Decide which packages or cursors to keep in memory.
- 2. Start up the database.
- 3. Call the DBMS SHARED POOL.KEEP package to pin the objects.

This procedure ensures that the system does not run out of shared memory before the pinned objects are loaded. Pinning the objects early in the life of the database instance prevents memory fragmentation that may result from keeping a large portion of memory in the middle of the shared pool.



Oracle Database PL/SQL Packages and Types Reference for information about using $DBMS_SHARED_POOL$ procedures

Configuring the Reserved Pool

Although Oracle Database breaks down very large requests for memory into smaller chunks, on some systems there may be a requirement to find a contiguous chunk of memory (such as over 5 KB, the default minimum reserved pool allocation is 4,400 bytes).

If there is not enough free space in the shared pool, then Oracle Database must search for and free enough memory to satisfy this request. This operation may hold the latch resource for significant periods of time, causing minor disruption to other concurrent attempts at memory allocation.

To avoid this, Oracle Database internally reserves a small memory area in the shared pool by default that the database can use if the shared pool does not have enough space. This reserved pool makes allocation of large chunks more efficient. The database can use this memory for operations such as PL/SQL and trigger compilation, or for temporary space while loading Java objects. After the memory allocated from the reserved pool is freed, it is returned to the reserved pool.

For large allocations, Oracle Database attempts to allocate space in the shared pool in the following order:

- From the unreserved part of the shared pool.
- 2. From the reserved pool.

If there is not enough space in the unreserved part of the shared pool, then Oracle Database checks whether the reserved pool has enough space.

3. From memory.

If there is not enough space in the unreserved and reserved parts of the shared pool, then Oracle Database attempts to free enough memory for the allocation. The database then retries the unreserved and reserved parts of the shared pool.

This section describes how to configure the reserved pool and contains the following topics:

- Sizing the Reserved Pool
- Increasing Memory Allocated to the Reserved Pool
- Reducing Memory Allocated to the Reserved Pool

Sizing the Reserved Pool

Typically, it is not necessary to change the default amount of space Oracle Database reserves for the reserved pool. However, there may be cases where you need to set aside space in the shared pool for unusually large allocations of memory.

You can set the reserved pool size by setting the value of the <code>SHARED_POOL_RESERVED_SIZE</code> initialization parameter. The default value for the <code>SHARED_POOL_RESERVED_SIZE</code> parameter is 5% of the <code>SHARED_POOL_SIZE</code> parameter.

If you set the value of the <code>SHARED_POOL_RESERVED_SIZE</code> parameter to more than half of the <code>SHARED_POOL_SIZE</code> parameter, then Oracle Database returns an error because the database does not allow you to reserve too much memory for the reserved pool. The amount of operating system memory available may also constrain the size of the shared pool. In general, set the value of the <code>SHARED_POOL_RESERVED_SIZE</code> parameter to no higher than 10% of the <code>SHARED_POOL_SIZE</code> parameter. On most systems, this value is sufficient if the shared pool is adequately tuned. If you increase this value, then the database takes additional memory from the shared pool and reduces the amount of unreserved shared pool memory available for smaller allocations.

When tuning these parameters, use statistics from the V\$SHARED_POOL_RESERVED view. On a system with ample free memory to increase the size of the SGA, the value of the REQUEST_MISSES statistic should equal zero. If the system is constrained by operating system memory, then the goal is to have the REQUEST_FAILURES statistic equal zero, or at least prevent its value from increasing. If you cannot achieve these target values, then increase the value of the SHARED_POOL_RESERVED_SIZE parameter. Also, increase the value of the SHARED_POOL_SIZE parameter by the same amount, because the reserved list is taken from the shared pool.

The V\$SHARED_POOL_RESERVED fixed view can also indicate when the value of the SHARED_POOL_SIZE parameter is too small. This can be the case if the REQUEST_FAILURES statistic is greater than zero and increasing. If the reserved list is enabled, then decrease the value of the SHARED_POOL_RESERVED_SIZE parameter. If the reserved list is not enabled, then increase the value of the SHARED_POOL_SIZE parameter, as described in "Increasing Memory Allocated to the Shared Pool".

Increasing Memory Allocated to the Reserved Pool

The reserved pool is too small if the value of the REQUEST_FAILURES statistic is higher than zero and increasing. In this case, increase the amount of memory available to the reserved pool.

Note:

Increasing the amount of memory available on the reserved list does not affect users who do not allocate memory from the reserved list.

To increase the size of the reserved pool:

• Increase the value of the SHARED_POOL_RESERVED_SIZE and SHARED_POOL_SIZE initialization parameters accordingly.

The values that you select for these parameters depend on the system's SGA size constraints, as described in "Sizing the Reserved Pool".

Reducing Memory Allocated to the Reserved Pool

The reserved pool is too large if the:

- REQUEST MISSES statistic is zero or not increasing
- FREE_SPACE statistic is greater than or equal to 50% of the SHARED_POOL_RESERVED_SIZE minimum

If either of these conditions is true, then reduce the amount of memory available to the reserved pool.

To reduce the size of the reserved pool:

Decrease the value of the SHARED_POOL_RESERVED_SIZE initialization parameter.

Configuring the Large Pool

Unlike the shared pool, the large pool does not have an LRU list. Oracle Database does not attempt to age objects out of the large pool. Consider configuring a large pool if the database instance uses any of the following Oracle Database features:

Shared server

In a shared server architecture, the session memory for each client process is included in the shared pool.

Parallel query

Parallel query uses shared pool memory to cache parallel execution message buffers.

Recovery Manager

Recovery Manager (RMAN) uses the shared pool to cache I/O buffers during backup and restore operations. For I/O server processes, backup, and restore operations, Oracle Database allocates buffers that are a few hundred kilobytes in size.

This section describes how to configure the large pool for the shared server architecture and contains the following topics:

- Configuring the Large Pool for Shared Server Architecture
- Configuring the Large Pool for Parallel Query
- Sizing the Large Pool
- Limiting Memory Use for User Sessions



Reducing Memory Use Using Three-Tier Connections

See Also:

- Oracle Database Concepts for information about the large pool
- Oracle Database Backup and Recovery User's Guide for information about sizing the large pool with Recovery Manager

Configuring the Large Pool for Shared Server Architecture

As Oracle Database allocates shared pool memory to shared server sessions, the amount of shared pool memory available for the library cache and data dictionary cache decreases. If you allocate the shared server session memory from a different pool, then the shared pool can be reserved for caching shared SQL.

Oracle recommends using the large pool to allocate the User Global Area (UGA) for the shared server architecture. Using the large pool instead of the shared pool decreases fragmentation of the shared pool and eliminates the performance overhead from shrinking the shared SQL cache.

By default, the large pool is not configured. If you do not configure the large pool, then Oracle Database uses the shared pool for shared server user session memory. If you do configure the large pool, Oracle Database still allocates a fixed amount of memory (about 10K) for each configured session from the shared pool when a shared server architecture is used. In either case, consider increasing the size of the shared pool accordingly.



Even though use of shared memory increases with shared servers, the total amount of memory use decreases. This is because there are fewer processes; therefore, Oracle Database uses less PGA memory with shared servers when compared to dedicated server environments.

0

Tip:

To specify the maximum number of concurrent shared server sessions that the database allows, use the CIRCUITS initialization parameter.



Tip:

For best performance with sort operations using shared servers, set the values of the <code>SORT_AREA_SIZE</code> and <code>SORT_AREA_RETAINED_SIZE</code> initialization parameters to the same value. This keeps the sort result in the large pool instead of writing it to disk.



Configuring the Large Pool for Parallel Query

Parallel query uses shared pool memory to cache parallel execution message buffers when Automatic Memory Management or Automatic Shared Memory Management is not enabled. Caching parallel execution message buffers in the shared pool increases its workload and may cause fragmentation.

To avoid possible negative impact to performance, Oracle recommends that you do not manage SGA memory manually when parallel query is used. Instead, you should enable Automatic Memory Management or Automatic Shared Memory Management to ensure that the large pool will be used to cache parallel execution memory buffers.

See Also:

- "Automatic Memory Management"
- "Automatic Shared Memory Management"
- Oracle Database VLDB and Partitioning Guide

Sizing the Large Pool

When storing shared server-related UGA in the large pool, the exact amount of UGA that Oracle Database uses depends on the application. Each application requires a different amount of memory for session information, and configuration of the large pool should reflect the memory requirement.

Oracle Database collects statistics on memory used by a session and stores them in the V\$SESSTAT view. Table 12-2 lists the statistics from this view that reflect session UGA memory.

Table 12-2 Memory Statistics in the V\$SESSTAT View

Statistic	Description
session UGA memory	Shows the amount of memory in bytes allocated to the session.
session UGA memory max	Shows the maximum amount of memory in bytes ever allocated to the session.

There are two methods to use this view to determine a correct size for the large pool. One method is to configure the size of the large pool based on the number of simultaneously active sessions. To do this, observe UGA memory usage for a typical user and multiply this amount by the estimated number of user sessions. For example, if the shared server requires 200K to 300K to store session information for a typical user session and you anticipate 100 active user sessions simultaneously, then configure the large pool to 30 MB.

Another method is to calculate the total and maximum memory being used by all user sessions. Example 12-5 shows two queries of the V\$SESSTAT and V\$STATNAME views to do this.

Example 12-5 Querying the V\$SESSTAT and V\$STATNAME Views

While the application is running, issue the following queries:

```
SELECT SUM(value) || ' bytes' "total memory for all sessions" FROM V$SESSTAT, V$STATNAME
```



```
WHERE name = 'session uga memory'
AND V$SESSTAT.STATISTIC# = V$STATNAME.STATISTIC#;

SELECT SUM(value) || ' bytes' "total max mem for all sessions"
FROM V$SESSTAT, V$STATNAME
WHERE name = 'session uga memory max'
AND V$SESSTAT.STATISTIC# = V$STATNAME.STATISTIC#;
```

These queries also select from the V\$STATNAME view to obtain internal identifiers for session memory and max session memory.

The output of these queries might look like the following:

The result of the first query shows that the memory currently allocated to all sessions is 157,125 bytes. This value is the total memory with a location that depends on how the sessions are connected to the database. If the sessions are connected to dedicated server processes, then this memory is part of the memories of the user processes. If the sessions are connected to shared server processes, then this memory is part of the shared pool.

The result of the second query shows that the sum of the maximum size of the memory for all sessions is 417,381 bytes. The second result is greater than the first because some sessions have deallocated memory since allocating their maximum amounts.

Use the result of either queries to determine the correct size for the shared pool. The first value is likely to be a better estimate than the second, unless if you expect all sessions to reach their maximum allocations simultaneously.

To size the large pool:

- 1. Verify the pool (shared pool or large pool) in which the memory for an object resides by checking the POOL column in the V\$SGASTAT view.
- 2. Set a value for the LARGE POOL SIZE initialization parameter.

The minimum value for this parameter is 300K.

Limiting Memory Use for User Sessions

To restrict the memory used by each client session from the SGA, set a resource limit using PRIVATE_SGA.

PRIVATE_SGA defines the number of bytes of memory used from the SGA by a session. However, this parameter is rarely used, because most DBAs do not limit SGA consumption on a user-by-user basis.



Oracle Database SQL Language Reference for information about setting the PRIVATE SGA resource limit

Reducing Memory Use Using Three-Tier Connections

If there is a high number of connected users, then consider reducing memory usage by implementing three-tier connections. Using a transaction process (TP) monitor is feasible only with pure transactional models because locks and uncommitted DML operations cannot be held between calls.

Using a shared server environment:

- Is much less restrictive of the application design than a TP monitor.
- Dramatically reduces operating system process count and context switches by enabling users to share a pool of servers.
- Substantially reduces overall memory usage, even though more SGA is used in shared server mode.



Tuning the Result Cache

This chapter describes how to tune the result cache and contains the following topics:

- About the Result Cache
- · Configuring the Result Cache
- Specifying Queries for Result Caching
- Monitoring the Result Cache

About the Result Cache

A result cache is an area of memory, either in the Shared Global Area (SGA) or client application memory, that stores the results of a database query or query block for reuse. The cached rows are shared across SQL statements and sessions unless they become stale.

Server Result Cache Concepts

The server result cache is a memory pool within the shared pool. This memory pool consists of the SQL query result cache—which stores results of SQL queries—and the PL/SQL function result cache, which stores values returned by PL/SQL functions.

This section describes the server result cache and contains the following topics:

- · Benefits of Using the Server Result Cache
- Understanding How the Server Result Cache Works

See Also:

- Oracle Database Concepts for information about the server result cache
- Oracle Database PL/SQL Language Reference for information about the PL/SQL function result cache

Benefits of Using the Server Result Cache

The benefits of using the server result cache depend on the application. OLAP applications can benefit significantly from its use. Good candidates for caching are queries that access a high number of rows but return a small number, such as those in a data warehouse. For example, you can use advanced query rewrite with equivalences to create materialized views that materialize queries in the result cache instead of using tables.



Oracle Database Data Warehousing Guide for information about using the result cache and advance query rewrite with equivalences

Understanding How the Server Result Cache Works

When a query executes, the database searches the cache memory to determine whether the result exists in the result cache. If the result exists, then the database retrieves the result from memory instead of executing the query. If the result is not cached, then the database executes the query, returns the result as output, and stores the result in the result cache.

When users execute queries and functions repeatedly, the database retrieves rows from the cache, decreasing response time. Cached results become invalid when data in dependent database objects is modified.

The following sections contains examples of how to retrieve results from the server result cache:

- How Results are Retrieved in a Query
- How Results are Retrieved in a View

How Results are Retrieved in a Query

The following example shows a query of hr.employees that uses the RESULT_CACHE hint to retrieve rows from the server result cache.

```
SELECT /*+ RESULT_CACHE */ department_id, AVG(salary)
FROM hr.employees
GROUP BY department id;
```

A portion of the execution plan of this guery might look like the following:

	Ic	d	Operation		Name	Rows
 	0 1 2 3	 - -	SELECT STATEMENT RESULT CACHE HASH GROUP BY TABLE ACCESS	İ	8fpza04gtwsfr6n595au15yj4y EMPLOYEES	11 11 107

In this example, the results are retrieved directly from the cache, as indicated in step 1 of the execution plan. The value in the Name column is the cache ID of the result.

The following example shows a query of the V\$RESULT_CACHE_OBJECTS view to retrieve detailed statistics about the cached result.

In this example, the value of CACHE_ID is the cache ID obtained from the explain plan in the earlier example. The output of this query might look like the following:

ID TY	YPE CREA	ATION_ BLOCK	_COUNT COLU	UMN_COUNT 1	PIN_COUNT	ROW_COUNT
2 Re	esult 06-1	NOV-11	1	2	0	12

How Results are Retrieved in a View

Example 13-1 shows a query that uses the RESULT CACHE hint within a WITH clause view.

Example 13-1 RESULT_CACHE Hint Specified in a WITH View

```
WITH summary AS
( SELECT /*+ RESULT_CACHE */ department_id, avg(salary) avg_sal
    FROM hr.employees
    GROUP BY department_id )
SELECT d.*, avg_sal
    FROM hr.departments d, summary s
WHERE d.department_id = s.department_id;
```

A portion of the execution plan of this query might look like the following:

Id Operatio	n	Name	 	Rows		Bytes		Cost	(%CPU)	Ti	me	
		8nknkh64ctmz94a5muf2tyb8r		11 11 11 11		517 517 286	 	7 7 4 4	(29) (25)	00	:00:01 :00:01 :00:01	
	LE ACCESS FULL ACCESS FULL	EMPLOYEES DEPARTMENTS	 	107 27		749 567		3 2	(0) (0)		:00:01 :00:01	

In this example, the <code>summary</code> view results are retrieved directly from the cache, as indicated in step 3 of the execution plan. The value in the <code>Name</code> column is the cache ID of the result.

Client Result Cache Concepts

The Oracle Call Interface (OCI) client result cache is a memory area inside a client process that caches SQL query result sets for OCI applications. This client cache exists for each client process and is shared by all sessions inside the process. Oracle recommends client result caching for queries of read-only or read-mostly tables.



The client result cache is distinct from the server result cache, which resides in the SGA. When client result caching is enabled, the query result set can be cached on the client, server, or both. Client caching can be enabled even if the server result cache is disabled.

This section describes the client result cache and contains the following topics:

- Benefits of Using the Client Result Cache
- Understanding How the Client Result Cache Works

Benefits of Using the Client Result Cache

OCI drivers, such as OCCI, the JDBC OCI driver, and ODP.NET, support client result caching. Performance benefits of using the client result cache include:

· Reduced query response time

When queries are executed repeatedly, the application retrieves results directly from the client cache memory, resulting in faster query response time.

More efficient use of database resources

The reduction in server round trips may result in substantial performance savings of server resources, such as server CPU and I/O. These resources are freed for other tasks, thereby making the server more scalable.

Reduced memory cost

The result cache uses client memory, which may be less expensive than server memory.

Understanding How the Client Result Cache Works

The client result cache stores the results of the outermost query, which are the columns defined by the OCI application. Subqueries and query blocks are not cached.

The following figure illustrates a client process with a database login session. This client process has one client result cache shared amongst multiple application sessions running in the client process. If the first application session runs a query, then it retrieves rows from the database and caches them in the client result cache. If other application sessions run the same query, then they also retrieve rows from the client result cache.

Client Server **Client Process Client Result Cache** SELECT department FROM departments Result Set Application Keeps Session Consistent 10,20,30,40,. **Database** Application Session SELECT department FROM departments

Figure 13-1 Client Result Cache

The client result cache transparently keeps the result set consistent with session state or database changes that affect it. When a transaction changes the data or metadata of database objects used to build the cached result, the database sends an invalidation to the OCI client on its next round trip to the server.



See Also:

Oracle Call Interface Programmer's Guide for details about the client result cache

Brownouts and Automatic Exclusion for the Result Cache

Brownouts are situations such as when an excess of broadcasted invalidation messages overwhelms the system.

Adaptive result cache object exclusion, or automatic exclusion, or simply block listing, is a feature in 23ai to detect potential brownout situations, identify the object, and exclude that object from the result cache.

Block listing of objects that are not worth caching is enabled with the init.ora parameter RESULT_CACHE_AUTO_BLOCKLIST. When enabled, the database periodically finds hot tables in the background and block list them.

To learn why an object was block listed, use the view V\$SQL and its RESULT CACHE REJECTION REASON column for the object's rejection code.

Note:

The section, Block Listing Reasons, has a table of rejection codes, their meaning, and recommended remedial actions. The same table is published in the *Oracle Database Reference* near the V SQL parameter list.

Block listing can be added to with DBMS_RESULT_CACHE.BLOCKLIST_ADD. Block listing can be overridden by adding the object to the ignore list with DBMS_RESULT_CACHE.IGNORE_LIST. Objects in the ignore lists are never considered for auto block listing. Objects can be individually added and removed from the ignore list, and the ignore list itself can be cleared.

The V_RESULT_CACHE_OBJECTS displays statistics for ${\tt TIME_SAVED}$ and ${\tt INVALIDATION_COST}$.

The following are some of the statistics collected and used to block list a given object.

- Invalidation message round trip time.
- Number of recent invalidations per object...
- Total time saved per object.
- Total cost of invalidations

Configuring the Result Cache

This section describes how to configure the server and client result cache and contains the following topics:

- Configuring the Server Result Cache
- Configuring the Client Result Cache



- Setting the Result Cache Mode
- Requirements for the Result Cache

Configuring the Server Result Cache

By default, on database startup, Oracle Database allocates memory to the server result cache in the shared pool. The memory size allocated depends on the memory size of the shared pool and the selected memory management system:

Adaptive result cache object exclusion

Block listing of objects that are not worth caching is enabled with the ora.init parameter RESULT_CACHE_AUTO_BLOCKLIST. When enabled, a background action group runs periodically to find hot tables and block list them.

Automatic shared memory management

If you are managing the size of the shared pool using the SGA_TARGET initialization parameter, Oracle Database allocates 0.50% of the value of the SGA_TARGET parameter to the result cache.

Manual shared memory management

If you are managing the size of the shared pool using the <code>SHARED_POOL_SIZE</code> initialization parameter, then Oracle Database allocates 1% of the shared pool size to the result cache.



Oracle Database will not allocate more than 75% of the shared pool to the server result cache.

The size of the server result cache grows until it reaches the maximum size. Query results larger than the available space in the cache are not cached, but can spill to temp if enabled. (Refer to Setting the Result Cache Mode.) The database employs a Least Recently Used (LRU) algorithm to age out cached results, but does not otherwise automatically release memory from the server result cache.

This section describes how to configure the server result cache in memory and contains the following topics:

- Sizing the Server Result Cache Using Initialization Parameters
- Managing the Server Result Cache Using DBMS_RESULT_CACHE

Sizing the Server Result Cache Using Initialization Parameters

Table 13-1 lists the database initialization parameters that control the server result cache.

Table 13-1 Server Result Cache Initialization Parameters related to Memory

Parameter	Description				
RESULT_CACHE_MAX_SIZE	Specifies the memory allocated to the server result cache. To disable the server result cache, set this parameter to 0.				



Table 13-1 (Cont.) Server Result Cache Initialization Parameters related to Memory

Parameter	Description
RESULT_CACHE_MAX_RESULT	Specifies the maximum amount of server result cache memory (in percent) that can be used for a single result. Valid values are between 1 and 100. The default value is 5%. You can set this parameter at the system or session level.
RESULT_CACHE_REMOTE_EXPIRATION	Specifies the expiration time (in minutes) for a result in the server result cache that depends on remote database objects. The default value is 0, which specifies that results using remote objects will not be cached. If a non-zero value is set for this parameter, DML on the remote database does not invalidate the server result cache.



Oracle Database Reference for more information about these initialization parameters

To change the memory allocated to the server result cache:

• Set the value of the RESULT CACHE MAX SIZE initialization parameter to the desired size.

In an Oracle Real Application Clusters (Oracle RAC) environment, the result cache is specific to each database instance and can be sized differently on each instance. However, invalidations work across instances. To disable the server result cache in a cluster, you must explicitly set this parameter to 0 for each instance startup.

Managing the Server Result Cache Using DBMS_RESULT_CACHE

The <code>DBMS_RESULT_CACHE</code> package provides statistics, information, and operators that enable you to manage memory allocation for the server result cache. Use the <code>DBMS_RESULT_CACHE</code> package to perform operations such as retrieving statistics on the cache memory usage and flushing the cache.

Viewing Memory Usage Statistics for the Server Result Cache

This section describes how to view memory allocation statistics for the result cache using the $\texttt{DBMS}_\texttt{RESULT}_\texttt{CACHE}$ package.

To view memory usage statistics for the result cache:

Execute the DBMS RESULT CACHE.MEMORY REPORT procedure.

Example 13-2 shows an execution of this procedure.

Example 13-2 Using the DBMS_RESULT_CACHE Package

SQL> SET SERVEROUTPUT ON SQL> EXECUTE DBMS_RESULT_CACHE.MEMORY_REPORT

The output of this command might look like the following:

```
R e s u l t C a c h e M e m o r y R e p o r t
[Parameters]
Block Size = 1024 bytes
Maximum Cache Size = 950272 bytes (928 blocks)
Maximum Result Size = 47104 bytes (46 blocks)
[Memory]
Total Memory = 46340 bytes [0.048% of the Shared Pool]
... Fixed Memory = 10696 bytes [0.011% of the Shared Pool]
... State Object Pool = 2852 bytes [0.003% of the Shared Pool]
... Cache Memory = 32792 bytes (32 blocks) [0.034% of the Shared Pool]
..... Unused Memory = 30 blocks
..... Used Memory = 2 blocks
...... Dependencies = 1 blocks
...... Results = 1 blocks
...... SQL = 1 blocks
```

PL/SQL procedure successfully completed.

Flushing the Server Result Cache

This section describes how to remove all existing results and purge the result cache memory using the DBMS RESULT CACHE package.

To flush the server result cache:

Execute the DBMS RESULT CACHE.FLUSH procedure.



Oracle Database PL/SQL Packages and Types Reference for information about the DBMS RESULT CACHE package

Configuring the Client Result Cache

Table 13-2 lists the database initialization parameters that enable or influence the behavior of the client result cache.

Table 13-2 Client Result Cache Initialization Parameters

Parameter	Description
CLIENT_RESULT_CACHE_SIZE	Specifies the maximum size of the client result cache for each client process. To enable the client result cache, set the size to 32768 bytes or greater. A lesser value, including the default of 0, disables the client result cache.
	Note: If the CLIENT_RESULT_CACHE_SIZE setting disables the client cache, then a client node cannot enable it. If the CLIENT_RESULT_CACHE_SIZE setting enables the client cache, however, then a client node can override the setting. For example, a client node can disable client result caching or increase the size of its cache.

Table 13-2 (Cont.) Client Result Cache Initialization Parameters

Parameter	Description
CLIENT_RESULT_CACHE_LAG	Specifies the amount of lag time (in milliseconds) for the client result cache. The default value is 3000 (3 seconds). If the OCI application does not perform any database calls for a period of time, then this setting forces the next statement execution call to check for validations.
	If the OCI application accesses the database infrequently, then setting this parameter to a low value results in more round trips from the OCI client to the database to keep the client result cache synchronized with the database.
COMPATIBLE	Specifies the release with which Oracle Database must maintain compatibility.

An optional client configuration file overrides client result cache initialization parameters set in the server parameter file.



The client result cache lag can only be set with the ${\tt CLIENT_RESULT_CACHE_LAG}$ initialization parameter.

See Also:

- Oracle Call Interface Programmer's Guide for information about the parameters that can be set in the client configuration file
- Oracle Database Reference for more information about these client result cache initialization parameters

Setting the Result Cache Mode

The result cache mode is a database setting that determines which queries are eligible to store result sets in the server and client result caches. If a query is eligible for caching, then the application checks the result cache to determine whether the query result set exists in the cache. If it exists, then the result is retrieved directly from the result cache. Otherwise, the database executes the query and returns the result as output and stores it in the result cache. Oracle recommends result caching for queries of read-only or read-mostly database objects.

To set the result cache mode:

• Set the value of the RESULT_CACHE_MODE initialization parameter to determine the behavior of the result cache.

You can set this parameter for the instance (ALTER SYSTEM), session (ALTER SESSION), or in the server parameter file.

Table 13-3 describes the values for this parameter.

• The RESULT_CACHE_MAX_TEMP_SIZE parameter controls the maximum amount of temporary tablespace that the result cache will consume in a database. Defaults to 10 times the default or initialized value of RESULT_CACHE_MAX_SIZE. This can only be modified on the system level, not the session.

In addition, this value cannot be modified to be lower than RESULT CACHE MAX TEMP RESULT.

• The RESULT_CACHE_MAX_TEMP_RESULT parameter controls the maximum size of temporary tablespace that one cached query result can consume. This is similar to RESULT_CACHE_MAX_RESULT. This value cannot be modified to be higher than RESULT_CACHE_MAX_TEMP_SIZE. The default is 5% of the default or initialized value of RESULT_CACHE_MAX_TEMP_SIZE. This can only be modified on the system level, not the session.

Note:

When any /+ result_cache */ hint is used, it overrides the value of result cache mode.

Do not over-zealously increase the value of RESULT_CACHE_MAX_TEMP. A large allocation of temporary space for result cache can increase the size of the temporary tablespace significantly and reduce the amount of temporary space for database operations such as hash joins, sorts, and user-created temporary tables.

The recommendation is to only use temporary segments for result caching through the usage of hints. Creating temporary segments on disk incur additional write operations that can be measurable in highly volatile environments. Using hints ensures this functionality is deployed for only queries that are known to be expensive to compute, reused often, and mostly non-volatile.

The result_cache and corresponding negative (no_result_cache) hints are query-block level. A set query (union, union-all, minus, etc.) has naturally a top-level query block. To perform a set query in any branch query block (such as the immediately-enclosing set query block), use the scope=current hint.

Specifying <code>scope=toplevel</code> moves the hint to the top-most query block, which is the query/ view level. For <code>result_cache</code> hints with <code>scope=toplevel</code> that are specified inside views (real schema-level views), the hints are applied to the top query block of the view and not to the top query block of the actual query invoking the view.

Example: Assume result_cache_mode=force, and we have a union-all query that we don't want to have cached.

```
select /*+ no_result_cache(scope=toplevel) */ * from dual union all
select * from dual;
```

In the above example, the no result cache hint can be any branch of the union-all.

Table 13-3 Values for the RESULT_CACHE_MODE Parameter

Value	Description
MANUAL	Query results can only be stored in the result cache by using a query hint or table annotation. This is the default and recommended value.



Table 13-3 (Cont.) Values for the RESULT_CACHE_MODE Parameter

Value	Description
FORCE	All results are stored in the result cache. If a query result is not in the cache, then the database executes the query and stores the result in the cache. Subsequent executions of the same SQL statement that include the hint /*+ RESULT_CACHE */ retrieve data from the cache. Sessions use these results if possible. To exclude query results from the cache, the /*+ NO_RESULT_CACHE */ query hint must be used.
	Note: FORCE mode is not recommended because the database and clients will attempt to cache all queries, which may create significant performance and latching overhead. Moreover, because queries that call non-deterministic PL/SQL functions are also cached, enabling the result cache in such a broad-based manner may cause material changes to the results.
MANUAL_TEMP	Query results can only be stored in the result cache by using a query hint or table annotation.
	All hinted queries are allowed to leverage temporary segments on disk unless explicitly prohibited by using the /*+ result_cache (temp=false) */ hint
FORCE_TEMP	All results are stored in the result cache. All queries are allowed to leverage temporary segments on disk unless explicitly prohibited by a hint.
SCOPE	The optional scope=current (default) parameter for hint allows a set query to be performed on the immediately-enclosing set query block.
	Specifying $scope=toplevel$ moves the hint to the top-most query block, which is the query/view level.

Note:

When the result cache is enabled, the database also caches queries that call non-deterministic PL/SQL functions. When caching <code>SELECT</code> statements that call such functions, the result cache tracks data dependencies for the PL/SQL functions and the database objects. However, if the function uses data that are not being tracked (such as sequences, <code>SYSDATE</code>, <code>SYS_CONTEXT</code>, and package variables), using the result cache on queries that call this function can produce stale results. In this regard, the behavior of the result cache is identical to caching PL/SQL functions. Therefore, always consider data accuracy, as well as performance, when choosing to enable the result cache.

✓ See Also:

Oracle Database Reference for information about the RESULT_CACHE_MODE initialization parameter.

Requirements for the Result Cache

Enabling the result cache does not *guarantee* that a specific result set will be included in the server or client result cache. In order for results to be cached, the following requirements must be met:

- Read Consistency Requirements
- Query Parameter Requirements
- · Restrictions for the Result Cache

Read Consistency Requirements

For a snapshot to be reusable, it must have read consistency. For a result set to be eligible for caching, at least one of the following conditions must be true:

- The read-consistent snapshot used to build the result must retrieve the most current, committed state of the data.
- The guery points to an explicit point in time using flashback guery.

If the current session has an active transaction referencing objects in a query, then the results from this query are not eligible for caching.

Query Parameter Requirements

Cache results can be reused if they are parameterized with variable values when queries are equivalent and the parameter values are the same. Different values or bind variable names may cause cache misses. Results are parameterized if any of the following constructs are used in the query:

- Bind variables
- The SQL functions DBTIMEZONE, SESSIONTIMEZONE, USERENV/SYS_CONTEXT (with constant variables), UID, and USER
- NLS parameters

Restrictions for the Result Cache

Results cannot be cached when the following objects or functions are in a query:

- Temporary tables and tables in the SYS or SYSTEM schemas
- Sequence CURRVAL and NEXTVAL pseudo columns
- SQL functions CURRENT_DATE, CURRENT_TIMESTAMP, LOCAL_TIMESTAMP, USERENV/ SYS_CONTEXT (with non-constant variables), SYS_GUID, SYSDATE, and SYSTIMESTAMP

The client result cache has additional restrictions for result caching.



Oracle Call Interface Programmer's Guide for information about additional restrictions for the client result cache

Specifying Queries for Result Caching

This section describes how to specify queries for result caching and contains the following topics:

Using SQL Result Cache Hints



Using Result Cache Table Annotations

Using SQL Result Cache Hints

Use result cache hints at the application level to control caching behavior. The SQL result cache hints take precedence over the result cache mode and result cache table annotations.



Oracle Database SQL Language Reference for information about the RESULT_CACHE and NO RESULT CACHE hints, and to specify whether or not the result can spill to disk.

Using the RESULT_CACHE Hint

When the result cache mode is MANUAL, the /*+ RESULT_CACHE */ hint instructs the database to cache the results of a query block and to use the cached results in future executions.

Example 13-3 shows a query that uses the RESULT CACHE hint.

Example 13-3 Using the RESULT_CACHE Hint

```
SELECT /*+ RESULT_CACHE (TEMP=true) */ prod_id, SUM(amount_sold)
  FROM sales
GROUP BY prod_id
ORDER BY prod id;
```

In this example, the query instructs the database to cache rows for a query of the sales table and allows the database to store the result on disk in the temporary tablespace.

Using the NO RESULT CACHE Hint

The /*+ NO_RESULT_CACHE */ hint instructs the database *not* to cache the results in either the server or client result caches.

Example 13-4 shows a query that uses the NO RESULT CACHE hint.

Example 13-4 Using the NO_RESULT_CACHE Hint

```
SELECT /*+ NO_RESULT_CACHE */ prod_id, SUM(amount_sold)
  FROM sales
GROUP BY prod_id
ORDER BY prod_id;
```

In this example, the query instructs the database not to cache rows for a query of the sales table.

Using the RESULT_CACHE Hint in Views

The RESULT_CACHE hint applies only to the query block in which the hint is specified. If the hint is specified only in a view, then only these results are cached. View caching has the following characteristics:

- The view must be one of the following types:
 - A standard view (a view created with the CREATE ... VIEW statement)

- An inline view specified in the FROM clause of a SELECT statement
- An inline view created with the WITH clause
- The result of a view query with a correlated column (a reference to an outer query block) cannot be cached.
- Query results are stored in the server result cache, not the client result cache.
- A caching view is not merged into its outer (or referring) query block.

Adding the RESULT_CACHE hint to inline views disables optimizations between the outer query and inline view to maximize reusability of the cached result.

The following example shows a query of the inline view view1.

In this example, the SELECT statement from view1 is the outer block, whereas the SELECT statement from employees is the inner block. Because the RESULT_CACHE hint is specified only in the inner block, the results of the inner query are stored in the server result cache, but the results of the outer query are not cached.

Assume that the same session run a query of the view view2 as shown in the following example.

```
WITH view2 AS
( SELECT /*+ RESULT_CACHE */ department_id, manager_id, count(*) count
   FROM hr.employees
   GROUP BY department_id, manager_id )
SELECT *
   FROM view2
WHERE count BETWEEN 1 and 5;
```

In this example, because the RESULT_CACHE hint is specified only in the query block in the WITH clause, the results of the employees query are eligible to be cached. Because these results are cached from the query in the first example, the SELECT statement in the WITH clause in the second example can retrieve the cached rows.

Using Result Cache Table Annotations

You can also use table annotations to control result caching. Table annotations affect the entire query, not query segments. The primary benefit of using table annotations is avoiding the necessity of adding result cache hints to queries at the application level. Because a table annotation has a lower precedence than a SQL result cache hint, you can override table and session settings by using hints at the query level.

Table 13-4 describes the valid values for the RESULT CACHE table annotation.

Table 13-4 Values for the RESULT_CACHE Table Annotation

Value	Description
DEFAULT	If at least one table in a query is set to <code>DEFAULT</code> , then result caching is <i>not</i> enabled at the table level for this query, unless if the <code>RESULT_CACHE_MODE</code> initialization parameter is set to <code>FORCE</code> or the <code>RESULT_CACHE</code> hint is specified. This is the default value.

Table 13-4 (Cont.) Values for the RESULT_CACHE Table Annotation

Value	Description
FORCE	If all the tables of a query are marked as FORCE, then the query result is considered for caching. The table annotation FORCE takes precedence over the RESULT_CACHE_MODE parameter value of MANUAL set at the session level.

Using the DEFAULT Table Annotation

The DEFAULT table annotation prevents the database from caching results at the table level.

Example 13-5 shows a CREATE TABLE statement that uses the DEFAULT table annotation to create a table sales and a query of this table.

Example 13-5 Using the DEFAULT Table Annotation

```
CREATE TABLE sales (...) RESULT_CACHE (MODE DEFAULT);

SELECT prod_id, SUM(amount_sold)

FROM sales

GROUP BY prod_id

ORDER BY prod id;
```

In this example, the sales table is created with a table annotation that disables result caching. The example also shows a query of the sales table, whose results are not considered for caching because of the table annotation.



Oracle Database SQL Language Reference for information about the CREATE TABLE statement and its syntax

Using the FORCE Table Annotation

The FORCE table annotation forces the database to cache results at the table level.

Using the sales table created in Example 13-5, assume that you decide to force result caching for this table, you can do so by using the FORCE table annotation.

Example 13-6 shows an ALTER TABLE statement that uses the FORCE table annotation on the sales table.

Example 13-6 Using the FORCE Table Annotation

```
ALTER TABLE sales RESULT_CACHE (MODE FORCE);

SELECT prod_id, SUM(amount_sold)
  FROM sales
  GROUP BY prod_id
HAVING prod_id=136;

SELECT /*+ NO_RESULT_CACHE */ *
  FROM sales
  ORDER BY time id DESC;
```

This example includes two queries of the sales table. The first query, which is frequently used and returns few rows, is eligible for caching because of the table annotation. The second query, which is a one-time query that returns many rows, uses a hint to prevent result caching.

Monitoring the Result Cache

The result cache is improved with automatic exclusion (block list) of objects that are not worth caching. When enabled with the <code>ora.init</code> parameter <code>RESULT_CACHE_AUTO_BLOCKLIST</code>, a background action group is run periodically to find *hot tables* and make them candidates for automatic *block listing*.

To learn why a table was considered a *hot table* and block listed, use the view V\$SQL and its RESULT CACHE REJECTION REASON column to learn an object's rejection reason code.



V_SQL topic has a table of rejection codes, their meaning, and recommended remedial actions. The same table is published in the *Oracle Database Reference* after the V_SQL parameter list.

To view information about the server and client result caches, query the relevant database views and tables.

Table 13-5 describes the most useful views and tables for monitoring the result cache.

Table 13-5 Views and Tables with Information About the Result Cache

View/Table	Description
V\$RESULT_CACHE_STATISTICS	Lists various server result cache settings and memory usage statistics.
V\$RESULT_CACHE_MEMORY	Lists all the memory blocks in the server result cache and their corresponding statistics.
V\$RESULT_CACHE_OBJECTS	Lists all the objects whose results are in the server result cache along with their attributes.
	In 23ai, the new columns for TIME_SAVED and INVALIDATION_COST are used for the purposes of automatically excluding (block list) objects that are not worth caching.
V\$RESULT_CACHE_DEPENDENCY	Lists the dependency details between the results in the server result cache and dependencies among these results.
V\$SQL	Lists statistics on shared SQL areas without the GROUP BY clause and contains one row for each child of the original SQL text entered. Statistics displayed in V\$SQL are normally updated at the end of query execution. However, for long running queries, they are updated every 5 seconds. This makes it easy to see the impact of long running SQL statements while they are still in progress.
	The RESULT_CACHE_REJECTION_REASON column provides the rejection reason code for an object that is block listed.



Table 13-5 (Cont.) Views and Tables with Information About the Result Cache

View/Table	Description
CLIENT_RESULT_CACHE_STATS\$	Stores cache settings and memory usage statistics for the client result caches obtained from the OCI client processes. This statistics table contains entries for each client process that uses result caching. After the client processes terminate, the database removes their entries from this table. The client table contains information similar to V\$RESULT_CACHE_STATISTICS.
DBA_TABLES, USER_TABLES, ALL_TABLES	Contains a RESULT_CACHE column that shows the result cache mode annotation for the table. If the table is not annotated, then this column shows DEFAULT. This column applies to both server and client result caches.



Oracle Database Reference for more information about these views and tables.

Block Listing Reasons

Table 13-6 Result Cache Rejection Reason Codes

Code	Reason for Result Cache Rejection	Recommended Remedial Action
0	Success	N/A
1000	Result size exceeds per-result limit for in- memory	Increase the value of one or both of the RESULT_CACHE_MAX_SIZE and RESULT_CACHE_MAX_RESULT initialization parameters, or enable spilling results to disk.
1001	Result size exceeds per-result limit on disk	Increase the value of one or both of the RESULT_CACHE_MAX_TEMP_SIZE and RESULT_CACHE_MAX_TEMP_RESULT initialization parameters.
2000	Unable to acquire enqueues	The system is busy; try again later.
2001	Unable to acquire a DML locks on a referenced object	The system is busy; try again later.
2002	Unable to acquire a row cache lock on referenced objects	The system is busy; try again later.
2003	Unable to allocate memory	The system is busy; try again later.
3000	Result cache ID for the query is blocklisted	Remove the cache ID from the blocklist by using the DBMS_RESULT_CACHE package. Exercise caution because the query may have been intentionally blocklisted.
3001	An object referenced in the query is blocklisted	Remove the object ID from the blocklist by using the DBMS_RESULT_CACHE package. Exercise caution because the object may have been intentionally blocklisted.



Table 13-6 (Cont.) Result Cache Rejection Reason Codes

Code	Reason for Result Cache Rejection	Recommended Remedial Action
4001	Result cache is bypassed	The result cache was bypassed with DBMS_RESULT_CACHE.BYPASS or by setting the RESULT_CACHE_MAX_SIZE initialization parameter to 0. Exercise caution because the bypass may have been intentional.
4002	Result cache is corrupt	Reset the result cache by using DBMS_RESULT_CACHE.FLUSH. If that does not resolve the problem, then restart the PDB or instance.
4003	RCBG process is down in Oracle RAC	Restart the Oracle RAC instance
4004	The query was only partially executed	The cursor should be fully executed. If the cursor is being executed as part of PL/SQL, then it should be fetched from until the NO_DATA_FOUND error is returned. If the cursor is being executed as part of a SQL query, then do not cancel it by terminating the process (such as typing CTRL+C).
4005	Temporary space cannot be allocated because the RCBG process is down	Restart the Oracle RAC instance
9999	Unknown reason	Unknown

Examples for Result Cache

The following example shows a query of the V\$RESULT_CACHE_STATISTICS view to monitor server result cache statistics.

```
COLUMN name FORMAT a20
SELECT name, value
FROM V$RESULT_CACHE_STATISTICS;
```

The output of this query might look like the following:

NAME	VALUE
Block Size (Bytes)	1024
Block Count Maximum	3136
Block Count Current	32
Result Size Maximum (Blocks)	156
Create Count Success	2
Create Count Failure	0
Find Count	0
Invalidation Count	0
Delete Count Invalid	0
Delete Count Valid	0

The following example shows a query of the <code>CLIENT_RESULT_CACHE_STATS\$</code> table to monitor the client result cache statistics.

```
SELECT stat_id, SUBSTR(name,1,20), value, cache_id
FROM CLIENT_RESULT_CACHE_STATS$
ORDER BY cache_id, stat_id;
```

The output of this query might look like the following:

STAT_ID	NAME OF STATISTICS	VALUE	CACHE_ID
======	=======================================	=====	=======
1	Block Size	256	124



2	Block Count Max	256	124
3	Block Count Current	128	124
4	Hash Bucket Count	1024	124
5	Create Count Success	10	124
6	Create Count Failure	0	124
7	Find Count	12	124
8	Invalidation Count	8	124
9	Delete Count Invalid	0	124
10	Delete Count Valid	0	124

The CLIENT_RESULT_CACHE_STATS\$ table contains statistics entries for each active client process performing client result caching. Every client process has a unique cache ID.

To find the client connection information for the sessions performing client caching:

- 1. Obtain the session IDs from the CLIENT_REGID column in the GV\$SESSION_CONNECT_INFO view that corresponds to the CACHE_ID column in the CLIENT_RESULT_CACHE_STATS\$ table.
- 2. Query the relevant columns from the GV\$SESSION_CONNECT_INFO and GV\$SESSION views.

For both server and client result cache statistics, a database that is optimized for result caching should show relatively low values for the Create Count Failure and Delete Count Valid statistics, while showing relatively high values for the Find Count Statistic.



Tuning the Process Global Area

This chapter describes how to tune the Process Global Area (PGA). If you are using automatic memory management to manage the database memory on your system, then you do not need to manually tune the PGA as described in this chapter.

This chapter contains the following topics:

- About the Program Global Area
- Sizing the Program Global Area Using Automatic Memory Management
- · Sizing the Program Global Area Using a Hard Limit

About the Process Global Area

The Process Global Area (PGA) is a private memory region that contains the data and control information for a server process. Only a server process can access the PGA. Oracle Database reads and writes information in the PGA on behalf of the server process. An example of such information is the run-time area of a cursor. Each time a cursor is executed, a new run-time area is created for that cursor in the PGA memory region of the server process executing that cursor.



Part of the run-time area can be located in the Shared Global Area (SGA) when using shared servers.

For complex queries (such as decision support queries), a big portion of the run-time area is dedicated to work areas allocated by memory intensive operators, including:

- Sort-based operators, such as ORDER BY, GROUP BY, ROLLUP, and window functions
- Hash-join
- Bitmap merge
- Bitmap create
- Write buffers used by bulk load operations

A sort operator uses a work area (the sort area) to perform the in-memory sorting of a set of rows. Similarly, a hash-join operator uses a work area (the hash area) to build a hash table from its left input.

Work Area Sizes

Oracle Database enables you to control and tune the sizes of work areas. Generally, bigger work areas can significantly improve the performance of a particular operator at the cost of higher memory consumption. The available work area sizes include:

Optimal

Optimal size is when the size of a work area is large enough that it can accommodate the input data and auxiliary memory structures allocated by its associated SQL operator. This is the ideal size for the work area.

One-pass

One-pass size is when the size of the work area is below optimal size and an extra pass is performed over part of the input data. With one-pass size, the response time is increased.

Multi-pass

Multi-pass size is when the size of the work area is below the one-pass threshold and multiple passes over the input data are needed. With multi-pass size, the response time is dramatically increased because the size of the work area is too small compared to the input data size.

For example, a serial sort operation that must sort 10 GB of data requires a little more than 10 GB to run as optimal size and at least 40 MB to run as one-pass size. If the work area is less than 40 MB, then the sort operation must perform several passes over the input data.

When sizing the work area, the goal is to have most work areas running with optimal size (more than 90%, or even 100% for pure OLTP systems), and only a small number of them running with one-pass size (less than 10%). Multi-pass executions should be avoided for the following reasons:

- Multi-pass executions can severely degrade performance.
 - A high number of multi-pass work areas has an exponentially adverse effect on the response time of its associated SQL operator.
- Running one-pass executions does not require a large amount of memory.

Only 22 MB is required to sort 1 GB of data in one-pass size.

Even for DSS systems running large sorts and hash-joins, the memory requirement for onepass executions is relatively small. A system configured with a reasonable amount of PGA memory should not need to perform multiple passes over the input data.

Sizing the Program Global Area Using Automatic Memory Management

Automatic PGA memory management simplifies and improves the way PGA memory is allocated. By default, PGA memory management is enabled. In this mode, Oracle Database automatically sizes the PGA by dynamically adjusting the portion of the PGA memory dedicated to work areas, based on 20% of the SGA memory size. The minimum value is 10MB.

Note:

For backward compatibility, automatic PGA memory management can be disabled by setting the value of the PGA_AGGREGATE_TARGET initialization parameter to 0. When automatic PGA memory management is disabled, the maximum size of a work area can be sized with the associated _AREA_SIZE parameter, such as the SORT_AREA_SIZE initialization parameter.



This section describes how to size the PGA using automatic PGA memory management and contains the following topics:

- Configuring Automatic PGA Memory Management
- Setting the Initial Value for PGA_AGGREGATE_TARGET
- Monitoring Automatic PGA Memory Management
- Tuning PGA_AGGREGATE_TARGET

Configuring Automatic PGA Memory Management

When running Oracle Database in automatic PGA memory management mode, sizing of work areas for all sessions is automatic, and the *_AREA_SIZE parameters are ignored by all sessions running in this mode. Oracle Database automatically derives the total amount of PGA memory available to active work areas from the PGA_AGGREGATE_TARGET initialization parameter. The amount of PGA memory is set to the value of PGA_AGGREGATE_TARGET minus the amount of PGA memory allocated to other components of the system (such as PGA memory allocated by sessions). Oracle Database then assigns the resulting PGA memory to individual active work areas based on their specific memory requirements.

Oracle Database attempts to adhere to the PGA_AGGREGATE_TARGET value set by the DBA by dynamically controlling the amount of PGA memory allotted to work areas. To accomplish this, Oracle Database first tries to maximize the number of optimal work areas for all memory-intensive SQL operations. The rest of the work areas are executed in one-pass mode, unless the PGA memory limit set by the DBA (using the PGA_AGGREGATE_TARGET parameter) is so low that multi-pass execution is required to reduce memory consumption to honor the PGA target limit.

When configuring a new database instance, it can be difficult to determine the appropriate setting for PGA_AGGREGATE_TARGET.

To configure automatic PGA memory management:

- 1. Make an initial estimate for the value of the PGA_AGGREGATE_TARGET parameter, as described in "Setting the Initial Value for PGA_AGGREGATE_TARGET".
- Run a representative workload on the database instance and monitor its performance, as described in "Monitoring Automatic PGA Memory Management".
- 3. Tune the value of the PGA_AGGREGATE_TARGET parameter using Oracle PGA advice statistics, as described in "Tuning PGA_AGGREGATE_TARGET".

See Also:

Oracle Database Reference for information about the PGA_AGGREGATE_TARGET initialization parameter

Setting the Initial Value for PGA_AGGREGATE_TARGET

Set the initial value of the PGA_AGGREGATE_TARGET initialization parameter based on the amount of available memory for the Oracle database instance. This value can then be tuned and dynamically modified at the instance level. By default, Oracle Database uses 20% of the SGA size for this value. However, this setting may be too low for a large DSS system.



To set the initial value for PGA AGGREGATE TARGET:

 Determine how much of the total physical memory to reserve for the operating system and other non-Oracle applications running on the same system.

For example, you might decide to reserve 20% of the total physical memory for the operating system and other non-Oracle applications, dedicating 80% of the memory on the system to the Oracle database instance.

- 2. Divide the remaining available memory between the SGA and the PGA:
 - For OLTP systems, the PGA memory typically makes up a small fraction of the available memory, leaving most of the remaining memory for the SGA.

Oracle recommends initially dedicating 20% of the available memory to the PGA, and 80% to the SGA. Therefore, the initial value of the PGA_AGGREGATE_TARGET parameter for an OLTP system can be calculated as:

PGA_AGGREGATE_TARGET = $(total_mem * 0.8) * 0.2$ where $total_mem$ is the total amount of physical memory available on the system.

 For DSS systems running large, memory-intensive queries, PGA memory can typically use up to 70% of the available memory.

Oracle recommends initially dedicating 50% of the available memory to the PGA, and 50% to the SGA. Therefore, the initial value of the PGA_AGGREGATE_TARGET parameter for a DSS system can be calculated as:

PGA_AGGREGATE_TARGET = $(total_mem * 0.8) * 0.5$ where $total_mem$ is the total amount of physical memory available on the system.

For example, if an Oracle database instance is configured to run on a system with 4 GB of physical memory, and if 80% (or 3.2 GB) of the memory is dedicated to the Oracle database instance, then initially set PGA_AGGREGATE_TARGET to 640 MB for an OLTP system, or 1,600 MB for a DSS system.

Monitoring Automatic PGA Memory Management

Before starting the tuning process, run a representative workload on the database instance and monitor its performance. PGA statistics collected by Oracle Database enable you to determine if the maximum PGA size is under-configured or over-configured. Monitoring these statistics enables you to assess the performance of automatic PGA memory management and tune the value of the PGA AGGREGATE TARGET parameter accordingly.

This section describes how to use performance views to monitor automatic PGA memory management and contains the following topics:

- Using the V\$PGASTAT View
- Using the V\$PROCESS View
- Using the V\$PROCESS_MEMORY View
- Using the V\$SQL WORKAREA HISTOGRAM View
- Using the V\$WORKAREA_ACTIVE View
- Using the V\$SQL_WORKAREA View



Using the V\$PGASTAT View

The V\$PGASTAT view provides instance-level statistics about PGA memory usage and the automatic PGA memory manager.

The following example shows a query of this view.

SELECT *
FROM V\$PGASTAT;

The output of this query might look like the following:

NAME	VALUE	UNIT
aggregate PGA target parameter	41156608	bytes
aggregate PGA auto target	21823488	bytes
global memory bound	2057216	bytes
total PGA inuse	16899072	bytes
total PGA allocated	35014656	bytes
maximum PGA allocated	136795136	bytes
total freeable PGA memory	524288	bytes
PGA memory freed back to OS	1713242112	bytes
total PGA used for auto workareas	0	bytes
maximum PGA used for auto workareas	2383872	bytes
total PGA used for manual workareas	0	bytes
maximum PGA used for manual workareas	8470528	bytes
over allocation count	291	
bytes processed	2124600320	bytes
extra bytes read/written	39949312	bytes
cache hit percentage	98.15	percent

Table 14-1 describes the main statistics shown in the V\$PGASTAT view.

Table 14-1 Statistics in the V\$PGASTAT View

Statistic	Description
aggregate PGA target parameter	This statistic shows the current value of the PGA_AGGREGATE_TARGET parameter. The default value is 20% of the SGA size. Setting this parameter to 0 disables automatic PGA memory management.
aggregate PGA auto target	This statistic shows the amount of PGA memory Oracle Database can use for work areas running in automatic mode. This amount is dynamically derived from the value of the PGA_AGGREGATE_TARGET parameter and the current work area workload. Hence, it is continuously adjusted by Oracle Database. If this value is small compared to the PGA_AGGREGATE_TARGET value, then most of PGA memory is used by other system components (such as PL/SQL or Java) and little is left for work areas. Ensure that enough PGA memory remains for work areas running in automatic mode.
global memory bound	This statistic shows the maximum size of a work area executed in automatic mode. This value is continuously adjusted by Oracle Database to reflect the current state of the work area workload. The global memory bound generally decreases when the number of active work areas increases in the system. As a rule of thumb, the value of the global bound should not decrease to less than 1 MB. If it does, increase the value of the PGA_AGGREGATE_TARGET parameter.



Table 14-1 (Cont.) Statistics in the V\$PGASTAT View

Statistic	Description
total PGA allocated	This statistic shows the current amount of PGA memory allocated by the database instance. Oracle Database tries to keep this number less than the PGA_AGGREGATE_TARGET value. However, if the work area workload is increasing rapidly or the PGA_AGGREGATE_TARGET parameter is set to a value that is too low, it is possible for the PGA allocated to exceed this value by a small percentage and for a short time.
total freeable PGA memory	This statistic indicates how much allocated PGA memory can be freed.
total PGA used for auto workareas	This statistic indicates how much PGA memory is currently consumed by work areas running in automatic mode. Use this number to determine how much memory is consumed by other consumers of the PGA memory (such as PL/SQL or Java):
	PGA other = total PGA allocated - total PGA used for auto workareas
over allocation count	This statistic is cumulative from instance startup. Over-allocating PGA memory can happen if the PGA_AGGREGATE_TARGET value is too small to accommodate the PGA other component and the minimum memory required to execute the work area workload. In this case, Oracle Database cannot honor the PGA_AGGREGATE_TARGET value, and extra PGA memory must be allocated. If over-allocation occurs, increase the value of the PGA_AGGREGATE_TARGET parameter using the information provided by the V\$PGA_TARGET_ADVICE view, as described in "Using the V\$PGA_TARGET_ADVICE View".
total bytes processed	This statistic indicates the number of bytes processed by memory-intensive SQL operators since instance startup. For example, the number of bytes processed is the input size for a sort operation. This number is used to compute the cache hit percentage metric.
extra bytes read/ written	When a work area cannot run optimally, one or more extra passes is performed over the input data. This statistic represents the number of bytes processed during these extra passes since instance startup. This number is also used to compute the cache hit percentage metric. Ideally, it should be small compared to total bytes processed.
cache hit percentage	This metric is computed by Oracle Database to reflect the performance of the PGA memory component. It is cumulative from instance startup. A value of 100% means that all work areas executed by the system since instance startup are using an optimal amount of PGA memory. This is ideal but rarely happens except for pure OLTP systems. Typically, some work areas run one-pass or even multi-pass, depending on the overall size of the PGA memory. When a work area cannot run optimally, one or more extra passes are performed over the input data. This reduces the cache hit percentage in proportion to the size of the input data and the number of extra passes performed. For an example of how this metric is calculated, see Example 14-1.

Example 14-1 shows how extra passes affect the cache hit percentage **metric**.

Example 14-1 Calculating Cache Hit Percentage

Four sort operations have been executed, three were small (1 MB of input data) and one was bigger (100 MB of input data). The total number of bytes processed (BP) by the four operations

is 103 MB. If one of the small sorts runs one-pass, an extra pass over 1 MB of input data is performed. This 1 MB value is the number of extra bytes read/written, or EBP.

The cache hit percentage is calculated using the following formula:

```
BP \times 100 / (BP + EBP)
```

In this example, the cache hit percentage is 99.03%. This value reflects that only one of the small sort operations performed an extra pass, while all other sort operations were able to run in optimal size. Therefore, the cache hit percentage is almost 100%, because the extra pass over 1 MB represents a tiny overhead. However, if the bigger sort operation runs in one-pass size, then the EBP is 100 MB instead of 1 MB, and the cache hit percentage falls to 50.73%, because the extra pass has a much bigger impact.

Using the V\$PROCESS View

The V\$PROCESS view contains one row for each Oracle process connected to the database instance. Use the following columns in this view to monitor the PGA memory usage of these processes:

- PGA USED MEM
- PGA_ALLOC_MEM
- PGA_FREEABLE_MEM
- PGA MAX MEM

Example 14-2 shows a query of this view.

Example 14-2 Querying the V\$PROCESS View

```
SELECT program, pga_used_mem, pga_alloc_mem, pga_freeable_mem, pga_max_mem
FROM V$PROCESS;
```

The output of this query might look like the following:

PROGRAM		PGA_USED_MEM	PGA_ALLOC_MEM	PGA_FREEABLE_MEM	PGA_MAX_MEM
PSEUDO		0	0	0	0
oracle@examp1690	(PMON)	314540	685860	0	685860
oracle@examp1690	(MMAN)	313992	685860	0	685860
oracle@examp1690	(DBW0)	696720	1063112	0	1063112
oracle@examp1690	(LGWR)	10835108	22967940	0	22967940
oracle@examp1690	(CKPT)	352716	710376	0	710376
oracle@examp1690	(SMON)	541508	948004	0	1603364
oracle@examp1690	(RECO)	323688	685860	0	816932
oracle@examp1690	(q001)	233508	585128	0	585128
oracle@examp1690	(QMNC)	314332	685860	0	685860
oracle@examp1690	(MMON)	885756	1996548	393216	1996548
oracle@examp1690	(MMNL)	315068	685860	0	685860
oracle@examp1690	(q000)	330872	716200	65536	716200
oracle@examp1690	(CJQ0)	533476	1013540	0	1144612

Using the V\$PROCESS MEMORY View

The V\$PROCESS_MEMORY view displays dynamic PGA memory usage by named component categories for each Oracle process. This view contains up to six rows for each Oracle process, one row for:

Each named component category:

- Java
- PL/SOL
- OLAP
- SQL
- Freeable

Memory that has been allocated to the process by the operating system, but not to a specific category

Other

Memory that has been allocated to a category, but not to a named category

Use the following columns in this view to dynamically monitor the PGA memory usage of Oracle processes for each of the six categories:

- CATEGORY
- ALLOCATED
- USED
- MAX ALLOCATED

Note:

The V\$PROCESS_MEMORY_DETAIL view displays dynamic PGA memory usage for the Oracle processes that exceed 500 MB of PGA usage. The V\$PROCESS_MEMORY_DETAIL view is available starting with Oracle Database 12c Release 2.

See Also:

Oracle Database Reference for more information about the $V\$PROCESS_MEMORY$ and $V\$PROCESS_MEMORY$ DETAIL views

Using the V\$SQL WORKAREA HISTOGRAM View

The V\$SQL_WORKAREA_HISTOGRAM view shows the number of work areas executed with optimal, one-pass, and multi-pass memory size since instance startup. Statistics in this view are divided into buckets. The buckets are defined by the optimal memory requirements of the work areas. Each bucket is identified by a range of optimal memory requirements, specified by the values in the LOW OPTIMAL SIZE and HIGH OPTIMAL SIZE columns.

For example, a sort operation may require 3 MB of memory to run in optimal size (cached). Statistics about the work area used by this sort operation are placed in the bucket defined by:

- LOW OPTIMAL SIZE = 2097152 (2 MB)
- HIGH OPTIMAL SIZE = 4194303 (4 MB minus 1 byte)

Statistics are segmented by work area size, because the performance impact of running a work area in optimal, one-pass or multi-pass size depends mainly on the size of the work area.



In this example, statistics about the work area are placed in this bucket because 3 MB lies within that range of optimal sizes.

Example 14-3 and Example 14-4 show two methods for querying this view.

Example 14-3 Querying the V\$SQL_WORKAREA_HISTOGRAM View: Non-Empty Buckets

The following query shows statistics for all non-empty buckets:

The result of the query might look like the following:

LOW_KB	HIGH_KB	OPTIMAL_EXECUTIONS	ONEPASS_EXECUTIONS	MULTIPASSES_EXECUTIONS
8	16	156255	0	0
16	32	150	0	0
32	64	89	0	0
64	128	13	0	0
128	256	60	0	0
256	512	8	0	0
512	1024	657	0	0
1024	2048	551	16	0
2048	4096	538	26	0
4096	8192	243	28	0
8192	16384	137	35	0
16384	32768	45	107	0
32768	65536	0	153	0
65536	131072	0	73	0
131072	262144	0	44	0
262144	524288	0	22	0

In this example, the output shows that—in the 1 MB to 2 MB bucket—551 work areas ran in optimal size, while 16 ran in one-pass size and none ran in multi-pass size. It also shows that all work areas under 1 MB were able to run in optimal size.

Example 14-4 Querying the V\$SQL_WORKAREA_HISTOGRAM View: Percent Optimal

The following query shows the percentage of times work areas are executed in optimal, onepass, or multi-pass size since startup. This query only considers work areas of a certain size, with an optimal memory requirement of at least 64 KB:

The output of this query might look like the following:

```
OPTIMAL_COUNT OPTIMAL_PERC ONEPASS_COUNT ONEPASS_PERC MULTIPASS_COUNT MULTIPASS_PERC

2239 81.63 504 18.37 0 0
```



In this example, the output shows that 81.63% of the work areas were able to run in optimal size. The remaining work areas (18.37%) ran in one-pass size and none of them ran in multipass size.

Using the V\$WORKAREA_ACTIVE View

The V\$WORKAREA_ACTIVE view displays the work areas that are active (or executing) in the database instance. Small, active sort operations (under 64 KB) are excluded from this view. Use this view to precisely monitor the size of all active work areas and to determine whether these active work areas spill to a temporary segment.

Example 14-5 shows a query of this view.

Example 14-5 Querying the V\$WORKAREA_ACTIVE View

The output of this query might look like the following:

SID	OPERATION	ESIZE	MEM	MAX MEM	PASS	TSIZE
8	GROUP BY (SORT)	315	280	904	0	
8	HASH-JOIN	2995	2377	2430	1	20000
9	GROUP BY (SORT)	34300	22688	22688	0	
11	HASH-JOIN	18044	54482	54482	0	
12	HASH-JOIN	18044	11406	21406	1	120000

In this example, the output shows that:

- Session 12 (SID column) is running a hash-join operation (OPERATION column) in a work area running in one-pass size (PASS column)
- The maximum amount of memory that the PGA memory manager expects this hash-join operation to use is 18044 KB (ESIZE column)
- The work area is currently using 11406 KB of memory (MEM column)
- The work area used up to 21406 KB of PGA memory (MAX MEM column) in the past
- The work area spilled to a temporary segment of 120000 KB (TSIZE column)

When the work area is deallocated—or when the execution of its associated SQL operator is complete—it is automatically removed from this view.

Using the V\$SQL WORKAREA View

Oracle Database maintains cumulative work area statistics for each loaded cursor whose execution plan uses one or more work areas. Each time a work area is deallocated, the V\$SQL WORKAREA view is updated with execution statistics for that work area.

You can join the V\$SQL_WORKAREA view with the V\$SQL view to relate a work area to a cursor, and with the V\$SQL_PLAN view to precisely determine which operator in the plan uses a work area.

Example 14-6 shows three queries of this view.

Example 14-6 Querying the V\$SQL_WORKAREA View

The following query finds the top 10 work areas that require the most cache memory:

The following query finds the cursors with one or more work areas that have been executed in one or multiple passes:

Using the hash value and address of a particular cursor, the following query displays the cursor execution plan, including information about the associated work areas:

The output of this query might look like the following:

OPERATION	OPTIONS	NAME	input(MB)	LAST_MEM	OPTIMAL_ME	ONEPASS_ME	O/1/M
SELECT STATE							
HASH	GROUP BY		4582	8	16	16	16/0/0
HASH JOIN	SEMI		4582	5976	5194	2187	16/0/0
TABLE ACCESS	FULL	ORDERS	51				
TABLE ACCESS	FUL	LINEITEM	1000				

You can get the address and hash value from the V\$SQL view by specifying a pattern in the query, as shown in the following query:

```
SELECT address, hash_value
FROM V$SQL
WHERE sql_text LIKE '%my_pattern%';
```



Tuning PGA_AGGREGATE_TARGET

To help you tune the value of the PGA_AGGREGATE_TARGET initialization parameter, Oracle Database provides two PGA performance advisory views: V\$PGA_TARGET_ADVICE and V\$PGA_TARGET_ADVICE_HISTOGRAM. By using these views, you do not need to use an empirical approach to tune the value of the PGA_AGGREGATE_TARGET parameter. Instead, you can use these views to predict how changing the value of the PGA_AGGREGATE_TARGET parameter will affect key PGA statistics.

This section describes how to tune the value of the PGA_AGGREGATE_TARGET initialization parameter and contains the following topics:

- Enabling Automatic Generation of PGA Performance Advisory Views
- Using the V\$PGA_TARGET_ADVICE View
- Using the V\$PGA_TARGET_ADVICE_HISTOGRAM View
- Using the V\$SYSSTAT and V\$SESSTAT Views
- Tutorial: How to Tune PGA_AGGREGATE_TARGET

Enabling Automatic Generation of PGA Performance Advisory Views

Oracle Database generates the V\$PGA_TARGET_ADVICE and V\$PGA_TARGET_ADVICE_HISTOGRAM views by recording the workload history, and then simulating this history for different values of the PGA_AGGREGATE_TARGET parameter. The values of the PGA_AGGREGATE_TARGET parameter are derived from fractions and multiples of its current value to assess possible higher and lower values. These values are used for the prediction and range from 10 MB to a maximum of 256 GB. The simulation process happens in the background and continuously updates the workload history to produce the simulation result. You can view the result at any time by querying these views.

To enable automatic generation of PGA performance advice views:

 Set the PGA_AGGREGATE_TARGET parameter to enable automatic PGA memory management.

Setting this parameter to 0 disables automatic PGA memory management and is not recommended. For information about setting this parameter, see "Setting the Initial Value for PGA_AGGREGATE_TARGET".

2. Set the STATISTICS LEVEL parameter to TYPICAL (the default) or ALL.

Setting this parameter to BASIC disables generation of the PGA performance advice views and is not recommended.



The contents of the PGA advice performance views are reset at instance startup or when the value of the PGA_AGGREGATE_TARGET parameter is changed.



Using the V\$PGA TARGET ADVICE View

The V\$PGA_TARGET_ADVICE view predicts how changing the value of the PGA_AGGREGATE_TARGET initialization parameter will affect the following statistics in the V\$PGASTAT view:

- cache hit percentage
- over allocation count

The following example shows a query of this view.

The output of this query might look like the following:

TARGET_MB	CACHE_HIT_PERC	ESTD_OVERALLOC_COUNT
63	23	367
125	24	30
250	30	3
375	39	0
500	58	0
600	59	0
700	59	0
800	60	0
900	60	0
1000	61	0
1500	67	0
2000	76	0
3000	83	0
4000	85	0

The following figure illustrates how the result of this query can be plotted.



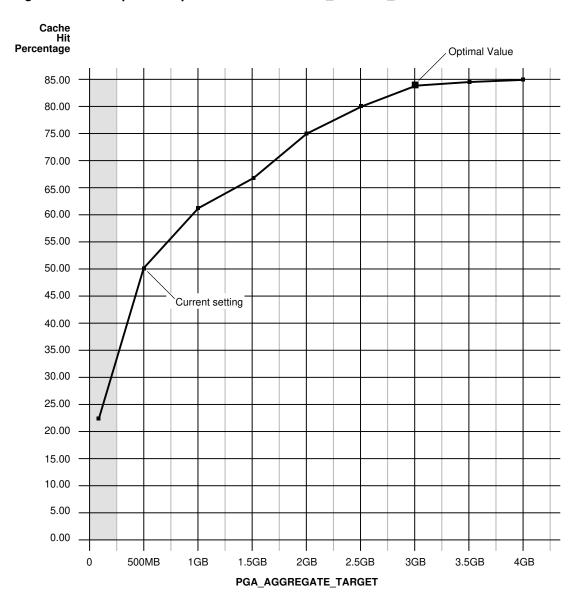


Figure 14-1 Graphical Representation of V\$PGA_TARGET_ADVICE

The curve shows how PGA cache hit percentage improves as the value of the PGA_AGGREGATE_TARGET parameter increases. The shaded zone in the graph represents the over allocation zone, where the value of the ESTD_OVERALLOCATION_COUNT column is non-zero. This area indicates that the value of the PGA_AGGREGATE_TARGET parameter is too small to meet the minimum PGA memory requirements. If the value of the PGA_AGGREGATE_TARGET parameter is set within the over allocation zone, then the memory manager will over-allocate memory and the actual PGA memory consumed will exceed the limit that was set. It is therefore meaningless to set a value of the PGA_AGGREGATE_TARGET parameter in that zone. In this particular example, the PGA_AGGREGATE_TARGET parameter should be set to at least 375 MB.

Beyond the over allocation zone, the value of the PGA cache hit percentage increases rapidly. This is due to an increase in the number of optimal or one-pass work areas and a decrease in the number of multi-pass executions. At some point, around 500 MB in this example, an inflection in the curve corresponds to the point where most (probably all) work

areas can run in optimal or at least one-pass size. Beyond this point, the cache hit percentage keeps increasing, though at a lower pace, up to the point where it starts to taper off and only slight improvement is achieved with increase in the value of the PGA_AGGREGATE_TARGET parameter. In the figure, this happens when PGA_AGGREGATE_TARGET reaches 3 GB. At this point, the cache hit percentage is 83% and only marginal improvement (by 2%) is achieved with one extra gigabyte of PGA memory. In this example, 3 GB is the optimal value for the PGA_AGGREGATE_TARGET parameter.

Note:

Although the theoretical maximum for the PGA cache hit percentage is 100%, a practical limit exists on the maximum size of a work area that may prevent this theoretical maximum from being reached, even when the value of the PGA_AGGREGATE_TARGET parameter is further increased. This should happen only in large DSS systems where the optimal memory requirement is large and may cause the value of the cache hit percentage to taper off at a lower percentage, such as 90%.

Ideally, the value of the PGA_AGGREGATE_TARGET parameter should be set to the optimal value, or at least to the maximum value possible in the region beyond the over allocation zone. As a rule of thumb, the PGA cache hit percentage should be higher than 60%, because at 60% the system is almost processing double the number of bytes it actually needs to process in an ideal situation. In this example, the value of the PGA_AGGREGATE_TARGET parameter should be set to at least 500 MB, and as close to 3 GB as possible. However, the correct setting for the PGA_AGGREGATE_TARGET parameter depends on how much memory can be dedicated to the PGA component. Generally, adding PGA memory requires reducing memory for some SGA components—like the shared pool or buffer cache—because the overall memory dedicated to the database instance is often bound by the amount of physical memory available on the system. Therefore, any decisions to increase PGA memory must be taken in the larger context of the available memory in the system and the performance of the various SGA components (which you can monitor with shared pool advisory and buffer cache advisory statistics). If you cannot reduce memory from the SGA components, consider adding more physical memory to the system.

Using the V\$PGA_TARGET_ADVICE_HISTOGRAM View

The V\$PGA_TARGET_ADVICE_HISTOGRAM view predicts how changing the value of the PGA_AGGREGATE_TARGET initialization parameter will affect the statistics in the V\$SQL_WORKAREA_HISTOGRAM view. Use this view to display detailed information about the predicted number of optimal, one-pass, and multi-pass work area executions for the PGA_AGGREGATE_TARGET values used for the prediction.

The V\$PGA_TARGET_ADVICE_HISTOGRAM view is identical to the V\$SQL_WORKAREA_HISTOGRAM view, with two additional columns to represent the PGA_AGGREGATE_TARGET values used for the prediction. Therefore, any query executed against the V\$SQL_WORKAREA_HISTOGRAM view can be used on this view, with an additional predicate to select the desired value of the PGA_AGGREGATE_TARGET parameter.

Example 14-7 shows a query of this view that displays the predicted content of the V\$SQL_WORKAREA_HISTOGRAM view for a value of the PGA_AGGREGATE_TARGET parameter set to twice its current value.

Example 14-7 Querying the V\$PGA_TARGET_ADVICE_HISTOGRAM View

The output of this query might look like the following:

LOW_KB	HIGH_KB	ESTD_OPTIMAL_CNT	ESTD_ONEPASS_CNT	ESTD_MPASS_CNT
8	16	156107	0	0
16	32	148	0	0
32	64	89	0	0
64	128	13	0	0
128	256	58	0	0
256	512	10	0	0
512	1024	653	0	0
1024	2048	530	0	0
2048	4096	509	0	0
4096	8192	227	0	0
8192	16384	176	0	0
16384	32768	133	16	0
32768	65536	66	103	0
65536	131072	15	47	0
131072	262144	0	48	0
262144	524288	0	23	0

In this example, the output shows that increasing the value of the PGA_AGGREGATE_TARGET parameter by a factor of 2 will enable all work areas under 16 MB to execute in optimal size.

Using the V\$SYSSTAT and V\$SESSTAT Views

Statistics in the V\$SYSSTAT and V\$SESSTAT views show the total number of work areas executed with optimal, one-pass, and multi-pass memory size. These statistics are cumulative since the instance or the session was started.

Example 14-8 shows a query of the V\$SYSSTAT view that displays the total number and the percentage of times work areas were executed in these three sizes since the instance was started:

Example 14-8 Querying the V\$SYSSTAT View

```
SELECT name profile, cnt, DECODE(total, 0, 0, ROUND(cnt*100/total)) percentage
FROM (SELECT name, value cnt, (SUM(value) over ()) total
FROM V$SYSSTAT
WHERE name
LIKE 'workarea exec%');
```

The output of this query might look like the following:

PROFILE			CNT	PERCENTAGE
workarea ex	ecutions -	optimal	5395	95
workarea ex	ecutions -	onepass	284	5
workarea ex	ecutions -	multipass	0	0



In this example, the output shows that 5,395 work area executions (or 95%) were executed in optimal size, and 284 work area executions (or 5%) were executed in one-pass size.

Tutorial: How to Tune PGA AGGREGATE TARGET

This tutorial provides a guideline for tuning the value of the PGA_AGGREGATE_TARGET parameter using the various views discussed in this chapter.

To tune PGA AGGREGATE TARGET:

- Set the value of the PGA AGGREGATE TARGET parameter to avoid memory over-allocation.
 - Use the V\$PGA_TARGET_ADVICE view to ensure that the PGA_AGGREGATE_TARGET value is not set within the over-allocation zone, as described in "Using the V\$PGA_TARGET_ADVICE View". In Example 16–8, the PGA_AGGREGATE_TARGET value should be set to at least 375 MB.
- Maximize the PGA cache hit percentage, based on response time requirements and memory constraints.
 - Use the V\$PGA_TARGET_ADVICE view to determine the optimal value for the PGA_AGGREGATE_TARGET parameter and set its value to the optimal value, or to the maximum value possible, as described in "Using the V\$PGA_TARGET_ADVICE View".

Assume a limit *X* on the memory that can be allocated to PGA:

- If limit X is higher than the optimal value, set the value of the PGA_AGGREGATE_TARGET parameter to the optimal value.
 - In Example 16–8, if you have 10 GB to dedicate to PGA, set the value of the $PGA_AGGREGATE_TARGET$ parameter to 3 GB and dedicate the remaining 7 GB to the SGA.
- If limit X is less than the optimal value, set the value of the PGA_AGGREGATE_TARGET parameter to X.
 - In Example 16–8, if you have only 2 GB to dedicate to PGA, set the value of the PGA_AGGREGATE_TARGET parameter to 2 GB and accept a cache hit percentage of 75%.
- 3. Verify that the new value of the PGA_AGGREGATE_TARGET parameter will result in the desired number of optimal and one-pass work area executions and avoid any multi-pass work area executions.
 - Use the V\$PGA_TARGET_ADVICE_HISTOGRAM view to predict the number of optimal, one-pass, and multi-pass work area executions, as described in "Using the V\$PGA_TARGET_ADVICE_HISTOGRAM View".
- **4.** If more PGA memory is required, then increase PGA memory by either reducing memory from SGA components or adding more physical memory to the system.
- At any time, ensure the number of optimal, one-pass, and multi-pass work area executions
 matches predictions and tune the value of the PGA_AGGREGATE_TARGET parameter if
 necessary.
 - Use the V\$SYSSTAT and V\$SESSTAT views to verify the total number of work areas executed with optimal, one-pass, and multi-pass memory size since instance or session startup, respectively, as described in "Using the V\$SYSSTAT and V\$SESSTAT Views".



Sizing the Program Global Area by Specifying an Absolute Limit

In automatic PGA memory management mode, Oracle Database attempts to adhere to the PGA_AGGREGATE_TARGET value by dynamically controlling the amount of PGA memory allotted to work areas. However, PGA memory usage may exceed the PGA_AGGREGATE_TARGET setting at times due to the following reasons:

- The PGA AGGREGATE TARGET setting acts as a target, and not a limit.
- PGA AGGREGATE TARGET only controls allocations of tunable memory.

Excessive PGA usage can lead to high rates of swapping. When this occurs, the system may become unresponsive and unstable. In that case, consider using any of the following methods to specify an absolute limit on the PGA memory usage:

 Use PGA_AGGREGATE_LIMIT parameter to set an absolute limit on the overall PGA memory usage.

See "Sizing the Program Global Area Using the PGA_AGGREGATE_LIMIT Parameter"

Use the Resource Manager procedure DBMS_RESOURCE_MANAGER.CREATE_PLAN_DIRECTIVE
to set an absolute limit on the PGA memory usage for each session in a particular
consumer group.

See "Sizing the Program Global Area Using the Resource Manager"

Sizing the Program Global Area Using the PGA_AGGREGATE_LIMIT Parameter

The PGA_AGGREGATE_LIMIT initialization parameter enables you to specify an absolute limit on the PGA memory usage. If the PGA_AGGREGATE_LIMIT value is exceeded, Oracle Database terminates the sessions or processes that are consuming the most untunable PGA memory in the following order:

- Calls for sessions that are consuming the most untunable PGA memory are canceled.
- If PGA memory usage is still over the PGA_AGGREGATE_LIMIT, then the sessions and processes that are consuming the most untunable PGA memory are terminated.

In determining the sessions and processes to terminate, Oracle Database treats parallel queries as a single unit.

By default, the PGA_AGGREGATE_LIMIT parameter is set to the greater of 2 GB, 200% of the PGA_AGGREGATE_TARGET value, or 3 MB times the value of the PROCESSES parameter. However, it will not exceed 120% of the physical memory size minus the total SGA size. The default value is printed into the alert log. A warning message is printed in the alert log if the amount of physical memory on the system cannot be determined.

To set PGA_AGGREGATE_LIMIT:

• Set the PGA_AGGREGATE_LIMIT initialization parameter to a desired value in number of bytes.

The value is expresses as a number followed by K (for kilobytes), M (for megabytes), or G (for gigabytes). Setting the value to 0 disables the hard limit on PGA memory.



See Also:

- Oracle Database Reference for information about the PGA_AGGREGATE_LIMIT initialization parameter
- Oracle Database Reference for information about the V\$PGASTAT view
- Oracle Database Administrator's Guide for information about Oracle Database Resource Manager and consumer groups

Sizing the Program Global Area Using the Resource Manager

You can set an absolute limit on the amount of PGA memory that can be allocated to each session in a particular consumer group using the SESSION_PGA_LIMIT parameter of the DBMS_RESOURCE_MANAGER.CREATE_PLAN_DIRECTIVE procedure of the Oracle Database Resource Manager. If a session exceeds the PGA memory limit set for its consumer group, then that session is terminated with the ORA-10260 error message.

See Also:

- Oracle Database Administration Guide topics:
 - "Program Global Area (PGA)" for more information about limiting the PGA memory for each session in a consumer group.
 - "Creating Resource Plan Directives" for more information about creating resource plan directives using the
 DBMS RESOURCE MANAGER.CREATE PLAN DIRECTIVE procedure.
- Oracle Database PL/SQL Packages and Types Reference for the syntax of the DBMS RESOURCE MANAGER.CREATE PLAN DIRECTIVE procedure.



Part IV

Managing System Resources

This part contains the following chapters:

- I/O Configuration and Design
- Managing Operating System Resources



I/O Configuration and Design

The I/O subsystem is a vital component of an Oracle database. This chapter introduces fundamental I/O concepts, discusses the I/O requirements of different parts of the database, and provides sample configurations for I/O subsystem design.

This chapter includes the following topics:

- About I/O
- I/O Configuration
- I/O Calibration Inside the Database
- I/O Calibration with the Oracle Orion Calibration Tool

About I/O

Every Oracle database reads or writes data on disk, thus generating **disk I/O**. The performance of many software applications is inherently limited by disk I/O. Applications that spend majority of their CPU time waiting for I/O activity to complete are said to be **I/O-bound**.

Oracle Database is designed so that if an application is well written, its performance should not be limited by I/O. Tuning I/O can enhance the performance of the application if the I/O system is operating at or near capacity and is not able to service the I/O requests within an acceptable time. However, tuning I/O cannot help performance if the application is not I/O-bound (for example, when CPU is the limiting factor).

Consider the following database requirements when designing an I/O system:

- Storage, such as minimum disk capacity
- Availability, such as continuous (24 x 7) or business hours only
- Performance, such as I/O throughput and application response times

Many I/O designs plan for storage and availability requirements with the assumption that performance will not be an issue. This is not always the case. Optimally, the number of disks and controllers to be configured should be determined by I/O throughput and redundancy requirements. The size of disks can then be determined by the storage requirements.

When developing an I/O design plan, consider using **Oracle Automatic Storage Management (Oracle ASM)**. Oracle ASM is an integrated, high-performance database file system and disk manager that is based on the principle that the database should manage storage instead of requiring an administrator to do it.

Oracle recommends that you use Oracle ASM for your database file storage, instead of raw devices or the operating system file system. Oracle ASM provides the following key benefits:

- Striping
- Mirroring
- Online storage reconfiguration and dynamic rebalancing
- Managed file creation and deletion

See Also:

Oracle Automatic Storage Management Administrator's Guide for additional information about Oracle ASM

I/O Configuration

This section describes the basic information to be gathered and decisions to be made when defining a system's I/O configuration. You want to keep the configuration as simple as possible, while maintaining the required availability, recoverability, and performance. The more complex a configuration becomes, the more difficult it is to administer, maintain, and tune.

This section contains the following topics:

- Lay Out the Files Using Operating System or Hardware Striping
- Manually Distributing I/O
- When to Separate Files
- Three Sample Configurations
- Oracle Managed Files
- Choosing Data Block Size

Lay Out the Files Using Operating System or Hardware Striping

If your operating system has LVM software or hardware-based striping, then it is possible to distribute I/O using these tools. Decisions to be made when using an LVM or hardware striping include **stripe depth** and **stripe width**.

- Stripe depth is the size of the stripe, sometimes called stripe unit.
- Stripe width is the product of the stripe depth and the number of drives in the striped set.

Choose these values wisely so that the system is capable of sustaining the required throughput. For an Oracle database, reasonable stripe depths range from 256 KB to 1 MB. Different types of applications benefit from different stripe depths. The optimal stripe depth and stripe width depend on the following:

- Requested I/O Size
- Concurrency of I/O Requests
- Alignment of Physical Stripe Boundaries with Block Size Boundaries
- · Manageability of the Proposed System

Requested I/O Size

Table 15-1 lists the Oracle Database and operating system parameters that you can use to set I/O size:



Table 15-1 Oracle Database and Operating System Operational Parameters

Parameter	Description
DB_BLOCK_SIZE	The size of single-block I/O requests. This parameter is also used in combination with multiblock parameters to determine multiblock I/O request size.
OS block size	Determines I/O size for redo log and archive log operations.
Maximum OS I/O size	Places an upper bound on the size of a single I/O request.
DB_FILE_MULTIBLOCK_READ _COUNT	The maximum I/O size for full table scans is computed by multiplying this parameter with $\mathtt{DB_BLOCK_SIZE}$. (the upper value is subject to operating system limits). If this value is not set explicitly (or is set to 0), the default value corresponds to the maximum I/O size that can be efficiently performed and is platform-dependent.
SORT_AREA_SIZE	Determines I/O sizes and concurrency for sort operations.
HASH_AREA_SIZE	Determines the I/O size for hash operations.

In addition to I/O size, the degree of concurrency also helps in determining the ideal stripe depth. Consider the following when choosing stripe width and stripe depth:

- On low-concurrency (sequential) systems, ensure that no single I/O visits the same disk twice. For example, assume that the stripe width is four disks, and the stripe depth is 32K. If a single 1MB I/O request (for example, for a full table scan) is issued by an Oracle server process, then each disk in the stripe must perform eight I/Os to return the requested data. To avoid this situation, the size of the average I/O should be smaller than the stripe width multiplied by the stripe depth. If this is not the case, then a single I/O request made by Oracle Database to the operating system results in multiple physical I/O requests to the same disk.
- On high-concurrency (random) systems, ensure that no single I/O request is broken up into multiple physical I/O calls. Failing to do this multiplies the number of physical I/O requests performed in your system, which in turn can severely degrade the I/O response times.

Concurrency of I/O Requests

In a system with a high degree of concurrent small I/O requests, such as in a traditional OLTP environment, it is beneficial to keep the stripe depth large. Using stripe depths larger than the I/O size is called **coarse grain striping**. In high-concurrency systems, the stripe depth can be as follows, where n > 1:

```
n * DB_BLOCK_SIZE
```

Coarse grain striping allows a disk in the array to service several I/O requests. In this way, a large number of concurrent I/O requests can be serviced by a set of striped disks with minimal I/O setup costs. Coarse grain striping strives to maximize overall I/O throughput. Multiblock reads, as in full table scans, will benefit when stripe depths are large and can be serviced from one drive. Parallel query in a data warehouse environment is also a candidate for coarse grain striping because many individual processes each issue separate I/Os. If coarse grain striping is used in systems that do not have high concurrent requests, then hot spots could result.

In a system with a few large I/O requests, such as in a traditional DSS environment or a low-concurrency OLTP system, then it is beneficial to keep the stripe depth small. This is called **fine grain striping**. In such systems, the stripe depth is as follows, where n is smaller than the multiblock read parameters, such as DB FILE MULTIBLOCK READ COUNT:

n * DB BLOCK SIZE

Fine grain striping allows a single I/O request to be serviced by multiple disks. Fine grain striping strives to maximize performance for individual I/O requests or response time.

Alignment of Physical Stripe Boundaries with Block Size Boundaries

On some Oracle Database ports, a database block boundary may not align with the stripe. If your stripe depth is the same size as the database block, then a single I/O issued by Oracle Database may result in two physical I/O operations.

This is not optimal in an OLTP environment. To ensure a higher probability of one logical I/O resulting in no more than one physical I/O, the minimum stripe depth should be at least twice the Oracle block size. Table 15-2 shows recommended minimum stripe depth for random access and for sequential reads.

Table 15-2 Minimum Stripe Depth

Disk Access	Minimum Stripe Depth
Random reads and writes	The minimum stripe depth is twice the Oracle block size.
Sequential reads	The minimum stripe depth is twice the value of DB_FILE_MULTIBLOCK_READ_COUNT, multiplied by the Oracle block size.



The specific documentation for your platform

Manageability of the Proposed System

With an LVM, the simplest configuration to manage is one with a single striped volume over all available disks. In this case, the stripe width encompasses all available disks. All database files reside within that volume, effectively distributing the load evenly. This single-volume layout provides adequate performance in most situations.

A single-volume configuration is viable only when used in conjunction with RAID technology that allows easy recoverability, such as RAID 1. Otherwise, losing a single disk means losing all files concurrently and, hence, performing a full database restore and recovery.

In addition to performance, there is a manageability concern: the design of the system must allow disks to be added simply, to allow for database growth. The challenge is to do so while keeping the load balanced evenly.

For example, an initial configuration can involve the creation of a single striped volume over 64 disks, each disk being 16 GB. This is total disk space of 1 terabyte (TB) for the primary data. Sometime after the system is operational, an additional 80 GB (that is, five disks) must be added to account for future database growth.

The options for making this space available to the database include creating a second volume that includes the five new disks. However, an I/O bottleneck might develop, if these new disks are unable to sustain the I/O throughput required for the files placed on them.

Another option is to increase the size of the original volume. LVMs are becoming sophisticated enough to allow dynamic reconfiguration of the stripe width, which allows disks to be added



while the system is online. This begins to make the placement of all files on a single striped volume feasible in a production environment.

If your LVM cannot support dynamically adding disks to the stripe, then it is likely that you need to choose a smaller, more manageable stripe width. Then, when new disks are added, the system can grow by a stripe width.

In the preceding example, eight disks might be a more manageable stripe width. This is only feasible if eight disks are capable of sustaining the required number of I/Os each second. Thus, when extra disk space is required, another eight-disk stripe can be added, keeping the I/O balanced across the volumes.

Note:

The smaller the stripe width becomes, the more likely it is that you will need to spend time distributing the files on the volumes, and the closer the procedure becomes to manually distributing I/O.

Manually Distributing I/O

If your system does not have an LVM or hardware striping, then I/O must be manually balanced across the available disks by distributing the files according to each file's I/O requirements. In order to make decisions on file placement, you should be familiar with the I/O requirements of the database files and the capabilities of the I/O system. If you are not familiar with this data and do not have a representative workload to analyze, you can make a first guess and then tune the layout as the usage becomes known.

To stripe disks manually, you need to relate a file's storage requirements to its I/O requirements.

- Evaluate database disk-storage requirements by checking the size of the files and the disks.
- 2. Identify the expected I/O throughput for each file. Determine which files have the highest I/O rate and which do not have many I/Os. Lay out the files on all the available disks so as to even out the I/O rate.

One popular approach to manual I/O distribution suggests separating a frequently used table from its index. This is not correct. During the course of a transaction, the index is read first, and then the table is read. Because these I/Os occur sequentially, the table and index can be stored on the same disk without contention. It is not sufficient to separate a data file simply because the data file contains indexes or table data. The decision to segregate a file should be made only when the I/O rate for that file affects database performance.

When to Separate Files

Regardless of whether you use operating system striping or manual I/O distribution, if the I/O system or I/O layout is not able to support the I/O rate required, then you need to separate files with high I/O rates from the remaining files. You can identify such files either at the planning stage or after the system is live.

The decision to segregate files should only be driven by I/O rates, recoverability concerns, or manageability issues. (For example, if your LVM does not support dynamic reconfiguration of stripe width, then you might need to create smaller stripe widths to be able to add n disks at a time to create a new stripe of identical configuration.)



Before segregating files, verify that the bottleneck is truly an I/O issue. The data produced from investigating the bottleneck identifies which files have the highest I/O rates.

The following sections describe how to segregate the following file types:

- Tables, Indexes, and TEMP Tablespaces
- Redo Log Files
- Archived Redo Logs

Tables, Indexes, and TEMP Tablespaces

If the files with high I/O are data files belonging to tablespaces that contain tables and indexes, then identify whether the I/O for those files can be reduced by tuning SQL or application code.

If the files with high-I/O are data files that belong to the TEMP tablespace, then investigate whether to tune the SQL statements performing disk sorts to avoid this activity, or to tune the sorting.

After the application has been tuned to avoid unnecessary I/O, if the I/O layout is still not able to sustain the required throughput, then consider segregating the high-I/O files.

Redo Log Files

If the high-I/O files are redo log files, then consider splitting the redo log files from the other files. Possible configurations can include the following:

- Placing all redo logs on one disk without any other files. Also consider availability; members of the same group should be on different physical disks and controllers for recoverability purposes.
- Placing each redo log group on a separate disk that does not store any other files.
- Striping the redo log files across several disks, using an operating system striping tool. (Manual striping is not possible in this situation.)
- Avoiding the use of RAID 5 for redo logs.

Redo log files are written sequentially by the Log Writer (LGWR) process. This operation can be made faster if there is no concurrent activity on the same disk. Dedicating a separate disk to redo log files usually ensures that LGWR runs smoothly with no further tuning necessary. If your system supports asynchronous I/O but this feature is not currently configured, then test to see if using this feature is beneficial. Performance bottlenecks related to LGWR are rare.

Archived Redo Logs

If the archiver is slow, then it might be prudent to prevent I/O contention between the archiver process and LGWR by ensuring that archiver reads and LGWR writes are separated. This is achieved by placing logs on alternating drives.

For example, suppose a system has four redo log groups, each group with two members. To create separate-disk access, the eight log files should be labeled 1a, 1b, 2a, 2b, 3a, 3b, 4a, and 4b. This requires at least four disks, plus one disk for archived files.

The following figure illustrates how redo members should be distributed across disks to minimize contention.



arch dest

1a 2a 1b 2b 4b

A 4a 3b

Figure 15-1 Distributing Redo Members Across Disks

lgwr

In this example, LGWR switches out of log group 1 (member 1a and 1b) and writes to log group 2 (2a and 2b). Concurrently, the archiver process reads from group 1 and writes to its archive destination. Note how the redo log files are isolated from contention.



Mirroring redo log files, or maintaining multiple copies of each redo log file on separate disks, does not slow LGWR considerably. LGWR writes to each disk in parallel and waits until each part of the parallel write is complete. Thus, a parallel write does not take longer than the longest possible single-disk write.

Because redo logs are written serially, drives dedicated to redo log activity generally require limited head movement. This significantly accelerates log writing.

Three Sample Configurations

This section contains three high-level examples of configuring I/O systems. These examples include sample calculations that define the disk topology, stripe depths, and so on:

- Stripe Everything Across Every Disk
- Move Archive Logs to Different Disks
- Move Redo Logs to Separate Disks

Stripe Everything Across Every Disk

The simplest approach to I/O configuration is to build one giant volume, striped across all available disks. To account for recoverability, the volume is mirrored (RAID 1). The striping unit for each disk should be larger than the maximum I/O size for the frequent I/O operations. This provides adequate performance for most cases.



Move Archive Logs to Different Disks

If archived redo log files are striped on the same set of disks as other files, then any I/O requests on those disks could suffer when the database is archiving the redo logs. Moving archived redo log files to separate disks provides the following benefits:

- The archive can be performed at very high rate (using sequential I/O).
- Nothing else is affected by the degraded response time on the archive destination disks.

The number of disks for archive logs is determined by the rate of archive log generation and the amount of archive storage required.

Move Redo Logs to Separate Disks

In high-update OLTP systems, the redo logs are write-intensive. Moving the redo log files to disks that are separate from other disks and from archived redo log files has the following benefits:

- Writing redo logs is performed at the highest possible rate. Hence, transaction processing performance is at its best.
- Writing of the redo logs is not impaired with any other I/O.

The number of disks for redo logs is mostly determined by the redo log size, which is generally small compared to current technology disk sizes. Typically, a configuration with two disks (possibly mirrored to four disks for fault tolerance) is adequate. In particular, by having the redo log files alternating on two disks, writing redo log information to one file does not interfere with reading a completed redo log for archiving.

Oracle Managed Files

When file systems can contain all Oracle Database data, database administration is simplified by using **Oracle Managed Files**. Oracle Database internally uses standard file system interfaces to create and delete files as needed for tablespaces, temp files, online logs, and control files. Administrators only specify the file system directory to be used for a particular type of file. You can specify one default location for data files and up to five multiplexed locations for the control and online redo log files.

Oracle Database ensures that a unique file is created and then deleted when it is no longer needed. This reduces corruption caused by administrators specifying the wrong file, reduces wasted disk space consumed by obsolete files, and simplifies creation of test and development databases. It also makes development of portable third-party tools easier, because it eliminates the need to put operating system-specific file names in SQL scripts.

New files can be created as Oracle Managed Files, while old ones are administered in the old way. Thus, a database can have a mixture of Oracle Managed Files and user-managed files.

Note:

Oracle Managed Files cannot be used with raw devices.

Several points should be considered when tuning Oracle Managed Files:



- Because Oracle Managed Files require the use of a file system, DBAs give up control over how the data is laid out. Therefore, it is important to correctly configure the file system.
- Build the file system for Oracle Managed Files on top of an LVM that supports striping. For load balancing and improved throughput, stripe the disks in the file system.
- Oracle Managed Files work best if used on an LVM that supports dynamically extensible logical volumes. Otherwise, configure the logical volumes as large as possible.
- Oracle Managed Files work best if the file system provides large extensible files.



Oracle Database Administrator's Guide for detailed information about using Oracle Managed Files

Choosing Data Block Size

One logical data block corresponds to a specific number of bytes of persistent storage. Data blocks are the smallest units of storage that Oracle Database can use or allocate.

A block size of 8 KB is optimal for most systems. However, OLTP systems occasionally use smaller block sizes and DSS systems occasionally use larger block sizes.

Traditionally, data files have been stored on magnetic disk. Alternative forms of non-volatile data storage are also supported, such as directly mapped buffer cache. [From a (database) software rather than a hardware perspective, the buffer cache contains the mapping of blocks in the form of hash chain entries, and some of those entries point to PMEM (persistent memory).] However the underlying data is physically stored, Oracle processes always read and write logical data blocks.

This section discusses considerations when choosing database block size for optimal performance and contains the following topics:

- Reads
- Writes
- Block Size Advantages and Disadvantages



The use of multiple block sizes in a single database instance is not encouraged because of manageability issues.

Reads

Regardless of the size of the data, the goal is to minimize the number of reads required to retrieve the desired data.

- If the rows are small and access is predominantly random, then choose a smaller block
- If the rows are small and access is predominantly sequential, then choose a larger block size.

- If the rows are small and access is both random and sequential, then it might be effective to choose a larger block size.
- If the rows are large, such as rows containing large object (LOB) data, then choose a larger block size.

Writes

For high-concurrency OLTP systems, consider appropriate values for INITRANS, MAXTRANS, and FREELISTS when using a larger block size. These parameters affect the degree of update concurrency allowed within a block. However, you do not need to specify the value for FREELISTS when using automatic segment-space management.

If you are uncertain about which block size to choose, then try a database block size of 8 KB for most systems that process a large number of transactions. This represents a good compromise and is usually effective. Only systems processing LOB data need more than 8 KB.



The Oracle Database installation documentation specific to your operating system for information about the minimum and maximum block size on your platform

Block Size Advantages and Disadvantages

Table 15-3 lists the advantages and disadvantages of different block sizes.

Table 15-3 Block Size Advantages and Disadvantages

Block Size	Advantages	Disadvantages
Smaller	Good for small rows with lots of random access.	Has relatively large space overhead due to metadata (that is, block header).
	Reduces block contention.	Not recommended for large rows. There might only be a few rows stored for each block, or worse, row chaining if a single row does not fit into a block,
Larger	Has lower overhead, so there is more room to store data.	Wastes space in the buffer cache, if you are doing random access to small rows and have a large block size. For
	Permits reading several rows into the buffer cache with a single I/O (depending on row size and block size).	example, with an 8 KB block size and 50 byte row size, you waste 7,950 bytes in the buffer cache when doing random access.
	Good for sequential access or very large rows (such as LOB data).	Not good for index blocks used in an OLTP environment, because they increase block contention on the index leaf blocks.

I/O Calibration Inside the Database

The I/O calibration feature of Oracle Database enables you to assess the performance of the storage subsystem, and determine whether I/O performance problems are caused by the database or the storage subsystem. Unlike other external I/O calibration tools that issue I/Os sequentially, the I/O calibration feature of Oracle Database issues I/Os randomly using Oracle data files to access the storage media, producing results that more closely match the actual performance of the database.



The section describes how to use the I/O calibration feature of Oracle Database and contains the following topics:

- Prerequisites for I/O Calibration
- Running I/O Calibration

Oracle Database also provides Orion, an I/O calibration tool. Orion is a tool for predicting the performance of an Oracle database without having to install Oracle or create a database. Unlike other I/O calibration tools, Oracle Orion is expressly designed for simulating Oracle database I/O workloads using the same I/O software stack as Oracle. Orion can also simulate the effect of striping performed by Oracle Automatic Storage Management. For more information, see "I/O Calibration with the Oracle Orion Calibration Tool".

Prerequisites for I/O Calibration

Before running I/O calibration, ensure that the following requirements are met:

- The user must be granted the SYSDBA privilege
- timed statistics must be set to TRUE
- Asynchronous I/O must be enabled

When using file systems, asynchronous I/O can be enabled by setting the ${\tt FILESYSTEMIO}$ OPTIONS initialization parameter to ${\tt SETALL}$.

Ensure that asynchronous I/O is enabled for data files by running the following query:

```
COL NAME FORMAT A50

SELECT NAME, ASYNCH_IO FROM V$DATAFILE F,V$IOSTAT_FILE I

WHERE F.FILE#=I.FILE_NO

AND FILETYPE NAME='Data File';
```

Additionally, only one calibration can be performed on a database instance at a time.

Running I/O Calibration

The I/O calibration feature of Oracle Database is accessed using the <code>DBMS_RESOURCE_MANAGER.CALIBRATE_IO</code> procedure. This procedure issues an I/O intensive read-only workload, made up of one megabyte of random of I/Os, to the database files to determine the maximum IOPS (I/O requests per second) and MBPS (megabytes of I/O per second) that can be sustained by the storage subsystem.

The I/O calibration occurs in two steps:

- In the first step of I/O calibration with the DBMS_RESOURCE_MANAGER.CALIBRATE_IO procedure, the procedure issues random database-block-sized reads, by default, 8 KB, to all data files from all database instances. This step provides the maximum IOPS, in the output parameter max_iops, that the database can sustain. The value max_iops is an important metric for OLTP databases. The output parameter actual_latency provides the average latency for this workload. When you need a specific target latency, you can specify the target latency with the input parameter max_latency (specifies the maximum tolerable latency in milliseconds for database-block-sized IO requests).
- The second step of calibration using the DBMS_RESOURCE_MANAGER.CALIBRATE_IO procedure issues random, 1 MB reads to all data files from all database instances. The second step yields the output parameter max_mbps, which specifies the maximum MBPS of I/O that the database can sustain. This step provides an important metric for data warehouses.

The calibration runs more efficiently if the user provides the <code>num_physical_disks</code> input parameter, which specifies the approximate number of physical disks in the database storage system.

Due to the overhead from running the I/O workload, I/O calibration should only be performed when the database is idle, or during off-peak hours, to minimize the impact of the I/O workload on the normal database workload.

To run I/O calibration and assess the I/O capability of the storage subsystem used by Oracle Database, use the DBMS RESOURCE MANAGER.CALIBRATE IO procedure:

```
SET SERVEROUTPUT ON
DECLARE
  lat NUMBER;
  iops INTEGER;
  mbps INTEGER;
BEGIN
-- DBMS_RESOURCE_MANAGER.CALIBRATE_IO (<DISKS>, <MAX_LATENCY>, iops, mbps, lat);
  DBMS_RESOURCE_MANAGER.CALIBRATE_IO (2, 10, iops, mbps, lat);
end;
//
```

When running the DBMS RESOURCE MANAGER.CALIBRATE IO procedure, consider the following:

- Only run one calibration at a time on databases that use the same storage subsystem. If you simultaneously run the calibration across separate databases that use the same storage subsystem, the calibration will fail.
- Quiesce the database to minimize I/O on the instance.
- For Oracle Real Application Clusters (Oracle RAC) configurations, ensure that all instances are opened to calibrate the storage subsystem across nodes.
- For an Oracle Real Application Clusters (Oracle RAC) database, the workload is simultaneously generated from all instances.
- The num_physical_disks input parameter is optional. By setting the num_physical_disks parameter to the approximate number of physical disks in the database's storage system, the calibration can be faster and more accurate.
- In some cases, asynchronous I/O is permitted for data files, but the I/O subsystem for submitting asynchronous I/O may be maximized, and I/O calibration cannot continue. In such cases, refer to the port-specific documentation for information about checking the maximum limit for asynchronous I/O on the system.

At any time during the I/O calibration process, you can query the calibration status in the V\$IO_CALIBRATION_STATUS view. After I/O calibration is successfully completed, you can view the results in the DBA RSRC IO CALIBRATE table.

See Also:

- Oracle Database PL/SQL Packages and Types Reference for more information about running the DBMS RESOURCE MANAGER.CALIBRATE IO procedure
- Oracle Database Reference for more information about the V\$IO_CALIBRATION_STATUS view and DBA_RSRC_IO_CALIBRATE table

I/O Calibration with the Oracle Orion Calibration Tool

This section describes the Oracle Orion Calibration Tool and includes the following sections:

- Introduction to the Oracle Orion Calibration Tool
- · Getting Started with Orion
- Orion Input Files
- Orion Parameters
- Orion Output Files
- Orion Troubleshooting

Introduction to the Oracle Orion Calibration Tool

Oracle Orion is a tool for predicting the performance of an Oracle database without having to install Oracle or create a database. Unlike other I/O calibration tools, Oracle Orion is expressly designed for simulating Oracle database I/O workloads using the same I/O software stack as Oracle. Orion can also simulate the effect of striping performed by Oracle Automatic Storage Management.

Table 15-4 lists the types of I/O workloads that Orion supports.

For each type of workload shown in Table 15-4, Orion can run tests using different I/O loads to measure performance metrics such as MBPS, IOPS, and I/O latency. Load is expressed in terms of the number of outstanding asynchronous I/Os. Internally, for each such load level, the Orion software keeps issuing I/O requests as fast as they complete to maintain the I/O load at that level. For random workloads, using either large or small sized I/Os, the load level is the number of outstanding I/Os. For large sequential workloads, the load level is a combination of the number of sequential streams and the number of outstanding I/Os per stream. Testing a given workload at a range of load levels can help you understand how performance is affected by load.

Note the following when you use Orion:

- Run Orion when the storage is idle (or pretty close to idle). Orion calibrates the
 performance of the storage based on the I/O load it generates; Orion is not able to properly
 assess the performance if non-Orion I/O workloads run simultaneously.
- If a database has been created on the storage, the storage can alternatively be calibrated using the PL/SQL routine dbms resource manager.calibrate io().

Table 15-4 Orion I/O Workload Support

Workload	Description
Small Random I/O	OLTP applications typically generate random reads and writes whose size is equivalent to the database block size, typically 8 KB. Such applications typically care about the throughput in I/Os Per Second (IOPS) and about the average latency (I/O turn-around time) per request. These parameters translate to the transaction rate and transaction turn-around time at the application layer.
	Orion simulates a random I/O workload with a given percentage of reads compared to writes, a given I/O size, and a given number of outstanding I/Os. In this Orion workload simulation, the I/Os are distributed across all disks.



	Table 15-4	(Cont.)	Orion I/O	Workload	Support
--	-------------------	---------	-----------	----------	---------

Workload	Description	
Large Sequential I/O	Data warehousing applications, data loads, backups, and restores generate sequential read and write streams composed of multiple outstanding 1 MB I/Os. Such applications are processing large amounts of data, such as a whole table or a whole database and they typically care about the overall data throughput in MegaBytes Per Second (MBPS).	
	Orion can simulate a given number of sequential read or write streams of a given I/O size with a given number of outstanding I/Os. Orion can optionally simulate Oracle Automatic Storage Management striping when testing sequential streams.	
Large Random I/O	A sequential stream typically accesses the disks concurrently with other database traffic. With striping, a sequential stream is spread across many disks. Consequently, at the disk level, multiple sequential streams are seen as random 1 MB I/Os.	
Mixed Workloads	Orion can simulate two simultaneous workloads: Small Random I/O and either Larg Sequential I/O or Large Random I/O. This workload type enables you to simulate, for example, an OLTP workload of 8 KB random reads and writes with a backup worklo four sequential read streams of 1 MB I/Os.	

Each Orion data point is a test for a specific mix of small and large I/O loads sustained for a duration. An Orion test consists of multiple data point tests. These data point tests can be represented as a two-dimensional matrix. Each column in the matrix represents data point tests with the same small I/O load, but varying large I/O loads. Each row represents data point tests with the same large I/O load, but varying small I/O loads. An Orion test can be for a single point, a single row, a single column, or for the whole matrix.

Orion Test Targets

You can use Orion to test any disk-based character device that supports asynchronous I/O. Orion has been tested on the following types of targets:

- DAS (direct-attached) storage: You can use Orion to test the performance of one or more local disks, volumes, or files on the local host.
- SAN (storage-area network) storage: Orion can be run on any host that has all or parts of
 the SAN storage mapped as character devices. The devices can correspond to striped or
 un-striped volumes exported by the storage array(s), or individual disks, or one or more
 whole arrays.
- NAS (network-attached storage): You can use Orion to test the performance on data files
 on NAS storage. In general, the performance results on NAS storage are dependent on the
 I/O patterns with which the data files have been created and updated. Therefore, you
 should initialize the data files appropriately before running Orion.

Orion for Oracle Administrators

Oracle administrators can use Orion to evaluate and compare different storage arrays, based on the expected workloads. Oracle administrators can also use Orion to determine the optimal number of network connections, storage arrays, storage array controllers, and disks for the expected peak workloads.

Getting Started with Orion

To get started using Orion, do the following:

- Select a test name to use with the Orion -testname parameter. This parameter specifies a
 unique identifier for your Orion run. For example, use the test name "mytest". For more
 information, see "Orion Parameters".
- 2. Create an Orion input file, based on the test name. For example, create a file named mytest.lun. In the input file list the raw volumes or files to test. Add one volume name per line. Do not put comments or anything else in the .lun file.

For example, an Orion input file could contain the following:

```
/dev/raw/raw1
/dev/raw/raw3
/dev/raw/raw4
/dev/raw/raw5
/dev/raw/raw6
/dev/raw/raw7
```

For more information, see "Orion Input Files".

3. Verify that the all volumes specified in the input file, for example mytest.lun, are accessible using the command dd or another equivalent file viewing utility. For example, for a typical sanity-check try the following on a Linux system:

```
$ dd if=/dev/raw/raw1 of=/dev/null bs=32k count=1024
```

Depending on your platform, the file viewing utility you use and its interface may be different.

- 4. Verify that your platform has the necessary libraries installed to do asynchronous I/Os. The Orion test is completely dependent on asynchronous I/O. On Linux and Solaris, the library libaio must be in the standard lib directories or accessible through the shell environment's library path variable (usually LD_LIBRARY_PATH or LIBPATH, depending on your shell). Windows has built-in asynchronous I/O libraries, so this issue does not apply.
- 5. As a first test with Orion, use -run with either the oltp or dss option. If the database is primarily OLTP, then use -run oltp. If the database is primarily for data warehousing or analytics, then use -run dss.

For example, use the following command to run an OLTP-like workload using the default input file name, orion.lun:

```
$ ./orion -run oltp
```

The I/O load levels generated by Orion take into account the number of disk spindles being tested (or specified with the <code>-num_disks</code> parameter). Keep in mind that the number of spindles *may or may not be* related to the number of volumes specified in the input file, depending on how these volumes are mapped.

6. The section "Orion Output Files" provides sample results showing the Orion output files. Using the sample file mytest_summary.txt is a good starting point for verifying the input parameters and analyzing the output. The sample files mytest_*.csv contain commadelimited values for several I/O performance measures.

Orion Input Files

When you specify the Orion -testname <testname> parameter, this sets the test name prefix for the Orion input and output filenames. The default value for the -testname option is "orion".

The Orion input file, <testname>.lun should contain a carriage-return-separated list of LUNs.

Orion Parameters

Use the Orion command parameters to specify the I/O workload type and to specify other Orion options.

Orion Required Parameter

The -run parameter is required with the Orion command. Table 15-5 describes the -run parameter.



Table 15-5 Required Orion Parameter

Option Description Default

-run *level*

Specifies the test run level to be *level*. This option provides the run level and allows complex commands to be specified at the advanced level. If not set as -run advanced, then setting any other parameter, besides -cache size or -verbose, results in an error.

normal

Except advanced, all of the -run level settings use a pre-specified set of parameters.

The level must be one of:

oltp

Tests with random small (8K) I/Os at increasing loads to determine the maximum IOPS.

This parameter corresponds to the following Orion invocation:

```
%> ./orion -run advanced \
   -num_large 0 -size_small 8 -type rand \
   -simulate concat -write 0 -duration 60 \
   -matrix row
```

dss

Tests with random large (1M) I/Os at increasing loads to determine the maximum throughput.

This parameter corresponds to the following Orion invocation:

```
%> ./orion -run advanced \
    -num_small 0 -size_large 1024 -type rand \
    -simulate concat -write 0 -duration 60 \
    -matrix column
```

simple

Generates the Small Random I/O and the Large Random I/O workloads for a range of load levels. In this option, small and large I/Os are tested in isolation. The only optional parameters that can be specified at this run level are -cache size and -verbose.

This parameter corresponds to the following Orion invocation:

```
%> ./orion -run advanced \
-size_small 8 -size_large 1024 -type rand \
-simulate concat -write 0 -duration 60 \
-matrix basic
```

normal

Same as simple, but also generates combinations of the small random I/O and large random I/O workloads for a range of loads. The only optional parameters that can be specified at this run level are -cache size and -verbose.

This parameter corresponds to the following Orion invocation:

```
%> ./orion -run advanced \
-size_small 8 -size_large 1024 -type rand \
-simulate concat -write 0 -duration 60 \
-matrix detailed
```

advanced

Tests the workload you specify with optional parameters. Any of the optional parameters can be specified at this run level.

Orion Optional Parameters

Table 15-6 Optional Orion Parameters

Option	Description	Default
-cache_size <i>num</i>	Size of the storage array's read or write cache (in MB). For Large Sequential I/O workloads, Orion warms the cache by doing random large I/Os before each data point. Orion uses the cache size to determine the duration for this cache warming operation. If set to 0, do not perform cache warming.	Default Value: If not specified, warming occurs for a default amount of time (two minutes). That is, issue two minutes of
	Unless this option is set to 0, Orion issues several unmeasured, random I/Os before each large sequential data point. These I/Os fill up the storage array's cache, if any, with random data so that I/Os from one data point do not result in cache hits for the next data point. Read tests are preceded with junk reads and write tests are preceded with junk writes. If specified, this 'cache warming' is performed until num MBs of I/O have been read or written.	unmeasured random I/Os before each data point.
-duration num_seconds	Set the duration to test each data point in seconds to the value <i>num_seconds</i> .	Default Value: 60
-help	Prints Orion help information. All other options are ignored with help set.	
-matrix <i>type</i>	Type of mixed workloads to test over a range of loads. An Orion test consists of multiple data point tests. The data point tests can be represented as a two-dimensional matrix.	Default Value: basic
	Each column in the matrix represents data point tests with the same small I/O load, but varying large I/O loads. Each row represents data point tests with the same large I/O load, but varying small I/O loads. An Orion test can be for a single point, a single row, a single column, or the whole matrix, depending on the matrix <i>type</i> :	
	 basic: No mixed workload. The Small Random and Large Random/Sequential workloads are tested separately. Test small I/Os only, then large I/Os only. 	
	 detailed: Small Random and Large Random/Sequential workloads are tested in combination. Test entire matrix. 	
	 point: A single data point with S outstanding Small Random I/Os and L outstanding Large Random I/Os or sequential streams. S is set by the -num_small parameter. L is set by the -num_large parameter. Test with -num_small small I/Os, -num_large large I/Os. 	
	 col: Large Random/Sequential workloads only. Test a varying large I/O load with -num_small small I/Os. 	
	 row: Small Random workloads only. Test a varying small I/O load with -num_large large I/Os. 	
	 max: Same as detailed, but only tests the workload at the maximum load, specified by the -num_small and -num_large parameters. Test varying loads up to the -num_small and - num_large limits. 	
-num_disks <i>value</i>	Specify the number of physical disks used by the test. Used to generate a range for the load. Specifies the number of disks (physical spindles). This number <i>value</i> is used to gauge the range of loads that Orion should test at. Increasing this parameter results in Orion using heavier I/O loads.	Default Value: the number of LUNs in < testname>.lun.



Table 15-6 (Cont.) Optional Orion Parameters

Option	Description	Default
-num_large <i>value</i>	Controls the large I/O load. Note, this option only applies when -matrix is specified as: row, point, or max.	Default Value: no default
	When the -type option is set to rand, the parameter argument <i>value</i> specifies the number of outstanding large I/Os.	
	When the -type option is set to seq, the parameter argument <i>value</i> specifies the number of sequential I/O streams.	
-num_small	Specify the maximum number of outstanding I/Os for the Small Random I/O workload.	Default Value: no default
	Note: this only applies when ${\operatorname{-matrix}}$ is specified as ${\operatorname{col}},$ ${\operatorname{point}},$ or ${\operatorname{max}}.$	
-num_streamIO <i>num</i>	Specify the number of concurrent I/Os per stream as <i>num</i> .	Default Value: 4
	Note: this parameter is only used if -type is seq.	
-simulate <i>type</i>	Data layout to simulate for Large Sequential I/O workload. Orion tests on a virtual LUN formed by combining specified LUNs in one of these ways. The <i>type</i> is one:	Default Value: concat
	 concat: A virtual volume is simulated by serially chaining the specified LUNs. A sequential test over this virtual volume will go from some point to the end of each one LUN, followed by the beginning to end of the next LUN, and so on. 	
	 raid0: A virtual volume is simulated by striping across the specified LUNs. Each sequential stream issues I/Os across all LUNs using raid0 striping. The stripe depth is 1M by default, to match the Oracle Automatic Storage Management stripe depth, and can be changed with the -stripe parameter. 	
	The offsets for I/Os are determined as follows:	
	For Small Random and Large Random workloads:	
	 The LUNs are concatenated into a single virtual LUN (VLUN) and random offsets are chosen within the VLUN. 	
	For Large Sequential workloads:	
	• With striping (-simulate raid0). The LUNs are used to create a single striped VLUN. With no concurrent Small Random workload, the sequential streams start at fixed offsets within the striped VLUN. For n streams, stream i start at offset VLUNsize * (i+1) / (n+1), unless n is 1, in which case the single stream start at offset 0. With a concurrent Small Random workload, streams start at random offsets within the striped VLUN.	
	 Without striping (-simulate CONCAT). The LUNs are concatenated into a single VLUN. The streams start at random offsets within the single VLUN. 	
	This parameter is typically only used if -type is seq.	
-size_large <i>num</i>	Specify the <i>num</i> , size of the I/Os (in KB) for the Large Random or Sequential I/O workload.	Default Value: 1024
-size_small <i>num</i>	Specify the num, size of the I/Os (in KB) for the Small Random I/O workload.	Default Value: 8



Table 15-6 (Cont.) Optional Orion Parameters

Option	Description	Default
-storax <i>type</i>	API to use for testing I/O workload.	Default Value: skgfr
	 skgfr: Use operating system I/O layer. 	
	 oss: Use OSS API for I/O with Cell server in an Exadata machine. 	
	 asmlib: Use ASMLIB disk devices based storage API for I/O. odmlib: Use Direct NFS storage based API for I/O. 	
-testname <i>tname</i>	Specify the <i>tname</i> identifier for the test run. When specified, the input file containing the LUN disk or file names must be named < <i>tname</i> >.lun.	Default Value: orion
	The output files are named with the prefix <tname></tname>	
-type [rand seq]	Type of the Large I/O workload.	Default Value: rand
	 rand: Randomly distributed large I/Os. 	
	• seq: Sequential streams of large I/Os.	
-verbose	Prints status and tracing information to standard output.	Default Value: option not set
-write <i>num_write</i>	Specify the percentage of I/Os that are writes to <i>num_write</i> ; the rest being reads.	Default Value: 0
	This parameter applies to both the Large and Small I/O workloads. For Large Sequential I/Os, each stream is either read-only or write-only; the parameter specifies the percentage of streams that are write-only. The data written to disk is garbage and unrelated to any existing data on the disk.	
	Caution: write tests obliterate all data on the specified LUNS.	



Write tests obliterate all data on the specified LUNS.

Orion Command Line Samples

The following provides sample Orion commands for different types of I/O workloads:

1. To evaluate storage for an OLTP database:

-run oltp

2. To evaluate storage for a data warehouse:

-run dss

3. For a basic set of data:

-run normal

4. To understand your storage performance with read-only, small and large random I/O workload:

\$ orion -run simple

To understand your storage performance with a mixed small and large random I/O workload:

```
$ orion -run normal
```

6. To generate combinations of 32KB and 1MB reads to random locations:

```
$ orion -run advanced -size_small 32 \
-size large 1024 -type rand -matrix detailed
```

7. To generate multiple sequential 1 MB write streams, simulating 1 MB RAID-0 stripes:

```
$ orion -run advanced -simulate raid0 \
-stripe 1024 -write 100 -type seq -matrix col -num small 0
```

8. To generate combinations of 32 KB and 1 MB reads to random locations:

```
-run advanced -size_small 32 -size_large 1024 -type rand -matrix detailed
```

9. To generate multiple sequential 1 MB write streams, simulating RAID0 striping:

```
-run advanced -simulate raid0 -write 100 -type seq -matrix col -num small 0
```

Orion Configuration for Exadata Storage Cells

The following steps configure Orion to run against Exadata storage cells.

1. Create \$OSSCONF/cellip.ora. Get the IP addresses using cellcli.

2. Check griddisk and create orion.lun file.

```
cellcli -e list griddisk
        orion FD 00 scaqaj04celadm08
                                       active
        orion FD 01 scaqaj04celadm08
                                      active
        orion FD 02 scaqaj04celadm08
                                      active
        orion FD 03 scaqaj04celadm08
                                       active
        orion FD 04 scaqaj04celadm08
                                      active
        orion FD 05 scaqaj04celadm08
                                     active
        orion FD 06 scaqaj04celadm08
                                       active
        orion FD 07 scaqaj04celadm08
                                        active
```

```
cat orion.lun o/192.168.0.31;192.168.0.32/orion_FD_00_scaqaj04celadm08 o/192.168.0.31;192.168.0.32/orion_FD_01_scaqaj04celadm08 o/192.168.0.31;192.168.0.32/orion_FD_02_scaqaj04celadm08 o/192.168.0.31;192.168.0.32/orion_FD_03_scaqaj04celadm08 o/192.168.0.31;192.168.0.32/orion_FD_04_scaqaj04celadm08 o/192.168.0.31;192.168.0.32/orion_FD_05_scaqaj04celadm08 o/192.168.0.31;192.168.0.32/orion_FD_06_scaqaj04celadm08 o/192.168.0.31;192.168.0.32/orion_FD_07_scaqaj04celadm08
```



Check cellinit.ora. If Orion is being run from compute node, the IP addresses will be different from the storage cells. This file should exist by default on a configured Exadata system.

```
cat $0SSCONF/cellinit.ora | grep ip
ipaddress2=192.168.0.32/20
ipaddress1=192.168.0.31/20
```

4. Run Orion.

```
which orion` -type rand -write 0 -run advanced -duration 30 -matrix point -num small 20 -num large 0 -testname orion
```

Orion Output Files

The output files for a test run are prefixed by <testname>_<date> where date is yyyymmdd_hhmm.

Table 15-7 lists the Orion output files.

Table 15-7 Orion Generated Output Files

Output File	Description
<testname>_<date>_hist.csv</date></testname>	Histogram of I/O latencies.
<testname>_<date>_iops.csv</date></testname>	Performance results of small I/Os in IOPS.
<testname>_<date>_lat.csv</date></testname>	Latency of small I/Os in microseconds.
<testname>_<date>_mbps.csv</date></testname>	Performance results of large I/Os in MBPS.
<testname>_<date>_summary.txt</date></testname>	Summary of the input parameters, along with the minimum small I/O latency (in secs), the maximum MBPS, and the maximum IOPS observed.
<testname>_<date>_trace.txt</date></testname>	Extended, unprocessed output.



If you are performing write tests, be prepared to lose any data stored on the LUNs.

Orion Sample Output Files

Orion creates several output files as specified in Table 15-7. For the sample "mytest" shown in the section, "Getting Started with Orion", the output files are:

- mytest_summary.txt: This file contains:
 - Input parameters
 - Maximum throughput observed for the Large Random/Sequential workload
 - Maximum I/O rate observed for the Small Random workload
 - Minimum latency observed for the Small Random workload
- mytest_mbps.csv: comma-delimited value file containing the data transfer rate (MBPS) results for the Large Random/Sequential workload. In the general case, this and all other

CSV files contains a two-dimensional table. Each row in the table corresponds to a large I/O load level and each column corresponds to a specific small I/O load level. Thus, the column headings are the number of outstanding small I/Os and the row headings are the number of outstanding large I/Os (for random large I/O tests) or the number of sequential streams (for sequential large I/O tests).

The following example shows the first few data points of the Orion MBPS output CSV file for "mytest". The simple mytest command-line does not test combinations of large and small I/Os. Hence, the MBPS file has just one column corresponding to 0 outstanding small I/Os. In this example, at a load level of 8 outstanding large reads and no small I/Os, the report data indicates a throughput of 103.06 MBPS.

The following graph shows a sample data transfer rate measured at different large I/O load levels. This graph can be generated by loading <code>mytest_mbps.csv</code> into a spreadsheet and graphing the data points. Orion does not directly generate such graphs. The x-axis corresponds to the number of outstanding large reads and the y-axis corresponds to the throughput observed.

The graph shows typical storage system behavior. As the number of outstanding I/O requests is increased, the throughput increases. However, at a certain point the throughput level stabilizes, indicating the storage system's maximum throughput value.

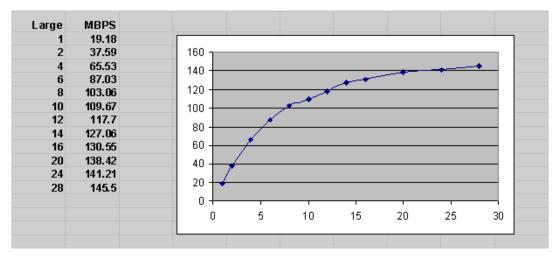


Figure 15-2 Sample I/O Load Levels

mytest_iops.csv: Comma-delimited value file containing the I/O throughput (in IOPS) results for the Small Random workload. Like in the MBPS file, the column headings are the number of outstanding small I/Os and the row headings are the number of outstanding large I/Os, when testing large random, or the number of sequential streams (for large sequential).

In the general case, a CSV file contains a two-dimensional table. However, for a simple test where you are not testing combinations of large and small I/Os the results file has just

one row. Hence, the IOPS results file just has one row with 0 large I/Os. As shown in the following example, data point with 12 outstanding small reads and no large I/Os provides a sample throughput of 951 IOPS.

```
Large/Small, 1, 2, 3, 6, 9, 12 . . . . 0, 105, 208, 309, 569, 782, 951 . . .
```

The following graph is generated by loading mytest_iops.csv into Excel and charting the data. This graph illustrates the IOPS throughput seen at different small I/O load levels.

The graph shows typical storage system behavior. As the number of outstanding I/O requests is increased, the throughput increases. However, at a certain point, the throughput level stabilizes, indicating the storage system reaches a maximum throughput value. At higher throughput levels, the latency for the I/O requests also increase significantly. Therefore, it is important to view this data with the latency data provided in the generated latency results in mytest lat.csv.

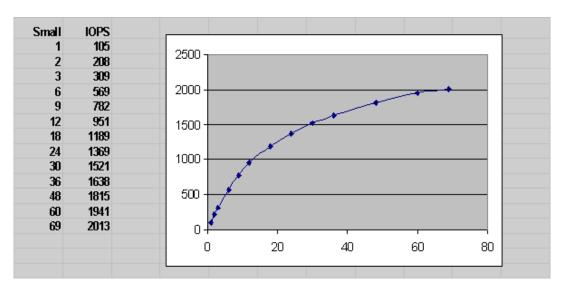


Figure 15-3 I/O Throughput at Different Small I/O Load Levels

 mytest_lat.csv: Comma-delimited value file containing the latency results for the Small Random workload. As with the MBPS and IOPS files, the column headings are the number of outstanding small I/Os and the row headings are the number of outstanding large I/Os (when testing large random I/Os) or the number of sequential streams.

In the general case, a CSV file contains a two-dimensional table. However, for a simple test where you are not testing combinations of large and small I/Os the results file has just one row. Hence, the IOPS results file just has one row with 0 large I/Os. In the following example, at a sustained load level of 12 outstanding small reads and no large I/Os, the generated results show an I/O turn-around latency of 22.25 milliseconds.

The following graph is generated by loading mytest_lat.csv into Excel and charting the data. This graph illustrates the small I/O latency at different small I/O load levels for mytest.



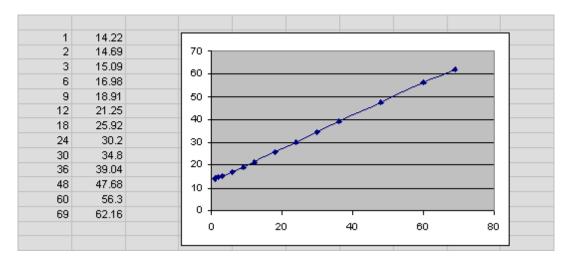


Figure 15-4 I/O Latency at Small I/O Load Levels

mytest_trace.txt: Contains the extended, unprocessed test output.

Note:

Orion reports errors that occur during a test on standard output.

Orion Troubleshooting

- If you are getting an I/O error on one or more of the volumes specified in the <testname>.lun file:
 - Verify that you can access the volume in the same mode as the test, read or write, using a file copy program such as dd.
 - Verify that your host operating system version can do asynchronous I/O.
 - On Linux and Solaris, the library libaio must be in the standard lib directories or accessible through the shell environment's library path variable (usually LD_LIBRARY_PATH or LIBPATH, depending on your shell).
- 2. If you run on NAS storage:
 - The file system must be properly mounted for Orion to run. Please consult your Oracle Installation Guide for directions (for example, the section, Appendix B "Using NAS Devices" in the Database Installation Guide for Linux x86).
 - The mytest.lun file should contain one or more paths of existing files. Orion does not work on directories or mount points. The file has to be large enough for a meaningful test. The size of this file should represent the eventual expected size of your datafiles (say, after a few years of use).
 - You may see poor performance doing asynchronous I/O over NFS on Linux (including 2.6 kernels).
 - If you are doing read tests and the reads are hitting untouched blocks of the file that
 were not initialized or previously written, some smart NAS systems may "fake" the read
 by returning zeroed-out blocks. When this occurs, you see unexpectedly good
 performance.



The workaround is to write all blocks, using a tool such as dd, before performing the read test.

- 3. If you run Orion on Windows: Testing on raw partitions requires temporarily mapping the partitions to drive letters and specifying these drive letters in the test.lun file.
- 4. If you run Orion 32-bit Linux/x86 binary on an x86_64 system: Please copy a 32-bit libaio.so file from a 32-bit computer running the same Linux version.
- 5. If you are testing with a lot of disks (num disks greater than around 30):
 - You should use the -duration option (see the optional parameters section for more details) to specify a long duration (like 120 seconds or more) for each data point. Since Orion tries to keep all the spindles running at a particular load level, each data point requires a ramp-up time, which implies a longer duration for the test.
 - You may get the following error message, instructing you to increase the duration value:

```
Specify a longer -duration value.
```

A duration of 2x the number of spindles seems to be a good rule of thumb. Depending on your disk technology, your platform may need more or less time.

- 6. If you get an error about libraries being used by Orion:
 - Linux/Solaris: See I/O error troubleshooting.
 - **NT-Only:** Do not move/remove the Oracle libraries included in the distribution. These must be in the same directory as orion.exe.
- 7. If you are seeing performance numbers that are "unbelievably good":
 - You may have a large read or write cache, or read and write cache somewhere between the Orion program and the disk spindles. Typically, the storage array controller has the biggest effect. Find out the size of this cache and use the cache_size advanced option to specify it to Orion (see the optional parameters section for more details).
 - The total size of your volumes may be really small compared to one or more caches along the way. Try to turn off the cache. This is needed if the other volumes sharing your storage show significant I/O activity in a production environment (and end up using large parts of the shared cache).
- 8. If Orion is reporting a long estimated run time:
 - The run time increases when -num_disks is high. Orion internally uses a linear formula to determine how long it takes to saturate the given number of disks.
 - The -cache_size parameter affects the run time, even when it is not specified. Orion does cache warming for two minutes per data point by default. If you have turned off the cache, specify -cache size 0.
 - The run time increases when a long -duration value is specified, as expected.



Managing Operating System Resources

This chapter explains how to tune the operating system for optimal performance of Oracle Database.

This chapter contains the following sections:

- Understanding Operating System Performance Issues
- Resolving Operating System Issues
- Understanding CPU
- Resolving CPU Issues

See Also:

- Your operating system documentation
- Your Oracle Database platform-specific documentation, which contains tuning information specific to your platform

Understanding Operating System Performance Issues

Operating system performance issues commonly involve process management, memory management, and scheduling. If you have tuned the Oracle database instance and still need to improve performance, verify your work or try to reduce system time. Ensure that there is enough I/O bandwidth, CPU power, and swap space. Do not expect, however, that further tuning of the operating system will have a significant effect on application performance. Changes in the Oracle Database configuration or in the application are likely to result in a more significant difference in operating system efficiency than simply tuning the operating system.

For example, if an application experiences excessive buffer busy waits, then the number of system calls increases. If you reduce the buffer busy waits by tuning the application, then the number of system calls decreases.

This section covers the following topics related to operating system performance issues:

- Using Operating System Caches
- Memory Usage
- Using Operating System Resource Managers

Using Operating System Caches

Operating systems and device controllers provide data caches that do not directly conflict with Oracle Database cache management. Nonetheless, these structures can consume resources while offering little or no performance benefit. This situation is most noticeable when database files are stored in a Linux or UNIX file system. By default, all database I/O goes through the file system cache.

On some Linux and UNIX systems, direct I/O is available to the filestore. This arrangement allows the database files to be accessed within the file system, bypassing the file system cache. Direct I/O saves CPU resources and allows the file system cache to be dedicated to non-database activity, such as program texts and spool files.

Note:

This problem does not occur on Windows. All file requests by the database bypass the caches in the file system.

Although the operating system cache is often redundant because the Oracle Database buffer cache buffers blocks, in some cases the database does not use the database buffer cache. In these cases, using direct I/O or raw devices may yield worse performance than using operating system buffering. Examples include:

- Reads or writes to the TEMP tablespace
- Data stored in NOCACHE LOBs
- Parallel execution servers reading data

Note:

In some cases the database can cache parallel query data in the database buffer cache instead of performing direct reads from disk into the PGA. This configuration may be appropriate when the database servers have a large amount of memory. See *Oracle Database VLDB and Partitioning Guide* to learn more using parallel execution.

You may want to cache but not all files at the operating system level.

Asynchronous I/O

With synchronous I/O, when an I/O request is submitted to the operating system, the writing process blocks until the write is confirmed as complete. It can then continue processing. With asynchronous I/O, processing continues while the I/O request is submitted and processed. Use asynchronous I/O when possible to avoid bottlenecks.

Some platforms support asynchronous I/O by default, others need special configuration, and some only support asynchronous I/O for certain underlying file system types.

FILESYSTEMIO_OPTIONS Initialization Parameter

You can use the <code>FILESYSTEMIO_OPTIONS</code> initialization parameter to enable or disable asynchronous I/O or direct I/O on file system files. This parameter is platform-specific and has a default value that is best for a particular platform.

FILESYSTEMIO OPTIONS can be set to one of the following values:

- ASYNCH: enable asynchronous I/O on file system files, which has no timing requirement for transmission.
- DIRECTIO: enable direct I/O on file system files, which bypasses the buffer cache.



- SETALL: enable both asynchronous and direct I/O on file system files.
- NONE: disable both asynchronous and direct I/O on file system files.



Your platform-specific documentation for more details

Limiting Asynchronous I/O in NFS Server Environments

In some Network File Storage (NFS) server environments, performance may be impaired if a large number of asynchronous I/O requests are made within a short period of time. In such cases, use the <code>DNFS_BATCH_SIZE</code> initialization parameter to improve performance and increase stability on your system by limiting the number of I/Os issued by an Oracle process.

The DNFS_BATCH_SIZE initialization parameter controls the number of asynchronous I/Os that can be queued by an Oracle foreground process when Direct NFS Client is enabled. In environments where the NFS server cannot handle a large number of outstanding asynchronous I/O requests, Oracle recommends setting this parameter to a value of 128. You can then increase or decrease its value based on the performance of your NFS server.

Note:

The default setting for the <code>DNFS_BATCH_SIZE</code> initialization parameter is 4096. The recommended value of 128 is only applicable on systems where the NFS server cannot handle a large number of asynchronous I/O requests and severe latency is detected.

See Also:

Oracle Database Reference for information about the DNFS_BATCH_SIZE initialization parameter

Improving I/O Performance Using Direct NFS Client

Direct NFS Client integrates the NFS client functionality directly in Oracle Database. Because Direct NFS Client is a specialized NFS client for Oracle Database, it is highly optimized. Direct NFS Client considerably improves database performance over NFS as compared to the traditional operating system NFS client.

Parallel NFS is an optional feature of Direct NFS Client that is introduced in NFS version 4.1 and is supported by Oracle Database 12c Release 2 (12.2) and later. Parallel NFS is a highly scalable distributed storage protocol, where clients, server, and storage devices are responsible for managing file access. In the NFS versions earlier to 4.1, only the server is responsible for managing file access. Thus, Parallel NFS enables highly scalable distributed NAS storage for better I/O performance.

Starting with Oracle Database 12c Release 2 (12.2), you can also use the Direct NFS dispatcher feature of Direct NFS Client. The Direct NFS dispatcher consolidates the TCP

connections that are created from a database instance to an NFS server. In large database deployments, using Direct NFS dispatcher improves scalability and network performance. Therefore, for a large number of TCP connections, Oracle recommends using Direct NFS dispatcher along with Parallel NFS for a Direct NFS Client deployment.

The Direct NFS Client supports the following NFS protocol versions: NFSv3, NFSv4, NFSv4.1, and pNFS. The default version is NFSv3. pNFS is used for Direct NFS with Parallel NFS. Direct NFS supports only the default sys security authentication with Parallel NFS. Direct NFS does not support Parallel NFS when combined with any of the Kerberos authentication parameters.

See Also:

- Oracle Database Installation Guide for information about enabling the Parallel NFS feature for Direct NFS Client by setting the value for the nfs_version parameter to pNFS in the Direct NFS configuration file oranfstab.
- Oracle Database Reference for information about enabling the Direct NFS dispatcher feature for the Direct NFS Client by setting the value for the ENABLE_DNFS_DISPATCHER initialization parameter to true.

Memory Usage

Memory usage is affected by both buffer cache limits and initialization parameters.

Buffer Cache Limits

The UNIX buffer cache consumes operating system memory resources. Although in some versions of UNIX, the UNIX buffer cache may be allocated a set amount of memory, it is common today for more sophisticated memory management mechanisms to be used. Typically, these will allow free memory pages to be used to cache I/O. In such systems, it is common for operating system reporting tools to show that there is no free memory, which is not generally a problem. If processes require more memory, the memory caching I/O data is usually released to allow the process memory to be allocated.

Parameters Affecting Memory Usage

The memory required by any one Oracle Database session depends on many factors. Typically the major contributing factors are:

- Number of open cursors
- Memory used by PL/SQL, such as PL/SQL tables
- SORT AREA SIZE initialization parameter

In Oracle Database, the PGA_AGGREGATE_TARGET initialization parameter gives greater control over a session's memory usage.

Using Operating System Resource Managers

Some platforms provide operating system resource managers. These are designed to reduce the impact of peak load use patterns by prioritizing access to system resources. They usually implement administrative policies that govern which resources users can access and how much of those resources each user is permitted to consume.



Operating system resource managers are different from domains or other similar facilities. Domains provide one or more completely separated environments within one system. Disk, CPU, memory, and all other resources are dedicated to each domain and cannot be accessed from any other domain. Other similar facilities completely separate just a portion of system resources into different areas, usually separate CPU or memory areas. Like domains, the separate resource areas are dedicated only to the processing assigned to that area; processes cannot migrate across boundaries. Unlike domains, all other resources (usually disk) are accessed by all partitions on a system.

Oracle Database runs within domains, and within these other less complete partitioning constructs, as long as the allocation of partitioned memory (RAM) resources is fixed, not dynamic.

Operating system resource managers prioritize resource allocation within a global pool of resources, usually a domain or an entire system. Processes are assigned to groups, which are in turn assigned resources anywhere within the resource pool.

Note:

- If you have multiple instances on a node, and you want to distribute resources among them, then each instance should be assigned to a dedicated operating-system resource manager group or managed entity. To run multiple instances in the managed entity, use instance caging to manage how the CPU resources within the managed entity should be distributed among the instances. When Oracle Database Resource Manager is managing CPU resources, it expects a fixed amount of CPU resources for the instance. Without instance caging, it expects the available CPU resources to be equal to the number of CPUs in the managed entity. With instance caging, it expects the available CPU resources to be equal to the value of the CPU_COUNT initialization parameter. If there are less CPU resources than expected, then Oracle Database Resource Manager is not as effective at enforcing the resource allocations in the resource plan.
- Oracle Database is not supported for use with any UNIX operating system
 resource manager's memory management and allocation facility. Oracle
 Database Resource Manager, which provides resource allocation capabilities
 within an Oracle database instance, cannot be used with any operating system
 resource manager.

For a complete list of operating system resource management and resource allocation and deallocation features that work with Oracle Database and Oracle Database Resource Manager, see your systems vendor and your Oracle representative. Oracle does not certify these system features for compatibility with specific release levels.

See Also:

- Oracle Database Administrator's Guide for information about Oracle Database Resource Manager
- Oracle Database Administrator's Guide for information about instance caging



Resolving Operating System Issues

This section provides hints for tuning various systems by explaining the following topics:

- Performance Hints on UNIX-Based Systems
- Performance Hints on Windows Systems
- Performance Hints on HP OpenVMS Systems

Familiarize yourself with platform-specific issues so that you know what performance options the operating system provides.



Your Oracle platform-specific documentation and your operating system vendor's documentation

Performance Hints on UNIX-Based Systems

On UNIX systems, try to establish a good ratio between the amount of time the operating system spends fulfilling system calls and doing process scheduling and the amount of time the application runs. The goal should be to run most of the time in application mode, also called user mode, rather than system mode.

The ratio of time spent in each mode is only a symptom of the underlying problem, which might involve the following:

- Paging or swapping
- Executing too many operating system calls
- · Running too many processes

If such conditions exist, then there is less time available for the application to run. The more time you can release from the operating system side, the more transactions an application can perform.

Performance Hints on Windows Systems

On Windows systems, as with UNIX-based systems, establish an appropriate ratio between time in application mode and time in system mode. You can easily monitor many factors with the Windows administrative performance tool: CPU, network, I/O, and memory are all displayed on the same graph to assist you in avoiding bottlenecks in any of these areas.

Performance Hints on HP OpenVMS Systems

Consider the paging parameters on a mainframe, and remember that Oracle Database can exploit a very large working set.

Free memory in HP OpenVMS environments is actually memory that is not mapped to any operating system process. On a busy system, free memory likely contains a page belonging to one or more currently active process. When that access occurs, a soft page fault takes place, and the page is included in the working set for the process. If the process cannot



expand its working set, then one of the pages currently mapped by the process must be moved to the free set.

Any number of processes might have pages of shared memory within their working sets. The sum of the sizes of the working sets can thus markedly exceed the available memory. When the Oracle server is running, the SGA, the Oracle Database kernel code, and the Oracle Forms run-time executable are normally all sharable and account for perhaps 80% or 90% of the pages accessed.

Understanding CPU

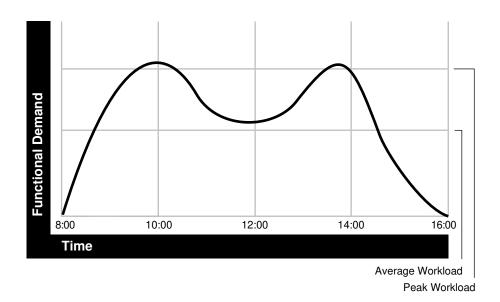
To address CPU problems, first establish appropriate expectations for the amount of CPU resources your system should be using. Then, determine whether sufficient CPU resources are available and recognize when your system is consuming too many resources. Begin by determining the amount of CPU resources the Oracle database instance utilizes with your system in the following three cases:

- System is idle, when little Oracle Database and non-Oracle activity exists
- System at average workloads
- System at peak workloads

You can capture various workload snapshots using the Automatic Workload Repository, Statspack, or the UTLBSTAT/UTLESTAT utility. Operating system utilities—such as vmstat, sar, and iostat on UNIX and the administrative performance monitoring tool on Windows—can be used along with the V\$OSSTAT or V\$SYSMETRIC_HISTORY view during the same time interval as Automatic Workload Repository, Statspack, or UTLBSTAT/UTLESTAT to provide a complimentary view of the overall statistics.

Workload is an important factor when evaluating your system's level of CPU utilization. During peak workload hours, 90% CPU utilization with 10% idle and waiting time can be acceptable. Even 30% utilization at a time of low workload can be understandable. However, if your system shows high utilization at normal workload, then there is no room for a peak workload. For example, The following figure illustrates workload over time for an application having peak periods at 10:00 AM and 2:00 PM.

Figure 16-1 Average Workload and Peak Workload





This example application has 100 users working 8 hours a day. Each user entering one transaction every 5 minutes translates into 9,600 transactions daily. Over an 8-hour period, the system must support 1,200 transactions an hour, which is an average of 20 transactions a minute. If the demand rate were constant, then you could build a system to meet this average workload.

However, usage patterns are not constant and in this context, 20 transactions a minute can be understood as merely a minimum requirement. If the peak rate you need to achieve is 120 transactions a minute, then you must configure a system that can support this peak workload.

For this example, assume that at peak workload, Oracle Database uses 90% of the CPU resource. For a period of average workload, then, Oracle Database uses no more than about 15% of the available CPU resource, as illustrated in the following equation:

```
20 tpm / 120 tpm * 90% = 15% of available CPU resource
```

where tpm is transactions a minute.

If the system requires 50% of the CPU resource to achieve 20 tpm, then a problem exists: the system cannot achieve 120 transactions a minute using 90% of the CPU. However, if you tuned this system so that it achieves 20 tpm using only 15% of the CPU, then, assuming linear scalability, the system might achieve 120 transactions a minute using 90% of the CPU resources.

As users are added to an application, the workload can rise to what had previously been peak levels. No further CPU capacity is then available for the new peak rate, which is actually higher than the previous.

Resolving CPU Issues

You can resolve CPU capacity issues by:

- Detecting and solving CPU problems from excessive consumption, as described in "Finding and Tuning CPU Utilization".
- Reducing the impact of peak load use patterns by prioritizing CPU resource allocation using Oracle Database Resource Manager, as described in "Managing CPU Resources Using Oracle Database Resource Manager".
- Using instance caging to limit the number of CPUs that a database instance can use simultaneously when running multiple database instances on a multi-CPU system, as described in "Managing CPU Resources Using Instance Caging".
- Increasing hardware capacity and improving the system architecture.

Finding and Tuning CPU Utilization

Every process running on your system affects the available CPU resources. Therefore, tuning non-database factors can also improve database performance.

Use the v\$osstat or vsysmetric_History$ view to monitor system utilization statistics from the operating system. Useful statistics contained in v\$osstat and vsysmetric_History$ include:

- Number of CPUs
- CPU utilization
- Load
- Paging



Physical memory



Oracle Database Reference for more information on ${\tt V\$OSSTAT}$ and ${\tt V\$SYSMETRIC_HISTORY}$

You can use operating system monitoring tools to determine which processes run on the system as a whole. If the system is too heavily loaded, check the memory, I/O, and process management areas described later in this section.

You can use tools such as <code>sar-u</code> on many UNIX-based systems to examine the level of CPU utilization on the system. In UNIX, statistics show user time, system time, idle time, and time waiting for I/O. A CPU problem exists if idle time and time waiting for I/O are both close to zero (less than 5%) at a normal or low workload.

On Windows, you can use the administrative performance tool to monitor CPU utilization. This utility provides statistics on processor time, user time, privileged time, interrupt time, and DPC time.

Related topics:

- Checking Memory Management
- Checking I/O Management
- · Checking Network Management
- Checking Process Management



This document describes how to check system CPU utilization on most UNIX-based and Windows systems. For other platforms, see your operating system documentation.

Checking Memory Management

Check the following memory management areas:

- Paging and Swapping
- Oversize Page Tables

Paging and Swapping

Use the V\$OSSTAT view, utilities such as sar or vmstat on UNIX, or the administrative performance tool on Windows, to investigate the cause of paging and swapping.

Oversize Page Tables

On UNIX, if the processing space becomes too large, then it can result in the page tables becoming too large. This is not an issue on Windows systems.



Checking I/O Management

Thrashing is an I/O management issue. Ensure that your workload fits into memory, so the computer is not thrashing (swapping and paging processes in and out of memory). The operating system allocates fixed portions of time during which CPU resources are available to your process. If the process wastes a large portion of each time period checking to ensure that it can run and ensuring that all necessary components are in the computer, then the process might be using only 50% of the time allotted to actually perform work.

Checking Network Management

Check client/server round trips. There is an overhead in processing messages. When an application generates many messages that need to be sent through the network, the latency of sending a message can result in CPU overload. To alleviate this problem, bundle multiple messages rather than perform lots of round trips. For example, you can use array inserts, array fetches, and so on.

Checking Process Management

Several process management issues discussed in this section should be checked.

- Scheduling and Switching
- Context Switching
- Starting New Operating System Processes

Scheduling and Switching

The operating system can spend excessive time scheduling and switching processes. Examine the way in which you are using the operating system, because it is possible that too many processes are in use. On Windows systems, do not overload the server with too many non-database processes.

Context Switching

Due to operating system specific characteristics, your system could be spending a lot of time in context switches. Context switching can be expensive, especially with a large SGA. Context switching is not an issue on Windows, which has only one process for each instance. All threads share the same page table.

Oracle Database has several features for context switching:

Post-wait driver

An Oracle process must be able to post another Oracle process (give it a message) and also must be able to wait to be posted. For example, a foreground process may need to post LGWR to tell it to write out all blocks up to a given point so that it can acknowledge a commit.

Often this post-wait mechanism is implemented through UNIX Semaphores, but these can be resource intensive. Therefore, some platforms supply a post-wait driver, typically a kernel device driver that is a lightweight method of implementing a post-wait interface.

Memory-mapped system timer

Oracle Database often needs to query the system time for timing information. This can involve an operating system call that incurs a relatively costly context switch. Some



platforms implement a memory-mapped timer that uses an address within the processes virtual address space to contain the current time information. Reading the time from this memory-mapped timer is less expensive than the overhead of a context switch for a system call.

List I/O interfaces to submit multiple asynchronous I/Os in One Call

List I/O is an application programming interface that allows several asynchronous I/O requests to be submitted in a single system call, rather than submitting several I/O requests through separate system calls. The main benefit of this feature is to reduce the number of context switches required.

Starting New Operating System Processes

There is a high cost in starting new operating system processes. Developers often create a single-purpose process, exit the process, and then create a new one. This technique recreates and destroys the process each time, consuming excessive amounts of CPU, especially in applications that have large SGAs. The CPU is needed to build the page tables each time. The problem is aggravated when you pin or lock shared memory because you must access every page.

For example, if you have a 1 gigabyte SGA, then you might have page table entries for every 4 KB, and a page table entry might be 8 bytes. You could end up with (1 GB / 4 KB) * 8 byte entries. This becomes expensive, because you need to continually ensure that the page table is loaded.

Managing CPU Resources Using Oracle Database Resource Manager

Oracle Database Resource Manager allocates and manages CPU resources among database users and applications in the following ways:

Preventing CPU saturation

If the CPUs run at 100%, then you can use Oracle Database Resource Manager to allocate a maximum amount of CPU to sessions in each consumer group. This feature can ensure that high-priority sessions can run immediately and lower the CPU consumption of low-priority sessions.

Limiting CPU usage for a consumer group

You can use the Resource Manager directive <code>max_utilization_limit</code> to place a hard limit on the percentage of CPU that a consumer group can use. This feature restricts the CPU consumption of low-priority sessions and can help provide more consistent performance for the workload in a consumer group.

Limiting damage from runaway queries

Starting with Oracle Database 11g Release 2 (11.2.0.2), Oracle Database Resource Manager can limit the damage from runaway queries by limiting the maximum execution time for a call, or by moving a long-running query to a lower-priority consumer group.

Limiting the parallel statement activity for a consumer group

Starting with Oracle Database 11*g* Release 2 (11.2.0.2), you can use the Resource Manager directive parallel_target_percentage to prevent one consumer group from monopolizing all parallel servers. The database queues parallel statements if they would cause this limit to be exceeded.

For example, assume that the target number of parallel servers is 64, and the consumer group ETL has this directive set to 50%. If consumer group ETL is using 30 parallel servers,

and if a new parallel statement needs 4 parallel servers, then the database would queue this statement.

See Also:

- Oracle Database Administrator's Guide to learn how to use Oracle Database Resource Manager
- Oracle Database VLDB and Partitioning Guide to learn how to use parallel query

Managing CPU Resources Using Instance Caging

When running multiple database instances on a single system, the instances compete for CPU resources. One resource-intensive database instance may significantly degrade the performance of the other instances. To avoid this problem, you can use instance caging to limit the number of CPUs that can used by each instance. Oracle Database Resource Manager then allocates CPU among the various database sessions according to the resource plan that you set for the instance, thereby minimizing the likelihood of the instance becoming CPU-bound.

See Also:

Oracle Database Administrator's Guide for information about using instance caging



Working With Automatic Materialized Views

Starting with Oracle Database Release 21c, materialized views can be created and maintained automatically.

The Oracle Database can automatically create and manage materialized views in order to optimize query performance. With very little or no interaction with the DBA, background tasks monitor and analyze workload characteristics and identifies where materialized views will improve SQL performance. The performance benefit of candidate materialized views is measured in the background (using workload queries) before they are made visible to the workload.

Note:

Automatic materialized views support partitioned and non-partitioned base tables. Incremental materialized view refresh is supported. In addition, for partitioned tables, there is support for Partition Change Tracking (PCT) view refresh. To be eligible for PCT-based refresh, partitioned base tables must use either range, list, or composite partitioning. If there is performance advantage, the automatic materialized view recommendations will include a partitioned automatic materialized view based on the partitioning of the base table of the materialized view. The partitioning type supported is auto-list partitioning, which will mirror the partitioning of the fact table.

The automatic materialized view maintenance module decides the type of refresh that is the most beneficial at the time of refresh, and will decide during run time whether to switch from incremental refresh to full refresh.

Overview of Automatic Materialized Views

The database automatically collects workload information, workload queries and query execution statistics. It also maintains and purges the history of the workload. This eliminates a time-consuming DBA task.

Although automatic materialized views can run with minimal DBA interaction, their behavior can be easily adjusted.

This is a summary of automatic materialized view functionality:

- Automatically detects and collects workload query execution statistics. These include buffer-gets, database time, estimated cost, and other statistics.
- Creates candidate materialized views hidden from the database workload and verifies that
 they will deliver the projected performance benefit. It does this by test executing a sample
 of workload queries in the background.
- Provides reports detailing performance test results and which materialized views have been implemented.
- Provides automatic materialized view refresh.

The database implements only automatic materialized views whose benefits far outweigh the cost of maintaining them. It does not implement those that provide marginal benefit.

Workload Information Provided by the Object Activity Tracking System

Automatic materialized views use workload information provided by the Object Activity Tracking System (OATS) as part of the automated decision-making processes.

Starting in Oracle Database 21c, the Object Activity Tracking System (OATS) tracks various activities associated with database objects. The automatic materialized view feature is one of the clients of this system. In the case of automatic materialized views, the usage data provided by OATS is one of the inputs into the analysis of cost versus benefit for creating or refreshing a materialized view, as well as in determining the best type of refresh and optimal refresh schedule.

OATS takes periodic snapshots of activity within any number of selected tables. The snapshot for each table captures the number of scans, loads, inserts/updates/deletes, truncations, and partition-related activity within the table from the beginning to the end of the snapshot interval. The DBA can use the DBMS_ACTIVITY PL/SQL package to set the OATS capture interval, snapshot retention period, and space limits.

For example, the DBA_ACTIVITY_TABLE view shows the usage data captured within each snaphot.

Data Dictionary Views That Provide Information About Automatic Materialized Views and OATS

As of Oracle Database 21c, the database includes data dictionary views that display information about automatic materialized views as well as OATS (Object Activity Tracking System).

Views for Monitoring Automatic Materialized Views

Use the following data dictionary views to check the automatic materialized view configuration and to examine various aspects of automatic materialized views activity:

DBA AUTO MV ANALYSIS ACTIONS

Displays information about analysis and tuning tasks, including actions, commands, advisor-specific flags, and command parameters.

- DBA AUTO MV ANALYSIS EXECUTIONS
 - Displays information about analysis and tuning executions, including concurrency, degree of parallelism (DOP) requested by the user and actual DOP upon execution finish, status, associated advisor, and informational or error message.
- DBA_AUTO_MV_ANALYSIS_RECOMMENDATIONS
 Displays recommendations associated with automatic materialized views.
- DBA_AUTO_MV_ANALYSIS_REPORT
 Reports on analyses and recommendations, including task and execution names, sequence number of the journal entry, and message entry in the journal.
- DBA AUTO MV ANALYSIS TASK



Displays analysis details associated with automatic materialized views, including task identifiers and task description, creation and last modification dates, execution data, parent task, status, and other information.

• DBA AUTO MV CONFIG

Displays the current automatic materialized view configuration.



The configuration parameters displayed in this view can be updated with CONFIGURE procedure of the DBMS AUTO MV package.

- DBA_AUTO_MV_MAINT_REPORT
 Displays the date, time, and message associated with automatic materialized view maintenance actions.
- DBA_AUTO_MV_REFRESH_HISTORY
 Displays the owner name, view name, date, start and end time, elapsed time, status, and error number (if an error occurred) for each automatic materialized view refresh.
- DBA_AUTO_MV_VERIFICATION_REPORT
 Displays the task name, execution name, and message associated with verifications.
- DBA_AUTO_MV_VERIFICATION_STATUS
 Displays the owner, start/end timestamps of verifications, SQL tuning sets used, and SQL Performance Analyzer tasks used in each verification.

Views for Monitoring OATS

DBA_ACTIVITY_CONFIG
 Displays the current value of the configuration parameters that control OATS.



The configuration parameters displayed in this view can be updated with CONFIGURE procedure of the DBMS ACTIVITY package.

- DBMS ACTIVITY TABLE
 - Describes table activity snapshots that were recently taken by OATS.
- DBA_ACTIVITY_SNAPSHOT_META
 Displays information about activity snapshots taken by OATS.
- DBA_ACTIVITY_MVIEW
 Describes materialized view activity snapshots that were recently taken by OATS.



Oracle Database Reference



The DBMS_AUTO_MV Package

This package contains procedures for controlling automatic materialized views.

DBMS_AUTO_MV.CONFIGURE

The DBA can use the <code>CONFIGURE</code> procedure of <code>DBMS_AUTO_MV</code> to start, stop, and configure automatic materialized views.

Table 17-1 Configure Procedure Parameters

Parameter	Description and Examples
AUTO_MV_MODE	IMPLEMENT: Implements automatic materialized views.
	OFF: Turns off automatic materialized views.
	REPORT ONLY: Report-only mode.
	<pre>exec dbms_auto_mv.configure('AUTO_MV_MODE', 'IMPLEMENT');</pre>
	<pre>exec dbms_auto_mv.configure('AUTO_MV_MODE', 'OFF'); exec dbms_auto_mv.configure('AUTO_MV_MODE', 'REPORT ONLY');</pre>
AUTO_MV_MAINT_TASK	ENABLE: Activates the task performing the maintenance (refreshes, validations, and cleanup).
	DISABLE: Deactivates the task performing the maintenance.
	CLEANUP_AND_DISABLE: Drops all automatic materialized views and deactivates the task. If automatic materialized views maintenance is in progress, then maintenance is allowed to finish before the task is deactivated.
	<pre>exec dbms_auto_mv.configure('AUTO_MV_MAINT_TASK', 'ENABLE'); exec dbms auto mv.configure('AUTO MV MAINT TASK',</pre>
	'DISABLE'); exec dbms_auto_mv.configure('AUTO_MV_MAINT_TASK', 'CLEANUP_AND_DISABLE');

Table 17-1 (Cont.) Configure Procedure Parameters

Parameter

Description and Examples

AUTO MV SPACE BUDGET

Specifies the percentage of space budgeted for implementing automatic materialized views within the tablespace where those views were created. This is a percentage of the total space used by all automatic materialized views and associated indexes within the tablespace.

A condition on the enforcement of AUTO MV SPACE BUDGET is the value of AUTO MV DEFAULT TABLESPACE:

- If AUTO MV DEFAULT TABLESPACE is not defined (NULL), then automatic materialized views are created on the tablespace of the view's parent object (which is the largest FACT table in the view's definition). In this case, the budget defined by AUTO MV SPACE BUDGET is enforced within that tablespace.
- If AUTO MV DEFAULT TABLESPACE is defined, then automatic materialized views are created in the designated default tablespace. In this case, the budget set by AUTO MV SPACE BUDGET is ignored.

If the budget is exceeded (possibly because of the growth of automatic materialized views), then the least-used automatic materialized view is dropped.

The value is an integer from 1 to 100. The default is 67 (67% of the total volume of the tablespace).

```
exec dbms auto mv.configure('AUTO MV SPACE BUDGET',
15);
```

AUTO MV DEFAULT TABLESPAC Specifies the default tablespace for the creation of automatic materialized views. Possible values are the name of a valid temporary tablespace or NULL (the default). In the case of NULL, new automatic materialized view is created in the default tablespace of the owner of the parent object. If the view has more than one parent object, such as materialized views defined on multiple base tables, then the default tablespace of the owner of largest base table is selected. If the value is changed dynamically, the change takes effect the next time automatic materialized view recommendations are implemented.

```
exec
dbms auto mv.configure('AUTO MV DEFAULT TABLESPACE',
'MYTABLESPACE');
exec
dbms auto mv.configure('AUTO MV DEFAULT TABLESPACE')
```

Table 17-1 (Cont.) Configure Procedure Parameters

Parameter	Description and Examples
AUTO_MV_TEMP_TABLESPACE	Specifies the temporary tablespace used for creation or refresh of automatic materialized views. Possible values are the name of a valid temporary tablespace or NULL. In the case of NULL, the tablepace is assigned to the owner of the largest parent object of the automatic materialized views. The default is NULL.
	<pre>exec dbms_auto_mv.configure('AUTO_MV_TEMP_TABLESPACE','TE MP2'); exec dbms_auto_mv.configure('AUTO_MV_TEMP_TABLESPACE');</pre>
AUTO_MV_RETENTION	Specifies the number of days automatic materialized views can continue to exist without being queried. If an automatic materialized view remains unused beyond this retention time, it is automatically dropped. Possible values are any integer between 1 and 373. The default is 33 days.
	<pre>exec dbms_auto_mv.configure('AUTO_MV_RETENTION', 365);</pre>
AUTO_MV_ANALYZE_REPORT_RETENTION	AUTO_MV_ANALYZE_REPORT_RETENTION Specifies the maximum number of days to retain analysis and recommendation history. Possible values are any integer from 0 to 90. A value of 0 means no history is maintained. The default is 31 days.
	<pre>exec dbms_auto_mv.configure('AUTO_MV_ANALYZE_REPORT_RETEN TION', 60);</pre>
AUTO_MV_VERIFY_REPORT_RET ENTION	Specifies the maximum number of days to retain verification history. Possible values are any integer from 0 to 90. The value 0 specifies that no verification history will be maintained. The default is 31 days.
	<pre>exec dbms_auto_mv.configure('AUTO_MV_VERIFY_REPORT_RETENT ION', 7);</pre>
AUTO MV_MAINT_REPORT_RETE NTION	Specifies the maximum number of days to retain history of automatic materialized view maintenance (refreshes) in the DBA_AUTO_MV_REFRESH_* dictionary tables. Possible values are any integer from 0 to 90. The value 0 specifies that no refresh history will be maintained. The default is 31 days.
	<pre>exec dbms_auto_mv.configure('AUTO_MV_MAINT_REPORT_RETENTI ON', 14);</pre>

Table 17-1 (Cont.) Configure Procedure Parameters

Parameter

Description and Examples

AUTO_MV_ANALYZE_WORKLOAD_WINDOW

Specifies the maximum number of hours to investigate queries from the latest snapshots and make recommendations. Possible values are any integer between from 1 to 8760. The default is 24 hours.

exec
dbms_auto_mv.configure('AUTO_MV_ANALYZE_WORKLOAD_WIN
DOW', 48);

AUTO_MV_ANALYZE_WORKLOAD_MIN TIME

Specifies the minimum time in seconds for a query to be considered for automatic materialized views recommendation. Queries below this threshold are not considered for recommendations. Possible values are any integer from 0 to 3600. The default is 120 seconds.

exec
dbms_auto_mv.configure('AUTO_MV_ANALYZE_WORKLOAD_MIN
_TIME', 1800);

AUTO MV SCHEMA

Specifies a schema to be either included or excluded during the creation of automatic materialized views. The schema is added to the *inclusion list* or the *exclusion list* in the configuration. Initially, both lists are empty and automatic materialized views can be created in all the schemas in a database where automatic materialized views are enabled. You can build the inclusion and exclusion lists by calling AUTO MV SCHEMA multiple times.

The boolean ALLOW determines if the schema is added to the inclusion list (TRUE) or to the exclusion list (FALSE). The default is TRUE. During workload processing, any query that does not contain a reference at least one table in a schema on the inclusion list is not analyzed and not auto tuned. It is not factored into recommendations and verifications. Likewise, if a query references a table in a schema on the exclusion list, that query is excluded from processing.

```
exec
dbms_auto_mv.configure('AUTO_MV_SCHEMA', 'SCHEMA_A');
exec
dbms_auto_mv.configure('AUTO_MV_SCHEMA', 'SCHEMA_B',
FALSE);
```

To enable or disable processing of all schemas, you can specific the schema as NULL. This either enables or disables all of them, depending on the value of ALLOW.

```
exec dbms_auto_mv.configure('AUTO_MV_SCHEMA','',
TRUE);
```

Table 17-1 (Cont.) Configure Procedure Parameters

Parameter	Description and Examples
AUTO_MV_APP_MODULE	Specifies application modules to include or exclude from the creation of automatic materialized views.
	<pre>exec dbms_auto_mv.configure('AUTO_MV_APP_MODULE', 'MODULE1', TRUE); exec dbms_auto_mv.configure('AUTO_MV_APP_MODULE', 'MODULE1', FALSE); exec dbms_auto_mv.configure('AUTO_MV_APP_MODULE', 'MODULE1');</pre>

DBMS_AUTO_MV.DROP_AUTO_MVS

This procedure drops an automatic materialized view. It can be executed only by users who have the DBA role.

Parameter	Description
OWNER	The name of the owner of the automatic materialized view.
MV_NAME	The name of the automatic materialized view.
ALLOW_RECREATE	Allow the materialized view to be recreated if necessary. Optional.

Note that if OWNER is specified and MV_NAME is set to NULL, then all automatic materialized views owned by OWNER are dropped.

```
exec dbms_auto_mv.drop_auto_mvs('SH', 'AUTO_MV$$_G2MKPB9SA1FB7', TRUE);
exec dbms_auto_mv.drop_auto_mvs('SH', 'AUTO_MV$$_G2MKPB9SA1FB7');
exec dbms_auto_mv.drop_auto_mvs('SH', '');
exec dbms_auto_mv.drop_auto_mvs('SH', '', TRUE);
```

DBMS_AUTO_MV.RECOMMEND

DBMS_AUTO_MV.RECOMMEND generates automatic materialized recommendations based on a given SQL tuning set. This API enables you to manually run automatic materialized view analysis and verification from a command line (instead of through an Automatic SQL Tuning task). You set the workload start and end time and determine whether this execution results in a report only, or an actual implementation. There is no default time limit for the workload window.

Execution of this API requires the DBA role.

Note:

Automatic materialized view configuration parameters can influence the results of <code>DBMS_AUTO_MV.RECOMMEND</code>. For example, the analysis and recommendations of this API are restricted to the schemas specified by the configuration parameter <code>AUTO_MV_SCHEMA</code>.

Parameter	Description
STS_OWNER	The name of the owner of the SQL tuning set. Default: SYS.
STS_NAME	The name of the SQL tuning set. Default: SYS_AUTO_STS.
WORKLOAD_START_TIME	Start time for the workload window. Default: WORKLOAD_END_TIME minus 24 hours.
WORKLOAD_END_TIME	End time for the workload window. Default: The current timestamp.
AUTO_MV_MODE	REPORT ONLY (recommendations only) or IMPLEMENT. Default: REPORT ONLY

Example:

Generate and report recommendations using SYS_AUTO_STS for the past 24 hours. Note that the default behavior is REPORT_ONLY, which means that no automatic materialized view will be implemented.

```
var exec_name varchar2(200);
begin
    :exec_name := dbms_auto_mv.recommend();
end;
SELECT * FROM DBA_AUTO_MV_ANALYSIS_RECOMMENDATIONS
    WHERE exec_name = :exec_name;
```

DBMS_AUTO_MV.REFRESH

The DBMS_AUTO_MV.RECOMMEND API enables you force a refresh of all stale automatic materialized views. The stale automatic materialized views are unconditionally refreshed in descending order, based on their verified query rewrite benefit values. There are no parameters. This routine can be executed only by users with the DBA role.

```
exec dbms auto mv.dbms auto refresh();
```

DBMS_AUTO_MV.REPORT_ACTIVITY

The DBMS_AUTO_MV.REPORT_ACTIVITY This API generates a report on automatic materialized view activities and usage within a specified time window. The report is returned as a CLOB.

Parameter	Description
ACTIVITY_START	The start of the time window. Default: SYSTIMESTAMP -1.
ACTIVITY_END	The end of the time window. Default: SYSTIMESTAMP.
TYPE	The format of the report. 'TEXT', 'HTML', and 'XML' are supported. Default: 'TEXT'.



Parameter	Description
SECTION	The section or sections covered by the report. The value can be any combination of: SUMMARY, MV_DETAILS, QUERY_DETAILS, VERIFICATION_DETAILS or ALL. Default: 'ALL'.
	Vse the "+" or "- " operator to concatenate a single string that includes or excludes sections of the report. This is shown in one of the examples below.
LEVEL	The level of detail in the report: BASIC, TYPICAL or

Examples:

Generate a report on all automatic materialized view activities. Output the report in HTML format:

ALL.

Default: 'TYPICAL'.

select dbms_auto_mv.report_activity(type => 'HTML') from dual;

Generate a report on all automatic materilalized view activities. Exclude the verification details. Output the report in XML format.

select dbms_auto_mv.report_activity(type => 'XML', section => 'ALL-VERIFICATION DETAILS') from dual;

DBMS_AUTO_MV.REPORT_LAST_ACTIVITY

The DBMS_AUTO_MV.REPORT_LAST_ACTIVITY API generates a report on the most recent automatic materialized view activities and usage.

Parameter	Description
TYPE	The format of the report. 'TEXT', 'HTML', and 'XML' are supported.
	Default: 'TEXT'.



Parameter	Description
SECTION	The section or sections covered by the report. The value can be any combination of: SUMMARY, MV_DETAILS, QUERY_DETAILS, VERIFICATION_DETAILS or ALL. Default: 'ALL'.
	Use the "+" or "- " operator to concatenate a single string that includes or excludes sections of the report. See the examples below.
LEVEL	The level of detail in the report: BASIC, TYPICAL of ALL. Default: 'TYPICAL'.

Examples:

Generate a comprehensive report of the most recent activity, at the typical level of detail. Output the report in text format (the default). Note that both of these statements return the same results.

```
select dbms_auto_mv.report_last_activity('TEXT', 'ALL', 'TYPICAL') from dual;
select dbms_auto_mv.report_last_activity() from dual;
```

Generate a report of the most recent activity that includes only the summary and the details of the materialized view. Report at the maximum level of detail. Output in XML format:

```
select dbms_auto_mv.report_last_activity('XML', 'SUMMARY+MV_DETAILS', 'ALL')
from dual;
```

Generate a report of the most recent activity at the basic level of detail. Exclude the verification details. Output in HTML format.

```
select dbms_auto_mv.report_last_activity('XML', 'ALL-VERIFICATION_DETAIL',
'BASIC') from dual;
```

For More Information



The Oracle Database PL/SQL Packages and Types Reference.

The DBMS_ACTIVITY Package

The DBMS_ACTIVITY PL/SQL package contains functions and procedures for configuring Object Activity Tracking System (OATS) information collection and management. Data collected by OATS is used in analyses performed by automatic materialized views.

DBAs can use the <code>DBMS_ACTIVITY.CONFIGURE</code> procedure to control three OATS parameters within a specific database.

ACTIVITY INTERVAL

The interval between snapshots.

```
exec dbms activity.configure('ACTIVITY INTERVAL MINUTES', 30)
```

• ACTIVITY_RETENTION_DAYS

How long snapshots are saved.

```
exec dbms activity.configure('ACTIVITY RETENTION DAYS',60)
```

ACTIVITY SPACE PERCENT

How much of available space is reserved for snapshots.

```
exec dbms activity.configure('ACTIVITY SPACE PERCENT', 10)
```

Note:

OATS is intended to be self-managing and the default configuration is recommended, particularly if the automatic materialized views feature is used.

See Also:

The PL/SQL Packages and Types Reference.



Glossary

ADO policy

A policy that specifies a rule and condition for Automatic Data Optimization (ADO). For example, an ADO policy may specify that an object is marked NOINMEMORY (action) 30 days after creation (condition). Specify ADO policies using the ILM clause of CREATE TABLE and ALTER TABLE statements.

Automatic Data Optimization (ADO)

A technology that creates policies, and automates actions based on those policies, to implement an Information Lifecycle Management (ILM) strategy.

Automatic In-Memory

A feature that automatically evicts cold (infrequently accessed) segments from the IM column store to ensure that the working data set is always populated.

availability

The degree to which an application, service, or function is accessible on demand.

Bloom filter

A low-memory data structure that tests membership in a set. The database uses Bloom filters to improve the performance of hash joins.

Column Compression Unit (CU)

Contiguous storage for a column in an In-Memory Compression Unit (IMCU).

columnar data pool

The subpool in the In-Memory Area that stores columnar data. It is also known as the 1 MB pool.



columnar format

The column-based format for objects that reside in the In-Memory Column Store. The columnar format contrasts with the row format used in data blocks.

common dictionary

A segment-level, instance-specific set of common dictionary codes, created from local dictionaries. A local dictionary is a sorted list of dictionary codes specific to a Column Compression Unit (CU). A join group uses a common dictionary to optimize joins.

compression tiering

The application of different levels of compression to data based on its access pattern. For example, administrators may compress inactive data at a higher rate of compression at the cost of slower access.

data flow operator (DFO)

The unit of work between data redistribution stages in a parallel query.

database buffer cache

The portion of the system global area (SGA) that holds copies of data blocks. All client processes concurrently connected to the database instance share access to the buffer cache.

database instance

The combination of the system global area (SGA) and background processes. An instance is associated with one and only one database. Every database instance is either a read/write database instance or a read-only database instance. In an Oracle Real Application Clusters configuration, multiple instances access a single database.

dense grouping key

A key that represents all grouping keys whose grouping columns come from a specific fact table or dimension.

dense join key

A key that represents all join keys whose join columns come from a particular fact table or dimension.

dense key

A numeric key that is stored as a native integer and has a range of values.

double buffering

A repopulation mechanism in which background processes create new In-Memory Compression Unit (IMCU) versions by combining the original rows with the latest modified rows. During repopulation, the stale IMCUs remain accessible for queries.

expression

A combination of one or more values, operators, and SQL functions that resolves to a value.

expression capture interval

The time interval within which the database considers IM expressions for possible capture.

expression capture window

An expression capture interval defined by invocation of the <code>IME_OPEN_CAPTURE_WINDOW</code> and <code>IME_OPEN_CAPTURE_WINDOW</code> procedures in the <code>DBMS_INMEMORY_ADMIN</code> package.

Expression Statistics Store (ESS)

A repository maintained by the optimizer to store statistics about expression evaluation. For each segment, the ESS monitors statistics such as frequency of execution, cost of evaluation, timestamp evaluation, and so on. The ESS is persistent in nature and has an SGA representation for fast lookup of expressions.

Heat Map

Heat Map shows the popularity of data blocks and rows. Automatic Data Optimization (ADO) to decide which segments are candidates for movement to a different storage tier.

home location

The database instance in which an IMCU resides. When auto DOP is enabled on Oracle RAC, the parallel query coordinator uses home location to determine where each IMCU is located, how large it is, and so on.

In-Memory hybrid scan

A query that scans both the IM column store and the row store. The optimizer considers an In-Memory hybrid scan automatically when all predicate columns have the INMEMORY attribute, and some columns in the SELECT list do not have the INMEMORY attribute.

hybrid partitioned table

A table in which some partitions are stored in data file segments and some are stored in external data source.



IM aggregation

An optimization that accelerates aggregation for queries that join from a single large table to multiple small tables. The transformation uses KEY VECTOR and VECTOR GROUP BY operators, which is why it is also known as VECTOR GROUP BY aggregation.

IM column store

An optional SGA area that stores copies of tables and partitions in a columnar format optimized for rapid scans.

IM dynamic scan

The use of lightweight threads to automatically parallelize In-Memory table scans.

IM expression

A SQL expression whose results are stored in the In-Memory Column Store. If $last_name$ is a column stored in the IM column store, then an IM expression might be $upper(last_name)$.

IMCU mirroring

In Oracle RAC, the duplication of an IMCU in multiple IM column stores. For example, the IM column stores on instance 1 and instance 2 are populated with the same sales table.

IMCU pruning

In a query of the In-Memory Column Store, the elimination of IMCUs based on the high and low values in each IMCU. For example, if a statements filters product IDs greater than 100, then the database avoids scanning IMCUs that contain values less than 100.

IM storage index

A data structure in an IMCU header that stores the minimum and maximum for all columns within the IMCU.

In-Memory Advisor

A downloadable PL/SQL package that analyzes the analytical processing workload in your database. This advisor recommends a size for the IM column store and a list of objects that would benefit from In-Memory population.

In-Memory Aggregation

See IM aggregation.



In-Memory Area

An optional SGA component that contains the IM column store.

In-Memory Column Store

See IM column store.

In-Memory Compression Unit (IMCU)

A storage unit in the In-Memory Column Store that is optimized for faster scans. The In-Memory Column Store stores each column in table separately and compresses it. Each IMCU contains all columns for a subset of rows in a specific table segment.

A one-to-many mapping exists between an IMCU and a set of database blocks. For example, if a table contains columns c1 and c2, and if its rows are stored in 100 database blocks on disk, then IMCU 1 might store the values for both columns for blocks 1-50, and IMCU 2 might store the values for both columns for blocks 51-100.

In-Memory Coordinator Process (IMCO)

A background process whose primary task is to initiate background population and repopulation of columnar data.

In-Memory Dynamic Scan

See IM dynamic scan.

In-Memory Expression

See IM expression.

In-Memory Expression Unit (IMEU)

A container that stores the computed result of an In-Memory Expression (IM expression). Each IMEU is linked to its own parent In-Memory Compression Unit (IMCU).

In-Memory FastStart

A feature that significantly reduces the time to populate data into the IM column store when a database instance restarts.

In-Memory population

See population.

In-Memory virtual column

A virtual column that is eligible to be populated in the In-Memory Column Store.

Information Lifecycle Management (ILM)

A set of processes and policies for managing data throughout its useful life.

join group

A user-defined object that specifies frequently joined columns from the same table or different tables. External tables are not supported.

A typical join group candidate is a set of columns used to join fact and dimension tables. Join groups are only supported when INMEMORY SIZE is a nonzero value.

key vector

A data structure that maps between dense join keys and dense grouping keys.

large pool

Optional area in the SGA that provides large memory allocations for backup and restore operations, I/O server processes, and session memory for the shared server and Oracle XA.

local dictionary

A sorted list of dictionary codes specific to a Column Compression Unit (CU).

lightweight thread

An execution entity used in an In-Memory Dynamic Scan. Lightweight threads help to parallelize scans of IMCUs.

metadata pool

A subpool of the In-Memory Area that stores metadata about the objects that reside in the IM column store. The metadata pool is also known as the *64 KB pool*.

memoptimize pool

An SGA pool that stores buffers and related structures for heap-organized tables specified as MEMOPTIMIZE FOR READ.

on-demand population

When INMEMORY PRIORITY is set to NONE, the IM column store *only* populates the object when it is accessed through a full scan. If the object is never accessed, or if it is accessed only through an index scan or fetch by rowid, then it is never populated.

operator

1. In memory management, operators control the flow of data. Examples include sort, hash join, and bitmap merge operators.

2. In SQL, an operator manipulates data items called *operands* or *arguments* and returns a result. Keywords or special characters represent the operators. For example, an asterisk (*) represents the multiplication operator.

OSON

Oracle's optimized binary JSON format. OSON enables fast queries and updates of the JSON data model in Oracle database server and Oracle database clients.

OZIP

A proprietary compression technique that offers extremely fast decompression. OZIP is tuned specifically for Oracle Database.

partition exchange load

A technique in which you create a table, load data into it, and then exchange an existing table partition with the table. This exchange process is a DDL operation with no actual data movement.

population

The operation of reading existing data blocks from data files, transforming the rows into columnar format, and then writing the columnar data to the IM column store. In contrast, *loading* refers to bringing new data into the database using DML or DDL.

priority-based population

When PRIORITY is set to a value other than NONE, Oracle Database adds the object to a prioritized population queue. The database populates objects based on their queue position, from CRITICAL to LOW. It is "priority-based" because the IM column store automatically populates objects using the prioritized list whenever the database re-opens. Unlike in ondemand population, objects do not require a full scan to be populated.

read-only database instance

A database instance that cannot process DML and does not support client connections.

read/write database instance

A database instance that can process DML and supports direct client connections. By default, a database instance is read/write.

repopulation

The automatic refresh of a currently populated In-Memory Compression Unit (IMCU) after its data has been significantly modified. In contrast, population is the initial creation of IMCUs in the IM column store.

service

The logical representation of an application workload that shares common attributes, performance thresholds, and priorities. A single service can be associated with one or more instances of an Oracle RAC database, and a single instance can support multiple services.

SGA

System global area. A group of shared memory structures that contain data and control information for one Oracle database instance.

SIMD

Single Instruction, Multiple Data. An instruction that processes data as a single unit, called a **vector**, rather than as separate instructions. SIMD processing is known as **vectorization**.

Snapshot Metadata Unit (SMU)

A storage unit in the In-Memory Area that contains metadata and transactional information for an associated In-Memory Compression Unit (IMCU).

Space Management Worker Process (Wnnn)

A process that populates or repopulates data in the IM column store on behalf of In-Memory Coordinator Process (IMCO).

staleness threshold

An internally set percentage of entries in the transaction journal for an IMCU that initiates repopulation.

storage tiering

The deployment of data on different tiers of storage depending on its level of access. For example, administrators migrate inactive data from high-performance, high-cost storage to low-cost storage.

system global area (SGA)

See SGA.



table scan process

A foreground or PQ process that coordinates an IM dynamic scan.

threshold-based repopulation

The automatic repopulation of an IMCU when the number of stale entries in an IMCU reaches an internal staleness threshold.

transaction journal

Metadata in a Snapshot Metadata Unit (SMU) that keeps the IM column store transactionally consistent.

trickle repopulation

A supplement to threshold-based repopulation. The In-Memory Coordinator Process (IMCO) may instigate trickle repopulation automatically for any IMCU in the IM column store that has stale entries but does not meet the staleness threshold.

vector aggregation

See IM aggregation.

virtual column

A column that is not stored on disk. The database derives the values in virtual columns on demand by computing a set of expressions or functions.

working data set

The subset of INMEMORY objects that is actively queried at a given time. Typically, the work working data set changes over time.



Index

ADDM enabling in a PDB, 7-6 allocation of memory, 9-1 applications deploying, 2-19 design principles, 2-9 development trends, 2-16 implementing, 2-14 Automatic Database Diagnostic Monitor actions and rationales of recommendations, 7-9 and DB time, 7-3 example report, 7-9 findings, 7-9 results, 7-9 setups, 7-10 types of problems considered, 7-3 types of recommendations, 7-9 automatic database diagnostic monitoring, 1-5 automatic segment-space management, 4-5, 15-10 Automatic Segment-space management, 10-1 automatic Segment-space management, 4-3 Automatic Workload Repository, 1-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 categorization of AWR statistics in a multitenant environment, 6-22 configuring, 6-7 setups, 7-10 types of problems considered, 7-3 types of recommendations, 7-9 automatic Workload Repository report, 6-37 Betree indexes, 2-11 benchmarking workloads, 2-17 big bang rollout strategy, 2-19 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
enabling in a PDB, 7-6 allocation of memory, 9-1 allocation of memory, 9-1 applications deploying, 2-19 design principles, 2-9 development trends, 2-16 implementing, 2-14 Automatic Database Diagnostic Monitor actions and rationales of recommendations, 7-9 analysis results example, 7-9 and DB time, 7-3 example report, 7-9 findings, 7-9 results, 7-9 setups, 7-10 types of problems considered, 7-3 types of problems considered, 7-3 automatic database diagnostic monitoring, 1-5 automatic segment-space management, 4-5, 15-10 Automatic SQL tuning, 1-5 automatic undo management, 4-3 Automatic SQL tuning, 1-5 automatic workload Repository, 1-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 categorization of AWR statistics in a multitenant environment, 6-22 configuring, 6-7
allocation of memory, 9-1 applications deploying, 2-19 design principles, 2-9 development trends, 2-16 implementing, 2-14 Automatic Database Diagnostic Monitor actions and rationales of recommendations, 7-9 analysis results example, 7-9 and DB time, 7-3 example report, 7-9 findings, 7-9 results, 7-9 setups, 7-10 types of problems considered, 7-3 types of recommendations, 7-9 automatic database diagnostic monitoring, 1-5 automatic segment-space management, 4-5, 15-10 Automatic SQL tuning, 1-5 automatic undo management, 4-3 Automatic Workload Repository report, 6-37 big bang rollout strategy, 2-19 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
applications deploying, 2-19 design principles, 2-9 development trends, 2-16 implementing, 2-14 Automatic Database Diagnostic Monitor actions and rationales of recommendations, 7-9 analysis results example, 7-9 and DB time, 7-3 example report, 7-9 findings, 7-9 results, 7-9 setups, 7-10 types of problems considered, 7-3 types of recommendations, 7-9 automatic database diagnostic monitoring, 1-5 automatic segment-space management, 4-5, 15-10 Automatic SQL tuning, 1-5 automatic SQL tuning, 1-5 automatic undo management, 4-3 Automatic Workload Repository, 1-5 automatic Workload Repository, 1-5 automatic SQL tuning, 1-5 automatic undo management, 4-3 Automatic Workload Repository, 1-5 automatic SQL tuning, 1-5 automatic workload Repository, 1-5 automatic workload Repository, 1-5 automatic workload Repository, 1-5 implementing, 2-14 automatic Workload Repository report, 6-37 B-tree indexes, 2-11 baselines, 1-2, 6-3 performance, 6-1 benchmarking workloads, 2-17 big bang rollout strategy, 2-19 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
deploying, 2-19 design principles, 2-9 development trends, 2-16 implementing, 2-14 Automatic Database Diagnostic Monitor actions and rationales of recommendations, 7-9 analysis results example, 7-9 and DB time, 7-3 example report, 7-9 findings, 7-9 results, 7-9 setups, 7-10 types of problems considered, 7-3 automatic database diagnostic monitoring, 1-5 automatic segment-space management, 4-5, 15-10 Automatic SQL tuning, 1-5 automatic undo management, 4-3 Automatic Workload Repository report, 6-37 Betree indexes, 2-11 baselines, 1-2, 6-3 performance, 6-1 benchmarking workloads, 2-17 big bang rollout strategy, 2-19 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache bit ratio, 11-4
design principles, 2-9 development trends, 2-16 implementing, 2-14 Automatic Database Diagnostic Monitor actions and rationales of recommendations, 7-9 analysis results example, 7-9 and DB time, 7-3 example report, 7-9 findings, 7-9 results, 7-9 setups, 7-10 types of problems considered, 7-3 types of recommendations, 7-9 automatic database diagnostic monitoring, 1-5 automatic sagment-space management, 4-5, 15-10 Automatic SQL tuning, 1-5 automatic Workload Repository, 1-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 configuring, 6-7 viewing remote snapshots for ADG standby databases, 6-35 views for accessing data, 6-20 Autonomous Data Warehouse ADW, 10-3 Autonomous Transaction Processing ATP, 10-3 awrrpt.sql Automatic Workload Repository report, 6-37 B-tree indexes, 2-11 baselines, 1-2, 6-3 performance, 6-1 benchmarking workloads, 2-17 big bang rollout strategy, 2-19 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
development trends, 2-16 implementing, 2-14 Automatic Database Diagnostic Monitor actions and rationales of recommendations, 7-9 analysis results example, 7-9 and DB time, 7-3 example report, 7-9 findings, 7-9 results, 7-9 setups, 7-10 types of problems considered, 7-3 types of recommendations, 7-9 automatic database diagnostic monitoring, 1-5 automatic segment-space management, 4-5, 15-10 Automatic shared memory management, 4-3 automatic SQL tuning, 1-5 automatic SQL tuning, 1-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 categorization of AWR statistics in a multitenant environment, 6-22 configuring, 6-7 databases, 6-35 views for accessing data, 6-20 Autonomous Data Warehouse ADW, 10-3 Autonomous Transaction Processing ATP, 10-3 awrrpt.sql Automatic Workload Repository report, 6-37 B B-tree indexes, 2-11 baselines, 1-2, 6-3 performance, 6-1 benchmarking workloads, 2-17 big bang rollout strategy, 2-19 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
implementing, 2-14 Automatic Database Diagnostic Monitor actions and rationales of recommendations, 7-9 analysis results example, 7-9 and DB time, 7-3 example report, 7-9 findings, 7-9 results, 7-9 setups, 7-10 types of problems considered, 7-3 types of recommendations, 7-9 automatic database diagnostic monitoring, 1-5 automatic segment-space management, 4-5, 15-10 Automatic sapred memory management, 4-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 configuring, 6-7 views for accessing data, 6-20 Autonomous Data Warehouse ADW, 10-3 Autonomous Data Warehouse
Automatic Database Diagnostic Monitor actions and rationales of recommendations, 7-9 analysis results example, 7-9 and DB time, 7-3 example report, 7-9 results, 7-9 results, 7-9 setups, 7-10 types of problems considered, 7-3 automatic database diagnostic monitoring, 1-5 automatic segment-space management, 4-5, 15-10 Automatic shared memory management, 4-5 automatic SQL tuning, 1-5 automatic workload Repository, 1-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 categorization of AWR statistics in a multitenant environment, 6-22 configuring, 6-7 Autonomous Transaction Processing ATP, 10-3 automous Data Warehouse ADW, 10-3 Autonomous Data Warehouse
actions and rationales of recommendations, 7-9 analysis results example, 7-9 and DB time, 7-3 example report, 7-9 findings, 7-9 results, 7-9 setups, 7-10 types of problems considered, 7-3 types of recommendations, 7-9 automatic database diagnostic monitoring, 1-5 automatic segment-space management, 4-5, 15-10 Automatic shared memory management, 4-5 automatic SQL tuning, 1-5 automatic Workload Repository, 1-5 automatic Workload Repository, 1-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 categorization of AWR statistics in a multitenant environment, 6-22 configuring, 6-7 ADW, 10-3 Autonomous Transaction Processing ATP, 10-3 awrrpt.sql Automatic Workload Repository report, 6-37 B-tree indexes, 2-11 baselines, 1-2, 6-3 performance, 6-1 benchmarking workloads, 2-17 big bang rollout strategy, 2-19 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
Autonomous Transaction Processing ATP, 10-3 and DB time, 7-3 example report, 7-9 findings, 7-9 results, 7-9 setups, 7-10 types of problems considered, 7-3 types of recommendations, 7-9 automatic database diagnostic monitoring, 1-5 automatic segment-space management, 4-5, 15-10 Automatic shared memory management, 4-3 Automatic SQL tuning, 1-5 automatic SQL tuning, 1-5 automatic Workload Repository report, 6-37 B-tree indexes, 2-11 baselines, 1-2, 6-3 performance, 6-1 benchmarking workloads, 2-17 big bang rollout strategy, 2-19 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
analysis results example, 7-9 and DB time, 7-3 example report, 7-9 findings, 7-9 results, 7-9 setups, 7-10 types of problems considered, 7-3 types of recommendations, 7-9 automatic database diagnostic monitoring, 1-5 automatic segment-space management, 4-5, 15-10 Automatic shared memory management, 10-1 automatic SQL tuning, 1-5 automatic Workload Repository report, 6-37 B -tree indexes, 2-11 baselines, 1-2, 6-3 performance, 6-1 benchmarking workloads, 2-17 big bang rollout strategy, 2-19 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
and DB time, 7-3 example report, 7-9 findings, 7-9 results, 7-9 setups, 7-10 types of problems considered, 7-3 automatic database diagnostic monitoring, 1-5 automatic memory management, 9-4 automatic segment-space management, 4-5, 15-10 Automatic SQL tuning, 1-5 automatic Workload Repository report, 6-37 B-tree indexes, 2-11 baselines, 1-2, 6-3 performance, 6-1 benchmarking workloads, 2-17 big bang rollout strategy, 2-19 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
example report, 7-9 findings, 7-9 results, 7-9 setups, 7-10 types of problems considered, 7-3 types of recommendations, 7-9 automatic database diagnostic monitoring, 1-5 automatic memory management, 9-4 automatic segment-space management, 4-5, 15-10 Automatic shared memory management, 10-1 automatic SQL tuning, 1-5 automatic Workload Repository report, 6-37 B-tree indexes, 2-11 baselines, 1-2, 6-3 performance, 6-1 benchmarking workloads, 2-17 big bang rollout strategy, 2-19 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
findings, 7-9 results, 7-9 setups, 7-10 types of problems considered, 7-3 types of recommendations, 7-9 automatic database diagnostic monitoring, 1-5 automatic segment-space management, 4-5, 15-10 Automatic shared memory management, 10-1 automatic SQL tuning, 1-5 automatic undo management, 4-3 Automatic Workload Repository, 1-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 categorization of AWR statistics in a multitenant environment, 6-22 configuring, 6-7 Betree indexes, 2-11 baselines, 1-2, 6-3 performance, 6-1 benchmarking workloads, 2-17 big bang rollout strategy, 2-19 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
results, 7-9 setups, 7-10 types of problems considered, 7-3 types of recommendations, 7-9 automatic database diagnostic monitoring, 1-5 automatic memory management, 9-4 automatic segment-space management, 4-5,
setups, 7-10 types of problems considered, 7-3 types of recommendations, 7-9 automatic database diagnostic monitoring, 1-5 automatic memory management, 9-4 automatic segment-space management, 4-5,
types of problems considered, 7-3 types of recommendations, 7-9 automatic database diagnostic monitoring, 1-5 automatic memory management, 9-4 automatic segment-space management, 4-5, 15-10 Automatic shared memory management, 10-1 automatic SQL tuning, 1-5 automatic undo management, 4-3 Automatic Workload Repository, 1-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 categorization of AWR statistics in a multitenant environment, 6-22 configuring, 6-7 B-tree indexes, 2-11 baselines, 1-2, 6-3 performance, 6-1 benchmarking workloads, 2-17 big bang rollout strategy, 2-19 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
types of recommendations, 7-9 automatic database diagnostic monitoring, 1-5 automatic memory management, 9-4 automatic segment-space management, 4-5,
automatic database diagnostic monitoring, 1-5 automatic memory management, 9-4 automatic segment-space management, 4-5,
automatic memory management, 9-4 automatic segment-space management, 4-5, 15-10 Automatic shared memory management, 10-1 automatic SQL tuning, 1-5 automatic undo management, 4-3 Automatic Workload Repository, 1-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 categorization of AWR statistics in a multitenant environment, 6-22 configuring, 6-7 benchmarking workloads, 2-17 big bang rollout strategy, 2-19 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
automatic segment-space management, 4-5, 15-10 Automatic shared memory management, 10-1 automatic SQL tuning, 1-5 automatic undo management, 4-3 Automatic Workload Repository, 1-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 categorization of AWR statistics in a multitenant environment, 6-22 configuring, 6-7 big bang rollout strategy, 2-19 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
automatic segment-space management, 4-5, 15-10 Automatic shared memory management, 10-1 automatic SQL tuning, 1-5 automatic undo management, 4-3 Automatic Workload Repository, 1-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 categorization of AWR statistics in a multitenant environment, 6-22 configuring, 6-7 big bang rollout strategy, 2-19 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
Automatic shared memory management, 10-1 automatic SQL tuning, 1-5 automatic undo management, 4-3 Automatic Workload Repository, 1-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 categorization of AWR statistics in a multitenant environment, 6-22 configuring, 6-7 bitmap indexes, 2-11 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
Automatic shared memory management, 10-1 automatic SQL tuning, 1-5 automatic undo management, 4-3 Automatic Workload Repository, 1-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 categorization of AWR statistics in a multitenant environment, 6-22 configuring, 6-7 block size choosing, 15-9 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
automatic undo management, 4-3 Automatic Workload Repository, 1-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 categorization of AWR statistics in a multitenant environment, 6-22 configuring, 6-7 automatic Workload Repository, 1-5 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
automatic undo management, 4-3 Automatic Workload Repository, 1-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 categorization of AWR statistics in a multitenant environment, 6-22 configuring, 6-7 optimal, 15-9 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
Automatic Workload Repository, 1-5 Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 categorization of AWR statistics in a multitenant environment, 6-22 configuring, 6-7 bottlenecks elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
Active Data Guard support, 6-26 AWR data storage and retrieval in a multitenant environment, 6-22 categorization of AWR statistics in a multitenant environment, 6-22 configuring, 6-7 elimination, 1-3 fixing, 3-1 identifying, 3-1 memory, 9-1 buffer cache hit ratio, 11-4
AWR data storage and retrieval in a multitenant environment, 6-22 identifying, 3-1 categorization of AWR statistics in a memory, 9-1 multitenant environment, 6-22 buffer cache configuring, 6-7 hit ratio, 11-4
multitenant environment, 6-22 identifying, 3-1 categorization of AWR statistics in a memory, 9-1 multitenant environment, 6-22 buffer cache configuring, 6-7 hit ratio, 11-4
categorization of AWR statistics in a memory, 9-1 multitenant environment, 6-22 buffer cache configuring, 6-7 hit ratio, 11-4
multitenant environment, 6-22 buffer cache configuring, 6-7 hit ratio, 11-4
configuring, 6-7 hit ratio, 11-4
111(10110) 22 1
default settings, 6-4 buffer pools
factors affecting space usage, 6-4 multiple, 11-7
managing snapshots in ADG standby business logic, 2-6, 2-14
databases, 6-32
minimizing space usage 6-4
modifying snapshot settings, 6-10
multitonant apvironment support 6.21
overview 6.2
recommendations for retention period 6.4
roporte 6.27
retention period, 6-4 Cloud Control, 1-4

Automatic Workload Repository (continued)



column order	DBMS_ADVISOR package (continued)
indexes, <i>2-12</i>	setups for ADDM, 7-10, 7-11
components	DBMS_SHARED_POOL package
hardware, 2-5	managing the shared pool, 12-20
software, 2-6	debugging designs, 2-18
conceptual modeling, 3-3	deploying applications, 2-19
consistent gets from cache statistic, 11-4	design principles, 2-9
contention	designs
memory, 9-1	debugging, 2-18
context switches, 16-10	testing, 2-18
CONTROL FILES initialization parameter, 4-1	validating, 2-18
<u> </u>	
CPUs, 2-5	development environments, 2-14
utilization, 16-8	diagnostic monitoring, 1-5, 7-1
CREATE INDEX statement	
PARALLEL clause, 4-8	E
CURSOR_SHARING initialization parameter,	
12-5	emergencies
CURSOR_SPACE_FOR_TIME initialization	performance, 3-5
parameter, 12-15	Emergency Performance Method, 3-6
cursors	End to End Application Tracing
accessing, 12-6	action and module names, 2-15
sharing, 12-6	estimating workloads, 2-17
	benchmarking, 2-17
D	extrapolating, 2-17
D	extrapolating workloads, 2-17
data	ontapolating normodus, 2 2.
and transactions, 2-7	_
cache, <i>16-1</i>	F
gathering, 5-1	FAILED AUTO TACKS DEPORT Superior 7.14
modeling, 2-10	FAILED_AUTO_TASKS_REPORT Function, 7-14
queries, 2-8	fast ingest, 10-9, 10-12
searches, 2-8	fast lookup, 10-17
	disabling for a table, 10-21
database caching modes	function-based indexes, 2-11
configuring, 11-14	
default database caching mode, 11-14	Н
determining which mode to use, 11-15	
force full database caching mode	hard parsing, 2-13
about, <i>11-15</i>	hardware
verifying, 11-16	components, 2-5
database monitoring, 1-5	limitations of components, 2-4
diagnostic, 7-1	sizing of components, 2-4
Database Resource Manager, 16-4, 16-11	HOLD_CURSOR clause, 12-7
databases	hours of service, 2-9
diagnosing and monitoring, 7-1	Hodro of Scrvice, 2 3
size, 2-9	
statistics, 5-1	
db block gets from cache statistic, 11-4	
DB time	indexes
metric, 7-2	adding columns, 2-11
statistic, 5-1	appending columns, 2-11
DB_BLOCK_SIZE initialization parameter, 15-4	B-tree, 2-11
DB_DOMAIN initialization parameter, 4-1	bitmap, <i>2-11</i>
-	column order, 2-12
DB_NAME initialization parameter, 4-1	costs, 2-12
DBA_OBJECTS view, 11-10	creating, 4-8
DBMS_ADVISOR package	design, <i>2-10</i>
setting DBIO_EXPECTED, 7-11	3 ,

indexes (continued)	memory
function-based, 2-11	hardware component, 2-5
partitioned, 2-11	memory allocation
placement on disk, 15-5	importance, 9-1
reducing I/O, 2-12	tuning, 10-7
reverse key, 2-11	memory management
selectivity, 2-12	automatic shared, 9-4
sequences in, 2-12	metrics, 6-1
serializing in, 2-12	mirroring
initialization parameters	redo logs, 15-7
CONTROL_FILES, 4-1	modeling
DB DOMAIN, 4-1	conceptual, 3-3
DB_NAME, 4-1	data, 2-10
OPEN_CURSORS, 4-1	workloads, 2-18
STREAMS_POOL_SIZE, 10-4	monitoring
instance caging, 16-8	diagnostic, 1-5
instance configuration	multiple buffer pools, 11-7
initialization files, 4-1	,
performance considerations, 4-1	N
Internet scalability, 2-3	N
y ,	NAMESPACE column
1	V\$LIBRARYCACHE view, 12-10
L	network
LARGE_POOL_SIZE initialization parameter,	hardware component, 2-5
12-26	speed, 2-8
latches	Speeu, 2-0
tuning, 1-2 linear scalability, 2-4	0
· · · · · · · · · · · · · · · · · · ·	
log writer processes	object-orientation, 2-16
log writer processes tuning, 15-6	OPEN_CURSORS initialization parameter, 4-1
log writer processes tuning, 15-6 LRU	OPEN_CURSORS initialization parameter, 4-1 operating system
log writer processes tuning, 15-6	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1
log writer processes tuning, 15-6 LRU	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization
log writer processes tuning, 15-6 LRU aging policy, 11-7	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1
log writer processes tuning, 15-6 LRU aging policy, 11-7	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization
log writer processes tuning, 15-6 LRU aging policy, 11-7	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4
log writer processes tuning, 15-6 LRU aging policy, 11-7	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer,
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4 advisors, 1-5
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7 MEMOPTIMIZE FOR WRITE clause, 10-14	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7 MEMOPTIMIZE FOR WRITE clause, 10-14 memoptimize pool, 10-18	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4 advisors, 1-5
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7 MEMOPTIMIZE FOR WRITE clause, 10-14 memoptimize pool, 10-18 MEMOPTIMIZE_POOL_SIZE initialization	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4 advisors, 1-5 Performance page, 1-5
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7 MEMOPTIMIZE FOR WRITE clause, 10-14 memoptimize pool, 10-18 MEMOPTIMIZE_POOL_SIZE initialization parameter, 10-18 MEMOPTIMIZE_WRITE hint, 10-14	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4 advisors, 1-5 Performance page, 1-5 Oracle Forms
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7 MEMOPTIMIZE FOR WRITE clause, 10-14 memoptimize pool, 10-18 MEMOPTIMIZE_POOL_SIZE initialization parameter, 10-18	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4 advisors, 1-5 Performance page, 1-5 Oracle Forms control of parsing and private SQL areas,
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7 MEMOPTIMIZE FOR WRITE clause, 10-14 memoptimize pool, 10-18 MEMOPTIMIZE_POOL_SIZE initialization parameter, 10-18 MEMOPTIMIZE_WRITE hint, 10-14 Memoptimized Rowstore about, 10-9	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4 advisors, 1-5 Performance page, 1-5 Oracle Forms control of parsing and private SQL areas, 12-8
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7 MEMOPTIMIZE FOR WRITE clause, 10-14 memoptimize pool, 10-18 MEMOPTIMIZE_POOL_SIZE initialization parameter, 10-18 MEMOPTIMIZE_WRITE hint, 10-14 Memoptimized Rowstore about, 10-9 fast ingest	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4 advisors, 1-5 Performance page, 1-5 Oracle Forms control of parsing and private SQL areas, 12-8 Oracle Managed Files, 15-8
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7 MEMOPTIMIZE FOR WRITE clause, 10-14 memoptimize pool, 10-18 MEMOPTIMIZE_POOL_SIZE initialization parameter, 10-18 MEMOPTIMIZE_WRITE hint, 10-14 Memoptimized Rowstore about, 10-9 fast ingest about, 10-9	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4 advisors, 1-5 Performance page, 1-5 Oracle Forms control of parsing and private SQL areas, 12-8 Oracle Managed Files, 15-8 Oracle Orion
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7 MEMOPTIMIZE FOR WRITE clause, 10-14 memoptimize pool, 10-18 MEMOPTIMIZE_POOL_SIZE initialization parameter, 10-18 MEMOPTIMIZE_WRITE hint, 10-14 Memoptimized Rowstore about, 10-9 fast ingest about, 10-9 disabling, 10-15	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4 advisors, 1-5 Performance page, 1-5 Oracle Forms control of parsing and private SQL areas,
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7 MEMOPTIMIZE FOR WRITE clause, 10-14 memoptimize pool, 10-18 MEMOPTIMIZE_POOL_SIZE initialization parameter, 10-18 MEMOPTIMIZE_WRITE hint, 10-14 Memoptimized Rowstore about, 10-9 fast ingest about, 10-9 disabling, 10-15 enabling, 10-14	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4 advisors, 1-5 Performance page, 1-5 Oracle Forms control of parsing and private SQL areas, 12-8 Oracle Managed Files, 15-8 Oracle Orion calibration tool parameters, 15-16 command-line options, 15-16
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7 MEMOPTIMIZE FOR WRITE clause, 10-14 memoptimize pool, 10-18 MEMOPTIMIZE_POOL_SIZE initialization parameter, 10-18 MEMOPTIMIZE_WRITE hint, 10-14 Memoptimized Rowstore about, 10-9 fast ingest about, 10-9 disabling, 10-15 enabling, 10-14 prerequisites, 10-12	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4 advisors, 1-5 Performance page, 1-5 Oracle Forms control of parsing and private SQL areas, 12-8 Oracle Managed Files, 15-8 Oracle Orion calibration tool parameters, 15-16 command-line options, 15-16 Oracle performance improvement method, 3-1
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7 MEMOPTIMIZE FOR WRITE clause, 10-14 memoptimize pool, 10-18 MEMOPTIMIZE_POOL_SIZE initialization parameter, 10-18 MEMOPTIMIZE_WRITE hint, 10-14 Memoptimized Rowstore about, 10-9 fast ingest about, 10-9 disabling, 10-15 enabling, 10-14 prerequisites, 10-12 using, 10-14	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4 advisors, 1-5 Performance page, 1-5 Oracle Forms control of parsing and private SQL areas,
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7 MEMOPTIMIZE FOR WRITE clause, 10-14 memoptimize pool, 10-18 MEMOPTIMIZE_POOL_SIZE initialization parameter, 10-18 MEMOPTIMIZE_WRITE hint, 10-14 Memoptimized Rowstore about, 10-9 fast ingest about, 10-9 disabling, 10-15 enabling, 10-14 prerequisites, 10-12 using, 10-14 fast lookup	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4 advisors, 1-5 Performance page, 1-5 Oracle Forms control of parsing and private SQL areas, 12-8 Oracle Managed Files, 15-8 Oracle Orion calibration tool parameters, 15-16 command-line options, 15-16 Oracle performance improvement method, 3-1
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7 MEMOPTIMIZE FOR WRITE clause, 10-14 memoptimize pool, 10-18 MEMOPTIMIZE_POOL_SIZE initialization parameter, 10-18 MEMOPTIMIZE_WRITE hint, 10-14 Memoptimized Rowstore about, 10-9 fast ingest about, 10-9 disabling, 10-15 enabling, 10-14 prerequisites, 10-12 using, 10-14 fast lookup about, 10-17	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4 advisors, 1-5 Performance page, 1-5 Oracle Forms control of parsing and private SQL areas, 12-8 Oracle Managed Files, 15-8 Oracle Orion calibration tool parameters, 15-16 command-line options, 15-16 Oracle performance improvement method, 3-1 steps, 3-2
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7 MEMOPTIMIZE FOR WRITE clause, 10-14 memoptimize pool, 10-18 MEMOPTIMIZE_POOL_SIZE initialization parameter, 10-18 MEMOPTIMIZE_WRITE hint, 10-14 Memoptimized Rowstore about, 10-9 fast ingest about, 10-9 disabling, 10-15 enabling, 10-14 prerequisites, 10-12 using, 10-14 fast lookup about, 10-17 enabling for a table, 10-20	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4 advisors, 1-5 Performance page, 1-5 Oracle Forms control of parsing and private SQL areas,
log writer processes tuning, 15-6 LRU aging policy, 11-7 M max session memory statistic, 12-25 MAXOPENCURSORS clause, 12-7 MEMOPTIMIZE FOR WRITE clause, 10-14 memoptimize pool, 10-18 MEMOPTIMIZE_POOL_SIZE initialization parameter, 10-18 MEMOPTIMIZE_WRITE hint, 10-14 Memoptimized Rowstore about, 10-9 fast ingest about, 10-9 disabling, 10-15 enabling, 10-14 prerequisites, 10-12 using, 10-14 fast lookup about, 10-17	OPEN_CURSORS initialization parameter, 4-1 operating system data cache, 16-1 optimization described, 1-4 optimizer, introduction, 1-4 query, 1-4 Oracle Enterprise Manager Cloud Control, 1-4 advisors, 1-5 Performance page, 1-5 Oracle Forms control of parsing and private SQL areas, 12-8 Oracle Managed Files, 15-8 Oracle Orion calibration tool parameters, 15-16 command-line options, 15-16 Oracle performance improvement method, 3-1 steps, 3-2

PARALLEL clause	reducing (continued)
CREATE INDEX statement, 4-8	paging and swapping, 10-6
parameter	REEXECUTE FAILED AUTO TASKS
RESULT_CACHE_MAX_TEMP_RESULT,	Procedure, 7-15
13-10	RELEASE_CURSOR clause, 12-7
RESULT_CACHE_MAX_TEMP_SIZE, 13-10	Remote Management Framework (RMF), 6-28
RESULT_CACHE_MODE, 13-9	resources
parsing	allocation, 2-6, 2-14
hard, 2-13	response time, 2-8
Oracle Forms, 12-8	reverse key indexes, 2-11
Oracle precompilers, 12-7	RMF, 6-28
reducing unnecessary calls, 12-6	rollout strategies
soft, 2-13	big bang approach, 2-19
partitioned indexes, <i>2-11</i>	trickle approach, 2-19
per-session PGA memory limit	trickie approach, 2 13
PGA, 14-19	
performance	S
emergencies, 3-5	and the little of O
improvement method, 3-1	scalability, 2-2
improvement method steps, 3-2	factors preventing, 2-4
mainframe, 16-6	Internet, 2-3
	linear, 2-4
monitoring memory on Windows, 16-9	selectivity
tools for diagnosing and tuning, 1-4	ordering columns in an index, 2-12
tools for performance tuning, 1-4	service hours, 2-9
UNIX-based systems, 16-6	session memory statistic, 12-25
Windows, 16-6	SGA_TARGET initialization parameter
Performance Hub active reports	automatic memory management, 10-1
about, 6-43	shared server
generating, 6-43, 6-45	performance issues, 4-9
PGA_AGGREGATE_TARGET initialization	reducing contention, 4-9
parameter, 4-1	tuning, 4-9
physical reads from cache statistic, 11-4	tuning memory, 12-24
proactive monitoring, 1-3	SHARED_POOL_SIZE initialization parameter,
processes	12-15
scheduling, 16-10	SHOW SGA statement, 10-7
Program Global Area	sizing redo logs, 4-4
PGA, 10-3	snapshots
program global area (PGA)	about, 6-2
shared servers, 12-24	soft parsing, 2-13
programming languages, 2-14	software
	components, 2-6
Q	sort areas
<u>~</u>	tuning, <i>14-1</i>
queries	SQL Tuning Advisor, 1-5
data, 2-8	statistics
query optimizer, 1-4	baselines, 6-1
See also optimizer	consistent gets from cache, 11-4
·	databases, 5-1
D	db block gets from cache, 11-4
R	gathering, 5-1
redo logs, 4-4	max session memory, 12-25
mirroring, 15-7	physical reads from cache, 11-4
placement on disk, 15-6	session memory, 12-25
sizing, 4-4	shared server processes, 4-11
· · · · · · · · · · · · · · · · · · ·	time model, 5-1
reducing contention with dispatchers, 4-10	unic model, J-1
CARLETONIA WOLL CLOUDING CO. 4° 10	

U
undo management, automatic mode, 4-3
Unified Memory, 10-1
UNIX system performance, 16-6
user global area (UGA)
shared servers, 4-9, 12-24
user interface, 2-6
users
interaction method, 2-8
interfaces, 2-14
location, 2-8
network speed, 2-8
number of, 2-8
requests, 2-14
response time, 2-8
V
<u> </u>
V\$BUFFER_POOL_STATISTICS view, 11-9
V\$DB_CACHE_ADVICE view, 11-2
V\$JAVA_LIBRARY_CACHE_MEMORY view,
12-13
V\$JAVA_POOL_ADVICE view, 12-13
V\$LIBRARY_CACHE_MEMORY view, 12-12
V\$LIBRARYCACHE view
NAMESPACE column, 12-10
V\$QUEUE view, 4-11
V\$ROWCACHE view
performance statistics, 12-13
V\$SESSTAT view, 12-25
V\$SHARED_POOL_ADVICE view, 12-12
V\$SHARED_POOL_RESERVED view, 12-22
V\$SYSSTAT view
redo buffer allocation, 11-13
using, 11-4
validating designs, 2-18
views, 2-13
vmstat UNIX command, 16-9
W
wait events, 5-3
classes, 5-3
Windows performance, 16-6
workloads, <i>2-17</i> , <i>2-18</i>

